

24.3-8

Let $G = (V, E)$ be a weighted, directed graph with weight function $w : E \rightarrow \{0, 1, \dots, W\}$ for some nonnegative integer W . Modify Dijkstra's algorithm to compute the shortest path weights from a given source vertex $s \in V$ in $O(WV + E)$ time.

DIJKSTRA_ORIGINAL(G, w, s)

Total Running Time = $O(V^2 + E)$ when *unmodified*

```

1  INITIALIZE-SINGLE-SOURCE(  $G, s$  )
2   $S = \emptyset$ 
3   $Q = G.V$ 
4  while  $Q \neq \emptyset$ 
5       $u = \text{EXTRACT-MIN}( Q )$ 
6       $S = S \cup \{ u \}$ 
7      for each vertex  $v \in G.\text{Adj}[ u ]$ 
8          RELAX(  $u, v, w$  )
```

The running time of Dijkstra's algorithm unmodified depends on how we implement the min-priority queue. Consider first the case in which we maintain the min-priority queue by taking advantage of the vertices being numbered 1 to $|V|$. We simply store $v.d$ in the v^{th} entry of an array. Each INSERT and DECREASE-KEY operation takes $O(1)$ time, and each EXTRACT-MIN operation takes $O(V)$ time (since we have to search through the entire array), for a total time of $O(V^2 + E)$.

However, if we provide a modification to this algorithm, we can achieve runtime $O(WV + E)$. As a result of our defined weight function, we know that each edge has at most weight W , so the maximum possible value of the longest shortest-path in the graph is $(V - 1)W$, and also an integer. For the priority queue, we will still prioritize vertices based off of their upper bounds on the weight of a shortest path from the source. Because our longest possible path is $(V - 1)W$ in length, we will need a queue of size $(V - 1)W + 2$ buckets. Vertex v is placed in the $v.d^{\text{th}}$ bucket of the queue. The source is found in the first bucket of the priority queue. INITIALIZE-SINGLE-SOURCE will have all vertices $v.d = \infty$, excluding the source, and will have the last bucket in the queue reserved for undiscovered vertices. In what will take $O(VW)$ time, we will need to scan the buckets in increasing order, until we encounter a non-empty bucket. The first vertex in the bucket is removed, using the EXTRACT-MIN method. All adjacent vertices are then relaxed. The relaxation still happens over E edges. Therefore our final run time is $O(WV + E)$.

24-1

Suppose that we order the edge relaxations in each pass of the Bellman-Ford algorithm as follows. Before the first pass, we assign an arbitrary linear order $v_1, v_2, \dots, v_{|V|}$ to the vertices of the input graph $G = (V, E)$. Then, we partition the edge set E into $E_f \cup E_b$, where $E_f = \{(v_i, v_j) \in E : i < j\}$ and $E_b = \{(v_i, v_j) \in E : i > j\}$. Assume that G contains no self-loops, so that every edge is in either E_f or E_b . Define $G_f = (V, E_f)$ and $G_b = (V, E_b)$.

24-1A

Proof by contradiction for G_f : Let us consider that G_f does indeed have a cycle. A cycle would indicate that at least one edge (u,v) would have u with a higher index than v . However, $[v_1, v_2, \dots, v_{|V|}]$ is a topological sort for G_f because from the definition of E_f , the edges are directed from smaller to larger indices. In other words, this edge would have $\{(u_i, v_j) \in E : i > j\}$ which is in contradiction to the definition of E_f which is the edge set for G_f . Therefore, we have a contradiction with the definition of G_f , so the cyclic edge that causes contradiction must not exist. Thus G_f is acyclic.

Proof by contradiction for G_b : Let us consider that G_b does indeed have a cycle. A cycle would indicate that at least one edge (u,v) would have u with a lower index than v . However, $[v_{|V|}, v_{|V|-1}, \dots, v_1]$ is a topological sort for G_b because from the definition of E_b , the edges are directed from smaller to larger indices. In other words, this edge would have $\{(u_i, v_j) \in E : i < j\}$ which is in contradiction to the definition of E_b which is the edge set for G_b . Therefore, we have a contradiction with the definition of G_b , so the cyclic edge that causes contradiction must not exist. Thus G_b is acyclic.

24-1B

Suppose that we implement each pass of the Bellman-Ford algorithm in the following way. We visit each vertex in the order $v_1, v_2, \dots, v_{|V|}$, relaxing edges of E_f that leave the vertex. We then visit each vertex in the order $v_{|V|}, v_{|V|-1}, \dots, v_1$, relaxing edges of E_b that leave the vertex.

Assume edges are partitioned into two directed cyclic graphs E_f and E_b . E_f consists of edges that go from lower number vertex to higher $[v_1, v_2, \dots, v_{|V|}]$ and E_b is vice versa $[v_{|V|}, v_{|V|-1}, \dots, v_1]$. First loop through the edges in E_f in topological ordering of E_f . Then loop through the edges in E_b in topological ordering of E_b . In the first part of the loop, any path that is part of the shortest path tree that starts at a vertex with correct labeling and only goes through edges of E_f becomes labeled correctly. In the second part of the loop, any path that starts at a vertex with correct labeling and only goes through edges of E_b becomes labeled correctly. Thus, total number of iterations is the number of E_f - E_b alternations on any given path. This value is at most $|V|/2$ and at this point $v.d = \delta(s,v)$ for all $v \in V$.

24-1C

We would expect to see no improvement to the asymptotic runtime of the Bellman-Ford algorithm attributed by this scheme because even with performing only an upper bound of $|V|/2$ passes, the improvement will still have $O(V)$ passes. The original Bellman-Ford algorithm of $|V|-1$ passes also takes $O(V)$ time to complete. Each pass still takes $\Theta(E)$ time, so the total asymptotic runtime for Yen's improvement and the Bellman-Ford algorithm remains at $O(VE)$.