

Return-Oriented Programming

Andy Podgurski

Electrical Engineering & Computer Science
Department
Case Western Reserve University

Overview

- ❑ Return-oriented programming (ROP) is a technique that allows an attacker to induce arbitrary behavior in a program whose control flow he has diverted.
- ❑ ROP does *not* involve injecting code.
- ❑ It *links* code snippets *already present* in memory.
 - Each snippet ends in a **ret** instruction.
- ❑ ROP *defeats* the **W \oplus X** protections deployed by Microsoft, Intel, and AMD.
- ❑ It is readily exploitable on *multiple architectures*.
- ❑ It *bypasses* security measures that seek to prevent execution of malicious code.

Background

- For 20+ years, security research has focused on preventing introduction and execution of new **malicious code**:
 - Techniques to guarantee the **integrity of control flow** in existing programs, e.g.,
 - type-safe languages, stack cookies, CFI
 - Techniques to **isolate “bad” code** introduced into the system, e.g.,
 - $W\oplus X$, memory tainting, virus scanners, trusted computing

Example: $W\oplus X$

- ❑ Memory is marked as either **writable** or **executable**, never both.
- ❑ Thus an adversary can't **inject data** into a process and then **jump** to that memory.
 - Raises a **processor exception**
- ❑ ROP **categorically evades** $W\oplus X$ protections.
 - The snippets are in memory marked **executable**.

Gadgets

- ❑ The organizational unit of a ROP attack is the **gadget**:
 - An arrangement of **words on the stack**:
 - ❑ **pointers to instruction sequences**
 - ❑ **immediate data words**
 - Accomplishes some **well-defined task**
 - ❑ e.g., load, xor, conditional branch
- ❑ By assembling a **Turing-complete** collection of gadgets, an attacker can synthesize **any** malicious behavior.
- ❑ Roemer et al show how to build such gadgets using snippets from the **Standard C Library** on Linux and Solaris.

Background: Return-to-libc Attacks

- ❑ Attackers responded to code injection defenses by **reusing code** already present in the target process image.
- ❑ The standard C library, **libc**, was the usual target.
 - libc contains **wrappers** for **system calls**.
 - This is called a **return-to-libc attack**.
 - McDonald showed how libc calls could be **chained** together.
 - This is **sufficient to defeat W \oplus X** on machines **without immutable memory protections**.
- ❑ In principle, **any available code** could be used.
- ❑ ROP **generalizes** return-to-libc to allow arbitrary computation **without** calling any functions.

ROP Principles (for x86)

- ❑ A RO program involves a particular **layout** of the **stack segment**. (Fig. 3)
- ❑ Each RO instruction is a **stack word**, pointing to an instruction sequence somewhere in memory.
- ❑ The **stack pointer %esp** governs what RO instruction sequence is **fetches next**:
 - Execution of **ret** induces **fetch-and-decode**:
 1. The target of **%esp** is read and used as the new value for **%eip** (instruction pointer).
 2. **%esp** is incremented to point to the next word on the stack.
 - If the next instruction sequence also ends in **ret**, this process is **repeated**.




Fig. 2. Layout of an ordinary program.




Fig. 3. Layout of a return-oriented program.




Fig. 4. Ordinary and return-oriented nop sleds.

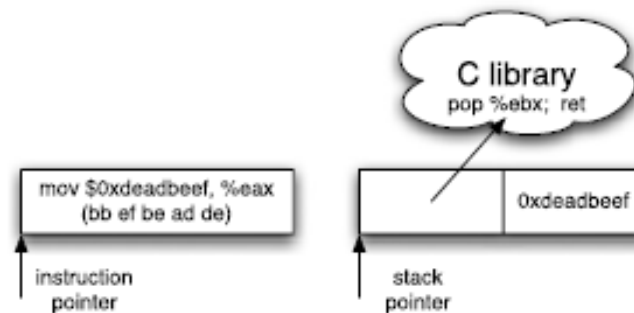


Fig. 5. Ordinary and return-oriented immediates.

No-op Instructions

- ❑ In ROP, a no-op is simply a stack word containing the address of a **ret**.
- ❑ These can be composed to form a **no-op sled**. (Fig. 4)

Encoding Immediate Constants

- In ROP, instructions encoding immediate constants can be approximated using a **pop *reg*** instruction
 - e.g., **pop %ebx; ret** will store the next stack word in **%ebx** and advance the stack pointer past it. (Fig. 5)

Control Flow

- ❑ In ROP, control flow is affected by perturbing the stack pointer **%esp**.
- ❑ Unconditional jumps: **pop %esp; ret** may be used
 - This is a form of immediate load. (Fig. 6)
- ❑ Conditional and indirect jumps are more difficult to implement.




Fig. 6. Ordinary and return-oriented direct jumps.




Fig. 7. A memory-load gadget.

Use of Gadgets

- ❑ Often, more than one instruction sequence is needed to encode a logical operation.
 - e.g., loading a value from memory may require first reading its address into a register, then reading memory.
- ❑ A gadget may include **multiple sequence pointers** and immediate values. (Fig. 7)
- ❑ They act like a return-oriented **instruction set**.

ROP Exploitation

- ❑ A **RO program** is one or more gadgets arranged to carry out an attack.
- ❑ The **payload** containing it is placed in the target's program's memory.
- ❑ The **stack pointer** must be **redirected** to point to the first gadget.
- ❑ The easiest way to accomplish these steps is a **BOF** on the stack.
 - The gadgets are placed so the first has overwritten the saved return address.
 - When the active function tries to return, the RO program is executed.
- ❑ The ROP payload could also be on the **heap**.
- ❑ The attacker could trigger its execution by **overwriting a function pointer** with the address of a snippet that sets **%esp** to the address of the first gadget and returns.

ROP Exploitation cont. (2)

- ❑ The gadgets need not be placed contiguously.
- ❑ With control flow gadgets, an attack can transfer control **between the stack and heap**.

Mitigations

- ❑ Address-space layout randomization (ASLR) defeats attacks that require knowledge of addresses in the target program image.
- ❑ It applies to code injection and code reuse attacks equally.
- ❑ Control-flow integrity systems also can prevent ROP attacks.

Sources

- ❑ R. Roemer et al, *Return-Oriented Programming: Systems, Languages, and Applications*, ACM TISSEC 15, 1, 2012.