

Chapter 3: Processes

Coverage:

- Process Concept
- Process Scheduling
- Operations on Processes
- Cooperating Processes
- Inter-process Communication

What is a Process?

- An operating system executes a variety of programs:
 - Batch system – jobs (different than processes)
 - Time-shared systems – user programs or tasks
- Process – a program in execution; process execution must progress in sequential fashion.
 - Textbook uses the terms *job* and *process* “almost” interchangeably.

What is a Process?

■ A process includes:

- instructions (code text section)
- program counter, other registers
- stack
- data section (static) and heap


```
#include <stdio.h>

main()
{
    FILE *file;
    int value;

    if ((file=fopen("abc", "r")) != NULL) exit(1);
    if (fread((char*)&value, sizeof(int), 1, file) == 1)
        printf("%d\n", value);
    fclose(file);
}
```

Opening and reading a file
via a stream. Read man pages [here](#)
and [here](#).

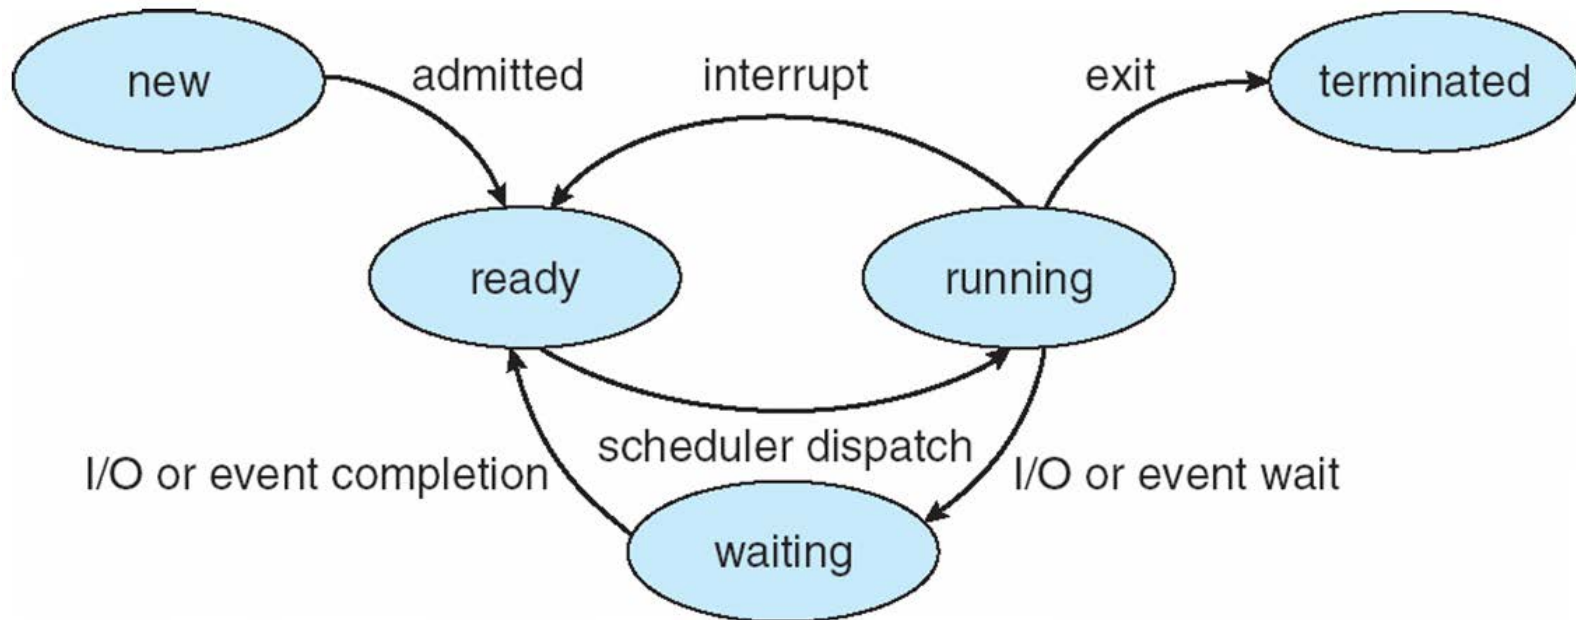
Process in Memory



Process State

- As a process executes, it changes **state**.
 - **new**: The process is being created.
 - **running**: Instructions are being executed.
 - **waiting**: The process is waiting for some event to occur.
 - **ready**: The process is waiting to be assigned to a processor.
 - **terminated**: The process has finished execution.

Diagram of Process State



Process Control Block (PCB)

Information associated with each process
(also called **Task Control Block (TCB)**):

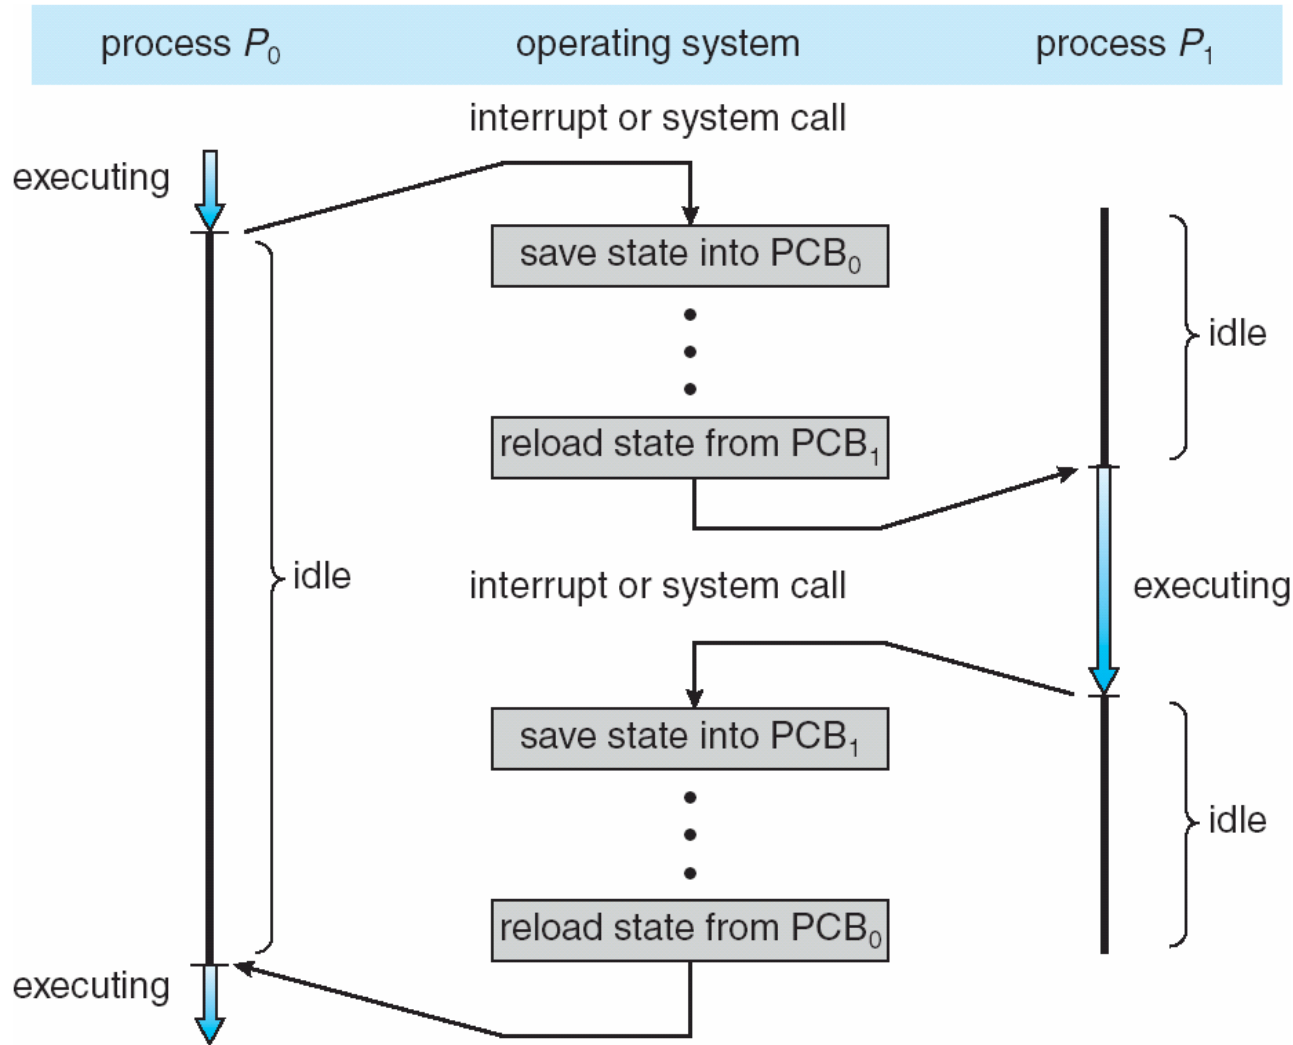
- Process state – running, waiting, etc.
- Program counter – location of instruction to next execute.
- CPU registers – contents of all process-centric registers.
- CPU scheduling information- priorities, scheduling queue pointers.
- Memory-management information – memory allocated to the process.
- Accounting information – CPU used, clock time elapsed since start, time limits.
- I/O status information – I/O devices allocated to process, list of open files.



Context Switch

- CPU switches to another process; OS kernel:
 - Save the state of the old process;
 - Load the saved state for the new process.
PCB represents the context.
- Context-switch time is overhead; no useful work while switching. Up to thousands of instructions run during a single process context switch.
 - Store/load registers
 - Update the PCB
 - Run the scheduling algorithm
 - Maintain the queues
 - Update process profiles, etc
- Time: dependent on hardware support;
Special hardware support for context switch.

CPU Switch From Process to Process



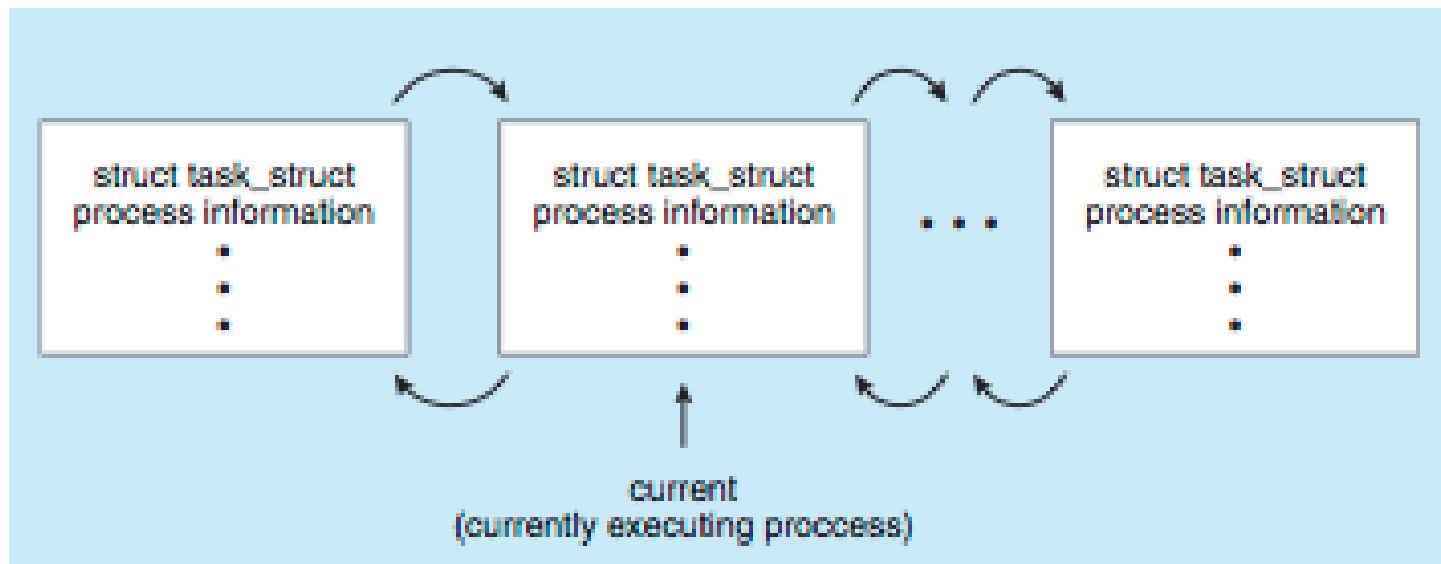
Threads

- So far, process has a single thread of execution.
- Consider having multiple program counters per process.
 - Multiple locations can execute at once.
 - ▶ Multiple threads of control -> **threads**
- Must then have storage for thread details, multiple program counters in PCB.
- See chapter 4.

Process Representation in Linux

- Represented by the C structure `task_struct`


```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



Process Scheduling

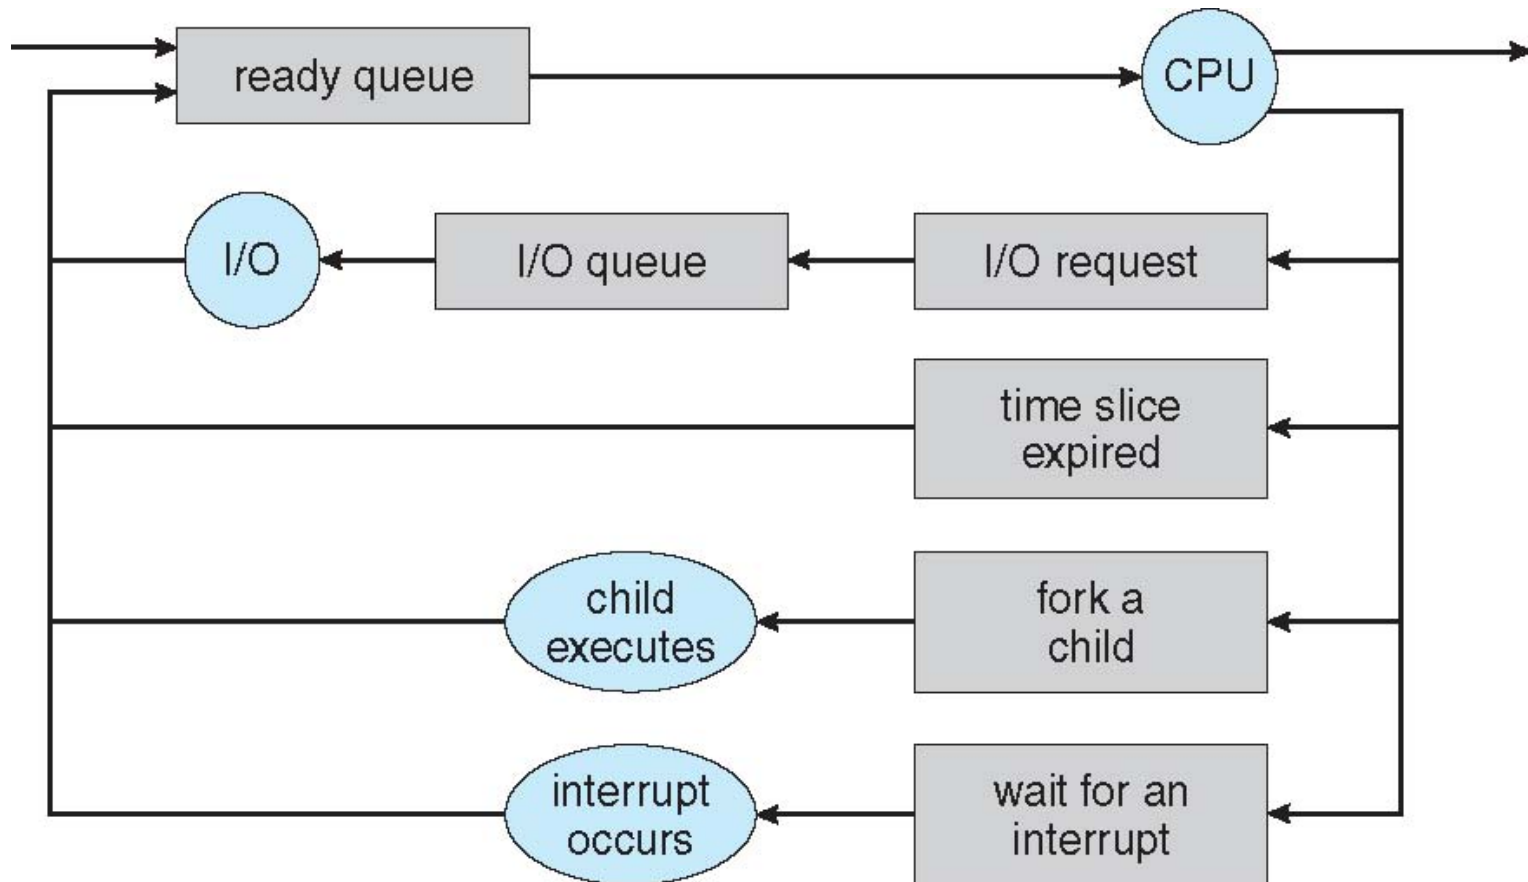
- Maximize CPU use, quickly switch processes onto CPU for time sharing.
- **Process scheduler** selects among available processes for next execution on CPU.
- Maintains **scheduling queues** of processes.
 - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute.
 - **Device queues** – set of processes waiting for an I/O device.
 - Processes migrate among the various queues.

Ready Queue And Various I/O Device Queues



Representation of Process Scheduling

Queueing diagram represents queues, resources, flows.



Schedulers

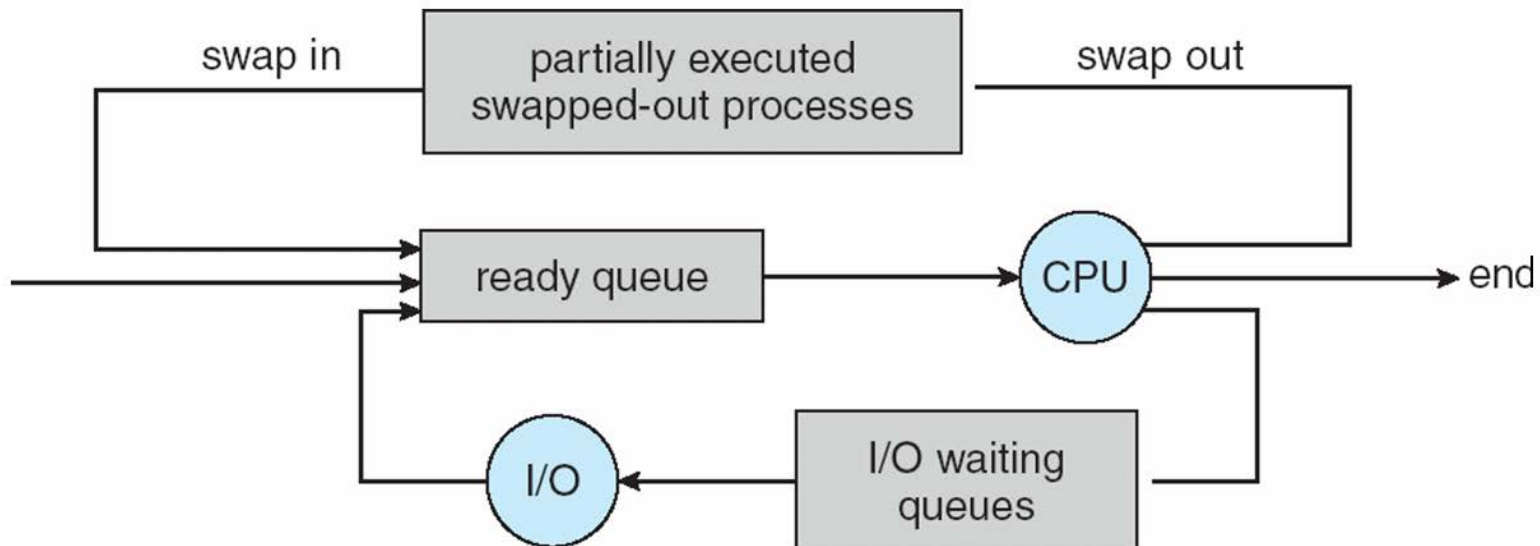
- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue.
 - Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow).
 - The long-term scheduler controls the *degree of multiprogramming*.
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU.
 - Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast).

A process can be described as either:

- **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts.
- **CPU-bound process** – spends more time doing computations; few very long CPU bursts.

Addition of Medium Term Scheduling

- **Medium-term scheduler** can be added if degree of multiple programming needs to decrease.
 - Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**.



Multitasking in Mobile Systems

- Early systems allowed only one process to run, others suspended.
- Due to screen real estate and user interface limits, iOS has a
 - Single **foreground** process- controlled via user interface.
 - Multiple **background** processes— in memory, running, but not on the display, and with limits.
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback.
- Android (Unix) runs foreground and background, with fewer limits.
 - Background process uses a **service** to perform tasks.
 - Service can keep running even if background process is suspended.
 - Service has no user interface, and small memory use.

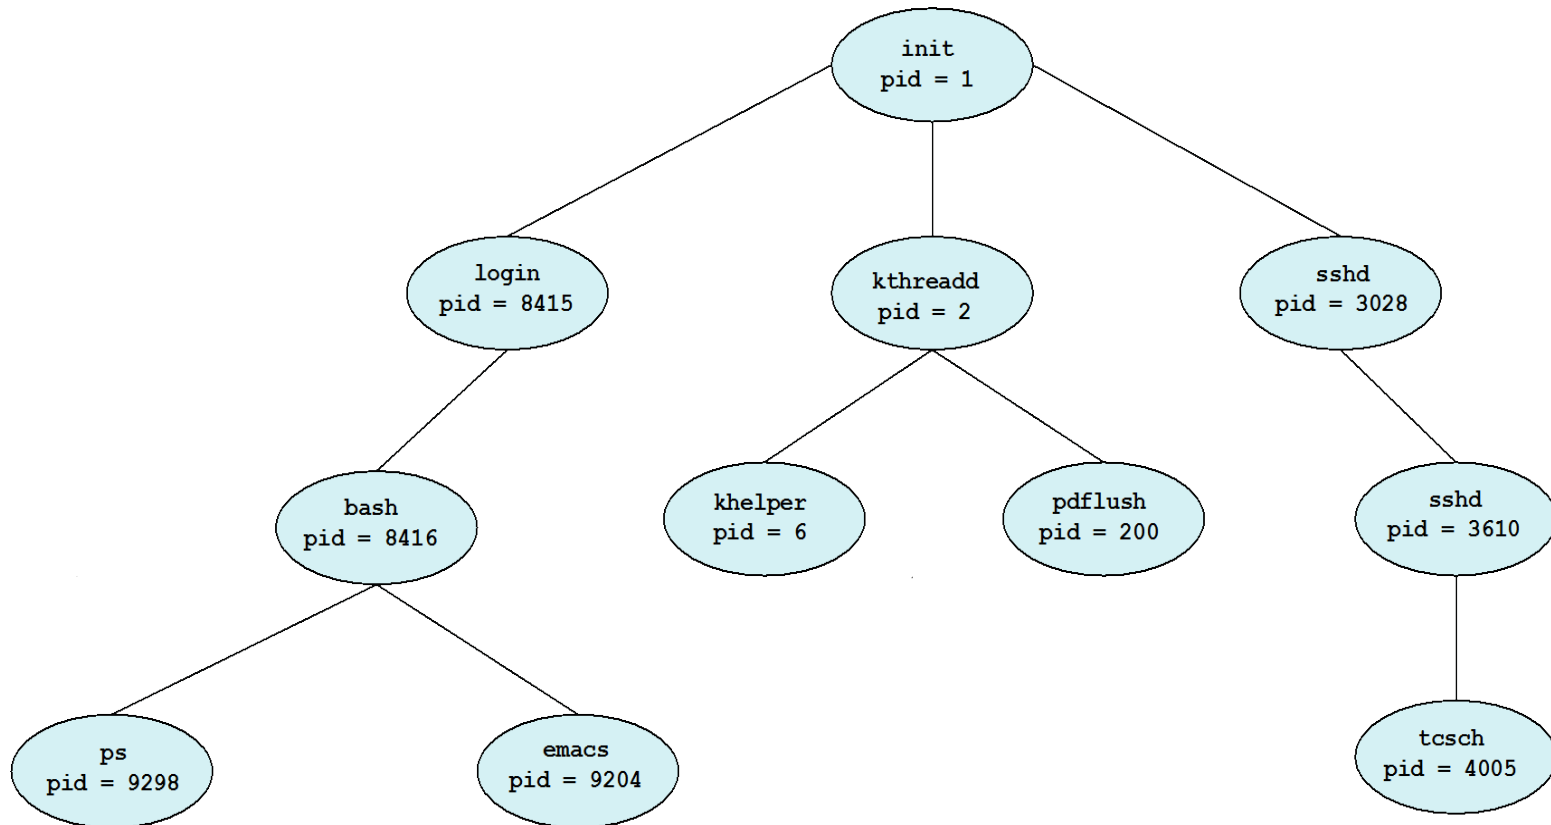
Operations on Processes

- System must provide mechanisms for process creation, termination, and so on as detailed next.

Process Creation

- **Parent** process creates **children** processes, which, in turn create other processes, forming a **tree** of processes.
- Generally, process identified and managed via a **process identifier (pid)**.
- Resource sharing options:
 - Parent and children share all resources.
 - Children share subset of parent's resources.
 - Parent and child share no resources.
- Execution options:
 - Parent and children execute concurrently.
 - Parent waits until children terminate.

A Tree of Processes in Linux



Process Tree vs Precedence Graph

Process Tree: Captures process creation (parent-child) history among processes.

Precedence Graph: Captures fork/join precedences among processes over a single address space.
Nodes: Compound statements.
Edges: Precedence constraints.

Process Creation

- Parent process creates children processes, which, in turn create other processes, forming a tree of processes--**process tree** (versus **precedence graph**).

- e.g., shell program

- Process ID

```
sxj63@volatile:~$ ps -l
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	5618	22326	22325	0	75	0	-	809	wait	pts/1	00:00:00	bash
0	R	5618	22332	22326	0	76	0	-	588	-	pts/1	00:00:00	ps

- Resource sharing

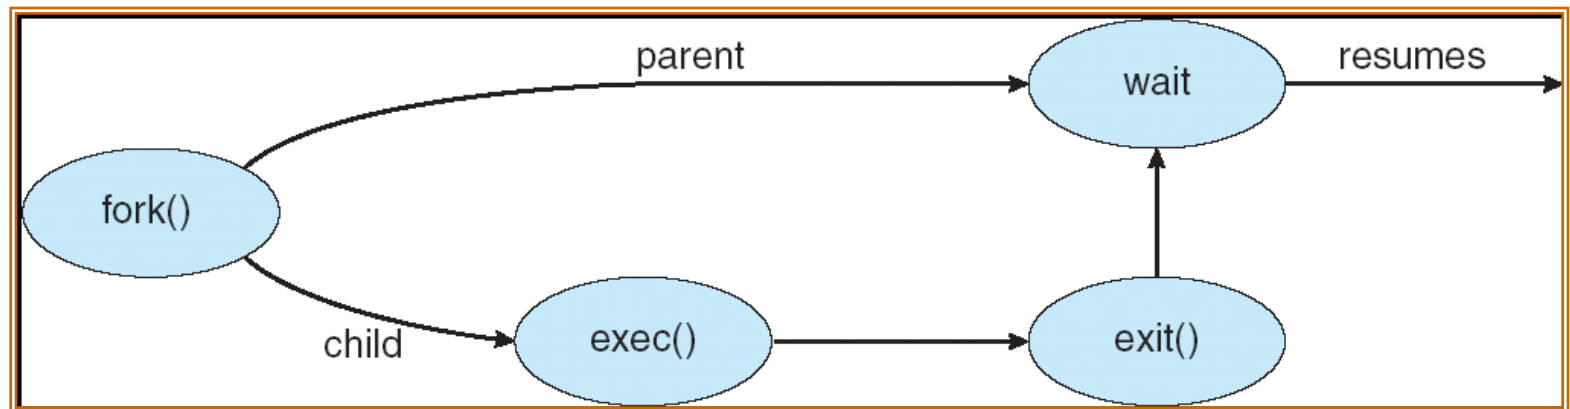
- Parent and children share all resources.
 - Children share subset of parent's resources.
 - Parent and child share no resources.

- Execution

- Parent and children execute concurrently.
 - Parent **wait()** until children terminate.

Process Creation (Cont.)

- Address space
 - Child duplicate of parent.
 - Child has a program loaded into it.
- UNIX examples
 - **fork()** system call creates new process. Alternative: vfork().
 - **exec()** system call (execl(), execlp(), etc) used after a **fork()** to replace the process' memory space with a new program.



An Example of Process Creation

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
```

```
main()
{
```

```
    int pid;
```

```
    pid = fork();
```

```
    if (pid < 0) { /* error occurs */
```

```
        perror("fork() failed");
```

```
        exit(1);
```

```
    } else if (pid == 0) { /* child process */
```

```
        sleep(10);
```

```
        execlp("/bin/sleep", "sleep", "10", NULL);
```

```
    } else { /* parent process */
```

```
        /* parent will wait for the child to complete */
```

```
        wait(NULL);
```

```
        printf("child complete\n");
```

```
        exit(0);
```

```
    }
```

```
}
```

```
sxj63@volatile:~$ ps -l
```

F	S	UID	PID	PPID	C	PRI	NI	ADDR	SZ	WCHAN	TTY	TIME	CMD
0	S	5618	18816	18815	0	75	0	-	811	wait	pts/1	00:00:00	bash
0	S	5618	18983	18816	0	76	0	-	350	wait	pts/1	00:00:00	a.out
0	S	5618	18984	18983	0	77	0	-	449	-	pts/1	00:00:00	sleep
0	R	5618	18985	18816	0	76	0	-	588	-	pts/1	00:00:00	ps

Changing Process Address Space

exec system call: The only way in which a program is executed in Unix is for an existing process to issue an exec system call.

exec() call does not change the pid.

The program invoked by exec() inherits: pid, ppid, cwd, file locks, real uid.

The new program can get a new “effective id”:

If the set-user-id bit of the program being exec'ed is set then the effective uid is changed to the user-id of the owner of the program file.

Execve(filename, argv, envp)

Variations: execve, execlp, execl, execvp, execl:

p: use current PATH for searching executable.

L: list of arguments

v: argv[] vector

e: pass own environment variable list.

Process Termination

- Process executes last instruction, and asks the operating system to terminate it (**exit()**).
 - Output data from child to parent (via **wait()**).
 - Process's resources are reclaimed by operating system.
- Parent may terminate child processes (e.g., by sending a signal to the child processes, signal KILL).
 - Child has exceeded allocated resources.
 - Task assigned to child is no longer required.
 - Parent is exiting.
 - ▶ An operating system may not allow a child to continue if its parent terminates.
 - ▶ Cascading termination.
- *Orphan and Zombie*
 - Orphan: A process whose parent has terminated and who has been adopted by the process **init(1)**.
 - Zombie (defunct): A process that has terminated and whose parent has not yet received notification of its termination.

Orphan and Zombie

```
#include <stdio.h>
#include <stdlib.h>

main()
{
    int pid;

    pid = fork();
    if (pid < 0) { /* error occurs */
        exit(1);
    }
    else if (pid == 0) { /* child process */
        sleep(10);
    }
    else { /* parent process */
        exit(0);
    }
}
```

```
#include <stdio.h>
#include <stdlib.h>

main()
{
    int pid;

    pid = fork();
    if (pid < 0) { /* error occurs */
        exit(1);
    }
    else if (pid == 0) { /* child process */
        exit(0);
    }
    else { /* parent process */
        sleep(10);
    }
}
```

C Program in Unix: Forking Separate Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```

Creating a Separate Process via Windows API

```
#include <stdio.h>
#include <windows.h>

int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;

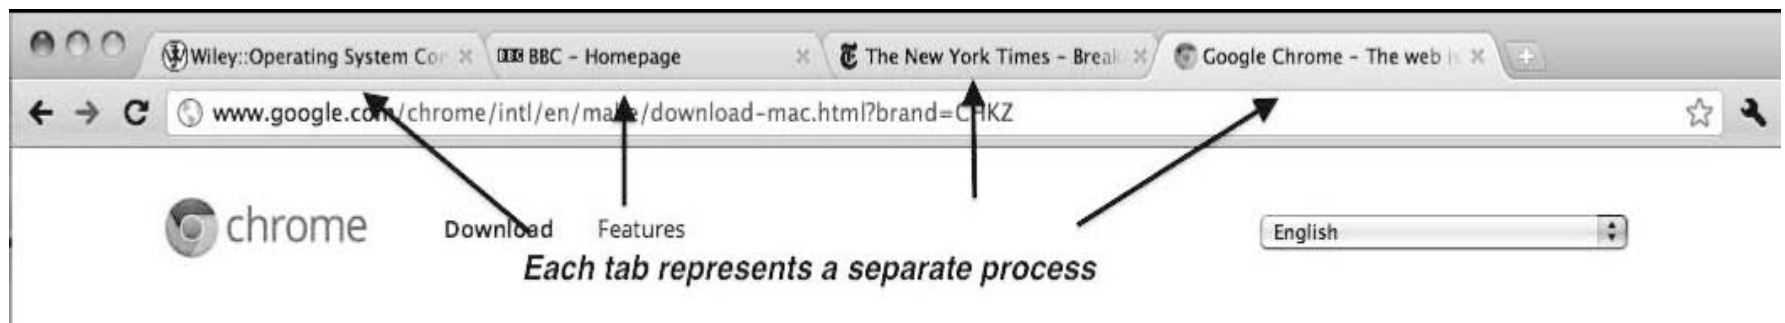
    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));

    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");

    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

Multiprocess Browser Architectures

- Many web browsers ran as single process (some still do).
 - If one web site causes trouble, entire browser can hang or crash.
- Google Chrome Browser is multi-process with 3 categories.
 - **Browser** process manages user interface, disk and network I/O.
 - **Renderer** process renders web pages, deals with HTML, Javascript, new one for each website opened.
 - ▶ Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits.
 - **Plug-in** process for each type of plug-in.



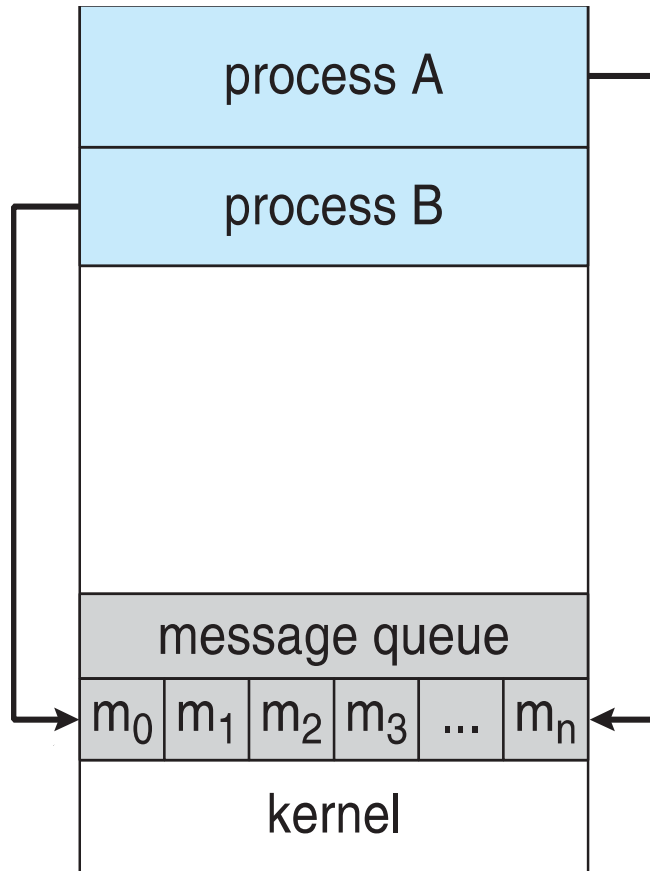
Cooperating Processes

- At the beginning of the course, we saw limited cooperation between parent and child processes
 - `wait()`, `signal()`, resource sharing on process creation
 - But these are not enough.
Example: backend database server process to provide answers to clients.


Interprocess Communication

- Processes within a system may be *independent* or *cooperating*.
- Cooperating process can affect or be affected by other processes, including sharing data.
- Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes need **interprocess communication (IPC)**.
- Two models of IPC
 - **Shared memory**
 - **Message passing**

Interprocess Communication Models



(a)



(b)

Producer-Consumer Problem

- We have seen one paradigm for cooperating processes:
Producer-Consumer model:
 - *producer* process produces information that is consumed by a *consumer* process.
 - **unbounded-buffer** places no practical limit on the size of the buffer.
 - **bounded-buffer** assumes that there is a fixed buffer size.

Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions.
- Message system – processes communicate with each other without resorting to shared variables.
- IPC facility provides two operations:
 - **send**(*message*) – message size fixed or variable
 - **receive**(*message*)

Interprocess Communication – Message Passing

- If P and Q wish to communicate, they need to:
 - establish a ***communication link*** between them.
 - exchange messages via send/receive.

- Implementation of communication link
 - physical (e.g., shared memory, hardware bus).
 - logical (e.g., direct or indirect, synchronous or asynchronous, automatic or explicit buffering).

Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message that the link can accommodate fixed or variable?
- Is a link unidirectional or bi-directional?

Synchronization

- Message passing may be either blocking or non-blocking.
- **Blocking** is considered **synchronous**.
 - **Blocking send** has the sender block until the message is received.
 - **Blocking receive** has the receiver block until a message is available.
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** has the sender send the message and continue.
 - **Non-blocking receive** has the receiver receive a valid message or null.

The rest of the slides are FYI only, and are already covered in depth in recitations.

Note: Your next two assignments will need this information.

Examples of IPC Systems - POSIX

■ POSIX Shared Memory

- Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR,  
0666);
```

- Also used to open an existing segment to share it

- Set the size of the object

```
ftruncate(shm_fd, 4096);
```

- Now the process could write to the shared memory

```
sprintf(shared_memory, "Writing to shared  
memory");
```

IPC POSIX Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0);
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);

    return 0;
}
```

IPC POSIX Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

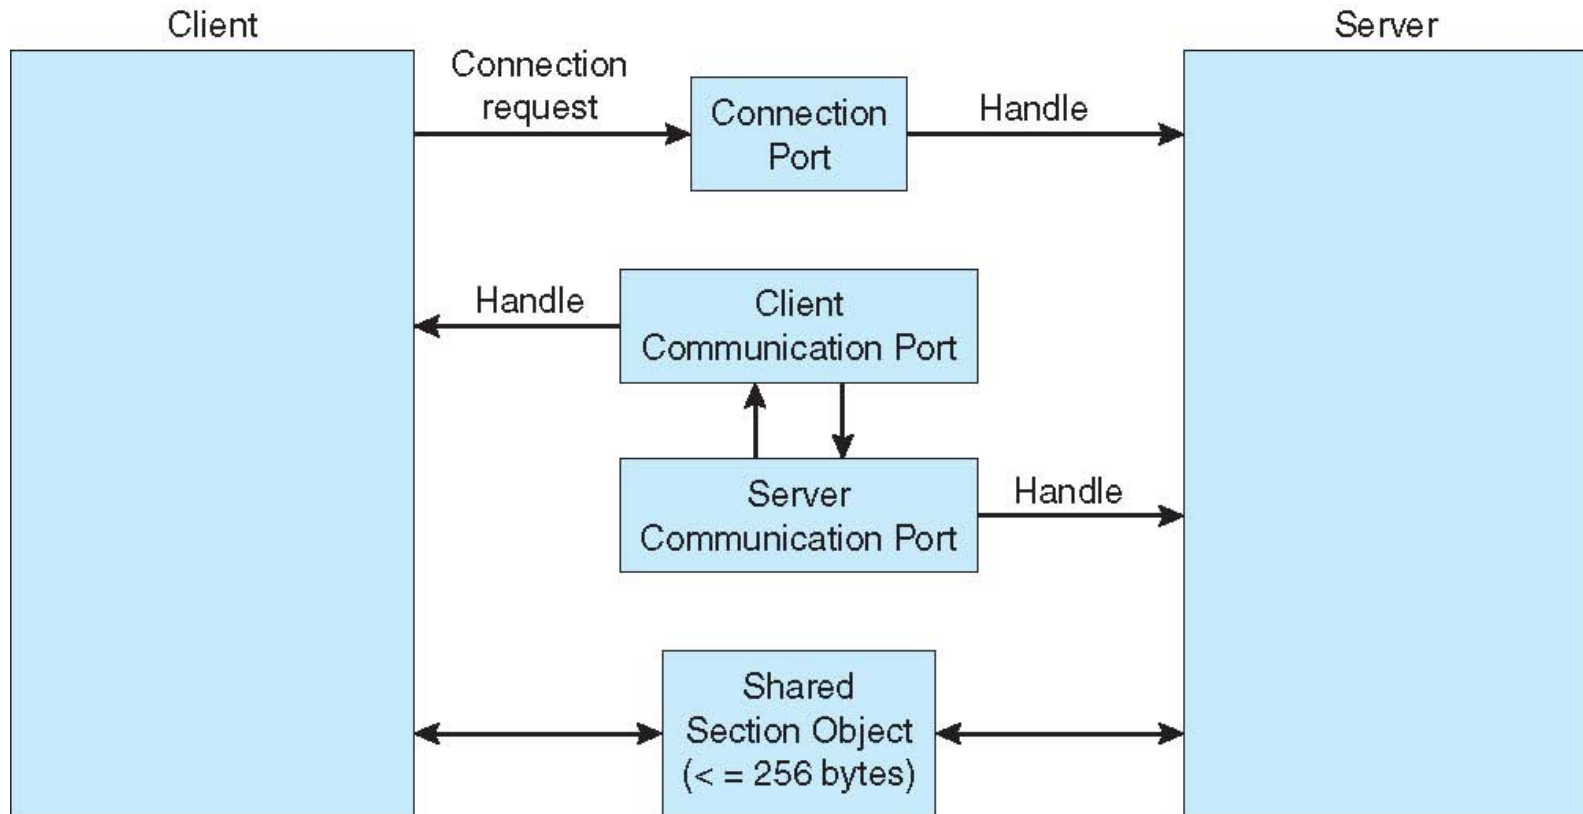
Examples of IPC Systems - Mach

- Mach communication is message-based.
 - Even system calls are messages.
 - Each task gets two mailboxes at creation- Kernel and Notify
 - Only three system calls needed for message transfer.
`msg_send()`, `msg_receive()`, `msg_rpc()`
 - Mailboxes needed for communication, created via
`port_allocate()`
 - Send and receive are flexible, for example four options if mailbox full:
 - ▶ Wait indefinitely.
 - ▶ Wait at most n milliseconds.
 - ▶ Return immediately.
 - ▶ Temporarily cache a message.

Examples of IPC Systems – Windows

- Message-passing centric via **advanced local procedure call (LPC)** facility.
 - Only works between processes on the same system.
 - Uses ports (like mailboxes) to establish and maintain communication channels.
 - Communication works as follows:
 - ▶ The client opens a handle to the subsystem's **connection port** object.
 - ▶ The client sends a connection request.
 - ▶ The server creates two private **communication ports** and returns the handle to one of them to the client.
 - ▶ The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.

Local Procedure Calls in Windows XP




Communications in Client-Server Systems

- Sockets (Used to be covered in this course; now covered in EECS 325).
- Remote Procedure Calls (will see in Week 12 approximately, and will have an assignment).
- Pipes (you have already seen in recitations; may have an assignment).
- Remote Method Invocation (Java).

Ordinary Pipes

- Ordinary Pipes allow communication in standard producer-consumer style
 - Producer writes to one end (the **write-end** of the pipe)
 - Consumer reads from the other end (the **read-end** of the pipe)
 - Ordinary pipes are therefore unidirectional
 - Require parent-child relationship between communicating processes.



- Windows calls these **anonymous pipes**
- See Unix and Windows code samples in textbook.

Named Pipes

- Named Pipes are more powerful than ordinary pipes
- Communication is bidirectional
- No parent-child relationship is necessary between the communicating processes
- Several processes can use the named pipe for communication
- Provided on both UNIX and Windows systems