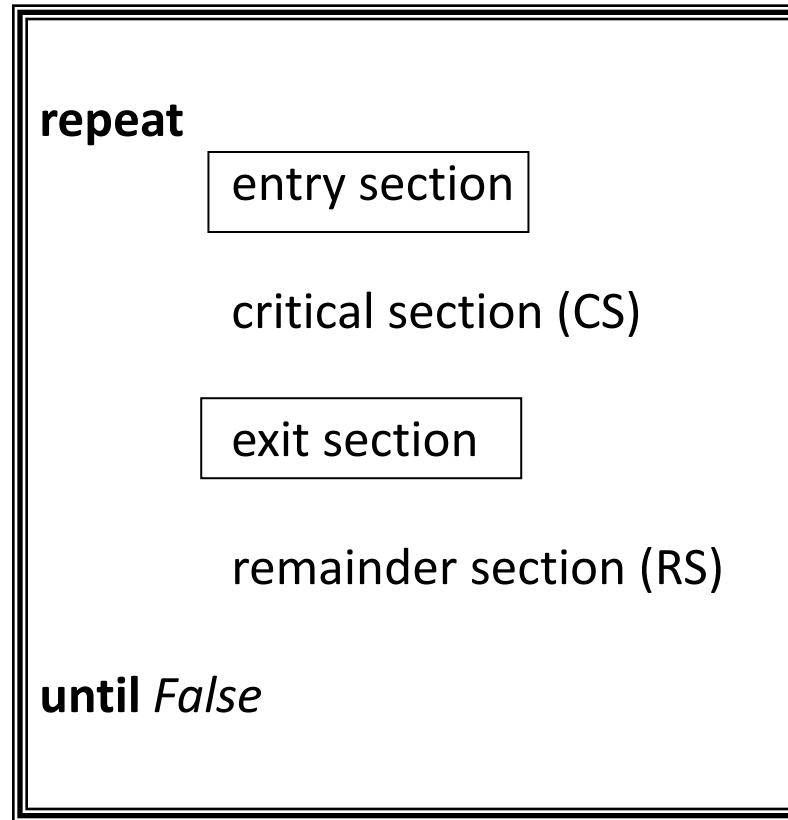


Critical Section in a Process: Environment

- n processes all competing to use some shared variables (data).
- Each process has multiple code segments, each called a ***critical section***, in which the shared data is accessed/modified.
- Problem – Make sure that the three critical section problem properties are always satisfied.

Program Structure (with a Single Critical Section) (simplification for the time being)



General structure of process P_i (*only in this part of slides*)

Two-Process (P_0 and P_1) Synchronization: Algorithm 1

```
int turn;          /* turn is a shared variable; initially, turn = 0
                   turn == i  $\Rightarrow$   $P_i$  can enter its critical section */

 $P_i$  : repeat

    while (turn != i) do no-op;          /* turn=0  $\Rightarrow$   $P_0$  can enter CS
                                         turn=1  $\Rightarrow$   $P_1$  can enter CS */
    critical section

    turn := j;                          /* j=1 - i, 0  $\leq$  i, j  $\leq$  1 */
    remainder section

until false;
```

Correct? Satisfaction of

- (1) mutual exclusion?
- (2) Progress?
- (3) Bounded Waiting?

Two-Process (P_0 and P_1) Synchronization: Algorithm 1

```
int turn:=0;  /* j=1 - i, 0 ≤ i, j ≤ 1
```

```
 $P_i$  : repeat
```

```
    while (turn != i) do no-op;
```

```
    CS
```

```
    turn := j;
```

```
    RS
```

```
until false;
```

```
int turn:=0;
```

```
 $P_j$  : repeat
```

```
    while (turn != j) do no-op;
```

```
    CS
```

```
    turn := i;
```

```
    RS
```

```
until false;
```

```

boolean flag[2]; /* shared variables; initially flag[0] = flag[1] = false
                  /* flag[i] == true  $\Rightarrow$  Pi can enter its critical section

Pi : repeat

    while ( flag[j] ) do no-op;
    flag[i] := True;

    critical section

    flag[i] := False;

    remainder section

until False;

```

5

Two-Process Synchronization: Algorithm 2

```
boolean flag[2]:=False;  
/* flag[i] == true  $\Rightarrow$   
/*      Pi can enter its CS.
```

```
Pi : repeat  
      while ( flag [ j ] ) do no-op;  
      flag[i] := True;  
      CS  
      flag[i] := False;  
      RS  
      until False;
```

```
boolean flag[2]:=False;  
/* flag[j] == true  $\Rightarrow$   
/*      Pj can enter its CS.
```

```
Pj : repeat  
      while ( flag [ i ] ) do no-op;  
      flag[j] := True;  
      CS  
      flag[j] := False;  
      RS  
      until False;
```

Two-Process Synchronization: Algorithm 3

boolean flag[2]; /* shared variables, initially flag[0] = flag[1] = false
/* flag[i] == true \Rightarrow P_i ready to enter its critical section.

P_i : repeat

flag[i] := *True*;
while (flag[j]) **do** *no-op*;

critical section

flag[i] := *False*;

remainder section

until *False*;

**Correct? Satisfaction of (1) Mutual Exclusion? (2) Progress?
(3) Bounded waiting?**

Two-Process Synchronization: Algorithm 3

```
boolean flag[2]:=False;  
/* flag[i] == true  $\Rightarrow$   
/*  $P_i$  ready to enter its CS.
```

```
 $P_i$  : repeat  
    flag[i] := True;  
    while ( flag[ j ] ) do no-op;  
    CS  
    flag[i] := False;  
    RS  
until False;
```

```
boolean flag[2]:=False;  
/* flag[j] == true  $\Rightarrow$   
/*  $P_j$  ready to enter its CS.
```

```
 $P_j$  : repeat  
    flag[ j ] := True;  
    while ( flag[ i ] ) do no-op;  
    CS  
    flag[ j ] := False;  
    RS  
until False;
```


Two-Process Synchronization: Algorithm 4

(Peterson's Algorithm)

```
boolean flag[2]; /* shared variables, initially flag[0] = flag[1] = false  
                  /* flag[i] == true  $\Rightarrow$  Pi ready to enter its critical section.
```

```
int turn;
```

```
Pi: repeat
```

```
    flag[i] := True;  
    turn := j;  
    while (flag[j] && turn == j) do no-op;
```

```
    critical section
```

```
    flag[i] := False;
```

```
    reminder section
```

```
until False;
```

Informal Proofs:

Mutual exclusion

Progress

Bounded waiting

Very fragile solution!

Meets all three requirements; solves the critical-section problem.

Two-Process Synchronization: Algorithm 4

```
boolean flag[2]:=False;
```

```
/* flag[i] == true  $\Rightarrow$ 
```

```
    Pi ready to enter its CS.
```

```
int turn;
```

```
Pi: repeat
```

```
    flag[i] := True;
```

```
    turn := i;
```

```
    while (flag[j] && turn == j)
```

```
        do no-op;
```

```
    CS
```

```
    flag[i] := False;
```

```
    RS
```

```
until False;
```

```
boolean flag[2]:=False;
```

```
/* flag[j] == true  $\Rightarrow$ 
```

```
    Pj ready to enter its CS.
```

```
int turn;
```

```
Pj: repeat
```

```
    flag[j] := True;
```

```
    turn := j;
```

```
    while (flag[i] && turn == i)
```

```
        do no-op;
```

```
    CS
```

```
    flag[j] := False;
```

```
    RS
```

```
until False;
```

Peterson's algorithm

Idea: Three-step solution:

1. Declare intention to enter into the CS.
2. Let others enter into their CSs.
3. Enter the CS.

This is an “example reasoning” for many critical section algorithms.

Where is the fragility of algorithm 4?

Another Algorithm:

Dekker's (much more complicated)

Shared variables: **int** turn; // 0 or 1
 boolean flag[2]; // initially *False*

Pi: **repeat**

```
flag[i] = True;
while (flag[j]) {
    if (turn == j) {
        flag[i] = False;
        while (turn == j);
        flag[i] = True;
    }
}
```

// I would like to enter.
// Does she?
// yes, and it is her turn.
// So I will yield
// waiting for her to leave.
// Finally! She is done.

critical section

```
turn = j;
flag[i] = False
```

// I had my turn.
// Done

remainder section

Proof: Mutual exclusion?
Progress? Bounded waiting?

until *False*;

Review:

So Far, Brute-Force Process Synchronization Algorithms--Review

(no synchronization constructs)

- So far we have seen:
Very simple concurrent program structure, and
two-process synchronization algorithms.
 - Algorithms 1 to 4. Correct one: fourth algorithm.
- Next: the same simple concurrent program structure, and
n-process synchronization algorithm.

N-process Synchronization: Bakery Algorithm (L. Lamport; 1974)

Program structure: As in the two process structure, except that there are now n processes, P_i , $0 \leq i \leq (n-1)$.

Solution Logic: Implement a bakery-ticket mechanism (sort of).

Before entering its critical section, a process gets a (ticket) number.

But, multiple processes *can* have the same ticket number!

Holder of the smallest number enters the critical section.

If two processes P_i and P_j receive the same number and $i < j$,
then P_i enters into its CS first; else P_j enters its CS first.

N-process Synchronization: Bakery Algorithm (L. Lamport; 1974)

P_i: repeat

```
    choosing[i] := True;  
    number[i] := max(number [0], number[1], ..., number [n - 1]) + 1;  
    choosing[i] := False;  
    for (j = 0; j < n; j++) {  
        while (choosing [ j ]) do no-op;  
        while (number [ j ] != 0 && (number [ j ], j) < (number [ i ], i))  
            do no-op ;  
    }
```

critical section

```
number [ i ] := 0;
```

remainder section

until *False*;

N-process Synchronization: Bakery Algorithm

Shared data:

boolean choosing[n];

int number[n];

Initially, all choosing values are *False*;

All number values are 0.

Notation: lexicographical order (ticket #, process id #)

$(a, b) < (c, d)$ **if** $a < c$ **or** $(a = c \text{ and } b < d)$

Bakery Algorithm: Proof of Mutual Exclusion

■ Observation: If

- (a) P_i is already in critical session, and
- (b) any P_k , $k \neq i$, has chosen its number (i.e., $\text{number}[k] \neq 0$),
then $(\text{number}[i], i) < (\text{number}[k], k)$

Proof:

- Case 1: P_k *hasn't started to choose* $\text{number}[k]$ when P_i executes `while(choosing[k])`. Obviously, $\text{number}[i] < \text{number}[k]$.
- Case 2: P_k *has chosen* $\text{number}[k] \neq 0$ before P_i executes `"while(choosing[k])..."`.
 P_i wins, implying $(\text{number}[i], i) < (\text{number}[k], k)$
- Case 3: P_k *is choosing* $\text{number}[k]$ when P_i executes `while(choosing[k])`.
 P_i will wait. We will reduce this to Case 2.

- ## ■ Because of this observation, if P_i is already in its CS, no other process can enter into its CS.

Why Do We Need choosing[]?

P_i: repeat

```
choosing[i] := True;  
number[i] := max(number[0], number[1], ..., number [n - 1]) + 1;  
choosing[i] := False;  
for (j = 0; j < n; j++) {  
    while (choosing[j]) do no-op;  
    while (number[j] != 0 && (number[j], j) < (number[i], i))  
        do no-op ;  
}
```

critical section

number[i] := 0;

remainder section

until *False*;

Why Do We Need choosing[]?

- Two processes, initially $\text{number}[0]=\text{number}[1]=0$.

P₀:

evaluate max () function;

... leave CPU...

P₁:

evaluate max () function;

number[1]:=1;

enter critical section

...leave CPU **while in CS...**

..start executing...

number[0] := 1;

while condition is false.

enter critical section → **violation of ME.**

Review: Brute-Force Process Synchronization Algorithms--Review (no synchronization constructs)

- So far, we have seen
 - Very simple program structure and two-process synchronization
 - ▶ Algorithms 1 to 4. Correct one: fourth algorithm.
 - Same simple program structure and n-process synchronization
 - ▶ Bakery Algorithm

Hardware Solutions to Synchronization

- Up until now, we have seen brute-force ways of providing a solution to the critical section (CS) problem—only when the program structure is very limited.
- Next, we will see hardware-based solutions to the CS problem.

“Single CPU Cycle” (atomic) function as part of the hardware instruction set:

```
function Test-and-Set (var target: bool): bool;  
  begin  
    Test-and-Set := target;  
    target := True;  
  end;
```

Mutual Exclusion with Test-and-Set

Shared data:

bool lock := *False*;

P_i: repeat

begin

while (*Test-and-Set* (lock)) **do** no-op;

critical section

lock := *False*;

remainder section

end

until *False*;

Satisfaction of: Mutual exclusion? Bounded waiting? Progress?

Another Hardware Sol'n: Swap

Also a “single CPU cycle” (atomic) procedure as part of the hardware instruction set:

```
Procedure Swap (var a, b: bool)
  var temp: bool;
  begin
    temp := a;
    a := b;
    b := temp;
  end;
```

Mutual Exclusion with Swap

Shared data:

bool lock := *False*;

P_i : repeat

key := *True*;

repeat

Swap(lock, key);

until key = *False*;

/* P_j has its own key variable.

critical section

lock := *False*;

remainder section

until *False*;

Satisfaction of: mutual exclusion? Bounded waiting? Progress?

Next: Three separate process synchronization constructs

- Semaphores
- Critical region or conditional critical region statements—CCRs
--Only briefly. No CCR-related exam/assignment questions.
- Monitors
 - Btw, every Java object is/has a monitor: more on this later.
- NOTE: NO MIXING OF THESE CONSTRUCTS in your algorithms!
 - Your concurrent process algorithm can only use ONE of the three constructs.
 - ▶ If your algorithm uses semaphores, it cannot use CCRs or monitors.
 - ▶ If your algorithm uses CCRs, it cannot use semaphores or monitors.
- NOTE: ALWAYS USE THE SYNTAX SPECIFIED IN SLIDES!!

Binary Semaphores

- Two *atomic operations* on a variable (semaphore) S:
- S below is a binary semaphore (takes on 0 or 1 as value):

wait (S): **while** $S \leq 0$ **do** *no-op*;
 $S := S - 1$;

signal (S): $S := S + 1$; Dijkstra: $V(S)$

- Semaphore S can **only** be **accessed via wait() and signal()**
--two indivisible (i.e., atomic) operations.
- Initialization of S: either 0 or 1. Normally, 1.
- For a nice intro to semaphores, see the [semaphores chapter](#) in a [online book](#).

Critical Section of n Processes

Shared data:

binary semaphore mutex; */* mutex is set to 1 initially.*

P_i: repeat

wait (mutex);

critical section

signal (mutex);

remainder section

Until *False;*

Mutual exclusion? Progress? Bounded waiting?

Problems with Binary Semaphores

Binary semaphore S:

wait (S): **while** $S \leq 0$ **do** *no-op*;
 $S := S - 1$;

signal (S): $S := S + 1$;

1. Busy waiting (i.e., no-op) above (in *wait*(S)) is wasteful.
2. Semaphore S above is binary.

Nonbinary (counting) semaphore operations form a **higher-level synchronization mechanism** that does not require busy waiting.

Eliminating Busy Waiting and Bounded Waiting

type *nonbinary semaphore* = **record**

value: integer;

/ value is usually*

L: **list of process**;

/ set to 1 initially*

end;

wait(S):

S.value := S.value - 1;

if (*S.value* < 0)

then begin

 add this process to *S.L*;

block;

end;

signal(S):

S.value := S.value + 1;

if (*S.value* ≤ 0)

then begin

 remove a process *P* from *S.L*;

wakeup(P);

end;

Two simple operations:

- *block* suspends the process that invokes it.
- *wakeup(P)* resumes the execution of a blocked process *P*.

Producer-Consumer (Bounded Buffer) Problem

Initialization:

binary semaphore *mutex*:=1;

nonbinary semaphore *full*:=0; *empty*:=*n*;

Producer:

repeat

...

Produce an item in nextp;

wait (*empty*); /* When $empty \leq 0$,

wait (*mutex*); /* block producer.

add nextp to buffer;

signal (*mutex*);

signal (*full*);

until *False*;

Consumer:

repeat

wait(*full*); /*When $full \leq 0$; block consumer.

wait (*mutex*);

remove an item from buffer to nextc;

signal(*mutex*);

signal(*empty*);

...

consume the item in nextc;

until *False*;

Proving Correctness:

Invariant I (single-producer and single-consumer):

$$(n - 2) \leq \text{empty} + \text{full} \leq n \text{ and } 0 \leq \text{mutex} \leq 1$$

Safety: Prove that $I \rightarrow \neg P_{\text{bad}}$

Producer-Consumer (Bounded Buffer) Problem

Initialization:

binary semaphore *mutex*:=1;

nonbinary semaphore *full*:=0; *empty*:=*n*;

Producer:

repeat

...

Produce an item in nextp;

wait (empty); /* When $empty \leq 0$,
wait (mutex); /* block producer.

add nextp to buffer;

signal (mutex);
signal (full);

until *False*;

Consumer:

repeat

wait(full); /*When $full \leq 0$; block consumer.
wait (mutex);

remove an item from buffer to nextc;

signal(mutex);
signal(empty);

....

consume the item in nextc;

until *False*;


Note the symmetrical use of empty and full semaphores in entry and exit sections of producer and consumer.

Dining Philosophers Problem (1965)

- 5 philosophers; 5 chopsticks; one bowl of rice.

nonbinary semaphore *chopstick* [0..4]:=1;

/* All chopsticks are initialized to 1.



Dining Philosophers Problem

- A philosopher eats or thinks!
 - Each philosopher needs two chopsticks to eat, and chopsticks are picked up one at a time.
 - After successfully picking up two chopsticks, a philosopher eats for a while, and, then, puts down the chopsticks, and thinks.

- Problem: Decide how philosophers pick the chopsticks so that
 - Philosophers (processes) do not starve!
 - There is no deadlock.

Dining Philosophers Problem

Nonbinary semaphore $chopstick[0..4]:=1;$

Philosopher i :

repeat

wait ($chopstick[i]$);
wait ($chopstick[(i+1) \bmod 5]$);

...

EAT

...

signal ($chopstick[i]$);
signal ($chopstick[(i+1) \bmod 5]$);

...

THINK

...

until *False*;

Problem: Deadlock

Dining Philosophers Problem

Invariant property characterization:

$up[i]$ = count of the times chopstick i has been picked up.

$down[i]$ = count of the times chopstick i has been put down.

A chopstick cannot be put down more times than it has been picked up.

Invariant property:

For all i : $[1 \leq i \leq 5 \rightarrow (down[i] \leq up[i] \leq (down[i] + 1))]$

Note that a chopstick can be picked up by at most one philosopher at a time.

Dining Philosophers Problem—No Deadlocks

nonbinary semaphore *chopstick*[0:4]:=1;

Philosopher *i*:

repeat

wait (*chopstick*[(*i*+1) mod 5]);
wait (*chopstick*[*i*]);

...

eat

...

signal (*chopstick*[*i*]);
signal (*chopstick*[(*i*+1) mod 5]);

...

think

...

until *False*;

**Switched these two
lines for philosopher 5,
and NO deadlocks!**

To Avoid Deadlocks-One Approach

- A number of processes (P_i) want to access an arbitrary number of resources (R_i), needing exclusive access. Processes must abide by the following rules:
 - **Circular wait elimination via ordered resource requests:** Processes request resources in increasing order of resource ids. A philosopher obtains a low-numbered fork first.
 - **Eliminating no preemption:** If a process has obtained a resource R_i and needs resource R_j , it must release R_i before making the request if $i > j$.

Readers-Writers Problem (1971)

Assumption: No reader waits unless a writer is writing! **Writers may starve!**

Initialization:

nonbinary semaphore *mutex:=1; wrt:=1;*

int *readcount:=0;*

Reader:

```
wait (mutex);  
    readcount := readcount + 1;  
    if readcount = 1 then wait (wrt);  
signal (mutex);
```

In CS: Reading!

```
wait (mutex);  
    readcount := readcount - 1;  
    if readcount = 0 then signal (wrt);  
signal (mutex);
```

Writer:

```
wait (wrt);
```

In CS: Writing!

```
signal (wrt);
```

**Soln is NOT
symmetric for
readers and
writers.
Not good.**

Readers-Writers Problem

Invariant: Use two auxiliary variables that characterize the state of concurrent execution.

writing[j] = 1 if writer j is in CS, 0 otherwise.

reading = 1 if **any** reader is in its CS, 0 otherwise.

Initially, reading=writing[j]=0 for all j, $1 \leq j \leq n$

Invariant I: reading + writing[1] + ... + writing[n] ≤ 1

Safety: Prove that $I \rightarrow \neg P_{\text{bad}}$

Readers-Writers Problem-Nobody Starves.

Reading/Writing Enforcement: FCFS

Initialization:

```
nonbinary semaphore mutex:=1; wrt:=1; rw:=1;  
int readcount:=0;
```

Reader:

```
wait(rw);  
wait(mutex);  
    readcount := readcount + 1;  
    if readcount = 1 then wait (wrt);  
signal (mutex);  
signal(rw);
```

...

In CS: Reading!

```
...  
wait (mutex);  
    readcount := readcount - 1;  
    if readcount = 0 then signal (wrt);  
signal (mutex);
```

Writer:

```
wait(rw);  
wait (wrt);
```

In CS: Writing!

```
signal (wrt);  
signal(rw);
```


Yet Another Readers-Writers Problem: Readers Starve!

Specifications:

- If a reader finds that there are writers waiting, the reader must wait until all waiting writers are finished writing.
- Even if writers arrive after the reader does, writers write before the reader reads, i.e., the reader can read only if the writers queue is empty.
- Readers can starve.

Readers-Writers Problem

Give a semaphore-based solution to the Readers-Writers problem where “readers starve” in the sense that a stream of writers can arrive and “write” one after another even when a reader has been waiting to read.

Global variables:

semaphore mutex: initially set to 1 rmutex, wtr, rdr : all initially set to 0
int nreaders = 0, nwriters = 0; boolean Busy = *False*; RBlocked = *False*;

Reader:

```
while True
{ wait (r-mutex);
  wait (mutex);
  if (nwriters > 0)
  then {RBlocked=True;
        signal(mutex);wait(rdr)}
  else {nreaders++;signal(mutex)
};
signal(r-mutex);
READ;
wait(mutex);
nreaders--;
if (nreaders = 0 and
      nwriters > 0)
then {Busy = True; signal(wrt)};
signal(mutex);
DO-SOMETHING;
}
```

Writer:

```
while True
{ wait(mutex);
  nwriters++;
  if ( Busy or nreaders > 0)
  then {signal(mutex); wait(wrt)}
  else { Busy = True;
        signal(mutex)};
  WRITE;
  wait(mutex);
  nwriters--;
  Busy = False;
  if (nwriters > 0)
  then {Busy = True; signal(wrt)}
  else if (RBlocked) then
    {RBlocked = False;
     nreaders++; signal(rdr)};
  signal(mutex);
  DO-SOMETHING;
}
```

Other Problems:

Sleeping Barber Problem

- A barbershop consists of a waiting room with N chairs. If there are no customers to be served, the barber goes to sleep. If a customer enters the barbershop and all chairs are occupied, then the customer leaves the shop. If the barber is busy, but chairs are available, then the customer sits in one of the free chairs. If the barber is asleep, the customer wakes up the barber. Write a program to coordinate the barber and the customers.

- Objectives
 - No starvation: a customer will have his hair cut once he sits down.
 - No deadlocks: both barber and customers could be waiting.

Cigarette-Smokers Problem

Consider a system with three smoker processes and one agent process. Each smoker continuously rolls a cigarette and then smokes it. But to roll and smoke a cigarette, the smoker needs three ingredients: tobacco, paper, and matches. One of the smoker processes has paper, another has tobacco, and the third has matches. The agent has an infinite supply of all three materials. The agent places two of the ingredients on the table. The smoker who has the remaining ingredient then makes and smokes a cigarette, signaling the agent on completion. The agent then puts out another two of the three ingredients, and the cycle repeats. Write a program to synchronize the agent and the smokers.

Exercises

- Three processes, P_0 , P_1 , and P_2 . Write code with semaphores to synchronize them such that P_2 executes only after both P_0 and P_1 execute.
- Two processes, P_0 and P_1 . Write code with semaphores to synchronize them such that P_0 executes at least twice as often as P_1 executes.