



Application Layer Part 4

Mark Allman
Case / ICSI

EECS 325/425
Fall 2018

*“I don't really want to stop the show,
But I thought that you might like to know,
That the singer's going to sing a song,
And he wants you all to sing along.”*

Many of these slides are more-or-less directly from the slide set developed by Jim Kurose and Keith Ross for their book “Computer Networking: A Top Down Approach, 5th edition”.

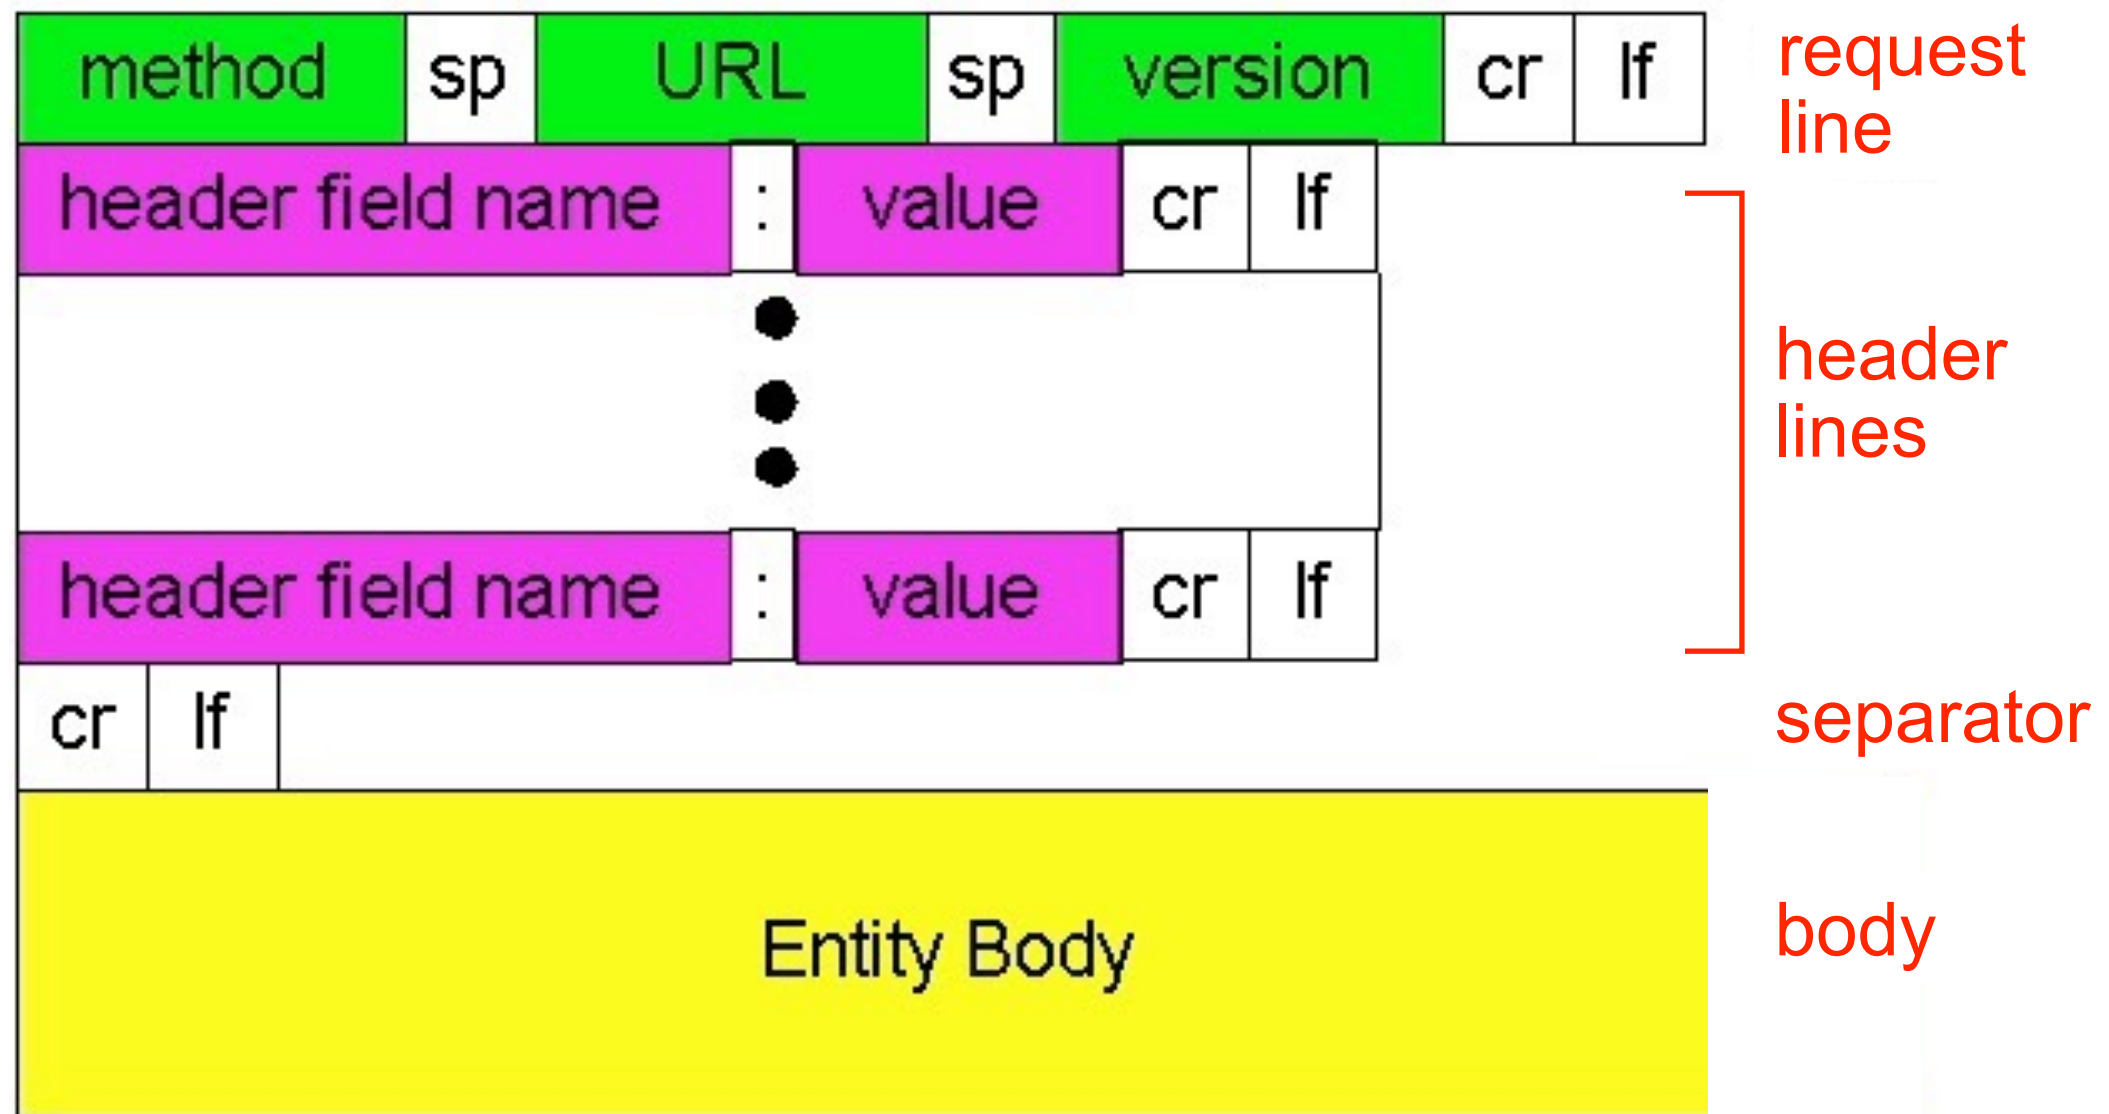
The slides have been lightly adapted for Mark Allman’s EECS 325/425 Computer Networks class at Case Western Reserve University.

All material copyright 1996-2010
J.F Kurose and K.W. Ross, All Rights Reserved

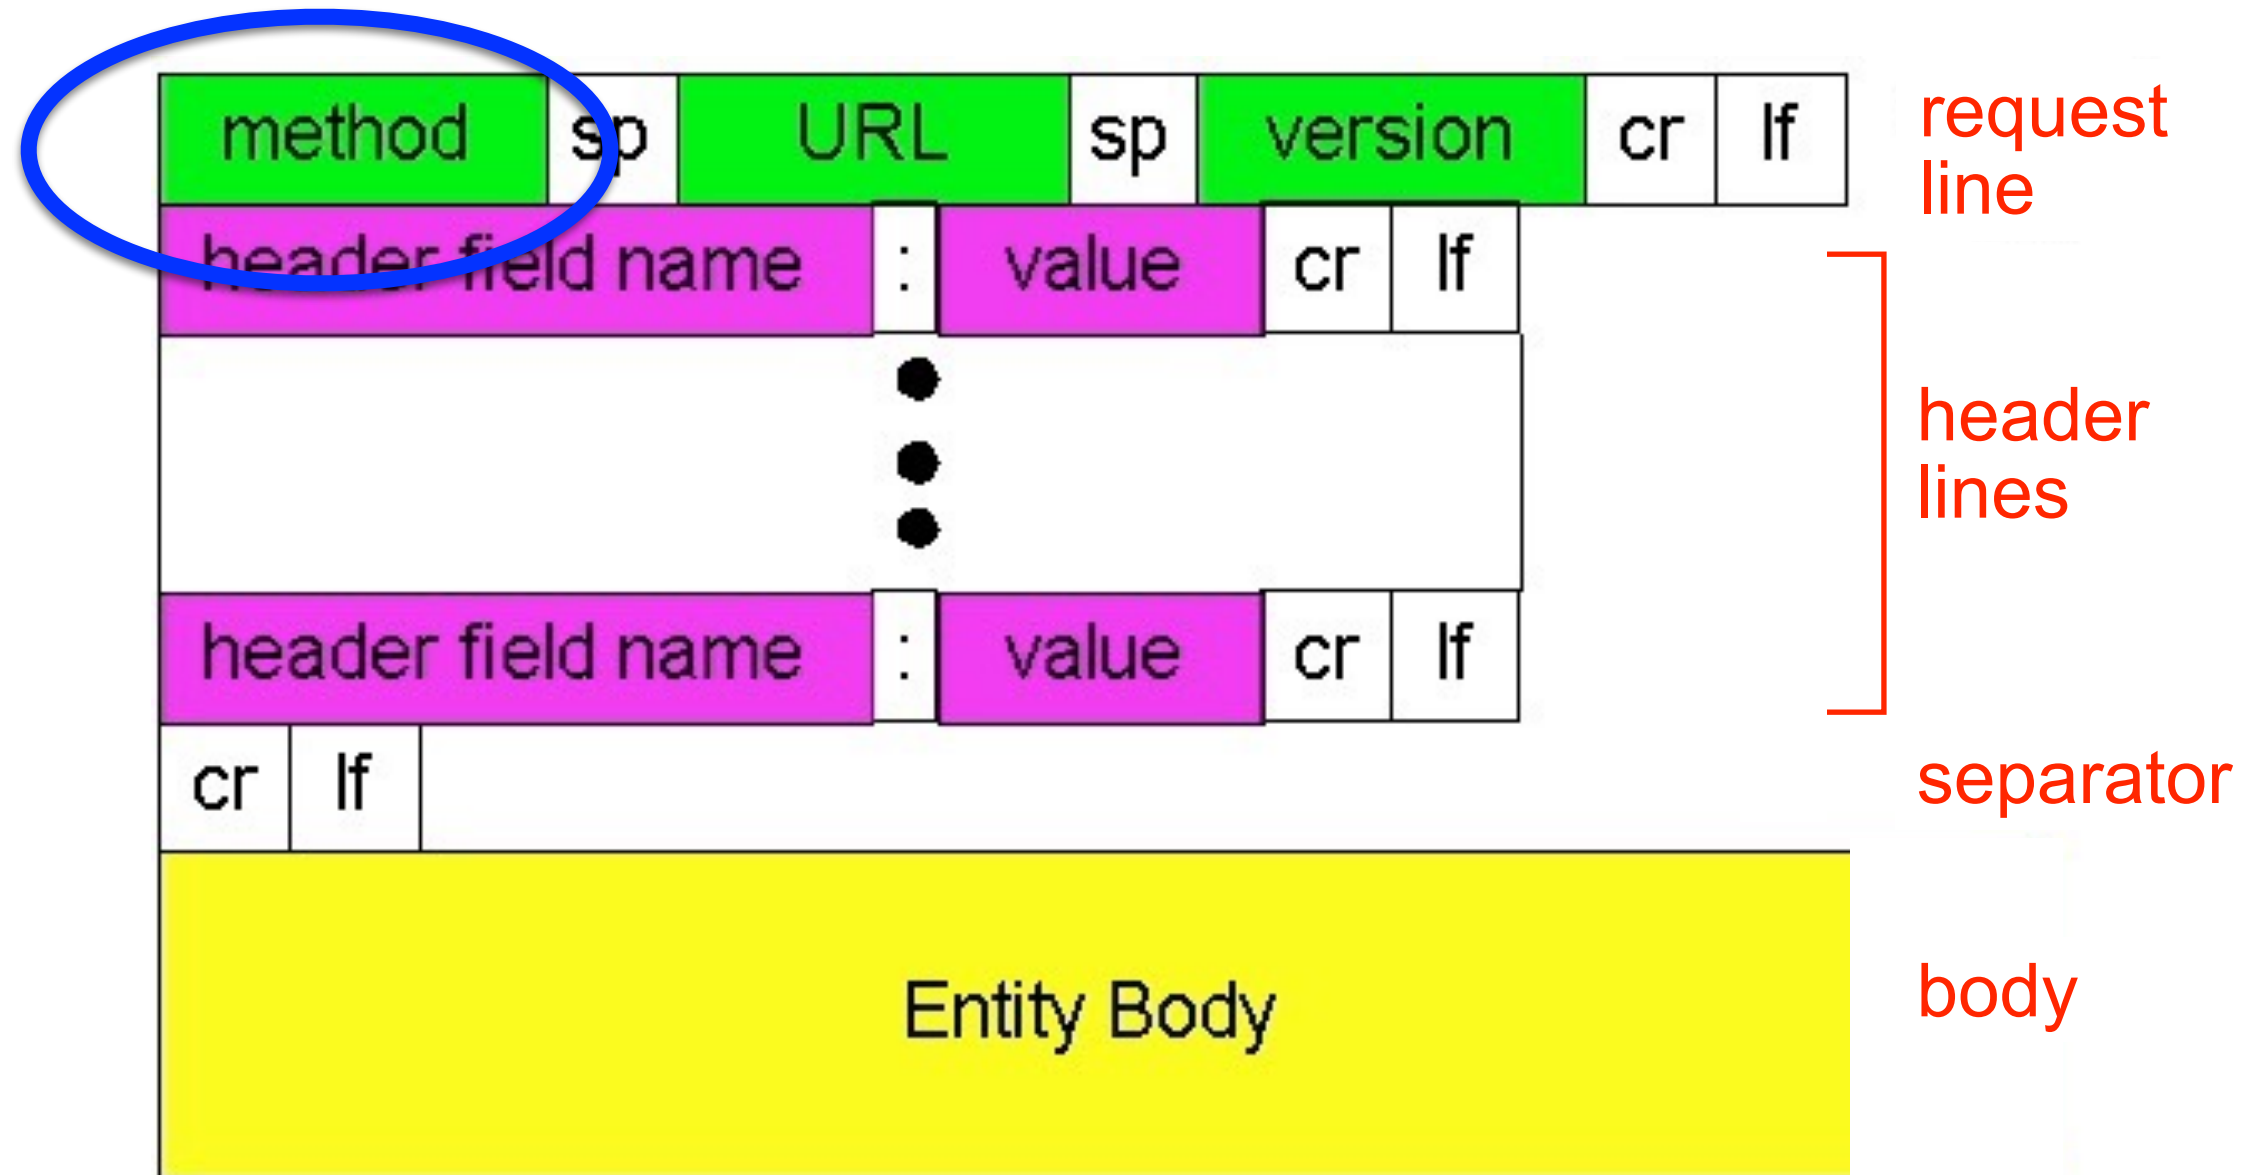
Beyond Basic HTTP

- HTTP is a *big* protocol / ecosystem with many extensions and variants
- Let's consider a few examples (not exhaustive in any way)

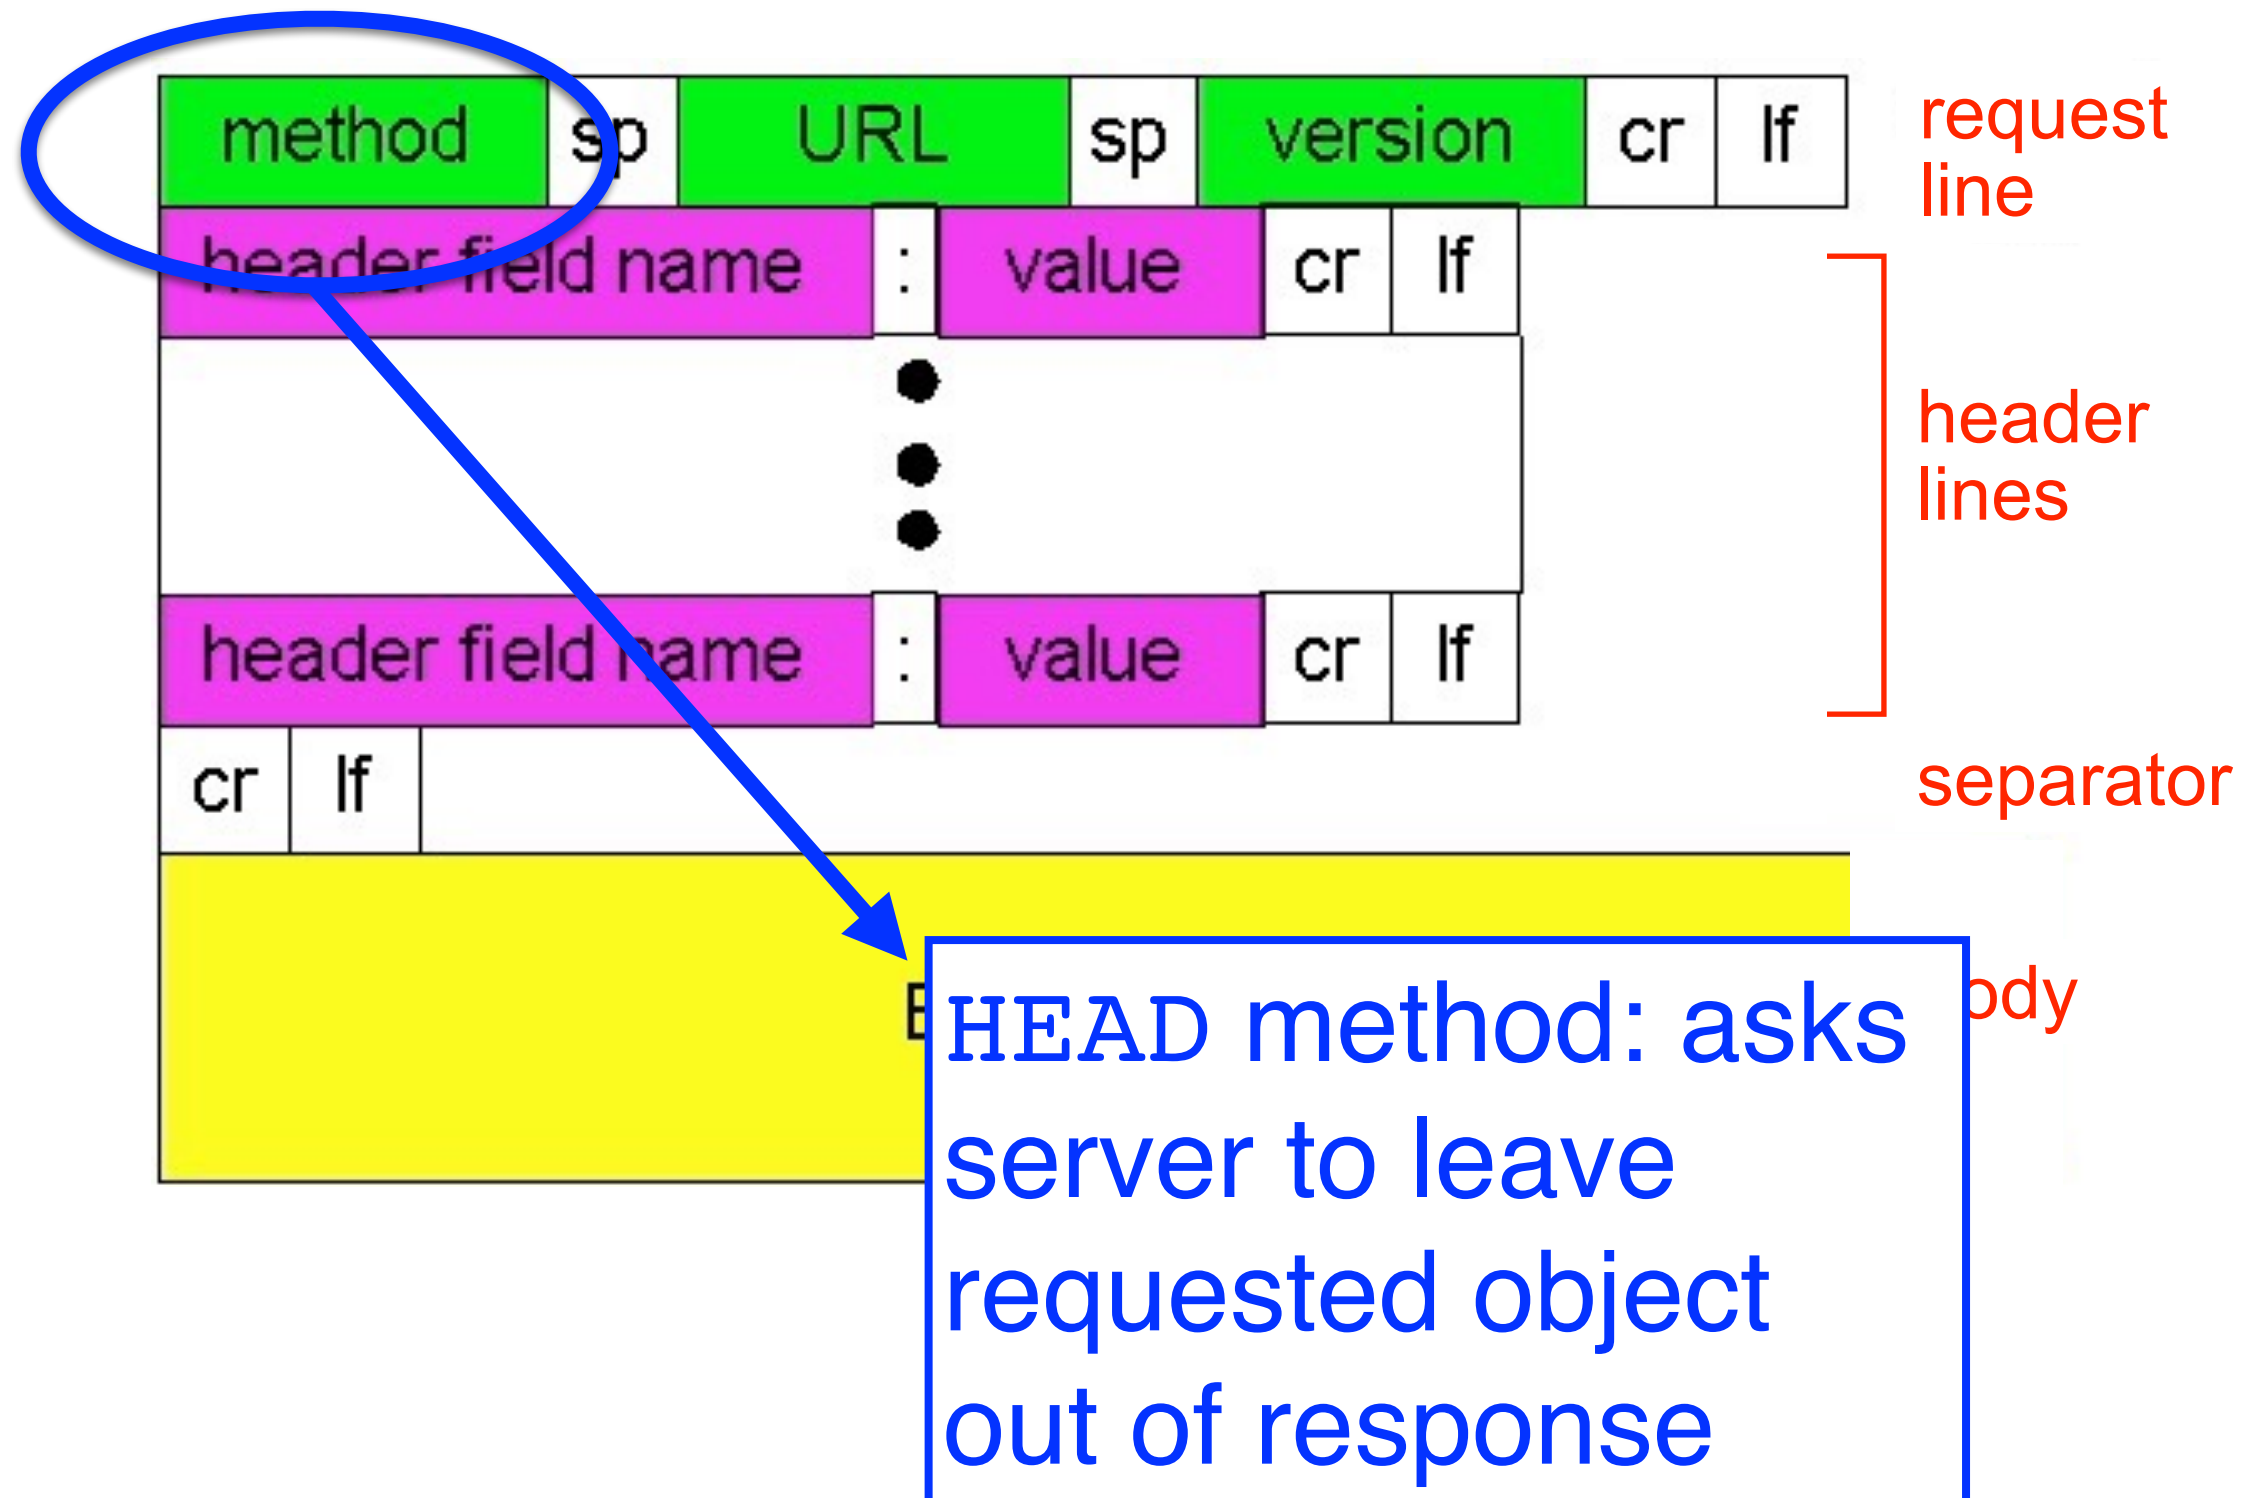
HTTP request message: general format



HTTP request message: general format



HTTP request message: general format



HTTP response message

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02 GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-1\r\n
\r\n
data data data data data ...
```

HTTP response message

```
HTTP/1.1 200 OK\r\n
Date: Sun, 26 Sep 2010 20:09:20 GMT\r\n
Server: Apache/2.0.52 (CentOS)\r\n
Last-Modified: Tue, 30 Oct 2007 17:00:02 GMT\r\n
ETag: "17dc6-a5c-bf716880"\r\n
Accept-Ranges: bytes\r\n
Content-Length: 2652\r\n
Keep-Alive: timeout=10, max=100\r\n
Connection: Keep-Alive\r\n
Content-Type: text/html; charset=ISO-8859-1\r\n
\r\n
data data data data data ...
```


Conditional GET

- ❖ **Goal:** don't send object if cache has up-to-date cached version
- ❖ cache: specify date of cached copy in HTTP request
 - `If-modified-since:`
`<date>`
- ❖ server: response contains no object if cached copy is up-to-date:
 - `HTTP/1.0 304 Not Modified`

Conditional GET

❖ **Goal:** don't send object if cache has up-to-date cached version

❖ cache: specify date of cached copy in HTTP request

`If-modified-since:
<date>`

❖ server: response contains no object if cached copy is up-to-date:

`HTTP/1.0 304 Not
Modified`

client

server

Conditional GET

- ❖ **Goal:** don't send object if cache has up-to-date cached version
- ❖ **cache:** specify date of cached copy in HTTP request
`If-modified-since:
<date>`
- ❖ **server:** response contains no object if cached copy is up-to-date:
`HTTP/1.0 304 Not
Modified`

client

server

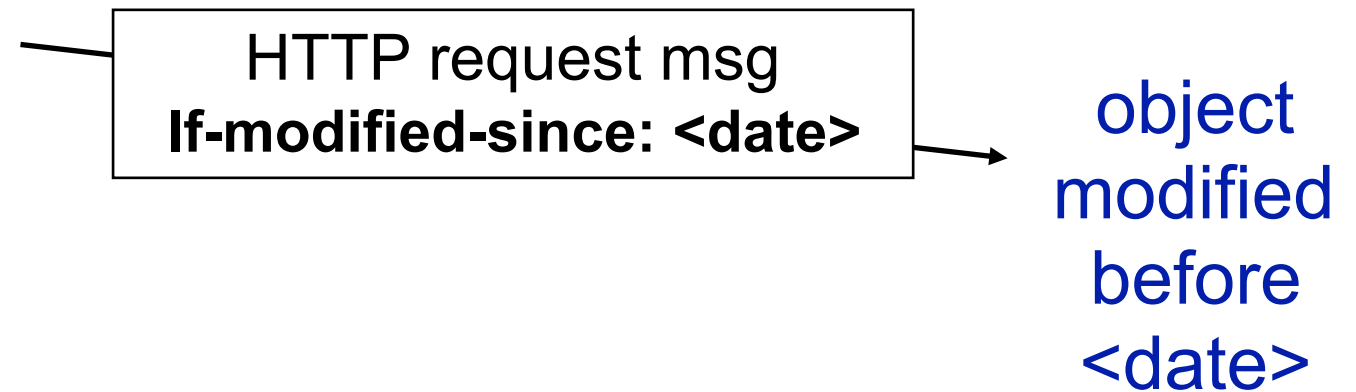


Conditional GET

- ❖ **Goal:** don't send object if cache has up-to-date cached version
- ❖ **cache:** specify date of cached copy in HTTP request
`If-modified-since: <date>`
- ❖ **server:** response contains no object if cached copy is up-to-date:
`HTTP/1.0 304 Not Modified`

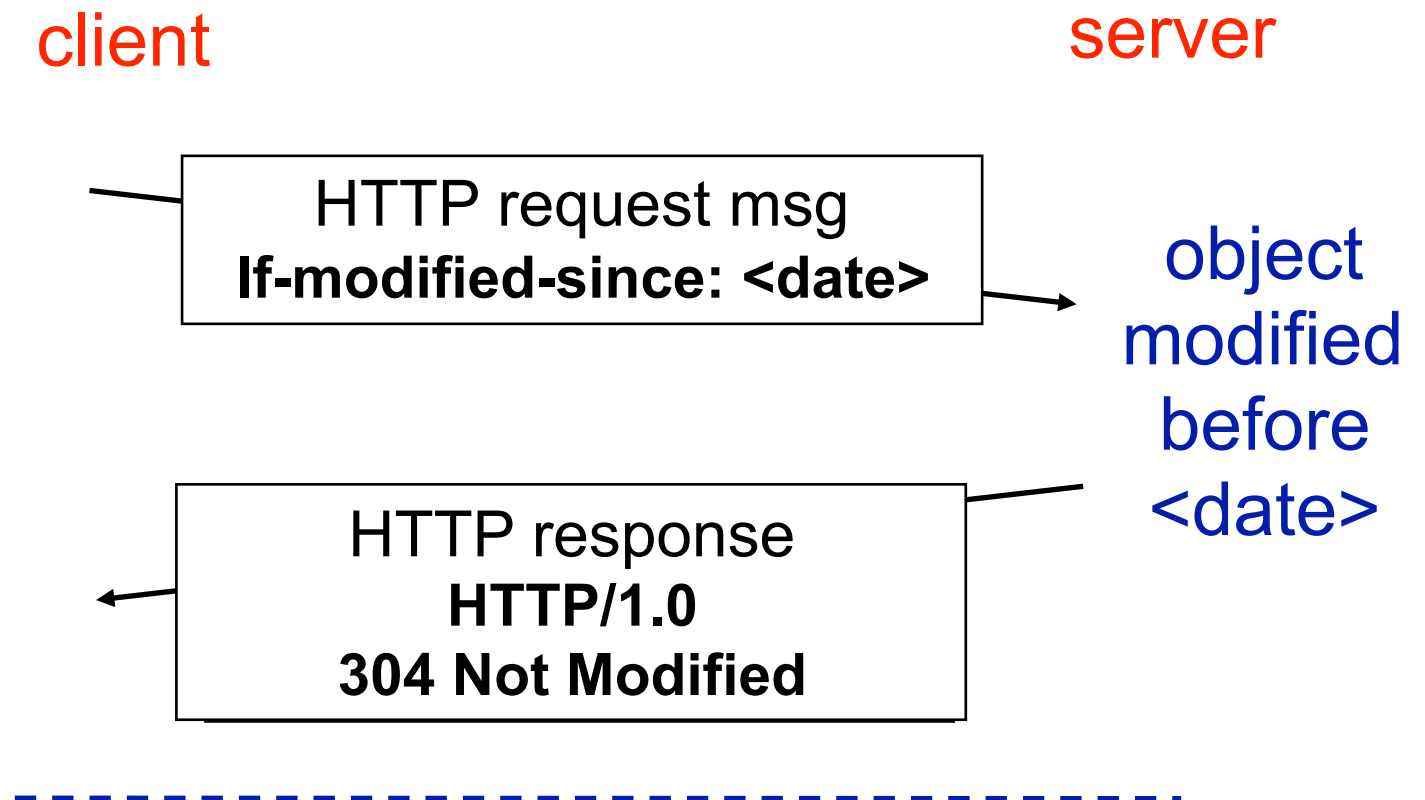
client

server



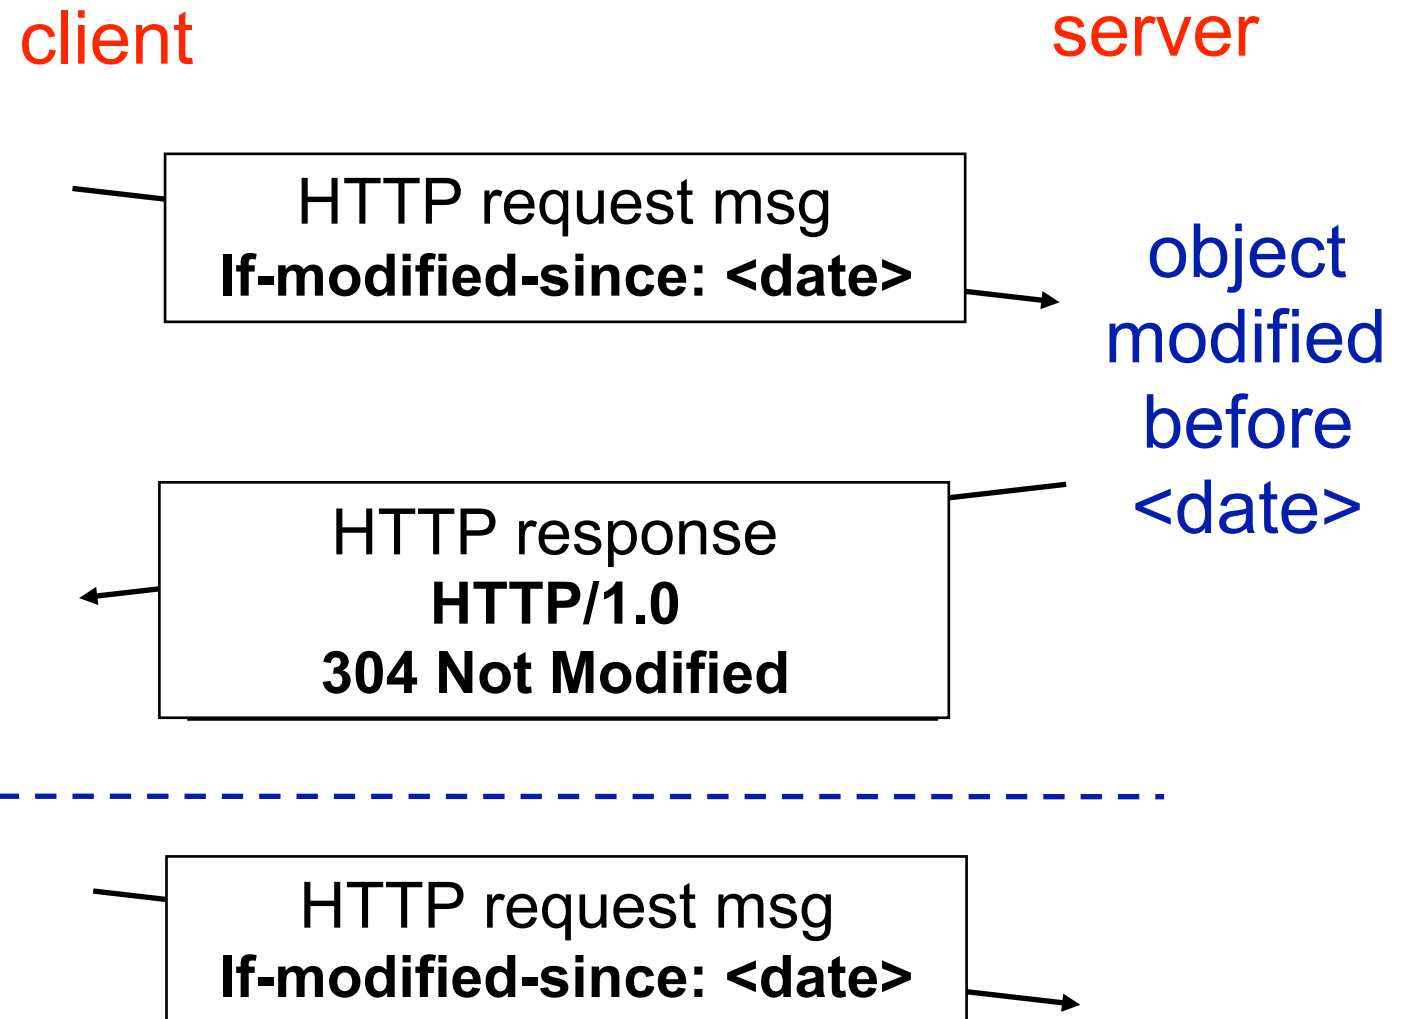
Conditional GET

- ❖ **Goal:** don't send object if cache has up-to-date cached version
- ❖ **cache:** specify date of cached copy in HTTP request
`If-modified-since: <date>`
- ❖ **server:** response contains no object if cached copy is up-to-date:
`HTTP/1.0 304 Not Modified`



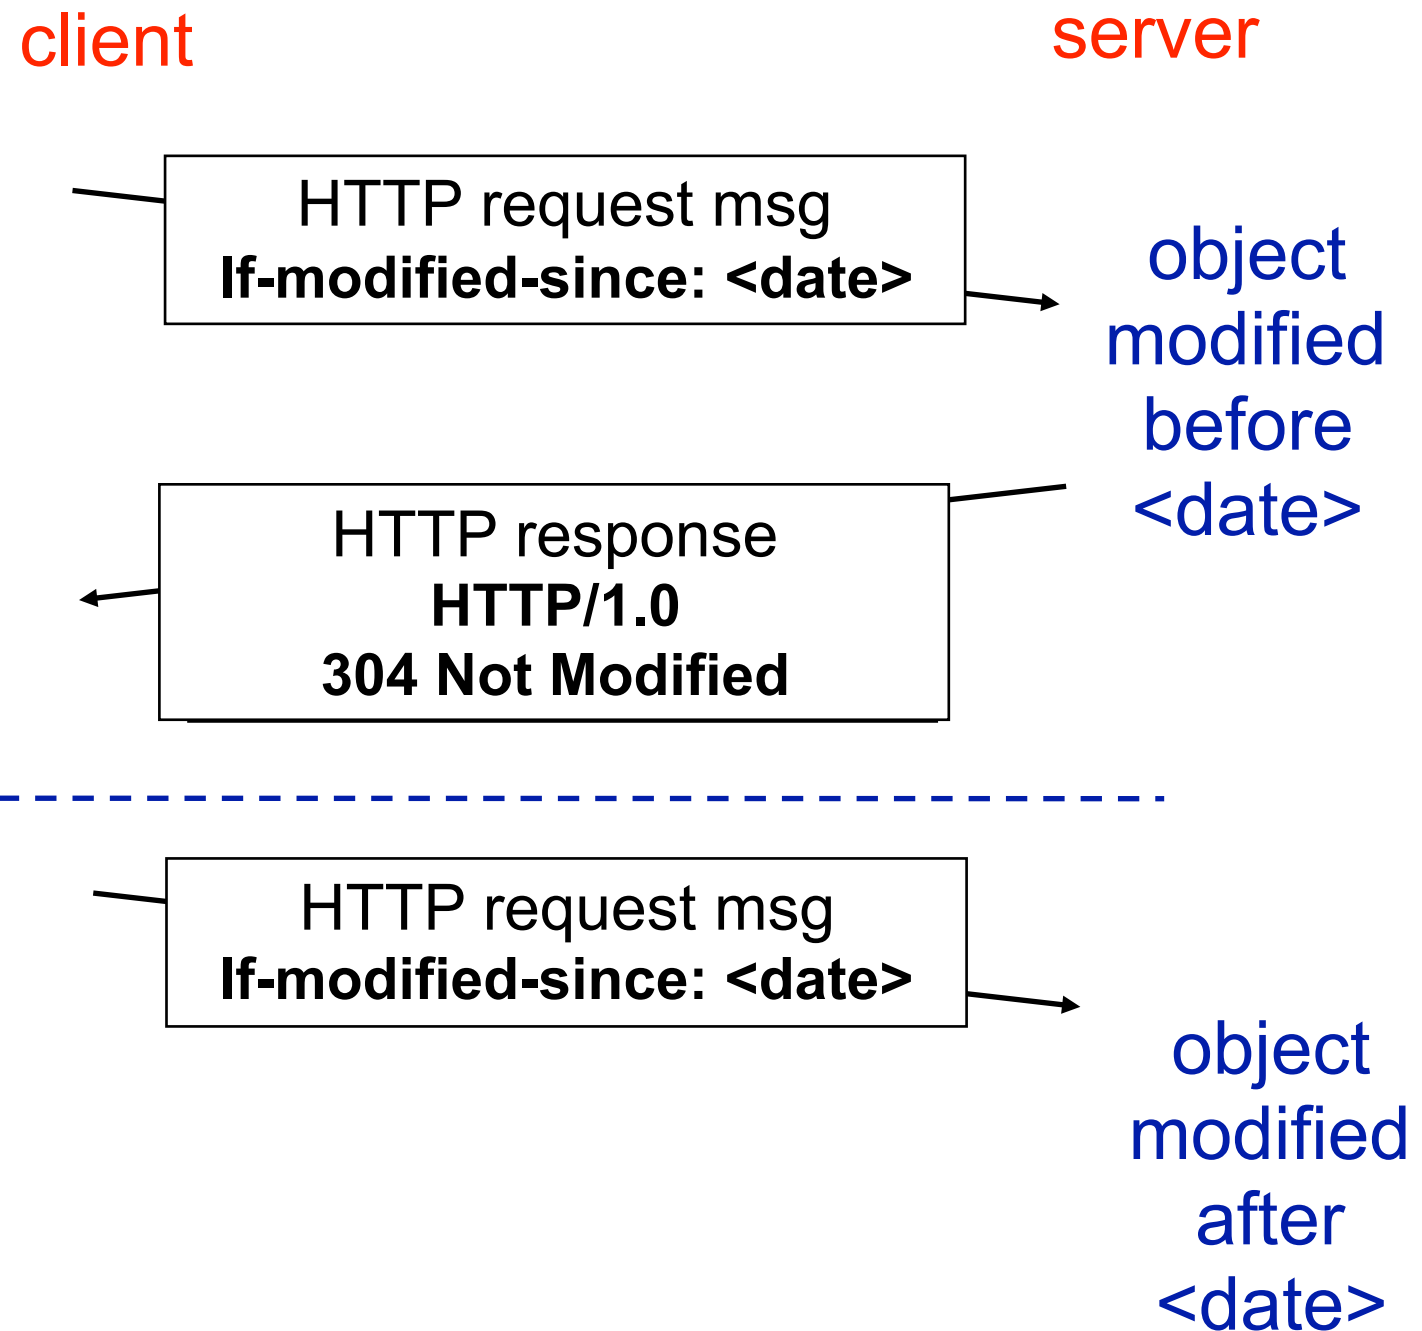
Conditional GET

- ❖ **Goal:** don't send object if cache has up-to-date cached version
- ❖ **cache:** specify date of cached copy in HTTP request
`If-modified-since: <date>`
- ❖ **server:** response contains no object if cached copy is up-to-date:
`HTTP/1.0 304 Not Modified`



Conditional GET

- ❖ **Goal:** don't send object if cache has up-to-date cached version
- ❖ **cache:** specify date of cached copy in HTTP request
`If-modified-since: <date>`
- ❖ **server:** response contains no object if cached copy is up-to-date:
`HTTP/1.0 304 Not Modified`

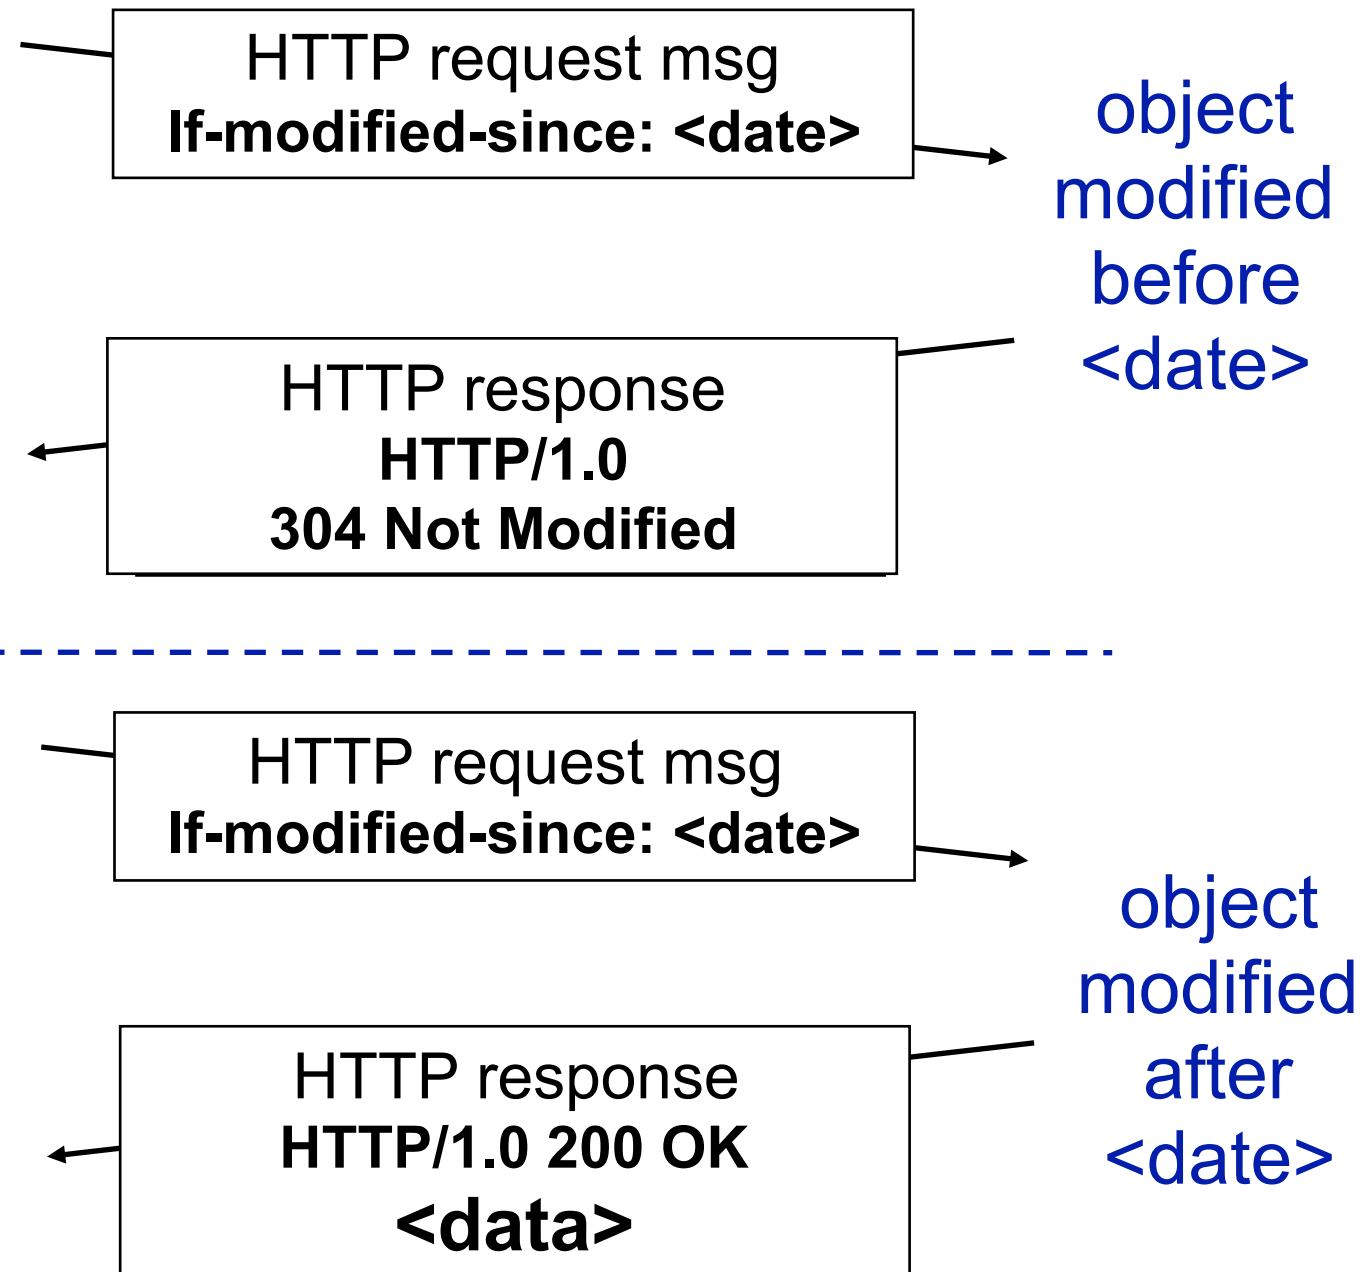


Conditional GET

- ❖ **Goal:** don't send object if cache has up-to-date cached version
- ❖ **cache:** specify date of cached copy in HTTP request
`If-modified-since: <date>`
- ❖ **server:** response contains no object if cached copy is up-to-date:
`HTTP/1.0 304 Not Modified`


client

server



Web caches (proxy server)

Goal: satisfy client request without involving origin server



Web caches (proxy server)


Goal: satisfy client request without involving origin server

- ❖ user sets browser:
Web accesses via
cache



client


Proxy
server



origin
server



client



origin
server


Web caches (proxy server)

Goal: satisfy client request without involving origin server

- ❖ user sets browser:
Web accesses via
cache
- ❖ browser sends all HTTP
requests to cache




Proxy
server



origin
server



client




origin
server


Web caches (proxy server)

Goal: satisfy client request without involving origin server

- ❖ user sets browser:
Web accesses via
cache
- ❖ browser sends all HTTP
requests to cache
 - object in cache: cache
returns object
 - else cache requests
object from origin server,
then returns object to
client



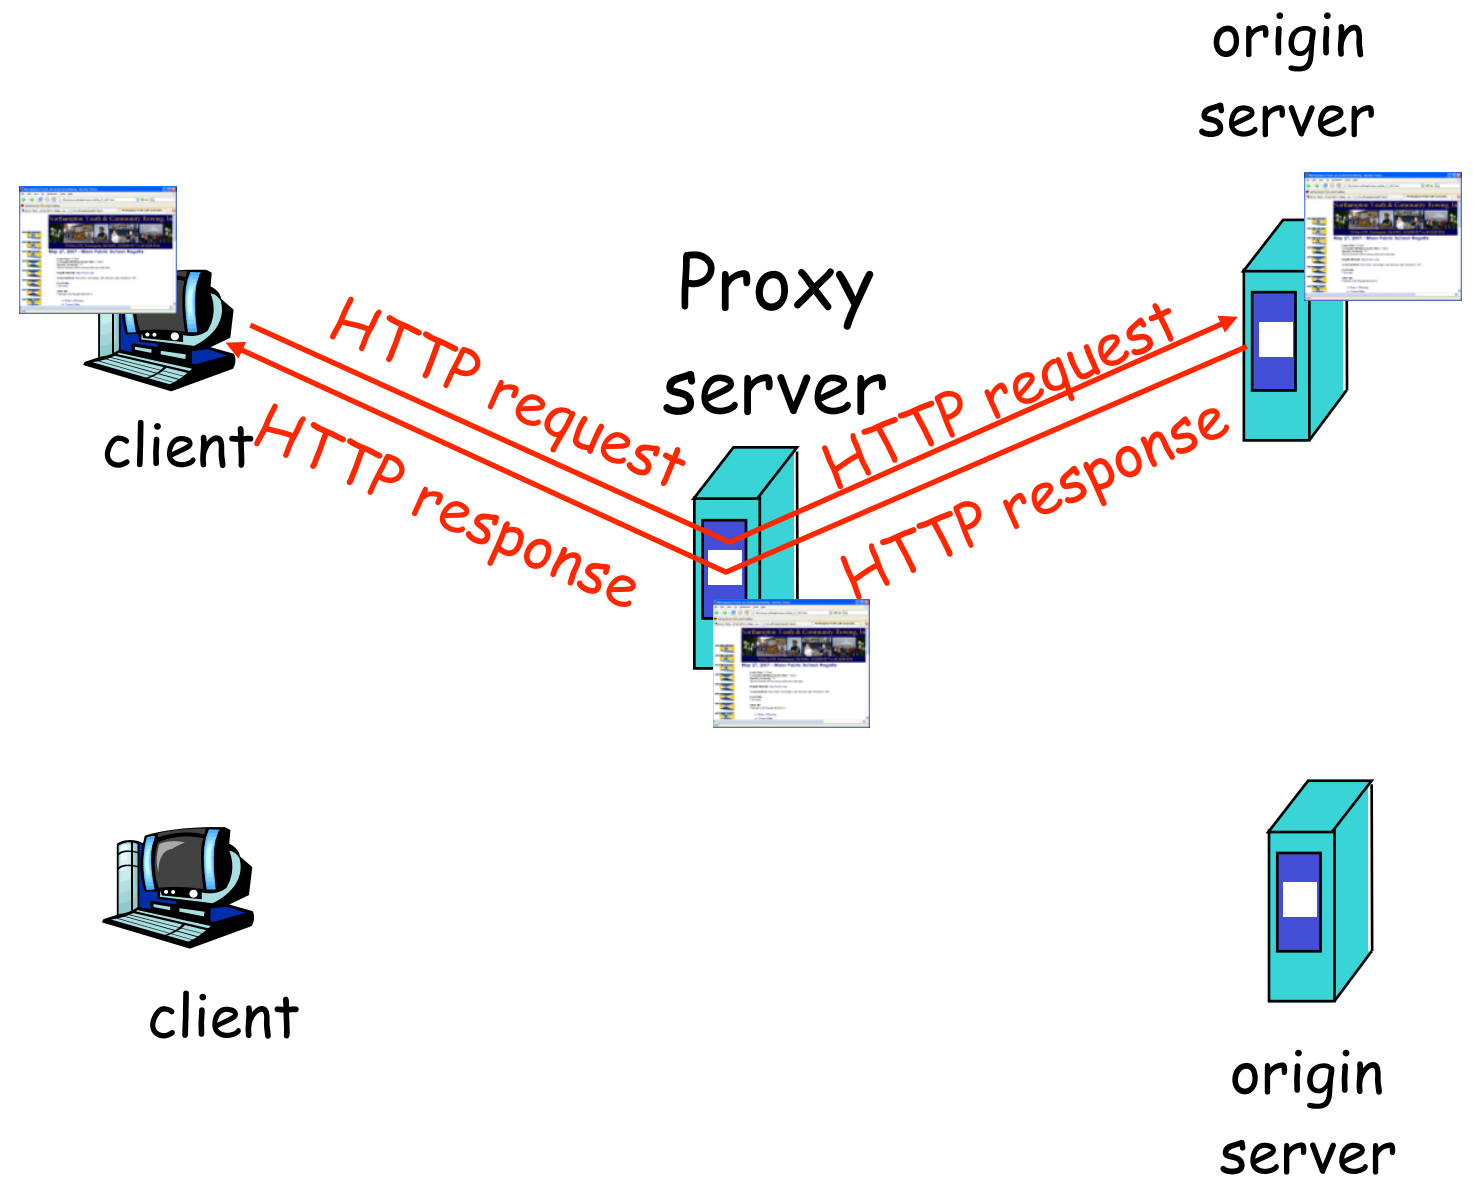
origin
server



Web caches (proxy server)

Goal: satisfy client request without involving origin server


- ❖ user sets browser:
Web accesses via cache
- ❖ browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client



Web caches (proxy server)

Goal: satisfy client request without involving origin server

- ❖ user sets browser:
Web accesses via cache
- ❖ browser sends all HTTP requests to cache
 - object in cache: cache returns object
 - else cache requests object from origin server, then returns object to client



More about Web caching

- ❖ cache acts as both client and server
- ❖ typically cache is installed by ISP (university, company, residential ISP)

More about Web caching

- ❖ cache acts as both client and server
- ❖ typically cache is installed by ISP (university, company, residential ISP)

why Web caching?

More about Web caching

- ❖ cache acts as both client and server
- ❖ typically cache is installed by ISP (university, company, residential ISP)

why Web caching?

- ❖ reduce response time for client request

More about Web caching

- ❖ cache acts as both client and server
- ❖ typically cache is installed by ISP (university, company, residential ISP)

why Web caching?

- ❖ reduce response time for client request
- ❖ reduce traffic on an institution's access link.


More about Web caching

- ❖ cache acts as both client and server
- ❖ typically cache is installed by ISP (university, company, residential ISP)

why Web caching?

- ❖ reduce response time for client request
- ❖ reduce traffic on an institution's access link.
- ❖ Internet dense with caches: enables "poor" content providers to effectively deliver content


Caching example



Caching example

assumptions

- ❖ average object size = 100,000 bits
- ❖ avg. request rate from institution's browsers to origin servers = 15/sec
- ❖ delay from institutional router to any origin server and back to router = 2 sec




Caching example

assumptions

- ❖ average object size = 100,000 bits
- ❖ avg. request rate from institution's browsers to origin servers = 15/sec
- ❖ delay from institutional router to any origin server and back to router = 2 sec

consequences




Caching example

assumptions

- ❖ average object size = 100,000 bits
- ❖ avg. request rate from institution's browsers to origin servers = 15/sec
- ❖ delay from institutional router to any origin server and back to router = 2 sec

consequences

- ❖ utilization on LAN = 15%




Caching example

assumptions

- ❖ average object size = 100,000 bits
- ❖ avg. request rate from institution's browsers to origin servers = 15/sec
- ❖ delay from institutional router to any origin server and back to router = 2 sec

consequences

- ❖ utilization on LAN = 15%
- ❖ utilization on access link = 100%




Caching example

assumptions


- ❖ average object size = 100,000 bits
- ❖ avg. request rate from institution's browsers to origin servers = 15/sec
- ❖ delay from institutional router to any origin server and back to router = 2 sec

consequences

- ❖ utilization on LAN = 15%
- ❖ utilization on access link = 100%
- ❖ total delay = Internet delay + access delay + LAN delay
= 2 sec + seconds + milliseconds




Caching example (cont)



Caching example (cont)

possible solution

- ❖ increase bandwidth of access link to, say, 10 Mbps




Caching example (cont)

possible solution

- ❖ increase bandwidth of access link to, say, 10 Mbps

consequence

- ❖ utilization on LAN = 15%




Caching example (cont)

possible solution

- ❖ increase bandwidth of access link to, say, 10 Mbps

consequence

- ❖ utilization on LAN = 15%
- ❖ utilization on access link = 15%




Caching example (cont)

possible solution

- ❖ increase bandwidth of access link to, say, 10 Mbps

consequence

- ❖ utilization on LAN = 15%
- ❖ utilization on access link = 15%
- ❖ Total delay = Internet delay + access delay + LAN delay
= 2 sec + msecs + msecs




Caching example (cont)

possible solution


- ❖ increase bandwidth of access link to, say, 10 Mbps

consequence

- ❖ utilization on LAN = 15%
- ❖ utilization on access link = 15%
- ❖ Total delay = Internet delay + access delay + LAN delay
= 2 sec + msecs + msecs
- ❖ often a costly upgrade




Caching example (cont)



Caching example (cont)

possible solution:

- ❖ install cache




Caching example (cont)

possible solution:

- ❖ install cache

consequence

- ❖ suppose hit rate is 0.4




Caching example (cont)

possible solution:

- ❖ install cache

consequence

- ❖ suppose hit rate is 0.4
 - 40% requests will be satisfied almost immediately
 - 60% requests satisfied by origin server




Caching example (cont)

possible solution:

- ❖ install cache

consequence

- ❖ suppose hit rate is 0.4
 - 40% requests will be satisfied almost immediately
 - 60% requests satisfied by origin server
- ❖ utilization of access link reduced to 60%, resulting in negligible delays (say 10 msec)




Caching example (cont)

possible solution:

- ❖ install cache

consequence

- ❖ suppose hit rate is 0.4
 - 40% requests will be satisfied almost immediately
 - 60% requests satisfied by origin server
- ❖ utilization of access link reduced to 60%, resulting in negligible delays (say 10 msec)
- ❖ total avg delay = Internet delay + access delay + LAN delay
$$= 0.6 * 2 \text{ secs} + 0.4 * \text{milliseconds} < 1.4 \text{ secs}$$



HTTP connections

basic HTTP

- ❖ one object sent over TCP connection

HTTP connections

basic HTTP

- ❖ one object sent over TCP connection

persistent HTTP

- ❖ multiple objects can be sent over single TCP connection between client, server.

Non-Persistent HTTP: Response time

Non-Persistent HTTP: Response time

definition of RTT: time for a
small packet to travel
from client to server and
back.

Non-Persistent HTTP: Response time

definition of RTT: time for a small packet to travel from client to server and back.

response time:

- ❖ one RTT to initiate TCP connection

Non-Persistent HTTP: Response time

definition of RTT: time for a small packet to travel from client to server and back.

response time:

- ❖ one RTT to initiate TCP connection
- ❖ one RTT for HTTP request and first few bytes of HTTP response to return

Non-Persistent HTTP: Response time

definition of RTT: time for a small packet to travel from client to server and back.

response time:

- ❖ one RTT to initiate TCP connection
- ❖ one RTT for HTTP request and first few bytes of HTTP response to return
- ❖ file transmission time

total = $2RTT + \text{transmit time}$

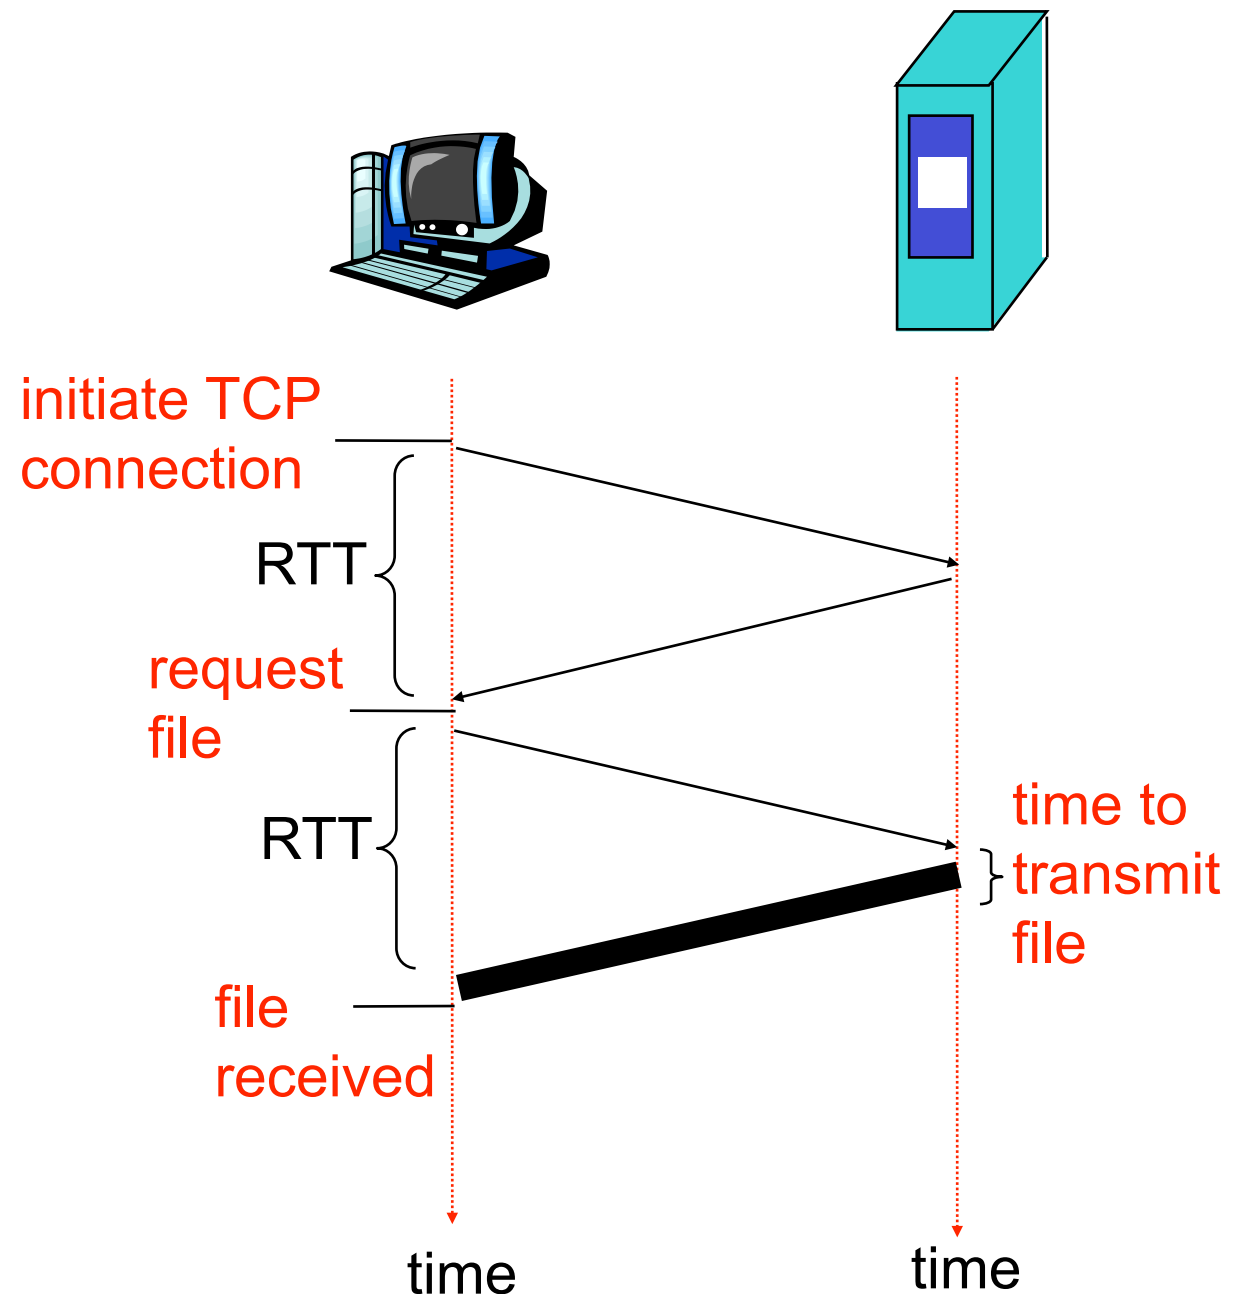
Non-Persistent HTTP: Response time

definition of RTT: time for a small packet to travel from client to server and back.

response time:

- ❖ one RTT to initiate TCP connection
- ❖ one RTT for HTTP request and first few bytes of HTTP response to return
- ❖ file transmission time

total = $2RTT + \text{transmit time}$



Persistent HTTP

non-persistent HTTP issues:

- ❖ requires 2 RTTs per object
- ❖ OS overhead for each TCP connection
- ❖ browsers often open parallel TCP connections to fetch referenced objects

Persistent HTTP

non-persistent HTTP issues:

- ❖ requires 2 RTTs per object
- ❖ OS overhead for each TCP connection
- ❖ browsers often open parallel TCP connections to fetch referenced objects

persistent HTTP

- ❖ server leaves connection open after sending response
- ❖ subsequent HTTP messages between same client/server sent over open connection
- ❖ client sends requests as soon as it encounters a referenced object
- ❖ as little as one RTT for all the referenced objects