



Application Layer Part 6

Mark Allman
Case / ICSI

EECS 325/425
Fall 2018


*“Now don’t you call James Bond or Secret Agent Man,
‘Cause they can’t do it Like I Can ...”*

Many of these slides are more-or-less directly from the slide set developed by Jim Kurose and Keith Ross for their book “Computer Networking: A Top Down Approach, 5th edition”.

The slides have been lightly adapted for Mark Allman’s EECS 325/425 Computer Networks class at Case Western Reserve University.

All material copyright 1996-2010
J.F Kurose and K.W. Ross, All Rights Reserved

Server-client vs. P2P: example



File distribution: BitTorrent


P2P file distribution

File distribution: BitTorrent


P2P file distribution

tracker: tracks peers
participating in torrent

torrent: group of
peers exchanging
chunks of a file




BitTorrent (1)




BitTorrent (1)

❖ file divided into 256KB **chunks**.




BitTorrent (1)

- ❖ file divided into 256KB **chunks**.
- ❖ peer joining torrent:
 - has no chunks, but will accumulate them over time
 - registers with tracker to get list of peers, connects to subset of peers ("neighbors")




BitTorrent (1)

- ❖ file divided into 256KB **chunks**.
- ❖ peer joining torrent:
 - has no chunks, but will accumulate them over time
 - registers with tracker to get list of peers, connects to subset of peers ("neighbors")
- ❖ while downloading, peer uploads chunks to other peers.




BitTorrent (1)

- ❖ file divided into 256KB **chunks**.
- ❖ peer joining torrent:
 - has no chunks, but will accumulate them over time
 - registers with tracker to get list of peers, connects to subset of peers ("neighbors")
- ❖ while downloading, peer uploads chunks to other peers.
- ❖ peers may come and go




BitTorrent (1)

- ❖ file divided into 256KB **chunks**.
- ❖ peer joining torrent:
 - has no chunks, but will accumulate them over time
 - registers with tracker to get list of peers, connects to subset of peers ("neighbors")
- ❖ while downloading, peer uploads chunks to other peers.
- ❖ peers may come and go
- ❖ once peer has entire file, ... ?



BitTorrent (1)

- ❖ file divided into 256KB **chunks**.
- ❖ peer joining torrent:
 - has no chunks, but will accumulate them over time
 - registers with tracker to get list of peers, connects to subset of peers ("neighbors")
- ❖ while downloading, peer uploads chunks to other peers.
- ❖ peers may come and go
- ❖ once peer has entire file, it may (selfishly) leave or (altruistically) remain



BitTorrent (2)

BitTorrent (2)

Pulling Chunks

BitTorrent (2)

Pulling Chunks

- ❖ at any given time,
different peers have
different subsets of file
chunks

BitTorrent (2)

Pulling Chunks

- ❖ at any given time, different peers have different subsets of file chunks
- ❖ periodically, a peer (Alice) asks each neighbor for list of chunks that they have.

BitTorrent (2)

Pulling Chunks

- ❖ at any given time, different peers have different subsets of file chunks
- ❖ periodically, a peer (Alice) asks each neighbor for list of chunks that they have.
- ❖ Alice sends requests for her missing chunks
 - rarest first

BitTorrent (2)

Pulling Chunks

- ❖ at any given time, different peers have different subsets of file chunks
- ❖ periodically, a peer (Alice) asks each neighbor for list of chunks that they have.
- ❖ Alice sends requests for her missing chunks
 - rarest first

Sending Chunks: tit-for-tat

BitTorrent (2)

Pulling Chunks

- ❖ at any given time, different peers have different subsets of file chunks
- ❖ periodically, a peer (Alice) asks each neighbor for list of chunks that they have.
- ❖ Alice sends requests for her missing chunks
 - rarest first

Sending Chunks: tit-for-tat

- ❖ Alice sends chunks to four neighbors currently sending her chunks at the highest rate
 - re-evaluate top 4 every 10 secs

BitTorrent (2)

Pulling Chunks

- ❖ at any given time, different peers have different subsets of file chunks
- ❖ periodically, a peer (Alice) asks each neighbor for list of chunks that they have.
- ❖ Alice sends requests for her missing chunks
 - rarest first

Sending Chunks: tit-for-tat

- ❖ Alice sends chunks to four neighbors currently sending her chunks at the highest rate
 - re-evaluate top 4 every 10 secs
- ❖ every 30 secs: randomly select another peer, starts sending chunks
 - newly chosen peer may join top 4
 - "optimistically unchoke"

BitTorrent (3)

❖ rough sketch

❖ why 30sec? why 256KB? why 4 neighbors?

Distributed Hash Table (DHT)

Distributed Hash Table (DHT)

❖ DHT: distributed (via P2P) data structure

Distributed Hash Table (DHT)

- ❖ DHT: distributed (via P2P) data structure
- ❖ data structure has (key, value) pairs;
 - key: ss number; value: human name

Distributed Hash Table (DHT)

- ❖ DHT: distributed (via P2P) data structure
- ❖ data structure has (key, value) pairs;
 - key: ss number; value: human name
- ❖ peers query data structure with key
 - value that match the key is returned

Distributed Hash Table (DHT)

- ❖ DHT: distributed (via P2P) data structure
- ❖ data structure has (key, value) pairs;
 - key: ss number; value: human name
- ❖ peers query data structure with key
 - value that match the key is returned
- ❖ peers can also insert (key, value) peers

DHT Identifiers

DHT Identifiers

- ❖ assign integer identifier to each peer in range $[0, 2^n - 1]$.
 - Each identifier can be represented by n bits.

DHT Identifiers

- ❖ assign integer identifier to each peer in range $[0, 2^n - 1]$.
 - Each identifier can be represented by n bits.
- ❖ require each key to be an integer in **same range**.

DHT Identifiers

- ❖ assign integer identifier to each peer in range $[0, 2^n - 1]$.
 - Each identifier can be represented by n bits.
- ❖ require each key to be an integer in **same range**.
- ❖ to get integer keys, hash original key.
 - e.g., $\text{key} = h(\text{"Led Zeppelin IV"})$
 - this is why they call it a distributed "hash" table

How to assign keys to peers?

How to assign keys to peers?

❖ central issue:

- assigning (key, value) pairs to peers.

How to assign keys to peers?

❖ central issue:

- assigning (key, value) pairs to peers.

❖ rule: assign key to the peer that has the **closest** ID.


How to assign keys to peers?

- ❖ central issue:
 - assigning (key, value) pairs to peers.
- ❖ rule: assign key to the peer that has the **closest** ID.
- ❖ convention in lecture: closest is the **immediate successor** of the key.

How to assign keys to peers?


- ❖ central issue:
 - assigning (key, value) pairs to peers.
- ❖ rule: assign key to the peer that has the **closest** ID.
- ❖ convention in lecture: closest is the **immediate successor** of the key.
- ❖ e.g.,: $n=4$; peers: 1,3,4,5,8,10,12,14;
 - key = 13, then successor peer = 14
 - key = 15, then successor peer = 1

Circular DHT (1)



- ❖ each peer only aware of immediate successor and predecessor.
- ❖ "overlay network"


Circular DHT (2)



Define closest
as closest
successor

Circular DHT (2)


$x = h(\text{"Led Zeppelin IV"})$
 $x = 1110$
Who's resp. for 1110?



Define closest
as closest
successor

Circular DHT (2)


$x = h(\text{"Led Zeppelin IV"})$
 $x = 1110$
Who's resp. for 1110?



Define closest
as closest
successor

Circular DHT (2)


$x = h(\text{"Led Zepplin IV"})$
 $x = 1110$
Who's resp. for 1110?



Define closest
as closest
successor

Circular DHT (2)


$x = h(\text{"Led Zeppelin IV"})$
 $x = 1110$
Who's resp. for 1110?



Define closest
as closest
successor

Circular DHT (2)


$x = h(\text{"Led Zeppelin IV"})$
 $x = 1110$
Who's resp. for 1110?



Define closest
as closest
successor

Circular DHT (2)


$x = h(\text{"Led Zeppelin IV"})$
 $x = 1110$
Who's resp. for 1110?



Define closest
as closest
successor


Circular DHT (2)

$x = h(\text{"Led Zepplin IV"})$
 $x = 1110$
Who's resp. for 1110?



Define closest
as closest
successor

Circular DHT (2)




$x = h(\text{"Led Zeppelin IV"})$
 $x = 1110$
Who's resp. for 1110?

Define closest
as closest
successor


Circular DHT (2)

$O(N)$ messages
on avg to resolve
query, when there
are N peers




Define closest
as closest
successor

Circular DHT with Shortcuts




Circular DHT with Shortcuts




- ❖ each peer keeps track of IP addresses of predecessor, successor, short cuts.

Circular DHT with Shortcuts




- ❖ each peer keeps track of IP addresses of predecessor, successor, short cuts.

Circular DHT with Shortcuts




- ❖ each peer keeps track of IP addresses of predecessor, successor, short cuts.

Circular DHT with Shortcuts




- ❖ each peer keeps track of IP addresses of predecessor, successor, short cuts.

Circular DHT with Shortcuts




- ❖ each peer keeps track of IP addresses of predecessor, successor, short cuts.
- ❖ reduced from 6 to 2 messages.

Circular DHT with Shortcuts




- ❖ each peer keeps track of IP addresses of predecessor, successor, short cuts.
- ❖ reduced from 6 to 2 messages.
- ❖ possible to design shortcuts so $O(\log N)$ neighbors, $O(\log N)$ messages in query

Peer Churn



- ❖ To handle peer churn, require each peer to know the IP address of its two successors.
- ❖ Each peer periodically pings its two successors to see if they are still alive.


Peer Churn



❖ peer 5 abruptly leaves

- ❖ To handle peer churn, require each peer to know the IP address of its two successors.
- ❖ Each peer periodically pings its two successors to see if they are still alive.

Peer Churn



- ❖ To handle peer churn, require each peer to know the IP address of its two successors.
- ❖ Each peer periodically pings its two successors to see if they are still alive.

- ❖ peer 5 abruptly leaves
- ❖ Peer 4 detects; makes 8 its immediate successor; asks 8 who its immediate successor is; makes 8's immediate successor its second successor.

Distributed Data Structures

Distributed Data Structures

- ❖ DHTs are broadly applicable to many applications

Distributed Data Structures

- ❖ DHTs are broadly applicable to many applications
- ❖ What is another distributed data structure?

Distributed Data Structures

- ❖ DHTs are broadly applicable to many applications
- ❖ What is another distributed data structure?
 - blockchain
 - essentially a cryptographically secure ledger that can be broadly shared with integrity guarantees

P2P Case study: Skype

- ❖ inherently P2P: pairs of users communicate.
- ❖ proprietary application-layer protocol (inferred via reverse engineering)
- ❖ hierarchical overlay with SuperNodes (SNs)
- ❖ Index maps usernames to IP addresses; distributed over SNs


P2P Case study: Skype

- ❖ inherently P2P: pairs of users communicate.
- ❖ proprietary application-layer protocol (inferred via reverse engineering)
- ❖ hierarchical overlay with SuperNodes (SNs)
- ❖ Index maps usernames to IP addresses; distributed over SNs




P2P Case study: Skype

- ❖ inherently P2P: pairs of users communicate.
- ❖ proprietary application-layer protocol (inferred via reverse engineering)
- ❖ hierarchical overlay with SuperNodes (SNs)
- ❖ Index maps usernames to IP addresses; distributed over SNs




P2P Case study: Skype

- ❖ inherently P2P: pairs of users communicate.
- ❖ proprietary application-layer protocol (inferred via reverse engineering)
- ❖ hierarchical overlay with SuperNodes (SNs)
- ❖ Index maps usernames to IP addresses; distributed over SNs




P2P Case study: Skype

- ❖ inherently P2P: pairs of users communicate.
- ❖ proprietary application-layer protocol (inferred via reverse engineering)
- ❖ hierarchical overlay with SuperNodes (SNs)
- ❖ Index maps usernames to IP addresses; distributed over SNs




P2P Case study: Skype

- ❖ inherently P2P: pairs of users communicate.
- ❖ proprietary application-layer protocol (inferred via reverse engineering)
- ❖ hierarchical overlay with SuperNodes (SNs)
- ❖ Index maps usernames to IP addresses; distributed over SNs




P2P Case study: Skype

- ❖ inherently P2P: pairs of users communicate.
- ❖ proprietary application-layer protocol (inferred via reverse engineering)
- ❖ hierarchical overlay with SuperNodes (SNs)
- ❖ Index maps usernames to IP addresses; distributed over SNs



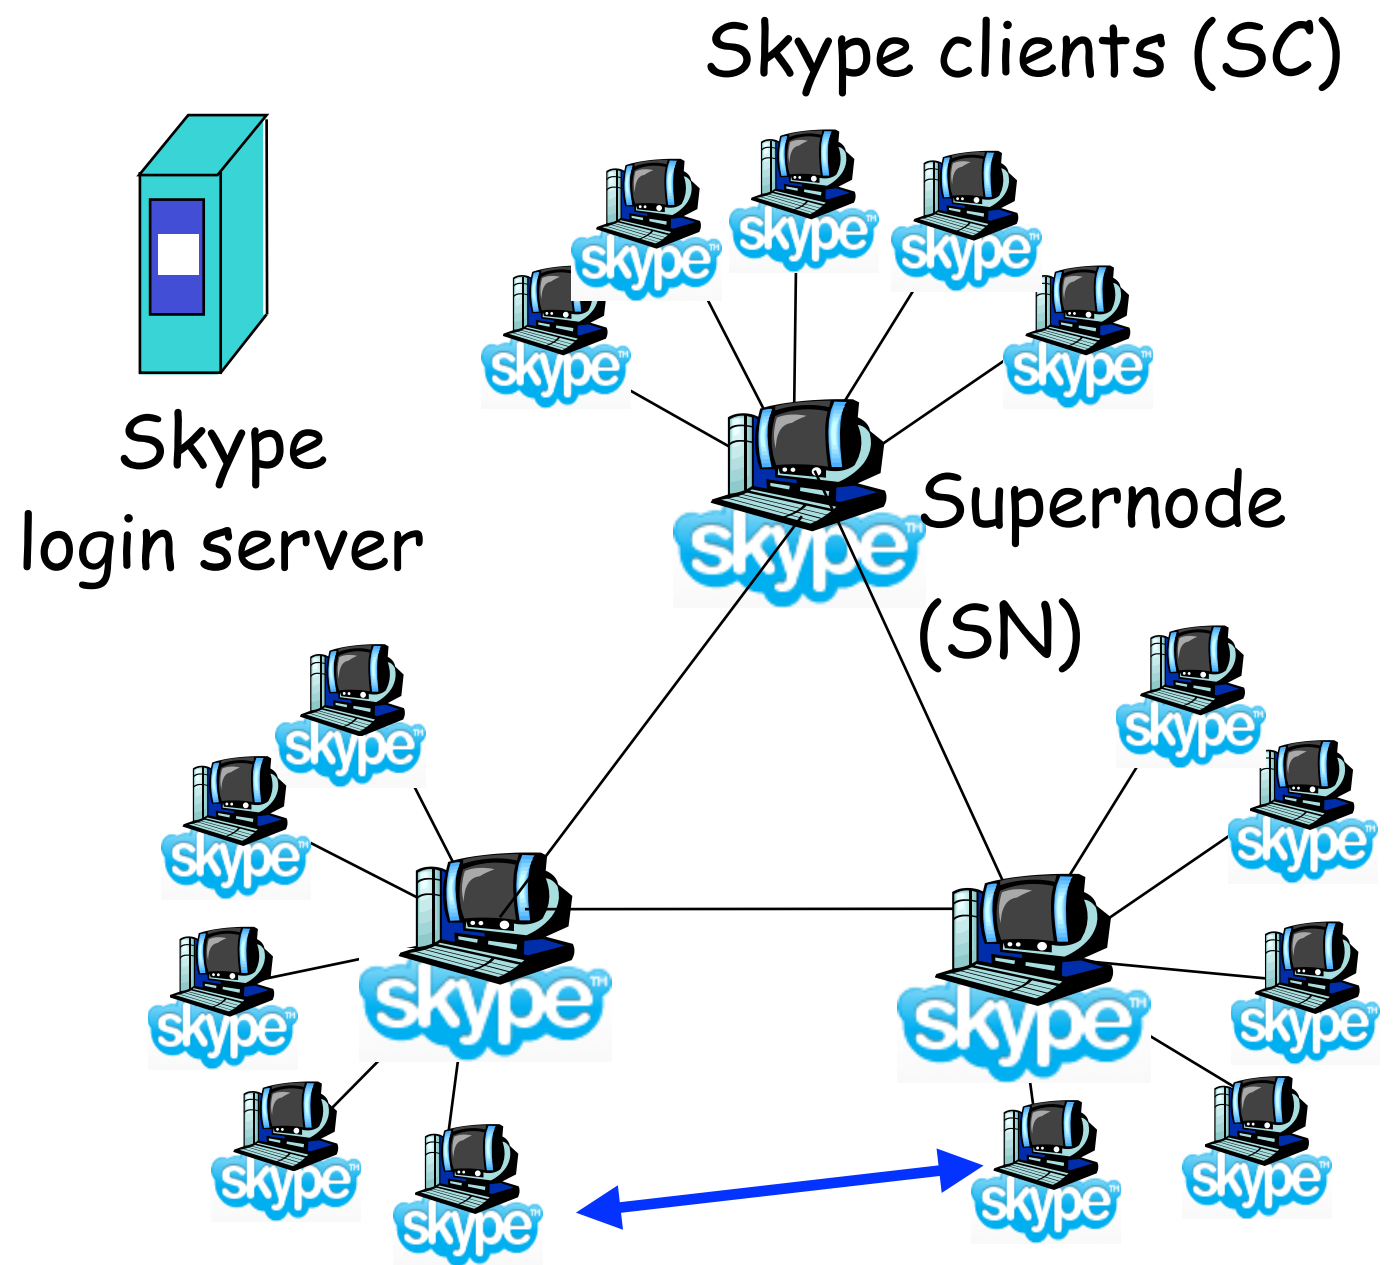
P2P Case study: Skype

- ❖ inherently P2P: pairs of users communicate.
- ❖ proprietary application-layer protocol (inferred via reverse engineering)
- ❖ hierarchical overlay with SuperNodes (SNs)
- ❖ Index maps usernames to IP addresses; distributed over SNs



P2P Case study: Skype

- ❖ inherently P2P: pairs of users communicate.
- ❖ proprietary application-layer protocol (inferred via reverse engineering)
- ❖ hierarchical overlay with SuperNodes (SNs)
- ❖ Index maps usernames to IP addresses; distributed over SNs



P2P Case study: Skype

- ❖ inherently P2P: pairs of users communicate.
- ❖ proprietary application-layer protocol (inferred via reverse engineering)
- ❖ hierarchical overlay with SuperNodes (SNs)
- ❖ Index maps usernames to IP addresses; distributed over SNs

