



Transport Layer Part 2

Mark Allman
mallman@case.edu

EECS 325/425
Fall 2018

*“Well, we busted outa class, had to get away from those fools,
We learned more from a three minute record than we ever learned in school.”*

These slides are more-or-less directly from the slide set developed by Jim Kurose and Keith Ross for their book “Computer Networking: A Top Down Approach, 5th edition”.

The slides have been lightly adapted for Mark Allman’s EECS 325/425 Computer Networks class at Case Western Reserve University.

All material copyright 1996-2010
J.F Kurose and K.W. Ross, All Rights Reserved

Connectionless demultiplexing

Connectionless demultiplexing

❖ create UDP sockets with host-local port numbers:

```
DatagramSocket mySocket1 = new DatagramSocket(12534);  
DatagramSocket mySocket2 = new DatagramSocket(12535);
```

Connectionless demultiplexing

- ❖ create UDP sockets with host-local port numbers:

```
DatagramSocket mySocket1 = new DatagramSocket(12534);
```

```
DatagramSocket mySocket2 = new DatagramSocket(12535);
```

- ❖ when creating datagram to send into UDP socket, must specify: (dest IP address, dest port number)

Connectionless demultiplexing

- ❖ create UDP sockets with host-local port numbers:

```
DatagramSocket mySocket1 = new DatagramSocket(12534);
```

```
DatagramSocket mySocket2 = new DatagramSocket(12535);
```

- ❖ when creating datagram to send into UDP socket, must specify: (dest IP address, dest port number)
- ❖ when host receives UDP segment:
 - checks destination port number in segment
 - directs UDP segment to socket with that port number

Connectionless demultiplexing

- ❖ create UDP sockets with host-local port numbers:


```
DatagramSocket mySocket1 = new DatagramSocket(12534);
```

```
DatagramSocket mySocket2 = new DatagramSocket(12535);
```


- ❖ when creating datagram to send into UDP socket, must specify: (dest IP address, dest port number)
- ❖ when host receives UDP segment:
 - checks destination port number in segment
 - directs UDP segment to socket with that port number
- ❖ IP datagrams with different source IP addresses and/or source port numbers directed to same socket
 - i.e., demuxing based on destination port only

Connectionless demux (cont)


```
DatagramSocket serverSocket = new DatagramSocket(6428);
```



client
IP: A



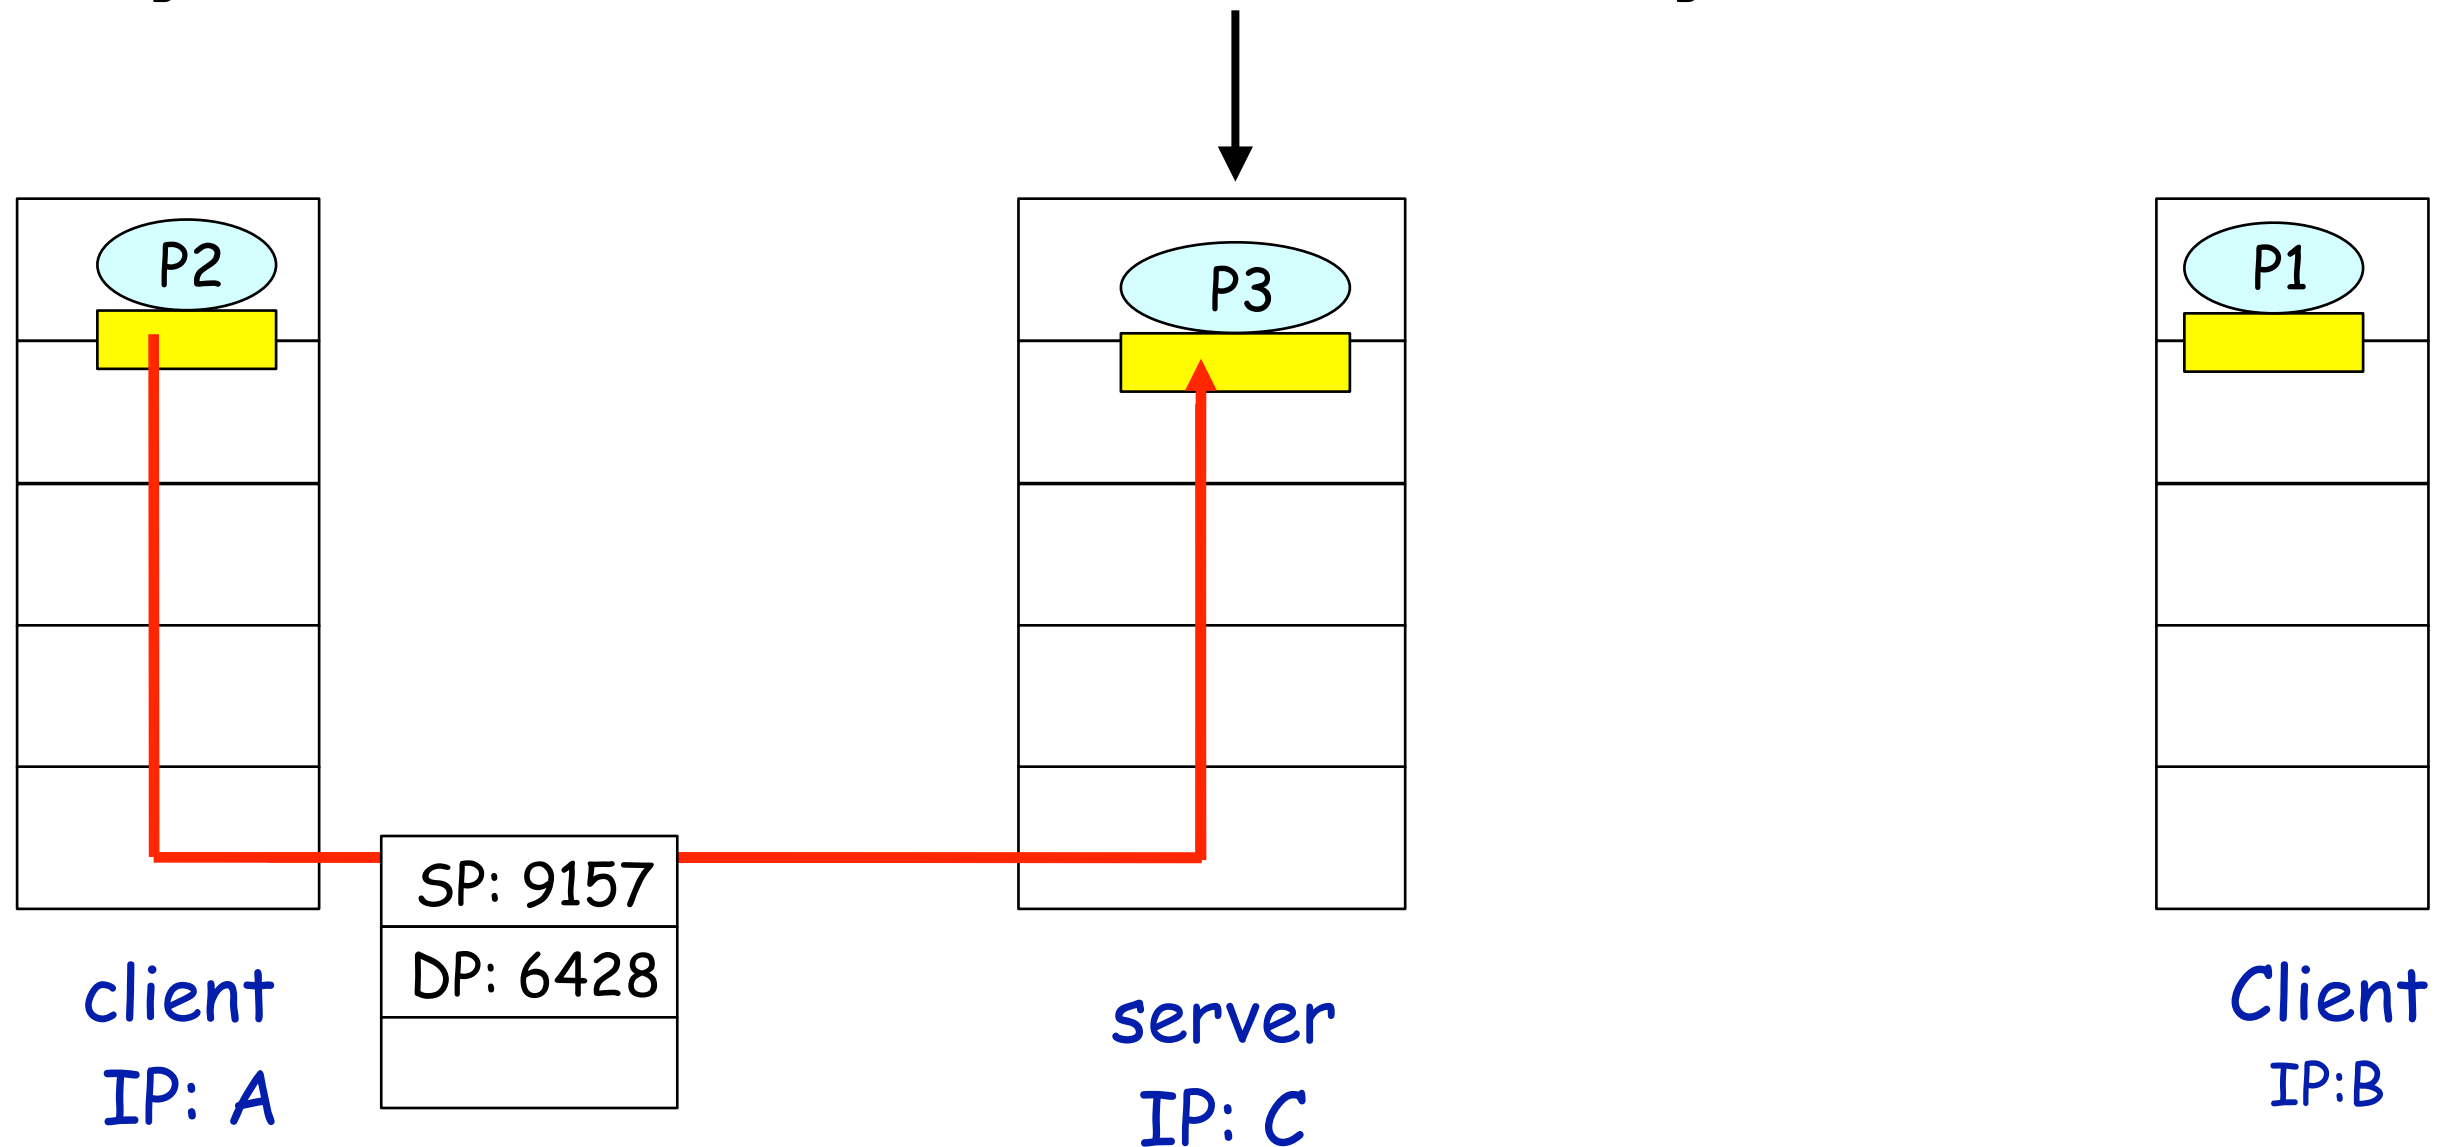
server
IP: C



Client
IP: B

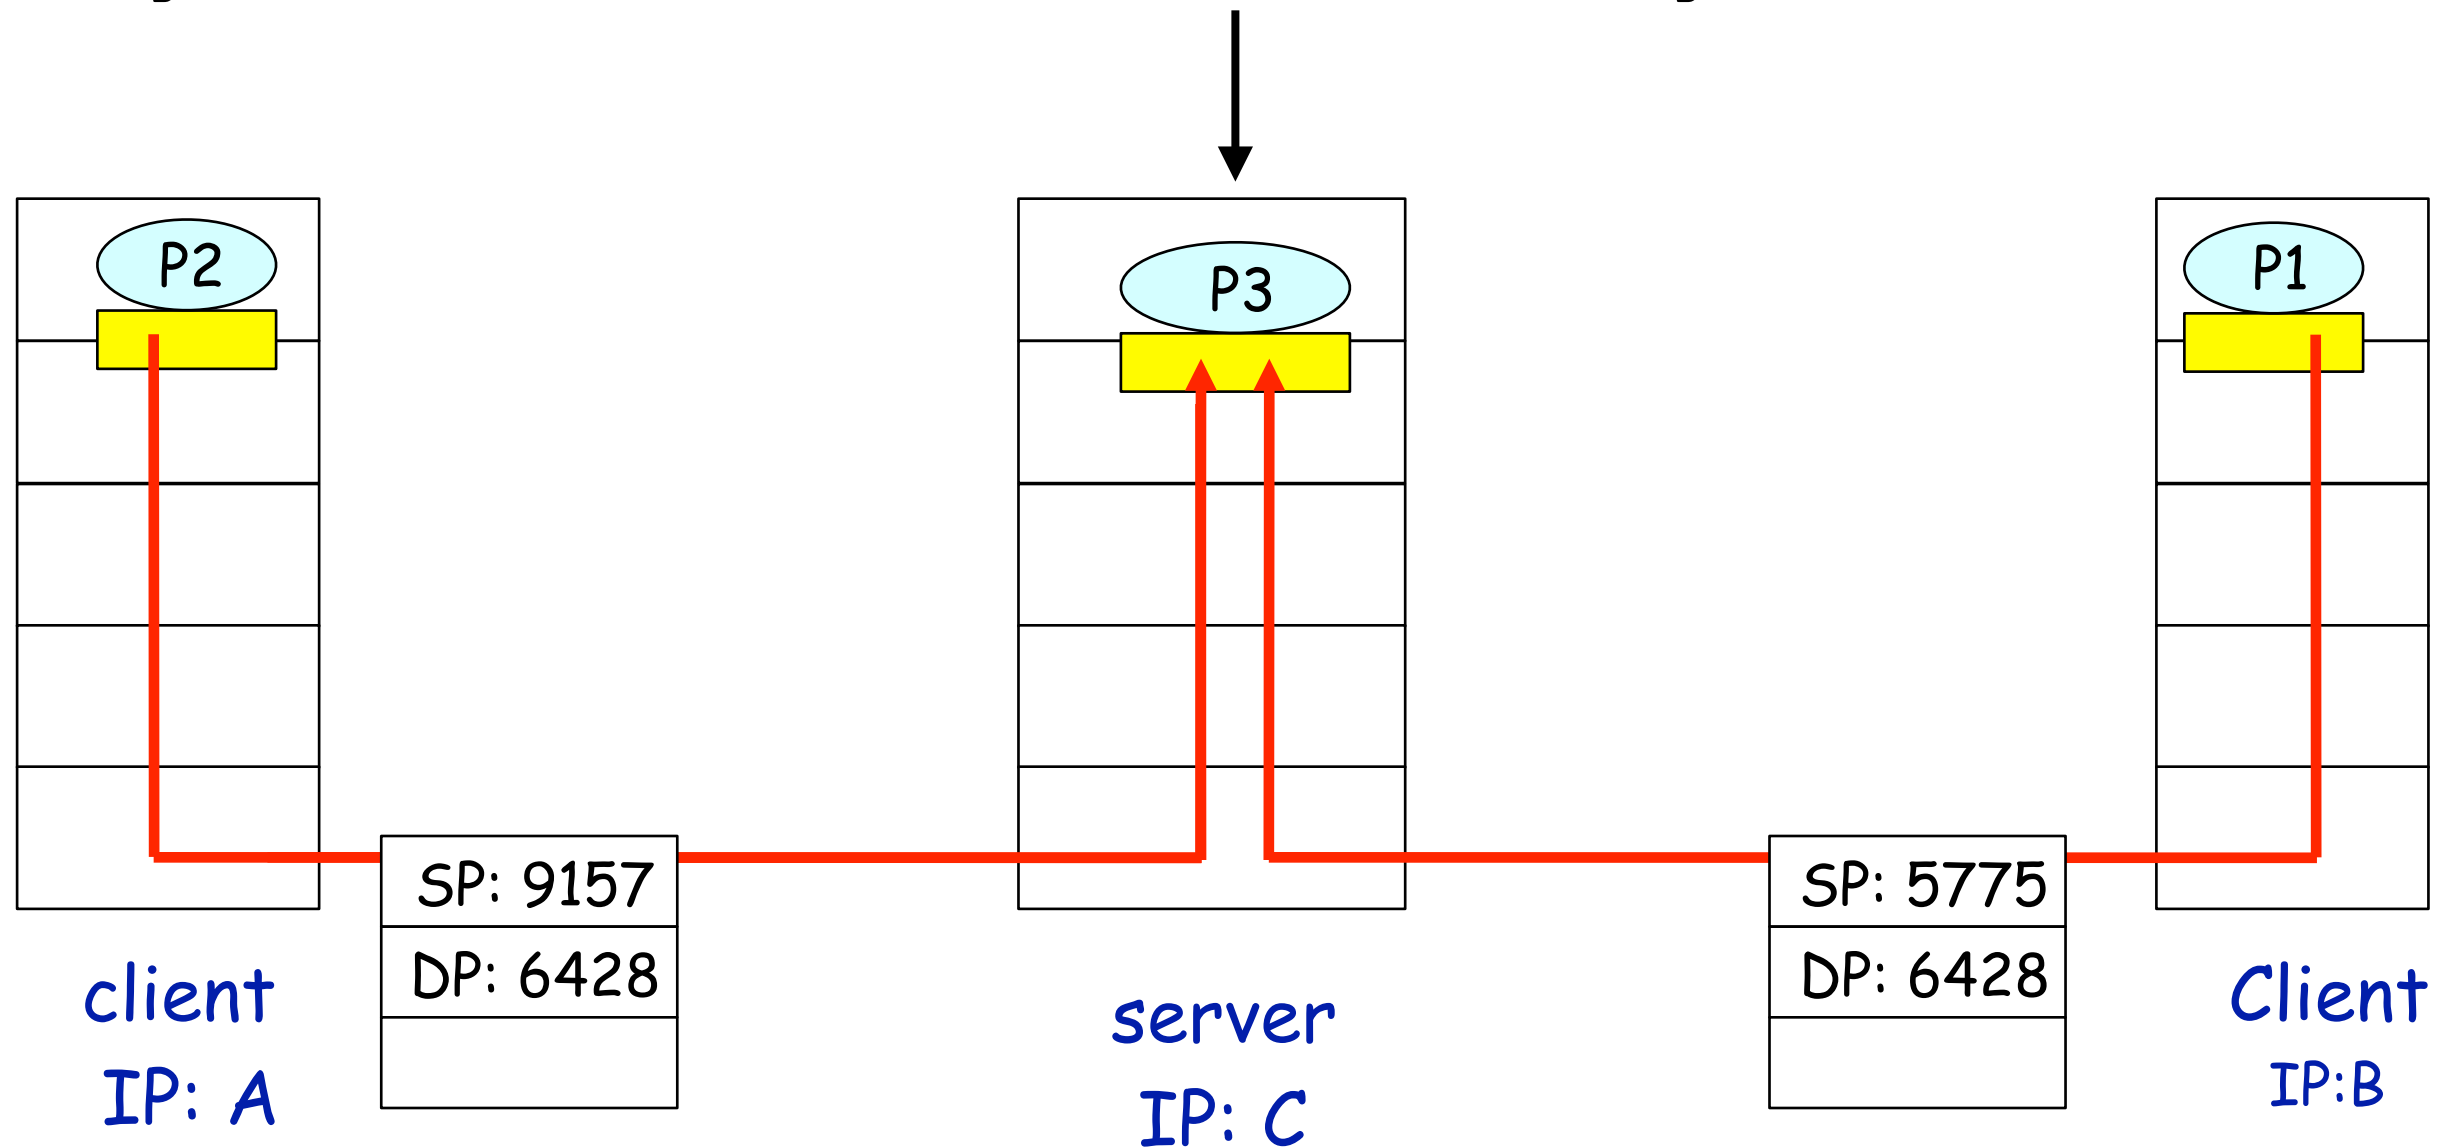
Connectionless demux (cont)

```
DatagramSocket serverSocket = new DatagramSocket(6428);
```



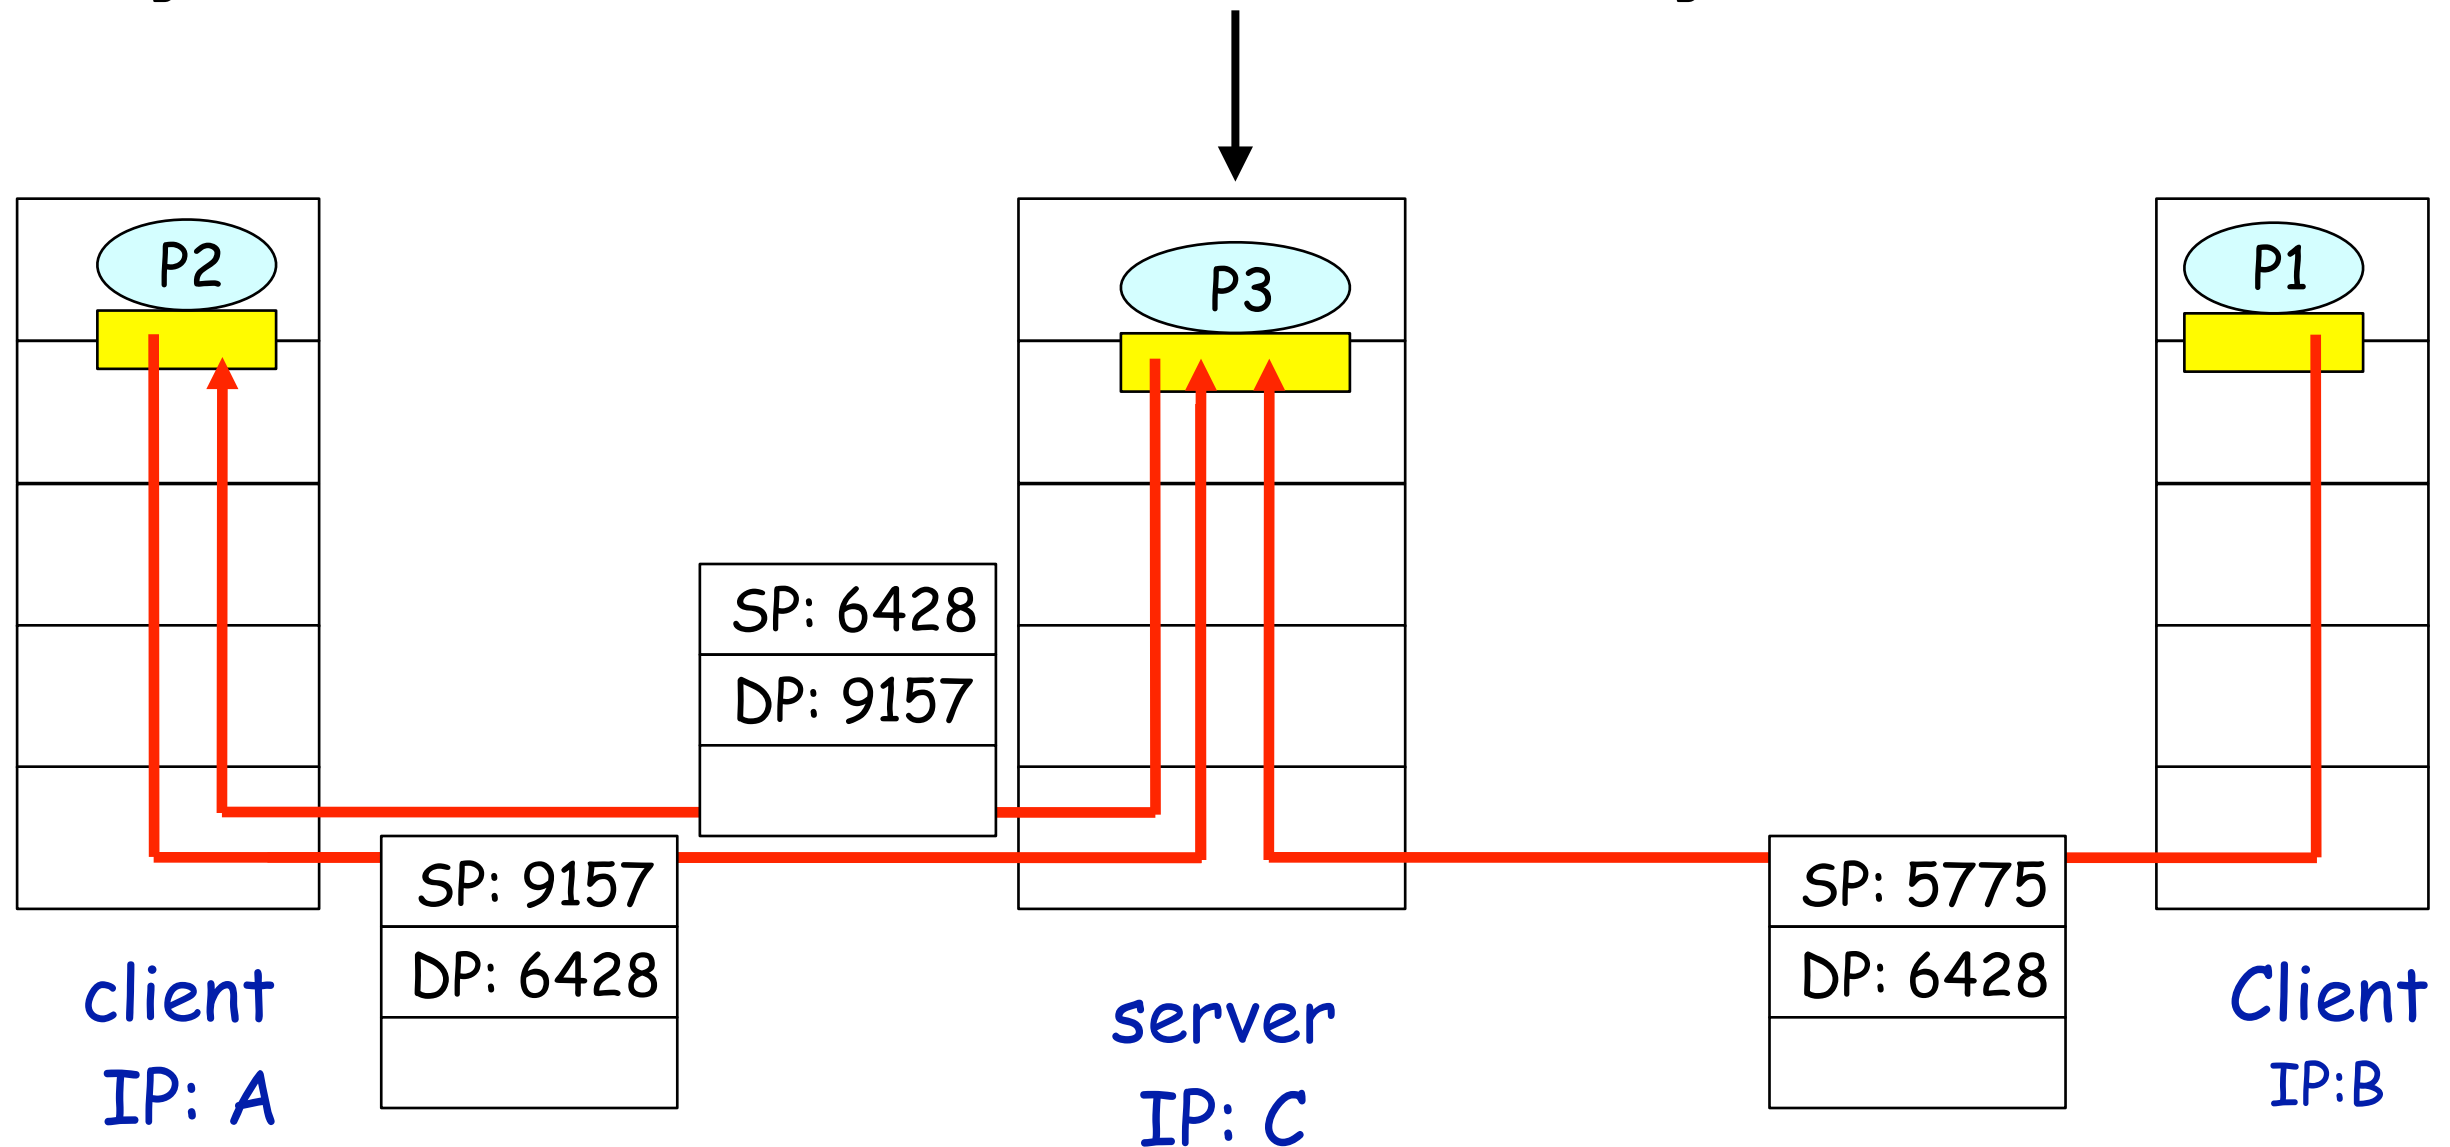
Connectionless demux (cont)

```
DatagramSocket serverSocket = new DatagramSocket(6428);
```




Connectionless demux (cont)

```
DatagramSocket serverSocket = new DatagramSocket(6428);
```



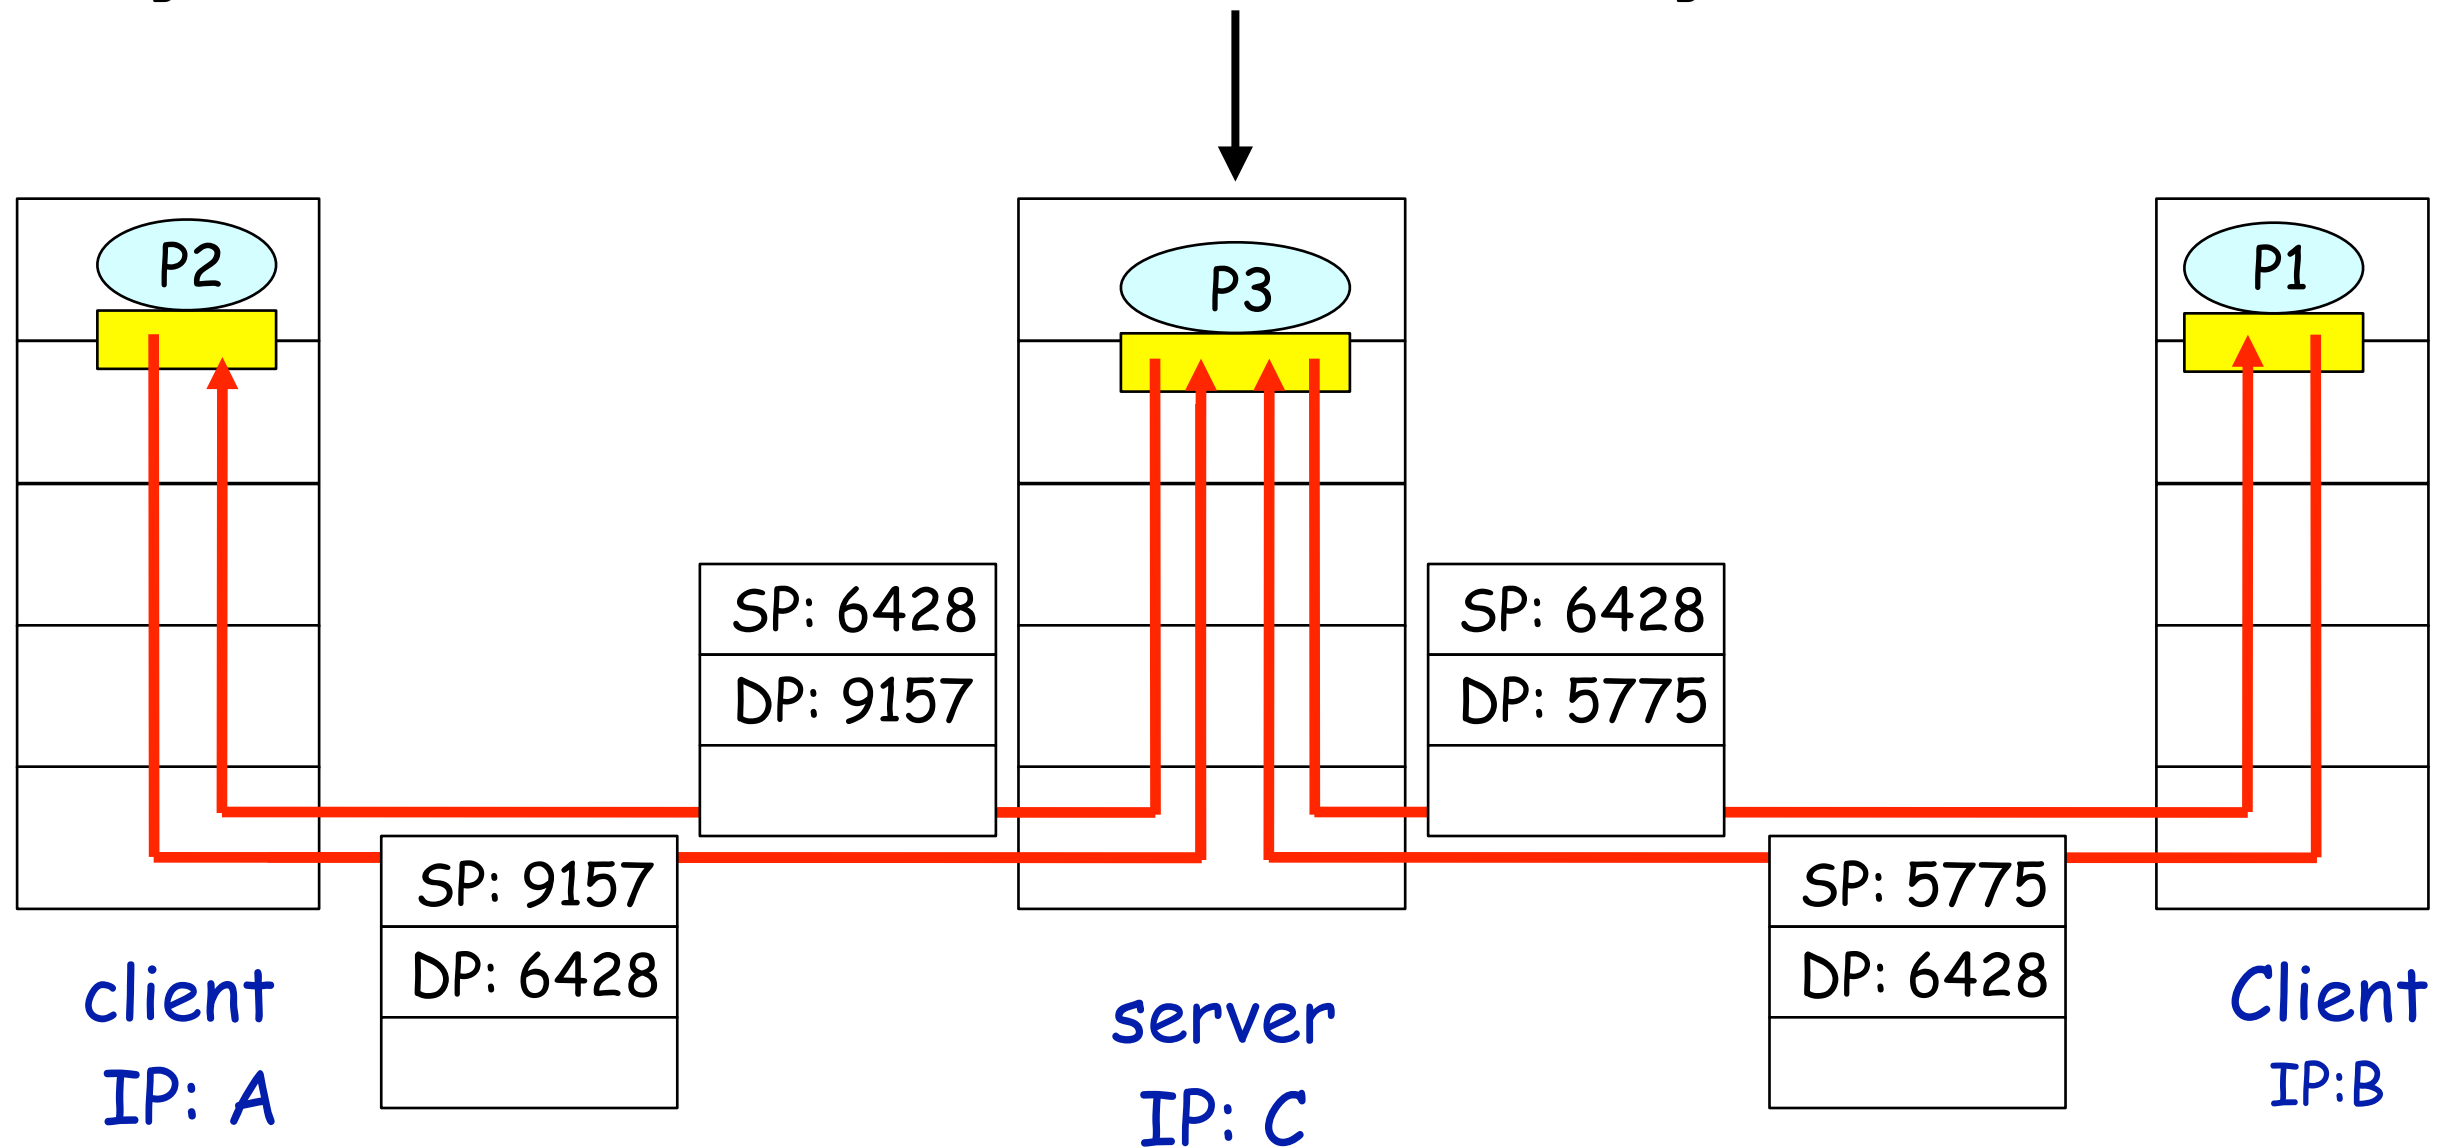
Connectionless demux (cont)

```
DatagramSocket serverSocket = new DatagramSocket(6428);
```



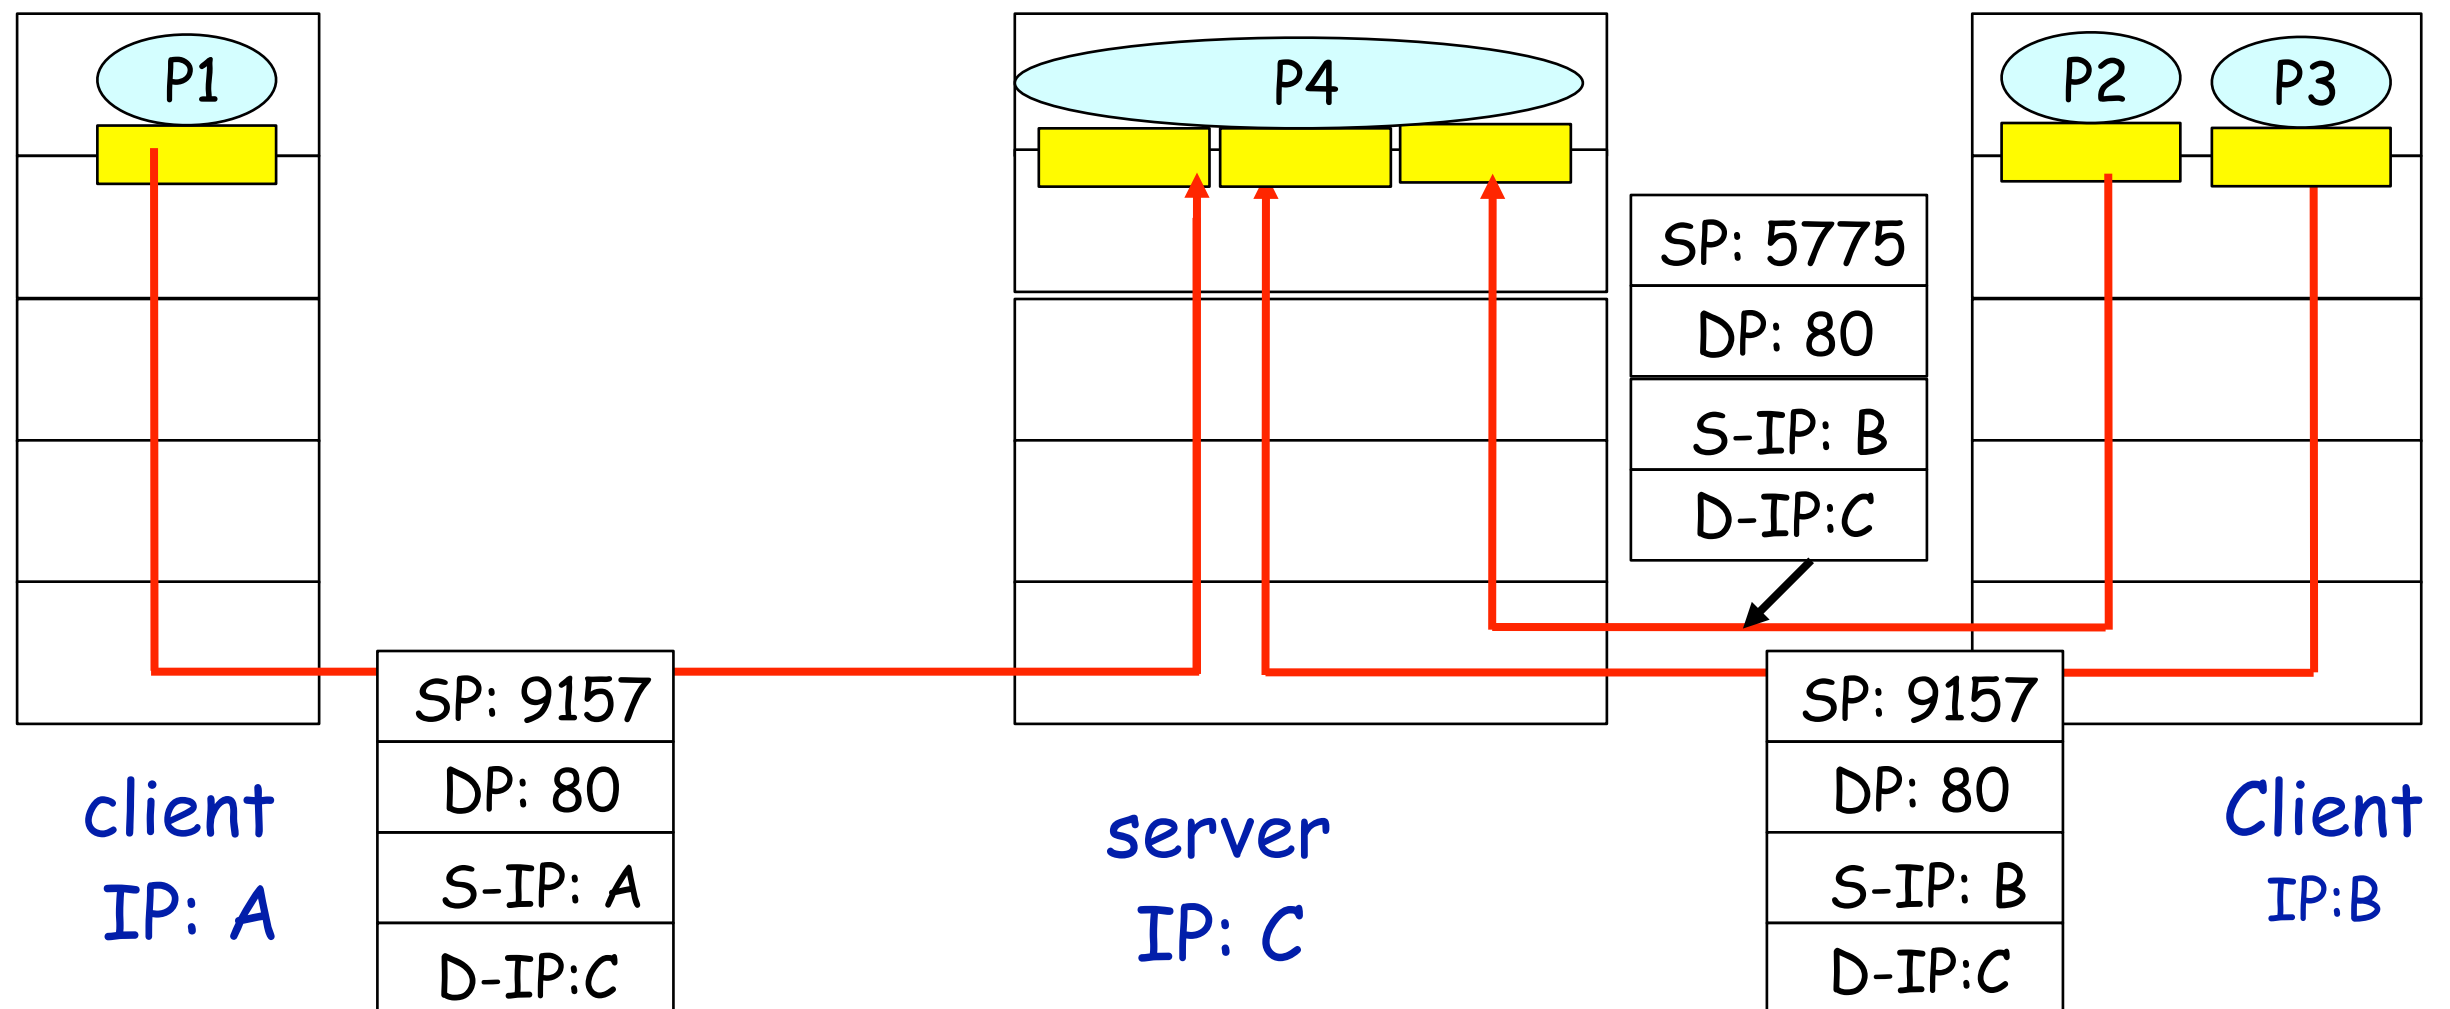
Connectionless demux (cont)

```
DatagramSocket serverSocket = new DatagramSocket(6428);
```

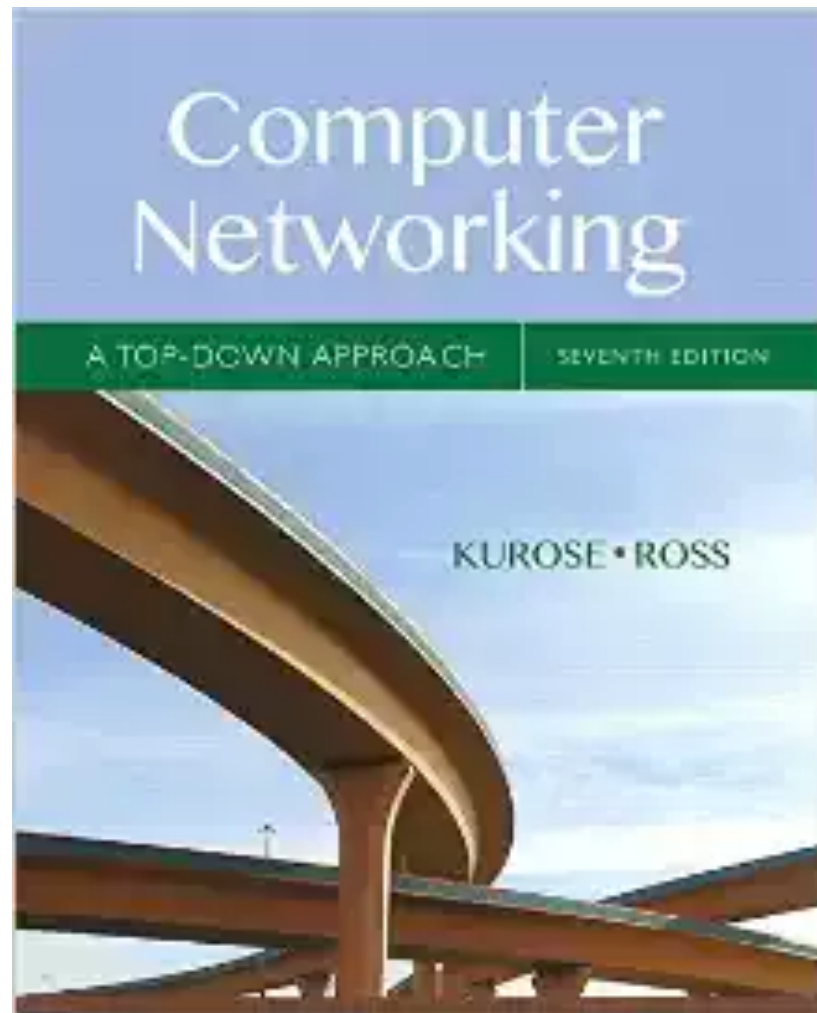


SP provides "return address"

More Complex Demux-ing



Reading Along ...



- 3.3: Connectionless Transport: UDP

UDP: User Datagram Protocol [RFC 768]

UDP: User Datagram Protocol [RFC 768]

❖ "no frills," "bare bones"

Internet transport
protocol

UDP: User Datagram Protocol [RFC 768]

- ❖ "no frills," "bare bones"
Internet transport
protocol
- ❖ "best effort" service, UDP
segments may be:

UDP: User Datagram Protocol [RFC 768]

- ❖ "no frills," "bare bones"
Internet transport
protocol
- ❖ "best effort" service, UDP
segments may be:
 - lost
 - delivered out of order
to app

UDP: User Datagram Protocol [RFC 768]

- ❖ "no frills," "bare bones"

Internet transport
protocol

- ❖ "best effort" service, UDP
segments may be:

- lost
- delivered out of order
to app

- ❖ **connectionless:**

- no handshaking between
UDP sender, receiver
- each UDP segment
handled independently
of others

UDP: User Datagram Protocol [RFC 768]

- ❖ "no frills," "bare bones"
Internet transport
protocol
- ❖ "best effort" service, UDP
segments may be:
 - lost
 - delivered out of order
to app
- ❖ **connectionless:**
 - no handshaking between
UDP sender, receiver
 - each UDP segment
handled independently
of others

Why is there a UDP?

UDP: User Datagram Protocol [RFC 768]

- ❖ "no frills," "bare bones"
Internet transport
protocol
- ❖ "best effort" service, UDP
segments may be:
 - lost
 - delivered out of order
to app
- ❖ **connectionless:**
 - no handshaking between
UDP sender, receiver
 - each UDP segment
handled independently
of others

Why is there a UDP?

- ❖ no connection establishment
(which can add delay)
- ❖ simple: no connection state
at sender, receiver
- ❖ small segment header
- ❖ no congestion control: UDP
can blast away as fast as
desired

UDP

❖ often used for streaming multimedia apps

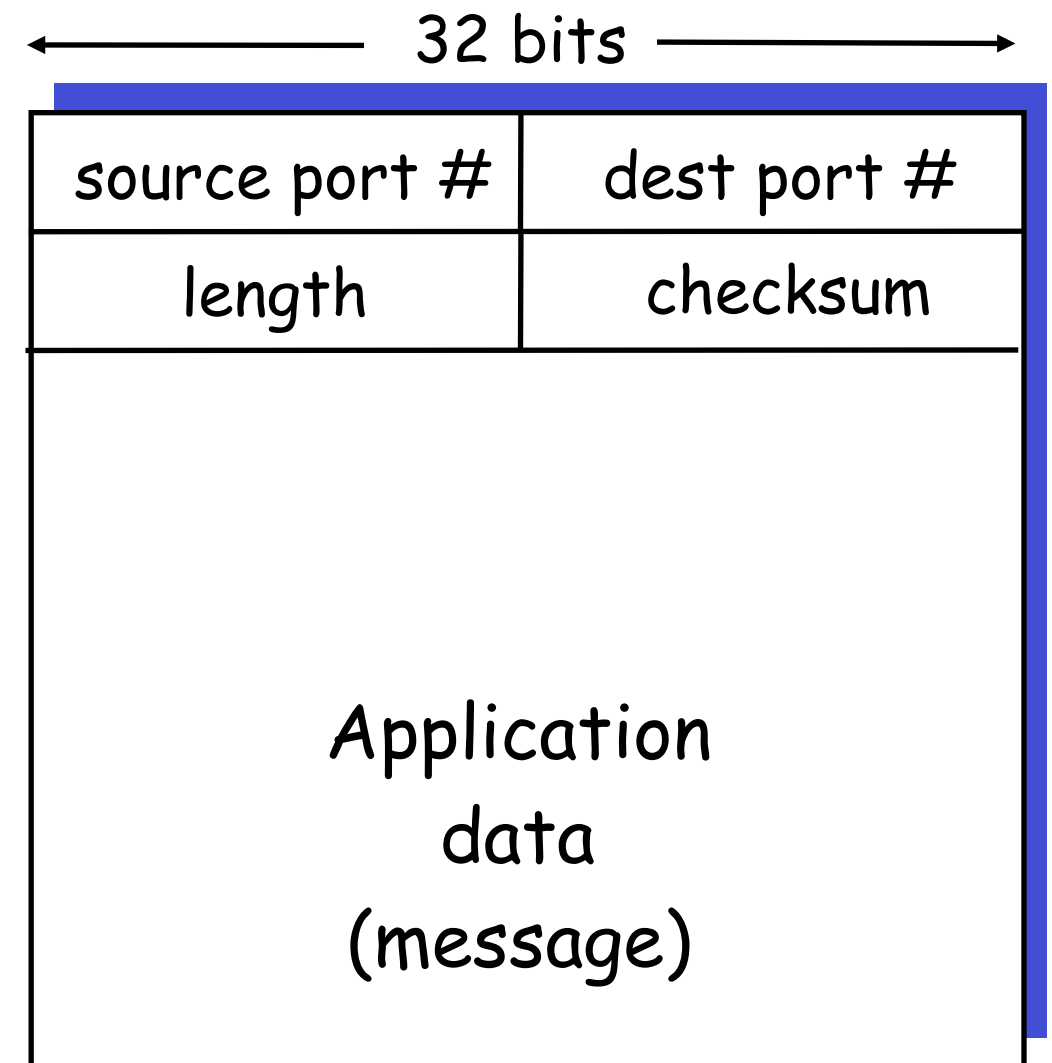
- loss tolerant
- rate controllable

❖ other UDP uses

- DNS
- SNMP

❖ reliable transfer over UDP: add reliability at application layer

- application-specific error recovery!



UDP segment format

UDP

❖ often used for streaming multimedia apps

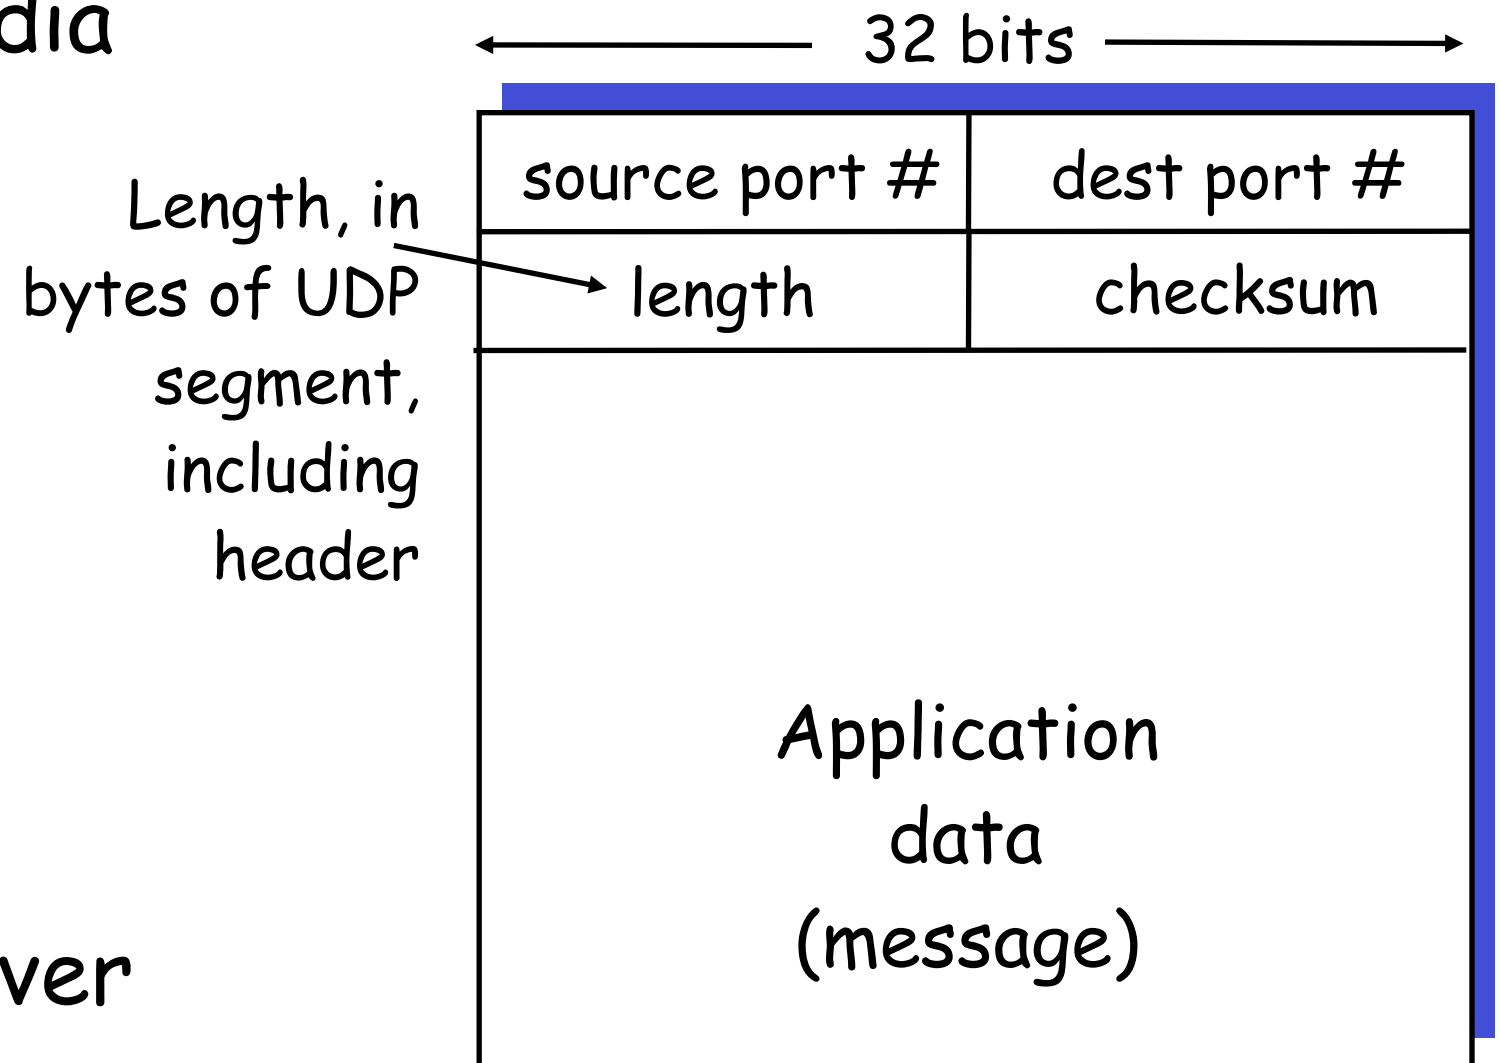
- loss tolerant
- rate controllable

❖ other UDP uses

- DNS
- SNMP

❖ reliable transfer over UDP: add reliability at application layer

- application-specific error recovery!



UDP segment format

UDP checksum

UDP checksum

Goal: detect "errors" (e.g., flipped bits) in transmitted segment

UDP checksum

Goal: detect "errors" (e.g., flipped bits) in transmitted segment

Sender:

- ❖ treat segment contents as sequence of 16-bit integers
- ❖ checksum: addition (1's complement sum) of segment contents
- ❖ sender puts checksum value into UDP checksum field

UDP checksum

Goal: detect "errors" (e.g., flipped bits) in transmitted segment

Sender:

- ❖ treat segment contents as sequence of 16-bit integers
- ❖ checksum: addition (1's complement sum) of segment contents
- ❖ sender puts checksum value into UDP checksum field

Receiver:

- ❖ compute checksum of received segment
- ❖ check if computed checksum equals checksum field value:
 - NO - error detected
 - YES - no error detected.

UDP checksum

Goal: detect "errors" (e.g., flipped bits) in transmitted segment

Sender:

- ❖ treat segment contents as sequence of 16-bit integers
- ❖ checksum: addition (1's complement sum) of segment contents
- ❖ sender puts checksum value into UDP checksum field

Receiver:

- ❖ compute checksum of received segment
- ❖ check if computed checksum equals checksum field value:
 - NO - error detected
 - YES - no error detected.
 - But maybe errors nonetheless?


UDP Packet Header

0	15	16	31
Source Port Number(16 bits)		Destination Port Number(16 bits)	
Length(UDP Header + Data)16 bits		UDP Checksum(16 bits)	
Application Data (Message)			

A Tangent...

❖ Binary number: 0111


Reading Along ...



- 3.4: Principles of reliable data transfer

Principles of Reliable data transfer


- ❖ important in app., transport, link layers
- ❖ top-10 list of important networking topics!



(a) provided service

Principles of Reliable data transfer

- ❖ important in app., transport, link layers
- ❖ top-10 list of important networking topics!



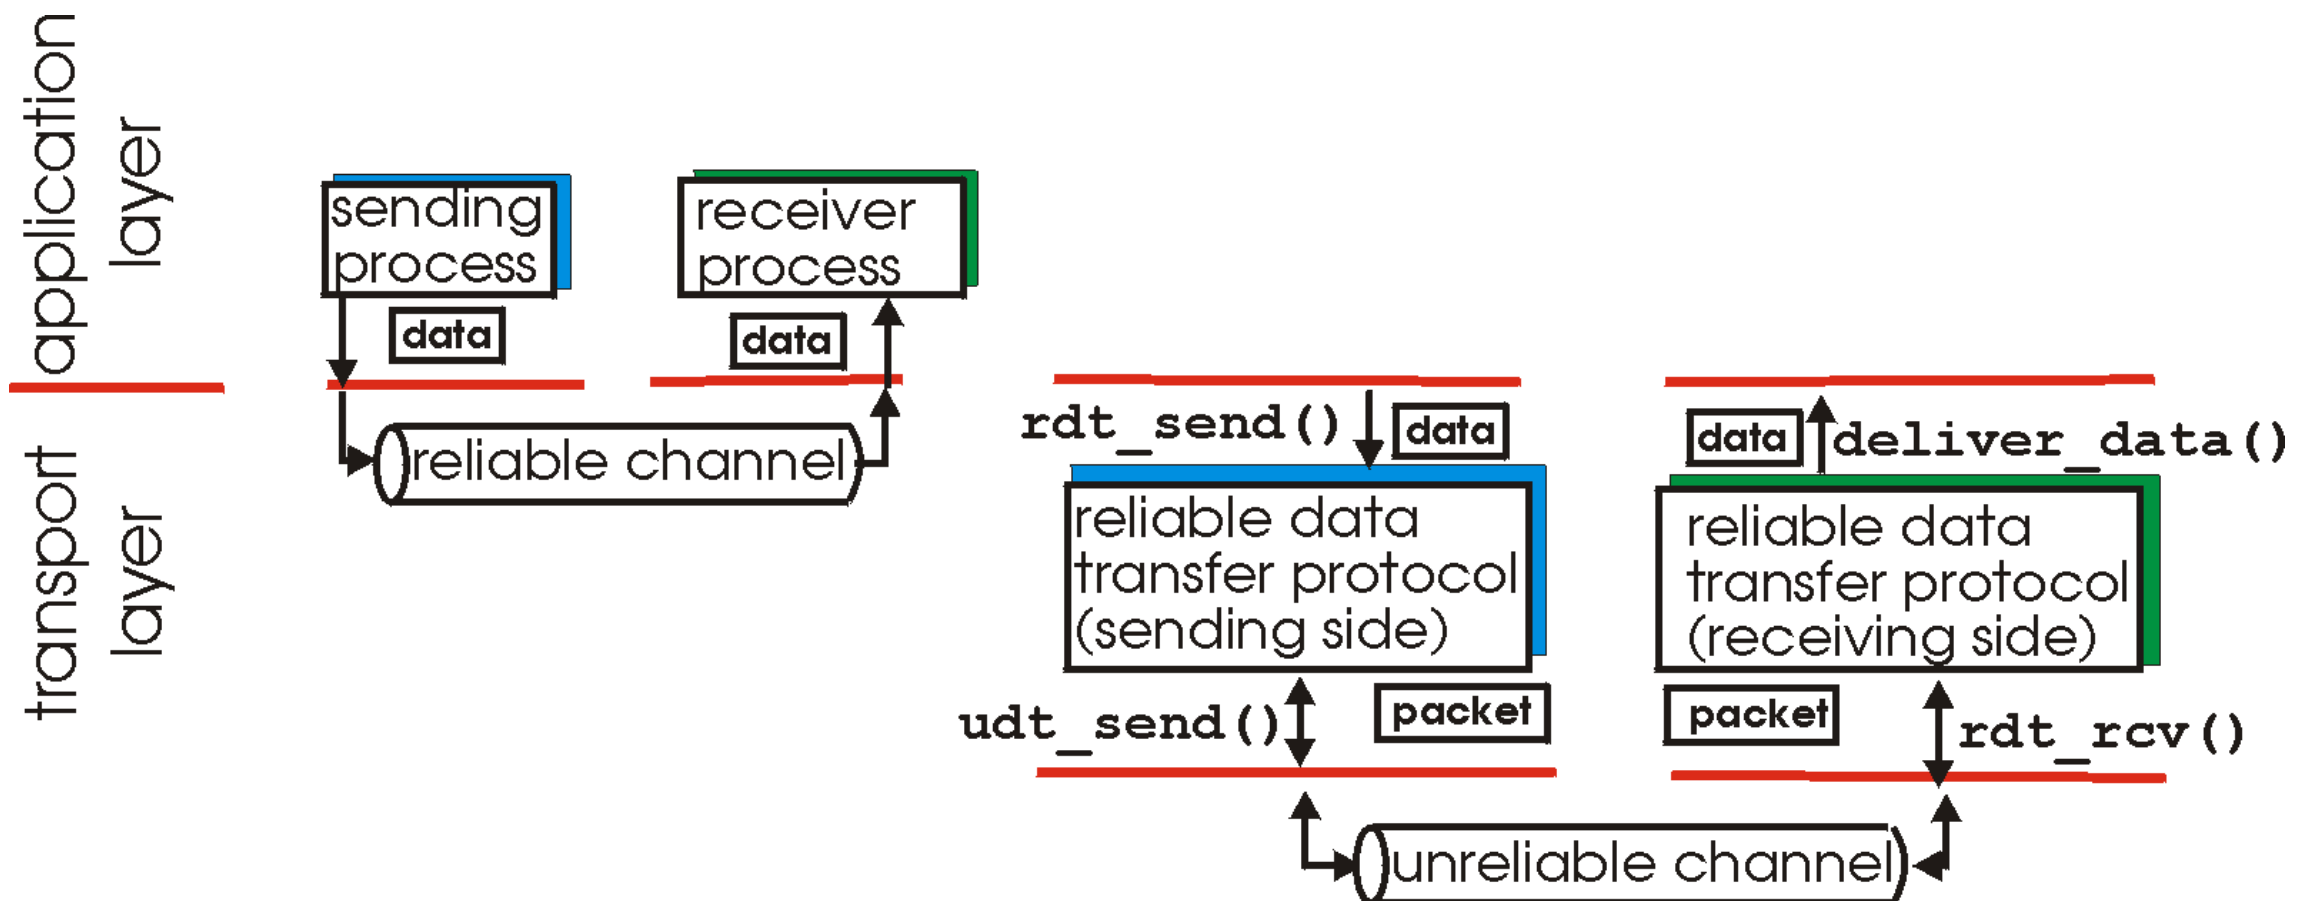
(a) provided service



(b) service implementation

Principles of Reliable data transfer

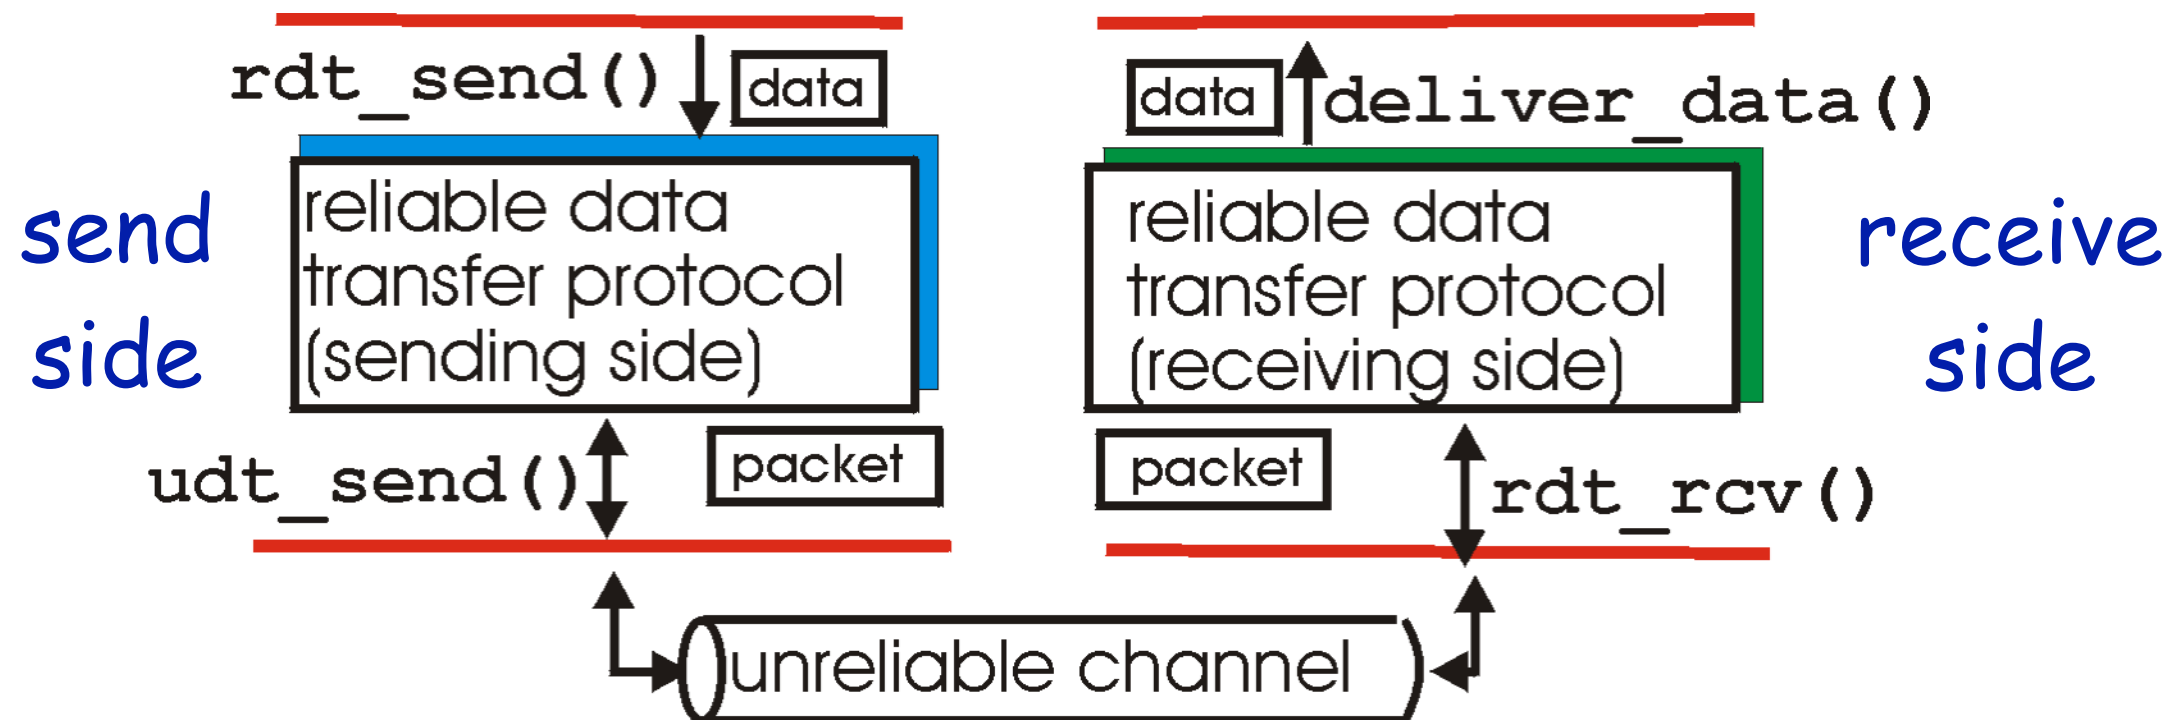
- ❖ important in app., transport, link layers
- ❖ top-10 list of important networking topics!



(a) provided service

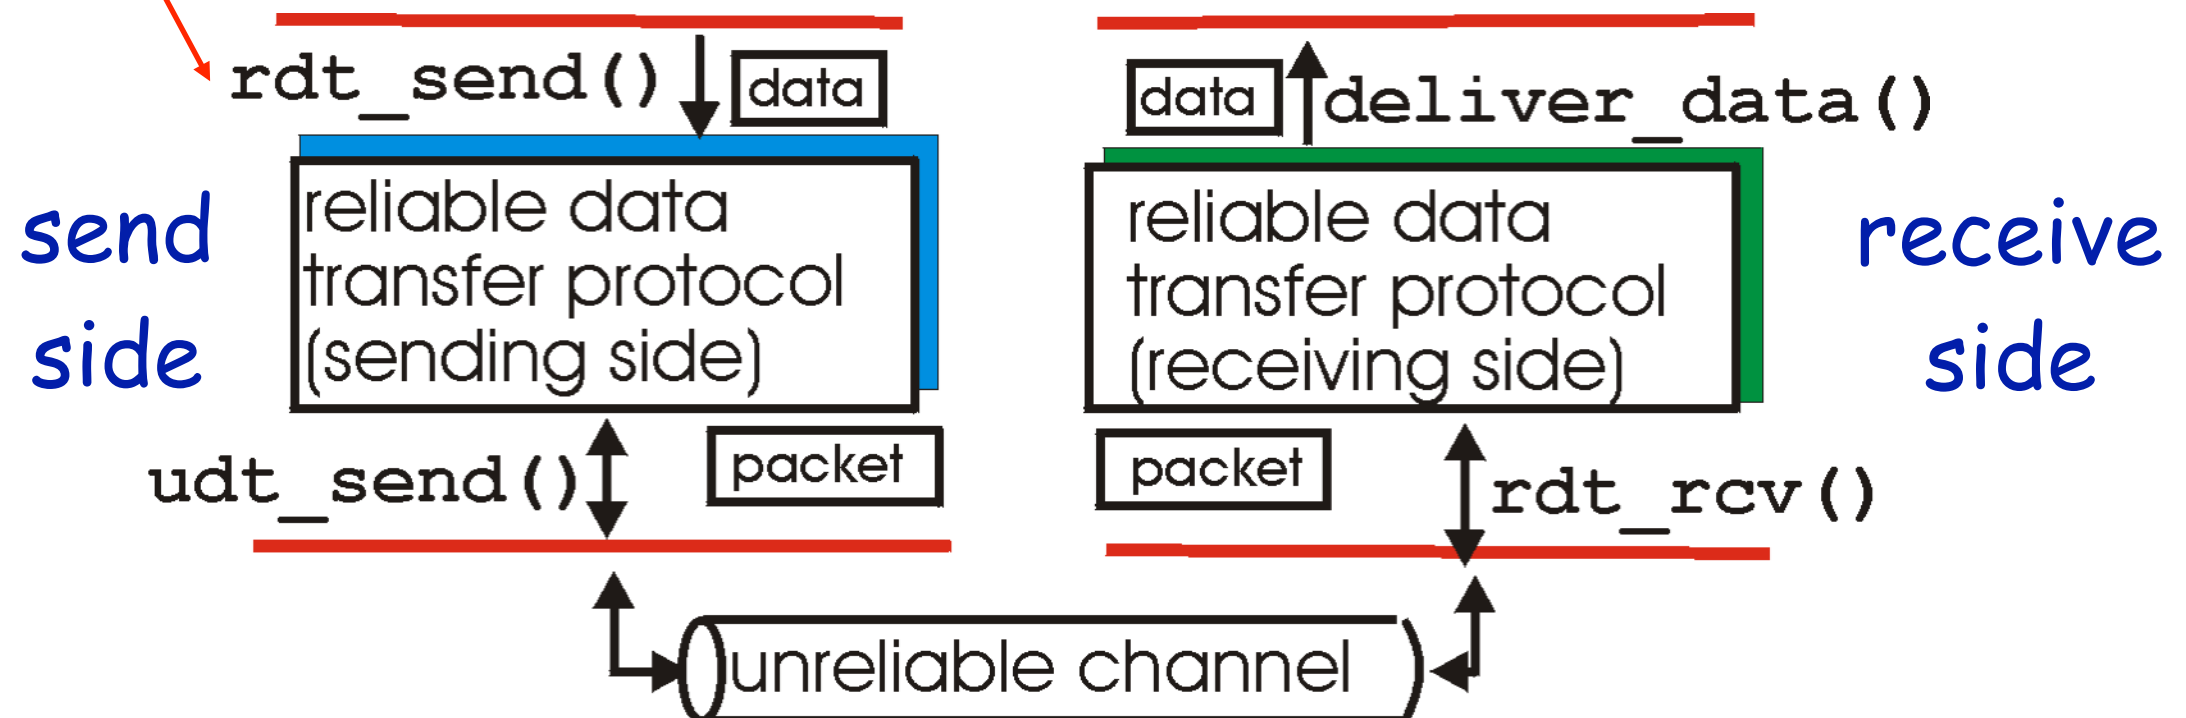
(b) service implementation

Reliable data transfer: getting started



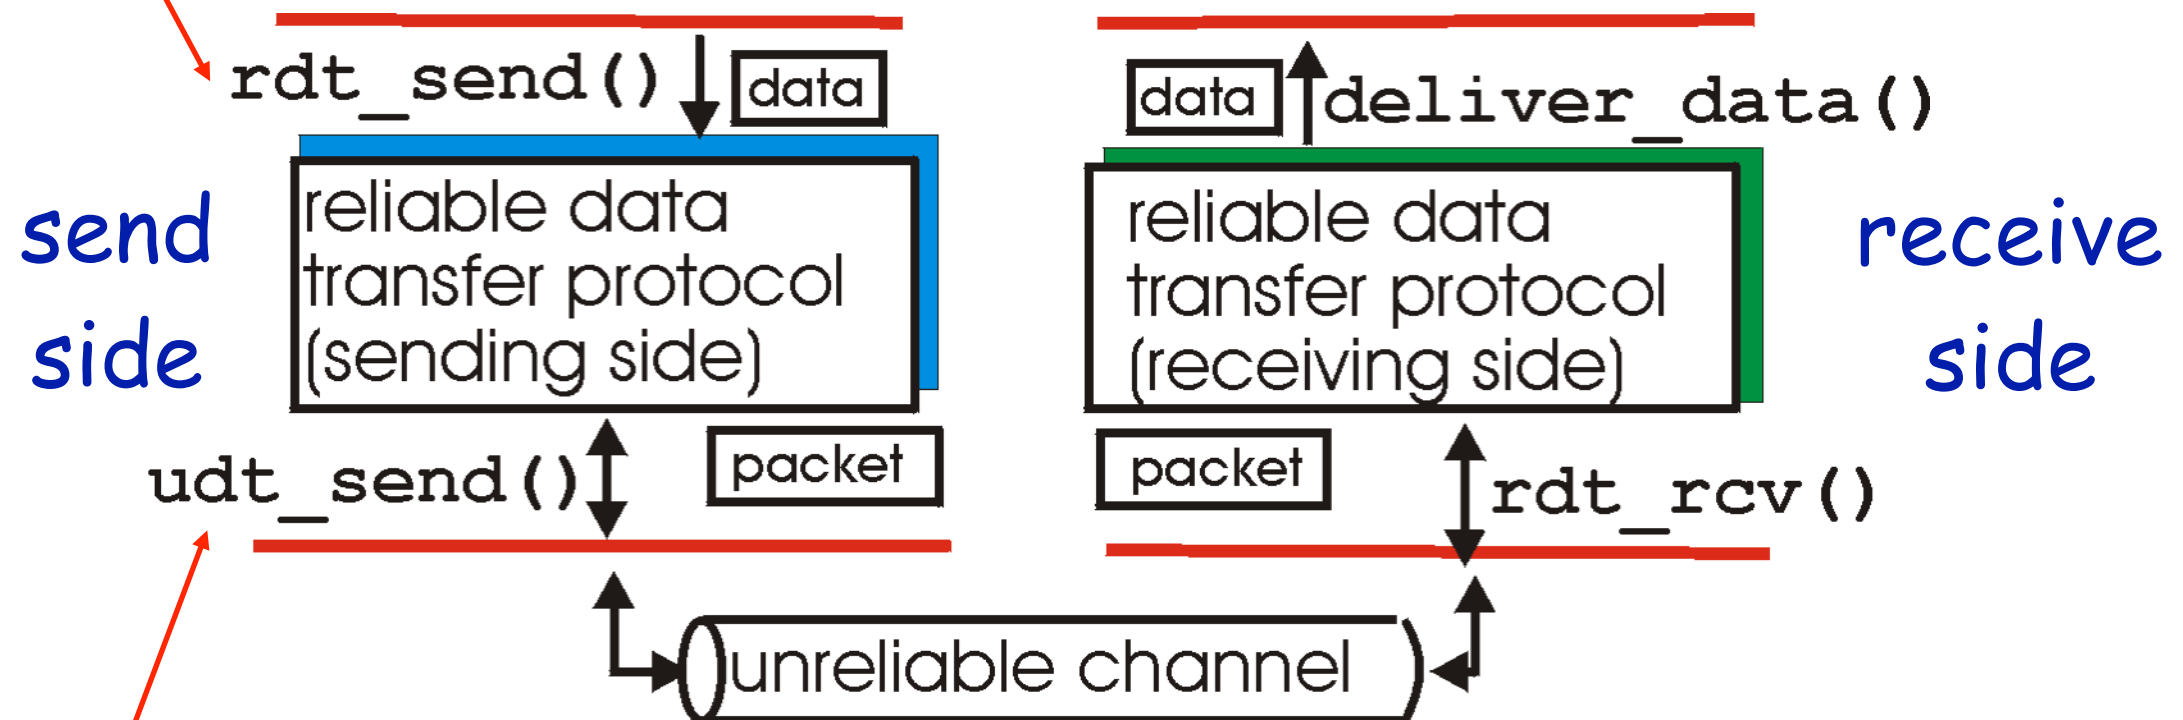
Reliable data transfer: getting started

rdt_send() : called from above,
(e.g., by app.). Passed data to
deliver to receiver upper layer



Reliable data transfer: getting started

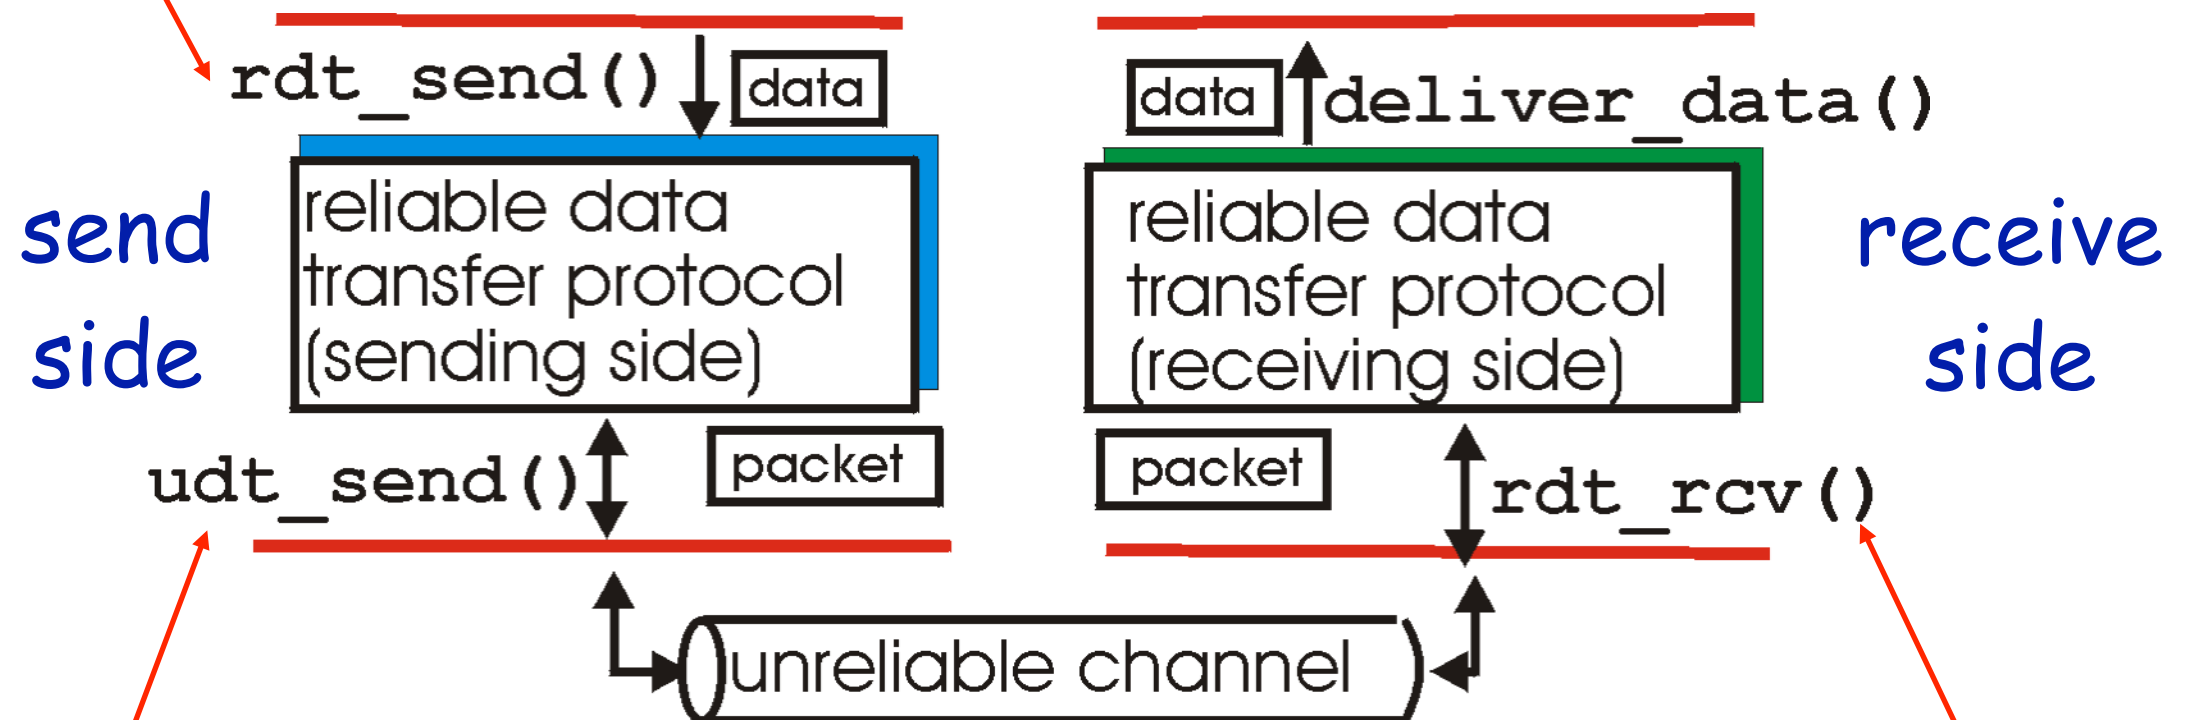
rdt_send() : called from above,
(e.g., by app.). Passed data to
deliver to receiver upper layer



udt_send() : called by rdt,
to transfer packet over
unreliable channel to receiver

Reliable data transfer: getting started

rdt_send() : called from above,
(e.g., by app.). Passed data to
deliver to receiver upper layer

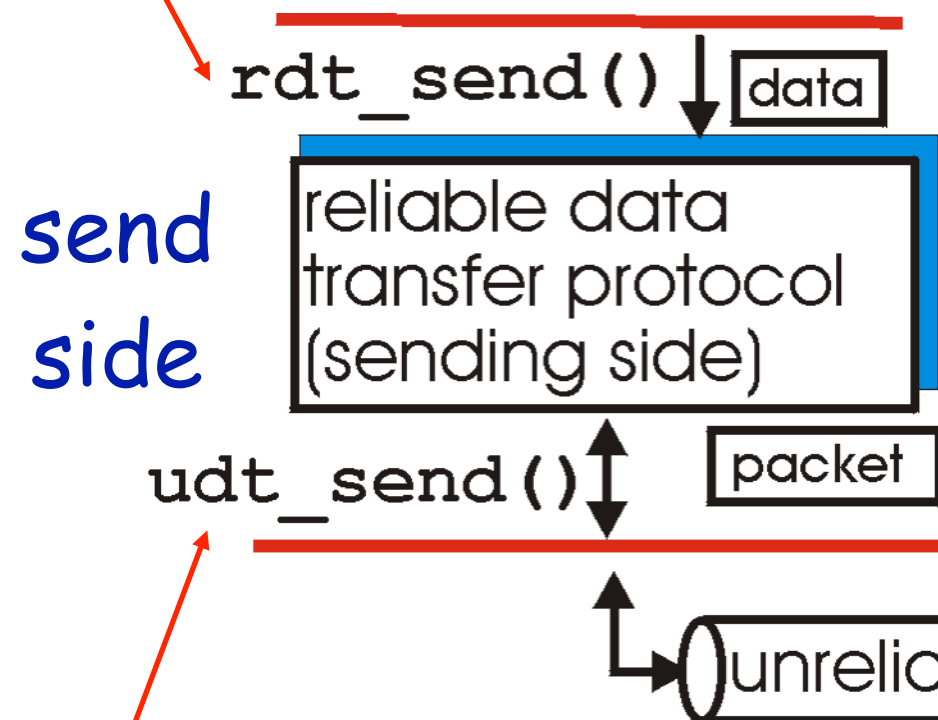


udt_send() : called by rdt,
to transfer packet over
unreliable channel to receiver

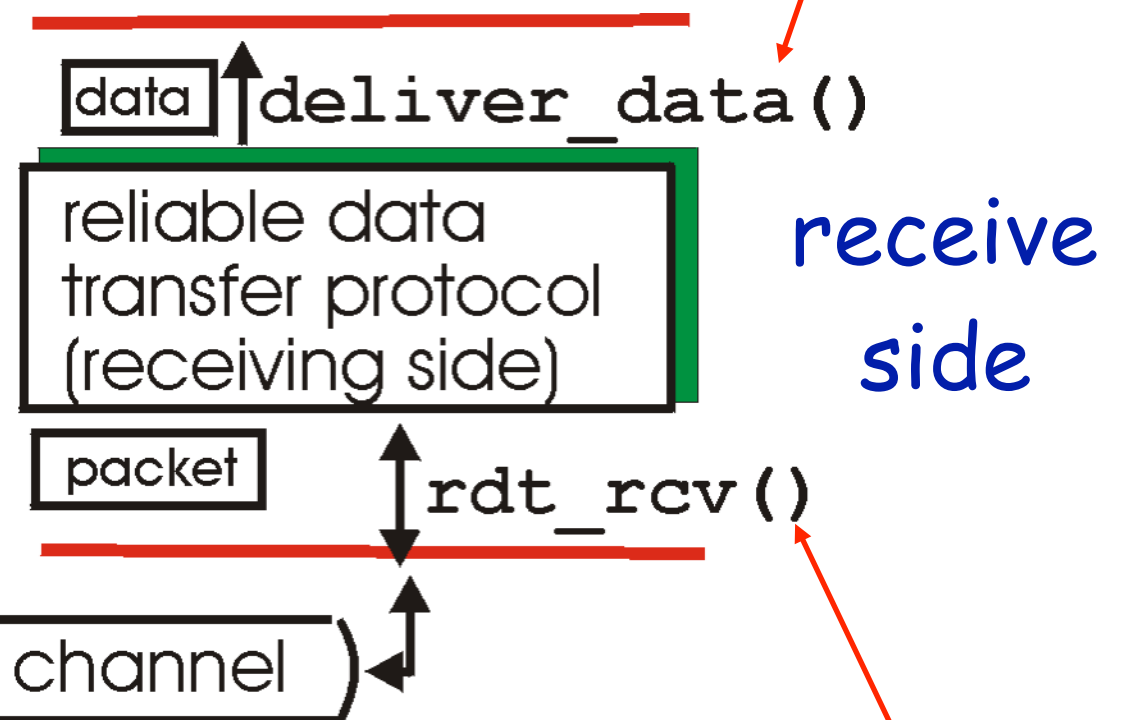
rdt_rcv() : called when packet
arrives on rcv-side of channel

Reliable data transfer: getting started

rdt_send() : called from above,
(e.g., by app.). Passed data to
deliver to receiver upper layer



deliver_data() : called by
rdt to deliver data to upper



udt_send() : called by rdt,
to transfer packet over
unreliable channel to receiver

rdt_rcv() : called when packet
arrives on rcv-side of channel

Reliable data transfer: getting started

Reliable data transfer: getting started

We'll:



- ❖ incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- ❖ consider only unidirectional data transfer
 - but control info will flow in both directions!
- ❖ use finite state machines (FSM) to specify sender, receiver

Reliable data transfer: getting started

We'll:

- ❖ incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- ❖ consider only unidirectional data transfer
 - but control info will flow in both directions!
- ❖ use finite state machines (FSM) to specify sender, receiver

state: when in this "state" next state uniquely determined by next event

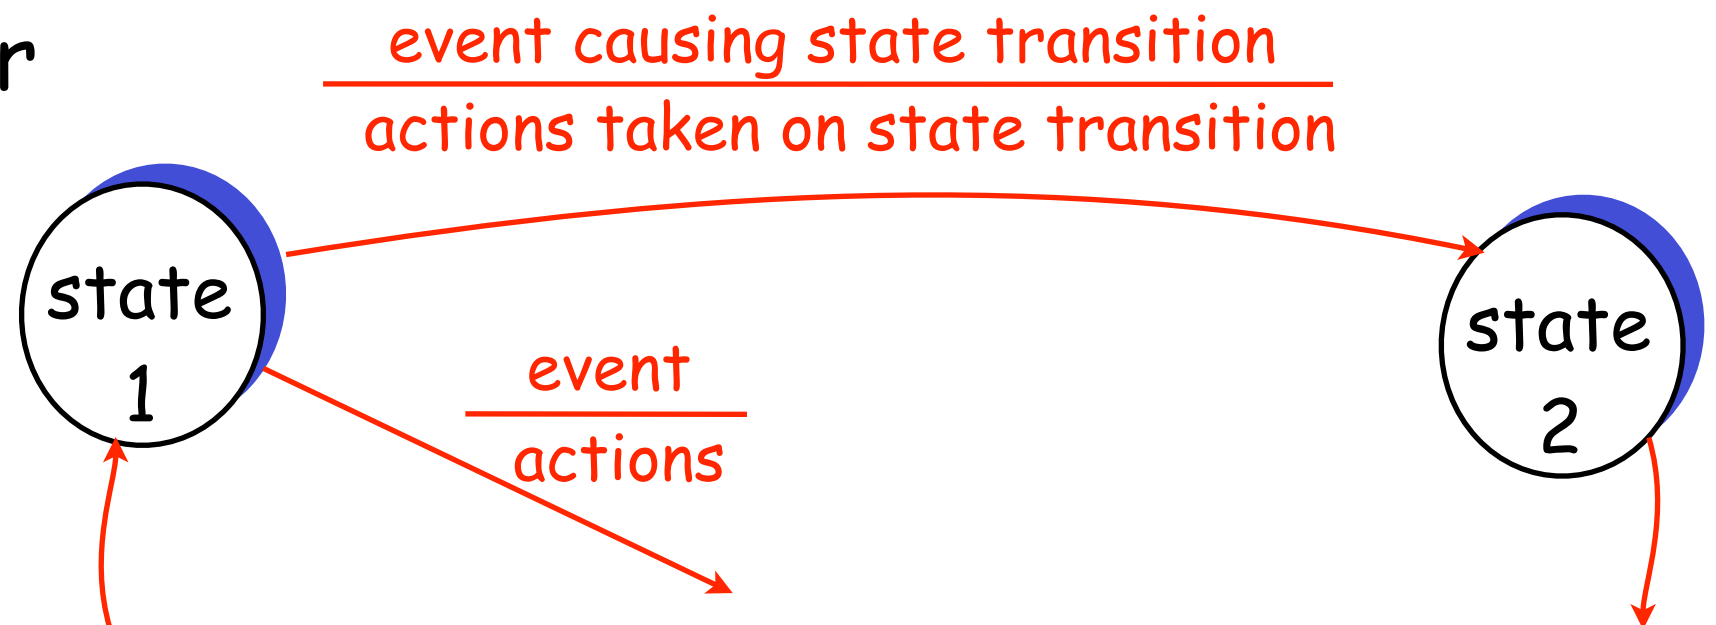


Reliable data transfer: getting started

We'll:

- ❖ incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- ❖ consider only unidirectional data transfer
 - but control info will flow in both directions!
- ❖ use finite state machines (FSM) to specify sender, receiver

state: when in this "state" next state uniquely determined by next event



Rdt1.0: reliable transfer over a reliable channel

Rdt1.0: reliable transfer over a reliable channel

❖ underlying channel perfectly reliable

- no bit errors
- no loss of packets

Rdt1.0: reliable transfer over a reliable channel

- ❖ underlying channel perfectly reliable
 - no bit errors
 - no loss of packets
- ❖ separate FSMs for sender, receiver:
 - sender sends data into underlying channel
 - receiver read data from underlying channel


Rdt1.0: reliable transfer over a reliable channel

❖ underlying channel perfectly reliable

- no bit errors
- no loss of packets

❖ separate FSMs for sender, receiver:

- sender sends data into underlying channel
- receiver read data from underlying channel



sender

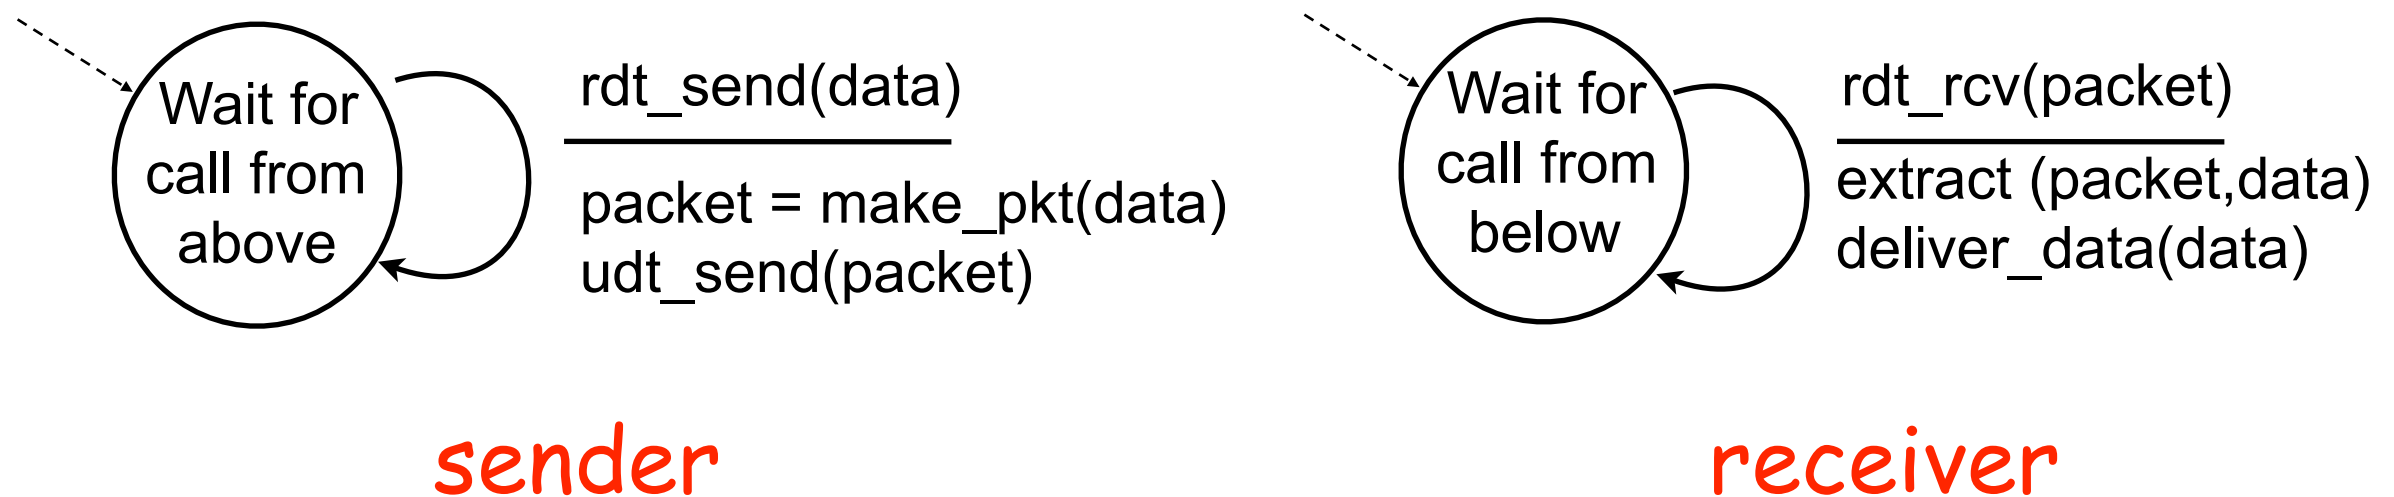
Rdt1.0: reliable transfer over a reliable channel

❖ underlying channel perfectly reliable

- no bit errors
- no loss of packets

❖ separate FSMs for sender, receiver:

- sender sends data into underlying channel
- receiver read data from underlying channel



Rdt 1.0

Rdt 1.0

- ❖ Obviously unrealistic assumptions!
- ❖ If that was all there was to it, there'd be little reason to even bring it up!

Rdt2.0: channel with bit errors

Rdt2.0: channel with bit errors

How do humans recover
from "errors"
during conversation?

Rdt2.0: channel with bit errors

Rdt2.0: channel with bit errors

- ❖ underlying channel may flip bits in packet
 - checksum to detect bit errors

Rdt2.0: channel with bit errors

- ❖ underlying channel may flip bits in packet
 - checksum to detect bit errors
- ❖ the question: how to recover from errors:

Rdt2.0: channel with bit errors

- ❖ underlying channel may flip bits in packet
 - checksum to detect bit errors
- ❖ the question: how to recover from errors:
 - **acknowledgements (ACKs)**: receiver explicitly tells sender that pkt received OK

Rdt2.0: channel with bit errors

- ❖ underlying channel may flip bits in packet
 - checksum to detect bit errors
- ❖ the question: how to recover from errors:
 - **acknowledgements (ACKs)**: receiver explicitly tells sender that pkt received OK
 - **negative acknowledgements (NAKs)**: receiver explicitly tells sender that pkt had errors

Rdt2.0: channel with bit errors

- ❖ underlying channel may flip bits in packet
 - checksum to detect bit errors
- ❖ the question: how to recover from errors:
 - **acknowledgements (ACKs)**: receiver explicitly tells sender that pkt received OK
 - **negative acknowledgements (NAKs)**: receiver explicitly tells sender that pkt had errors
 - sender retransmits pkt on receipt of NAK


Rdt2.0: channel with bit errors

- ❖ underlying channel may flip bits in packet
 - checksum to detect bit errors
- ❖ the question: how to recover from errors:
 - **acknowledgements (ACKs)**: receiver explicitly tells sender that pkt received OK
 - **negative acknowledgements (NAKs)**: receiver explicitly tells sender that pkt had errors
 - sender retransmits pkt on receipt of NAK
- ❖ new mechanisms in rdt2.0 (beyond rdt1.0):
 - error detection
 - receiver feedback: control msgs (ACK,NAK) rcvr->sender

rdt2.0: FSM specification

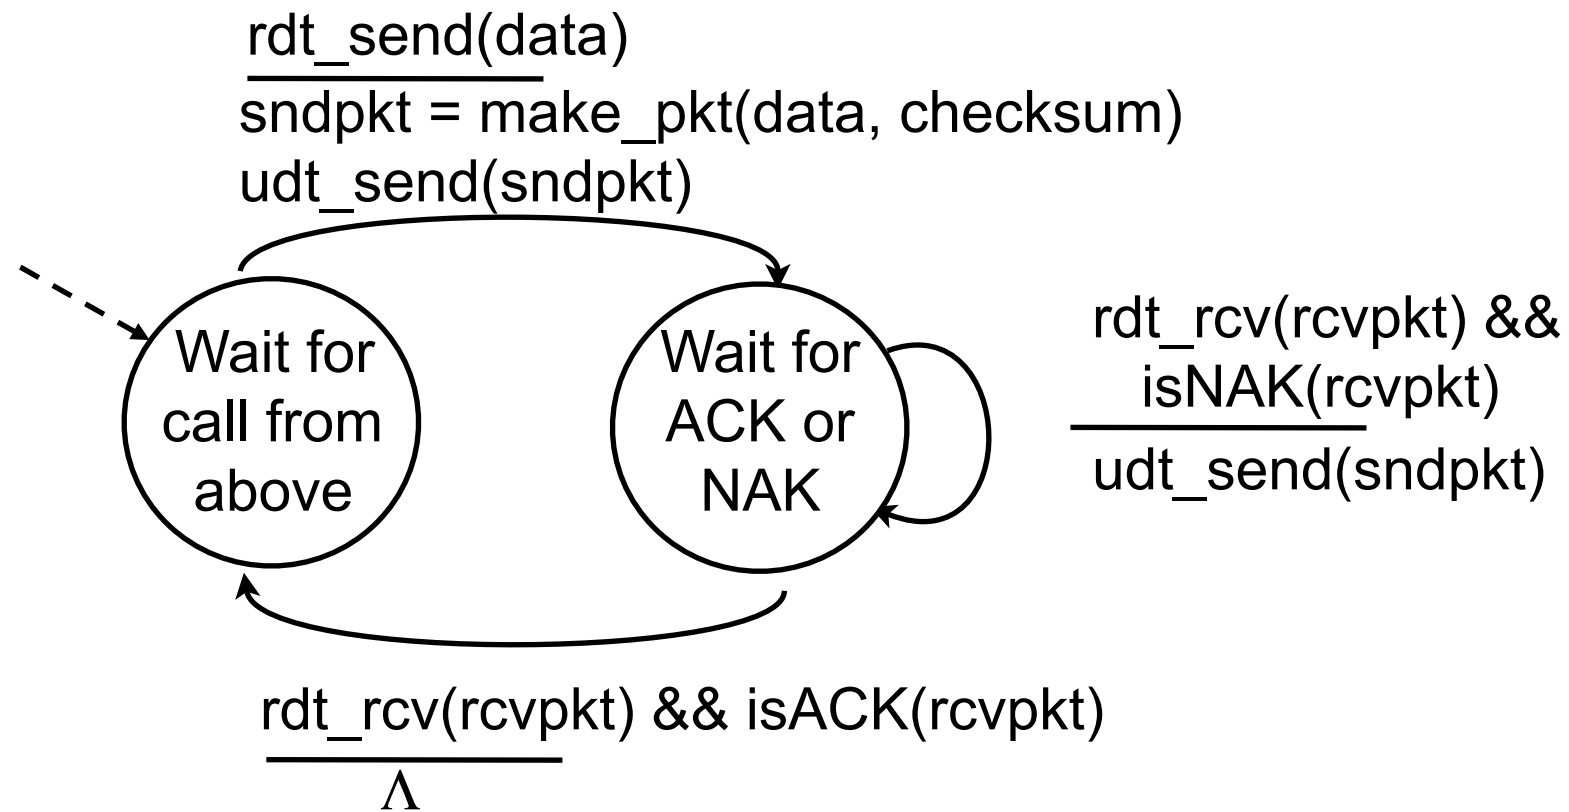
receiver

rdt_rcv(rcvpkt) &&
corrupt(rcvpkt)
udt_send(NAK)




rdt_rcv(rcvpkt) &&
notcorrupt(rcvpkt)
extract(rcvpkt,data)
deliver_data(data)
udt_send(ACK)

rdt2.0: FSM specification

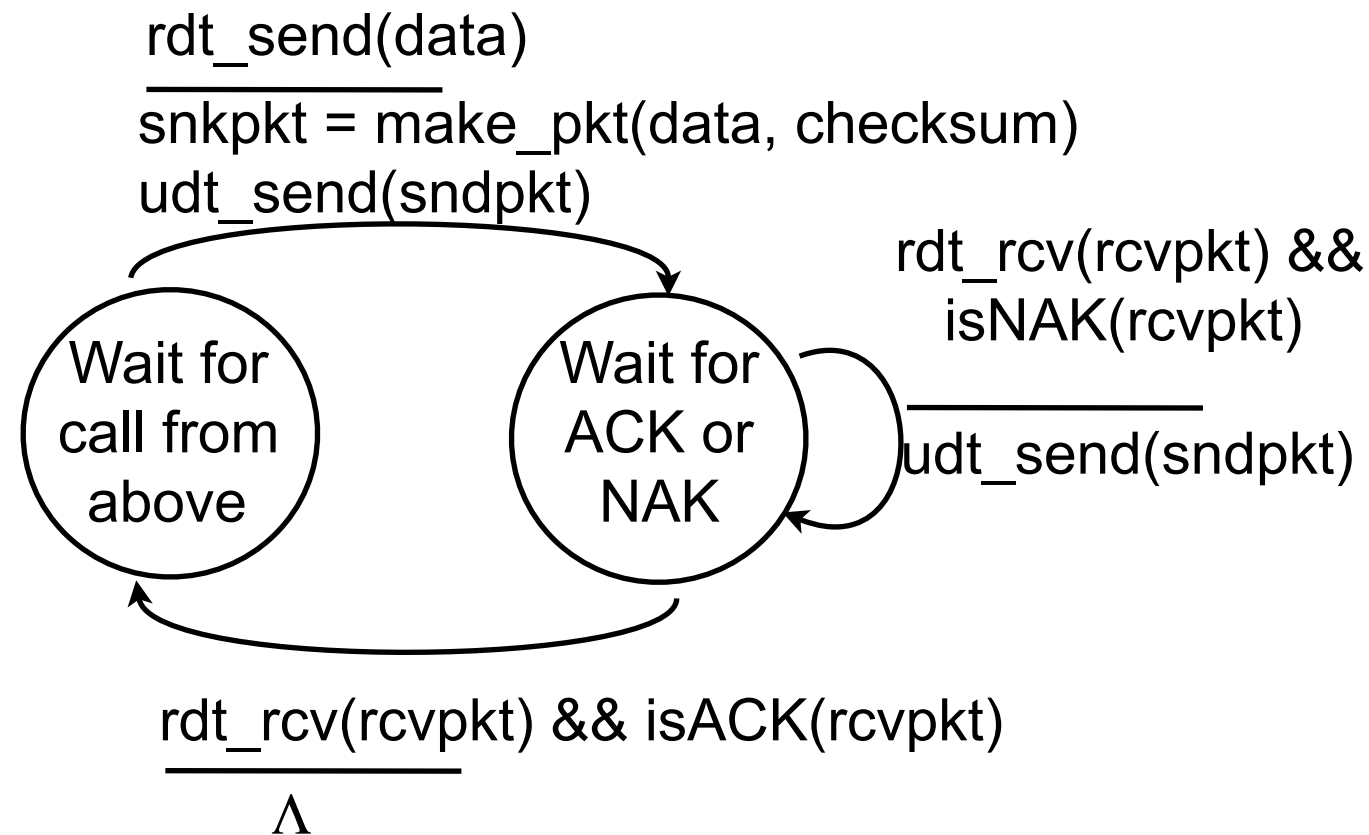


sender

receiver

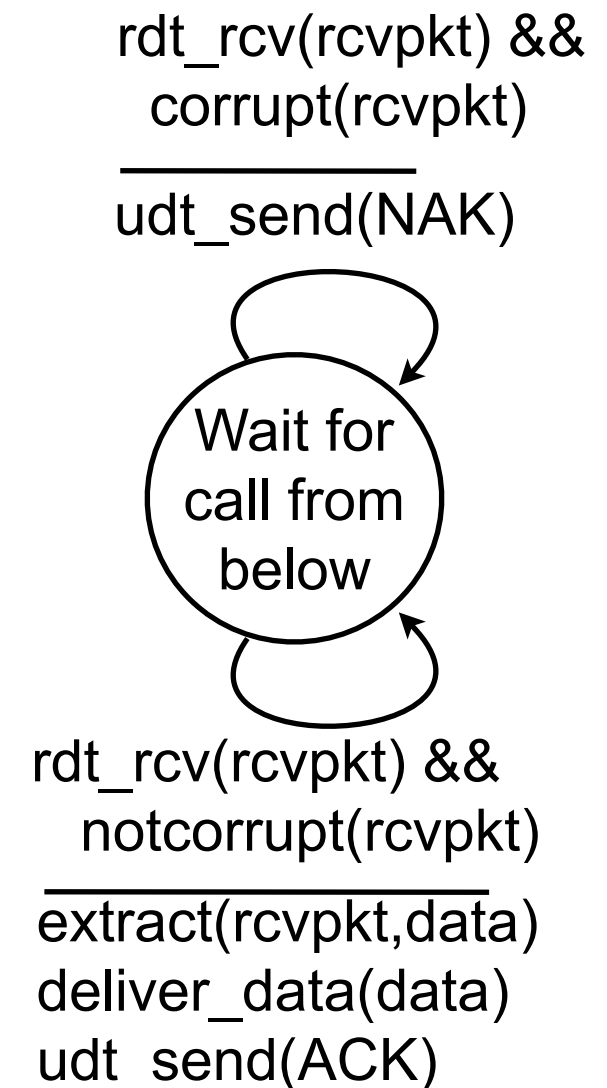


rdt2.0: operation with no errors


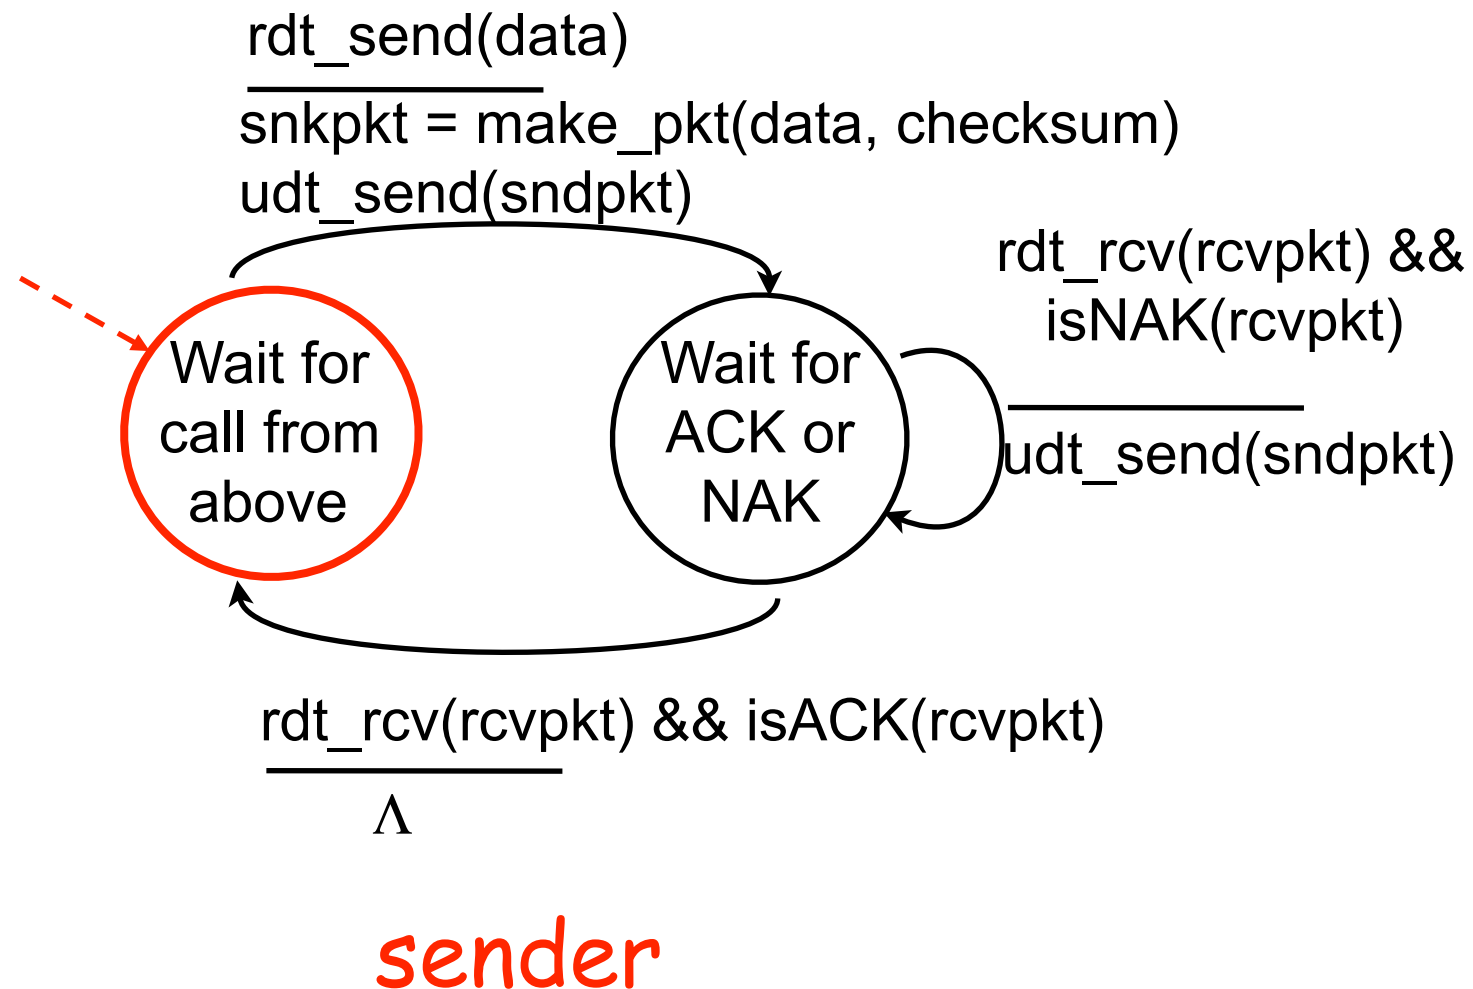


sender


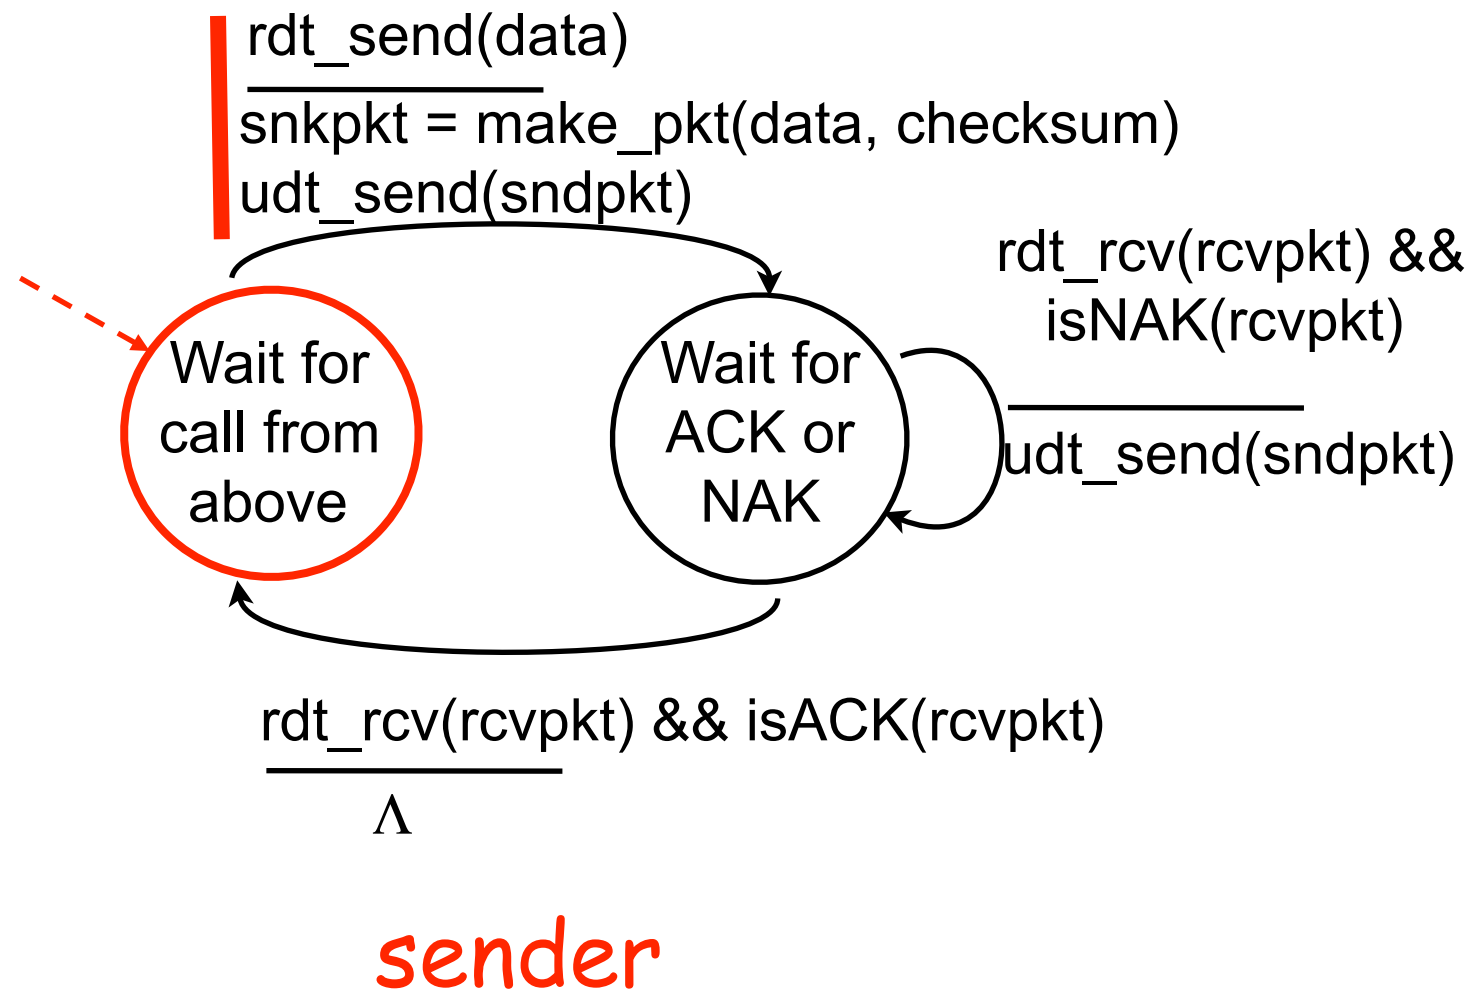
receiver



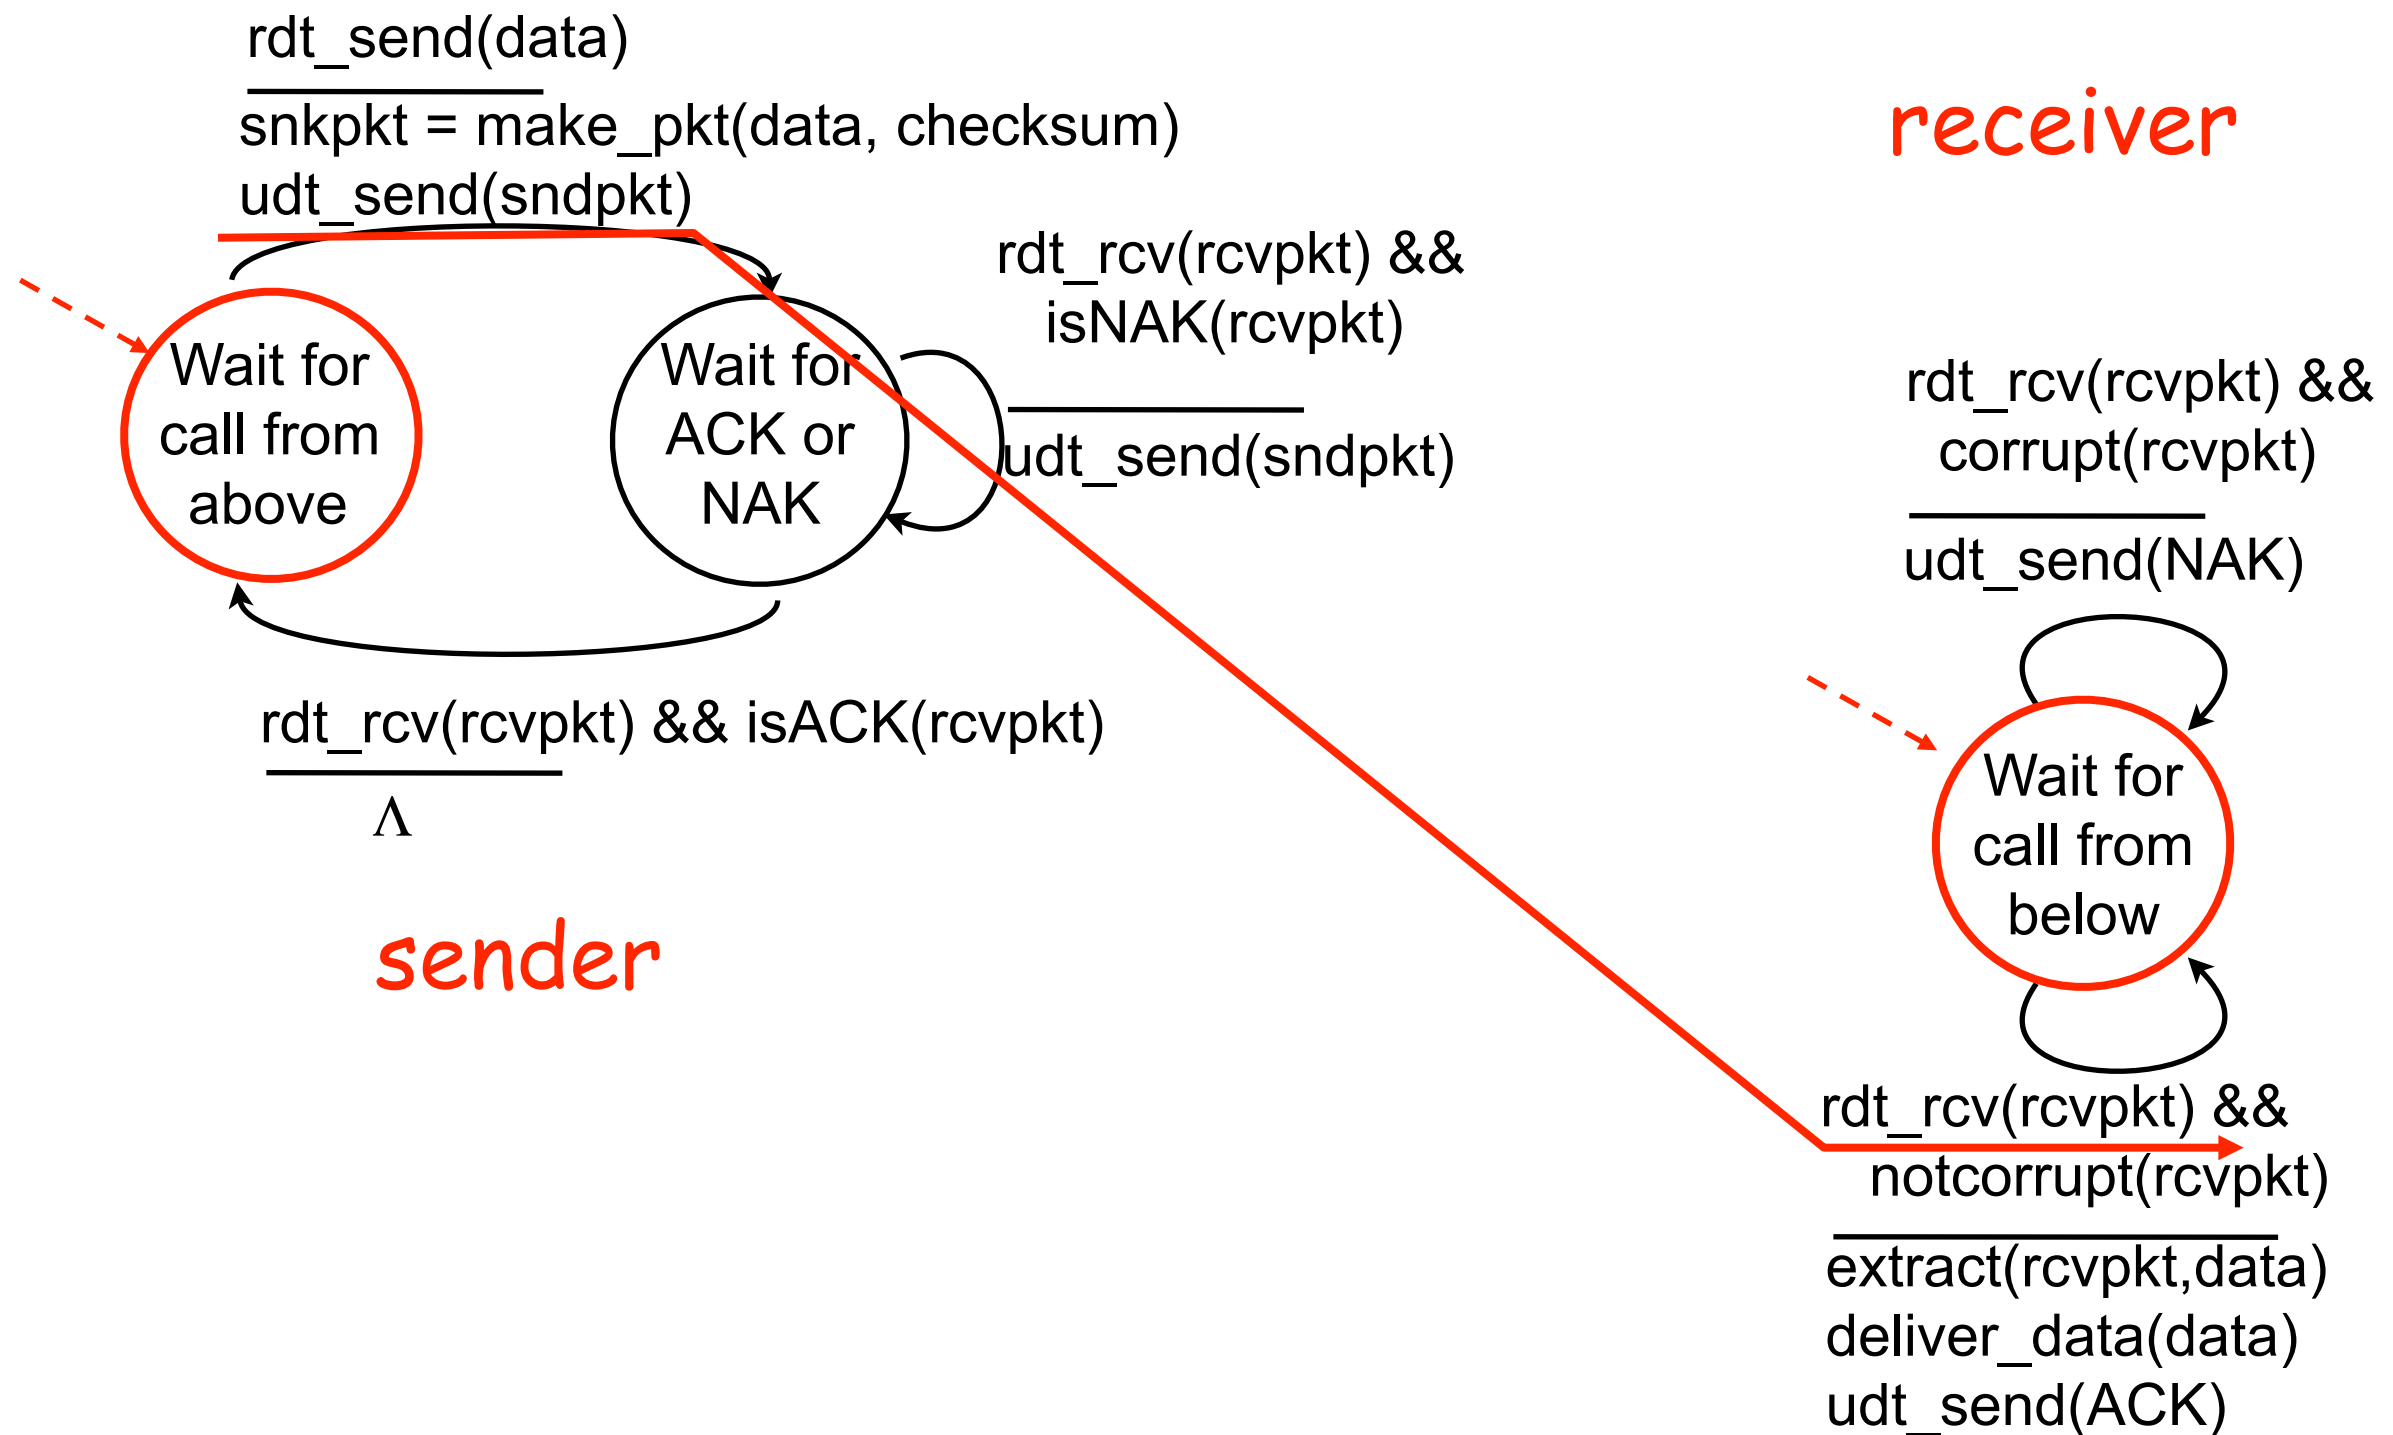
rdt2.0: operation with no errors




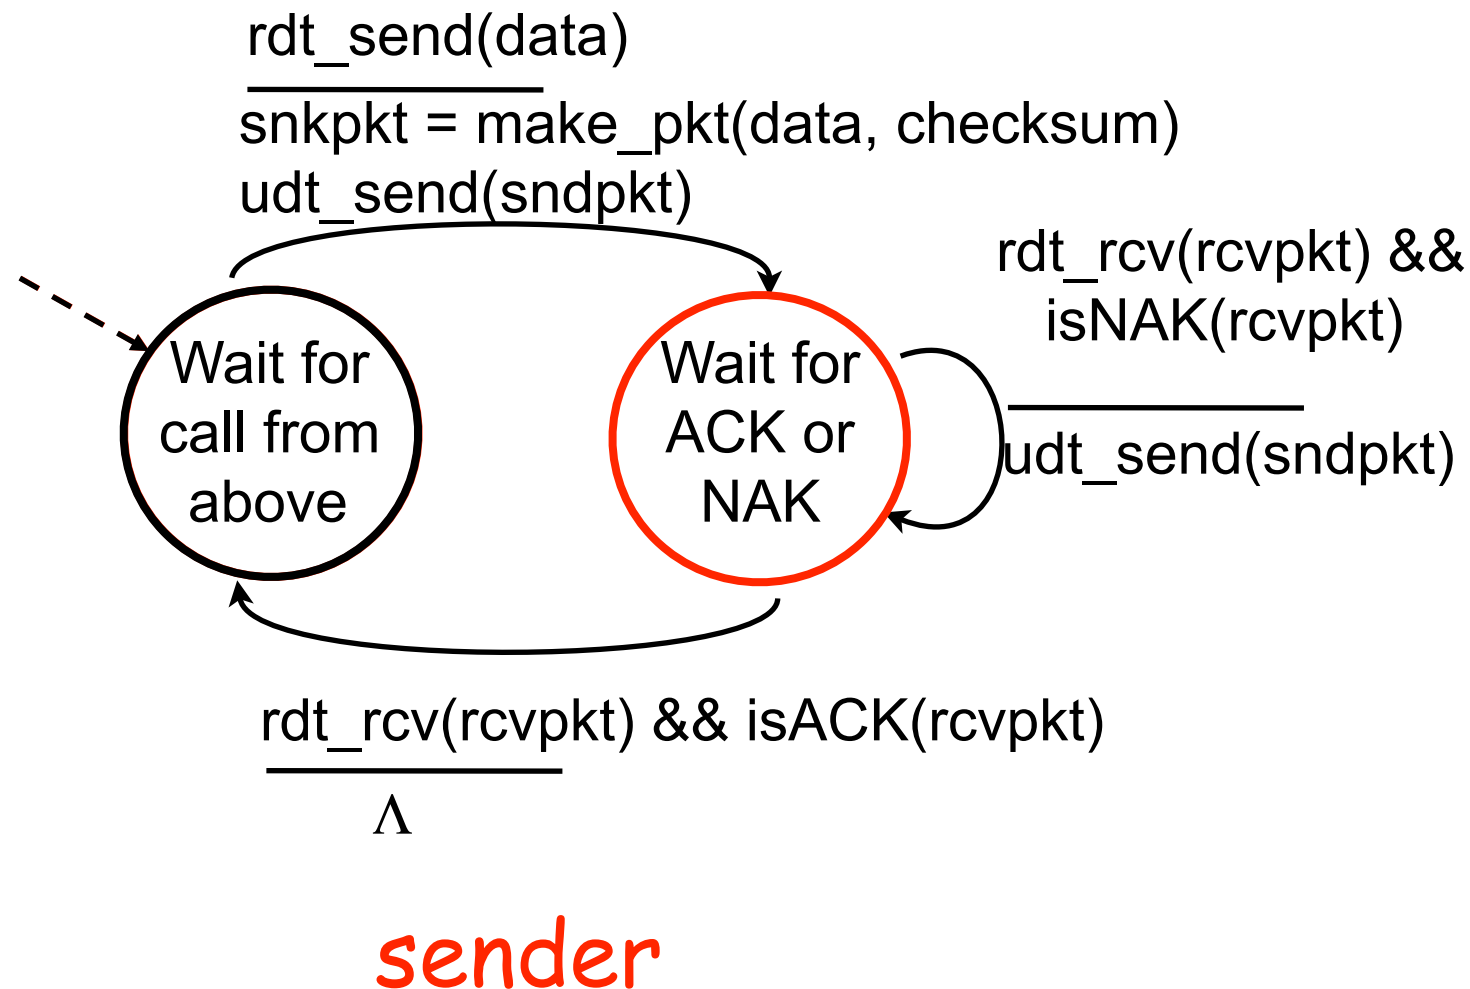
rdt2.0: operation with no errors



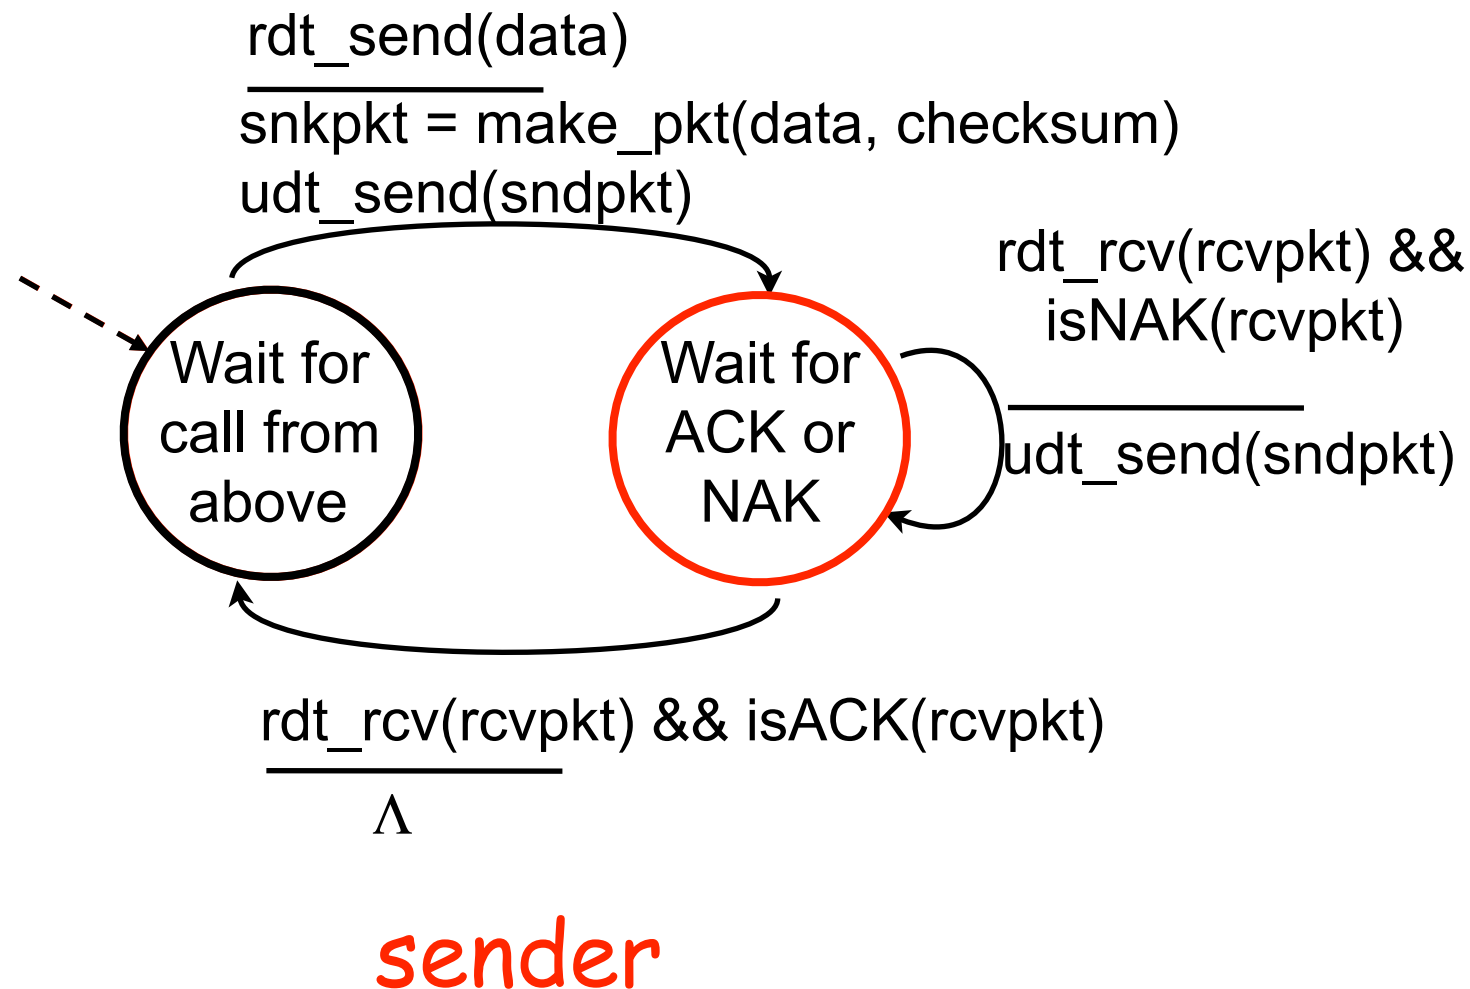
rdt2.0: operation with no errors



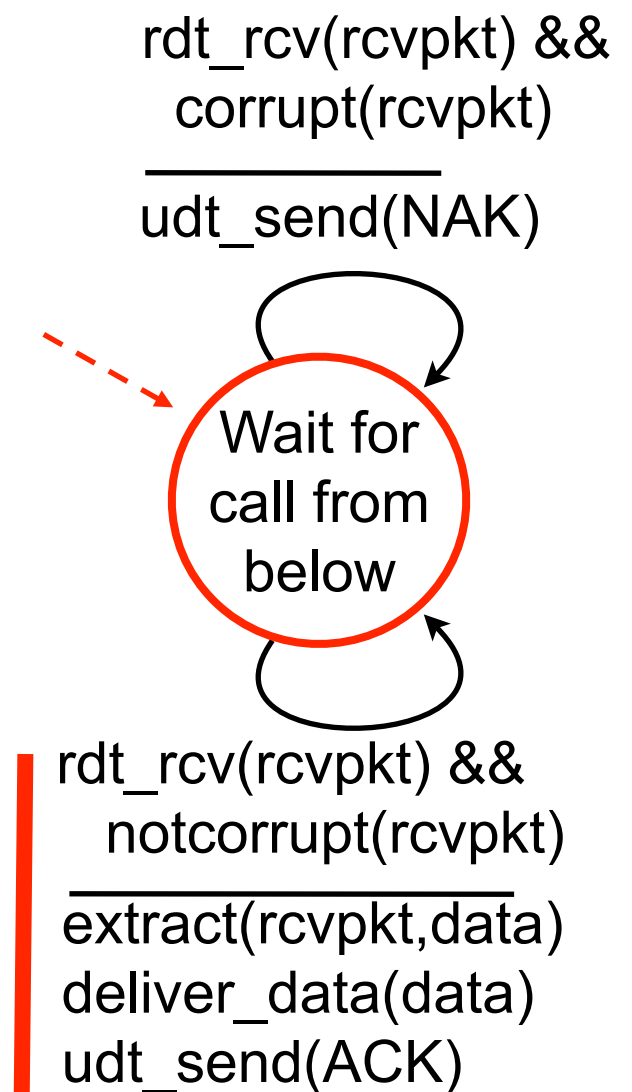
rdt2.0: operation with no errors



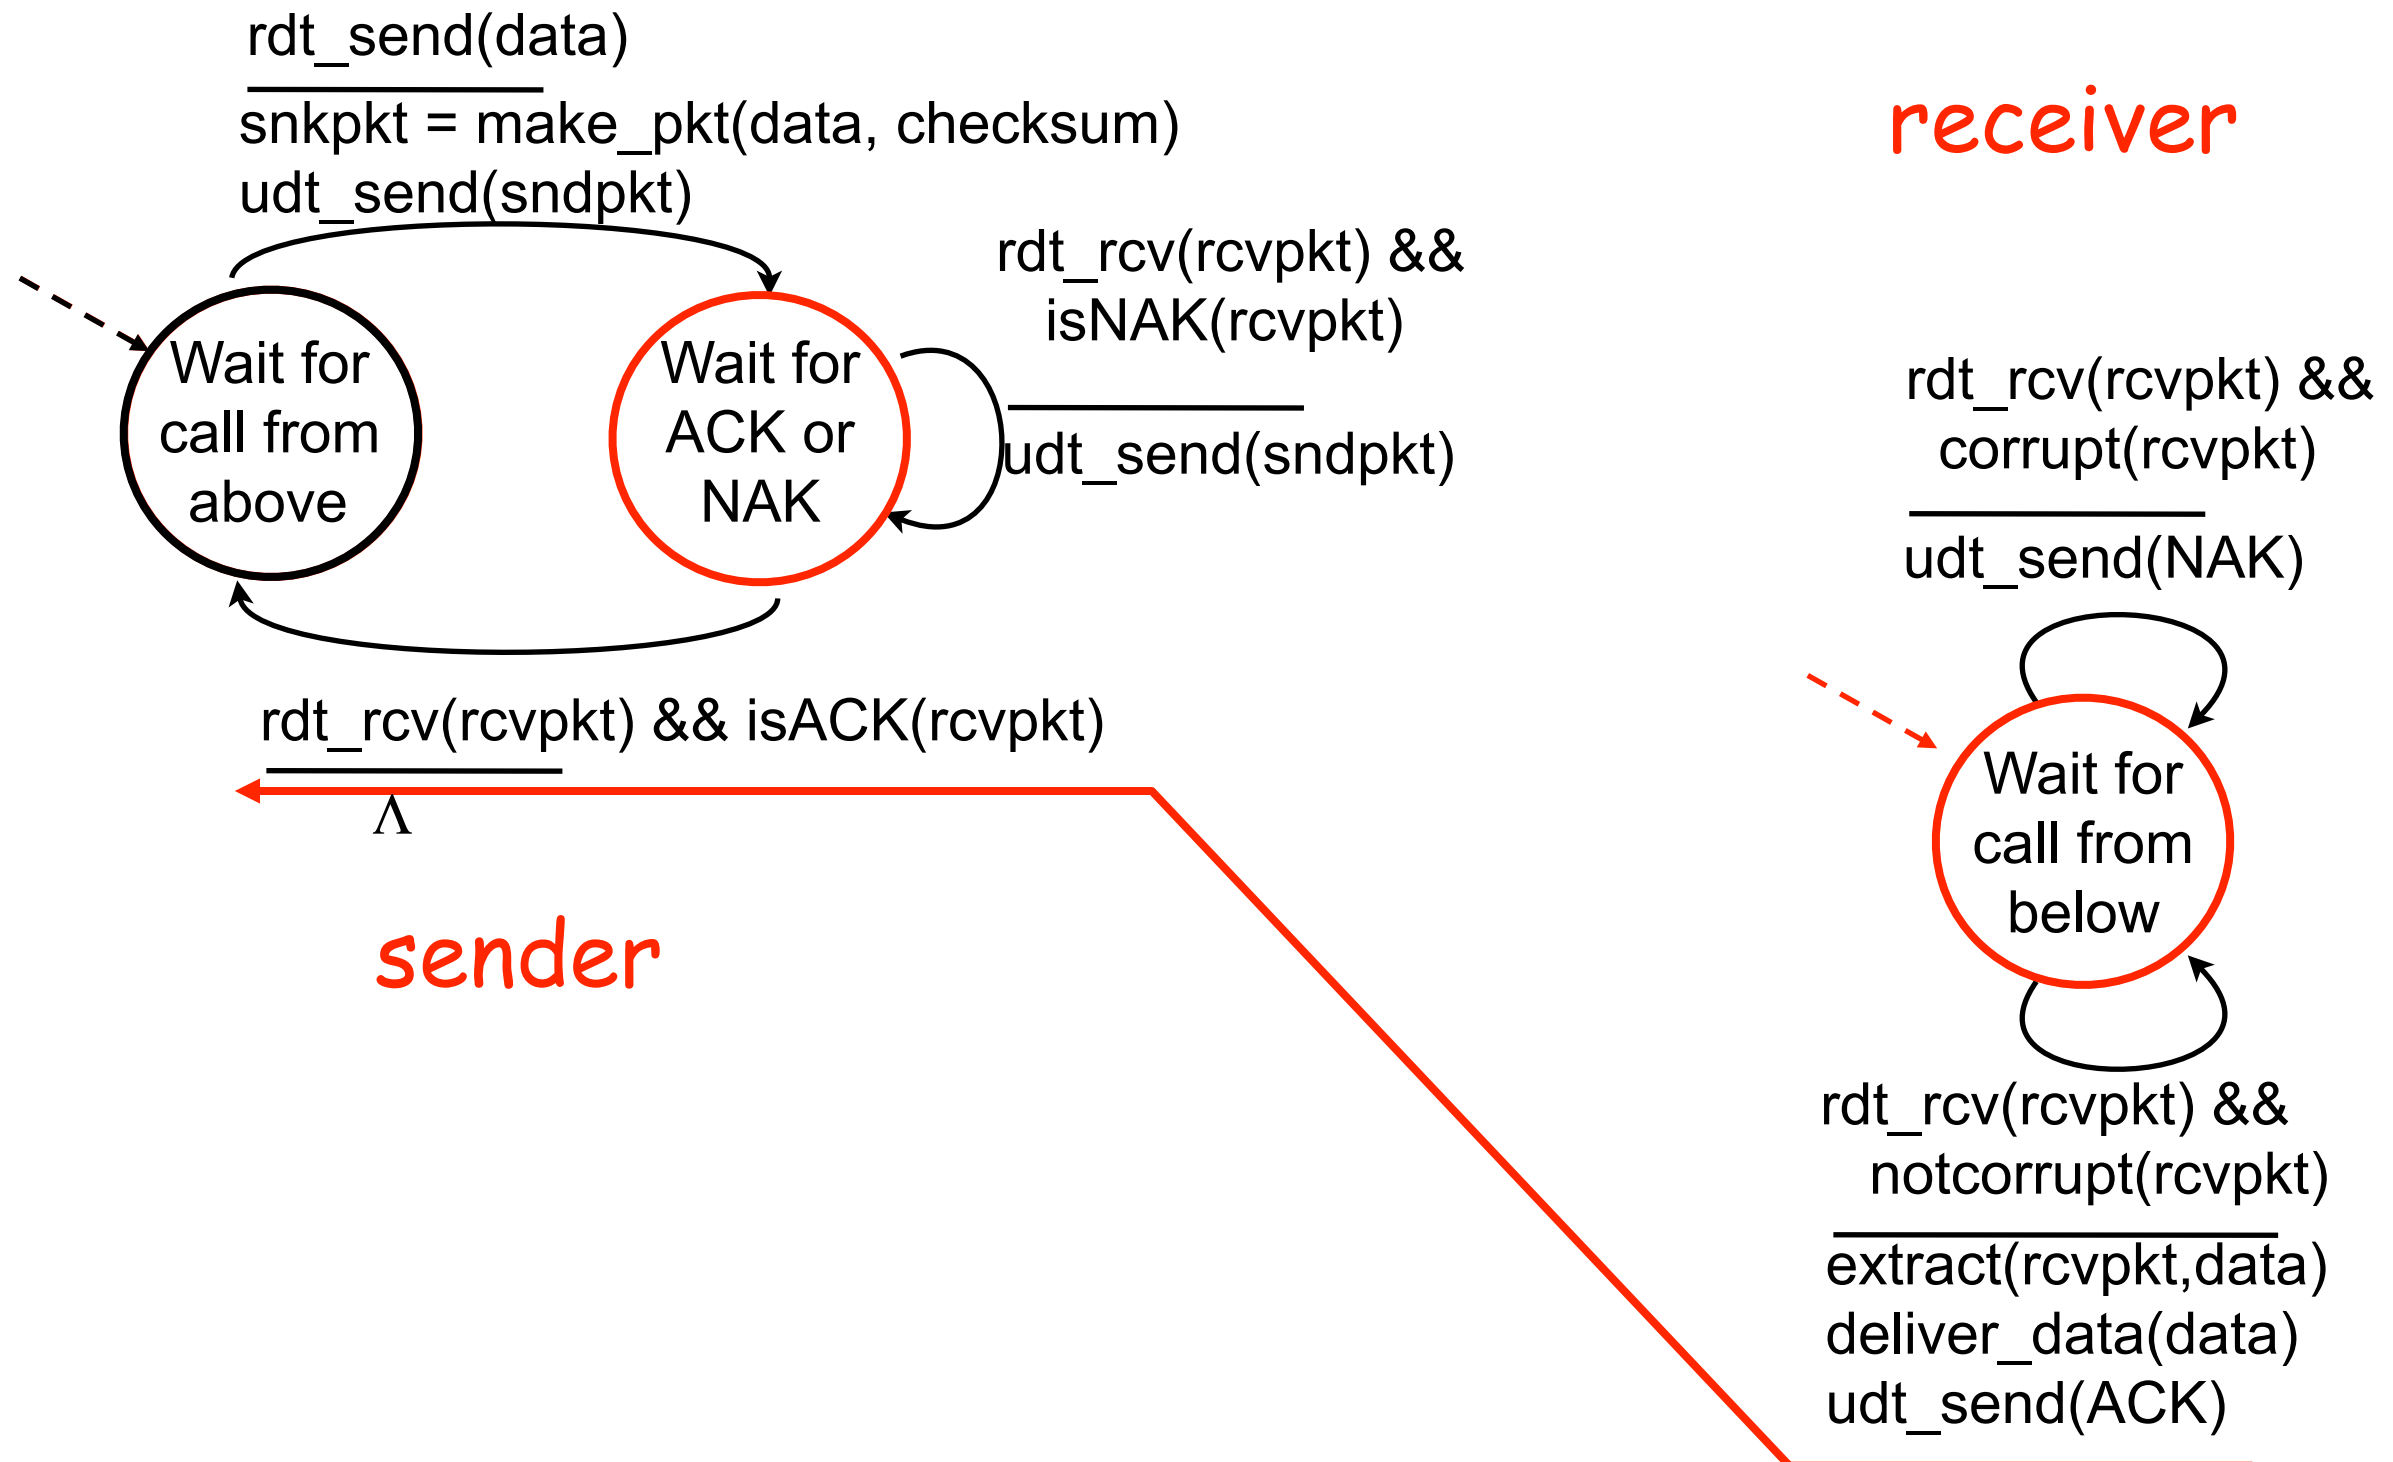
rdt2.0: operation with no errors



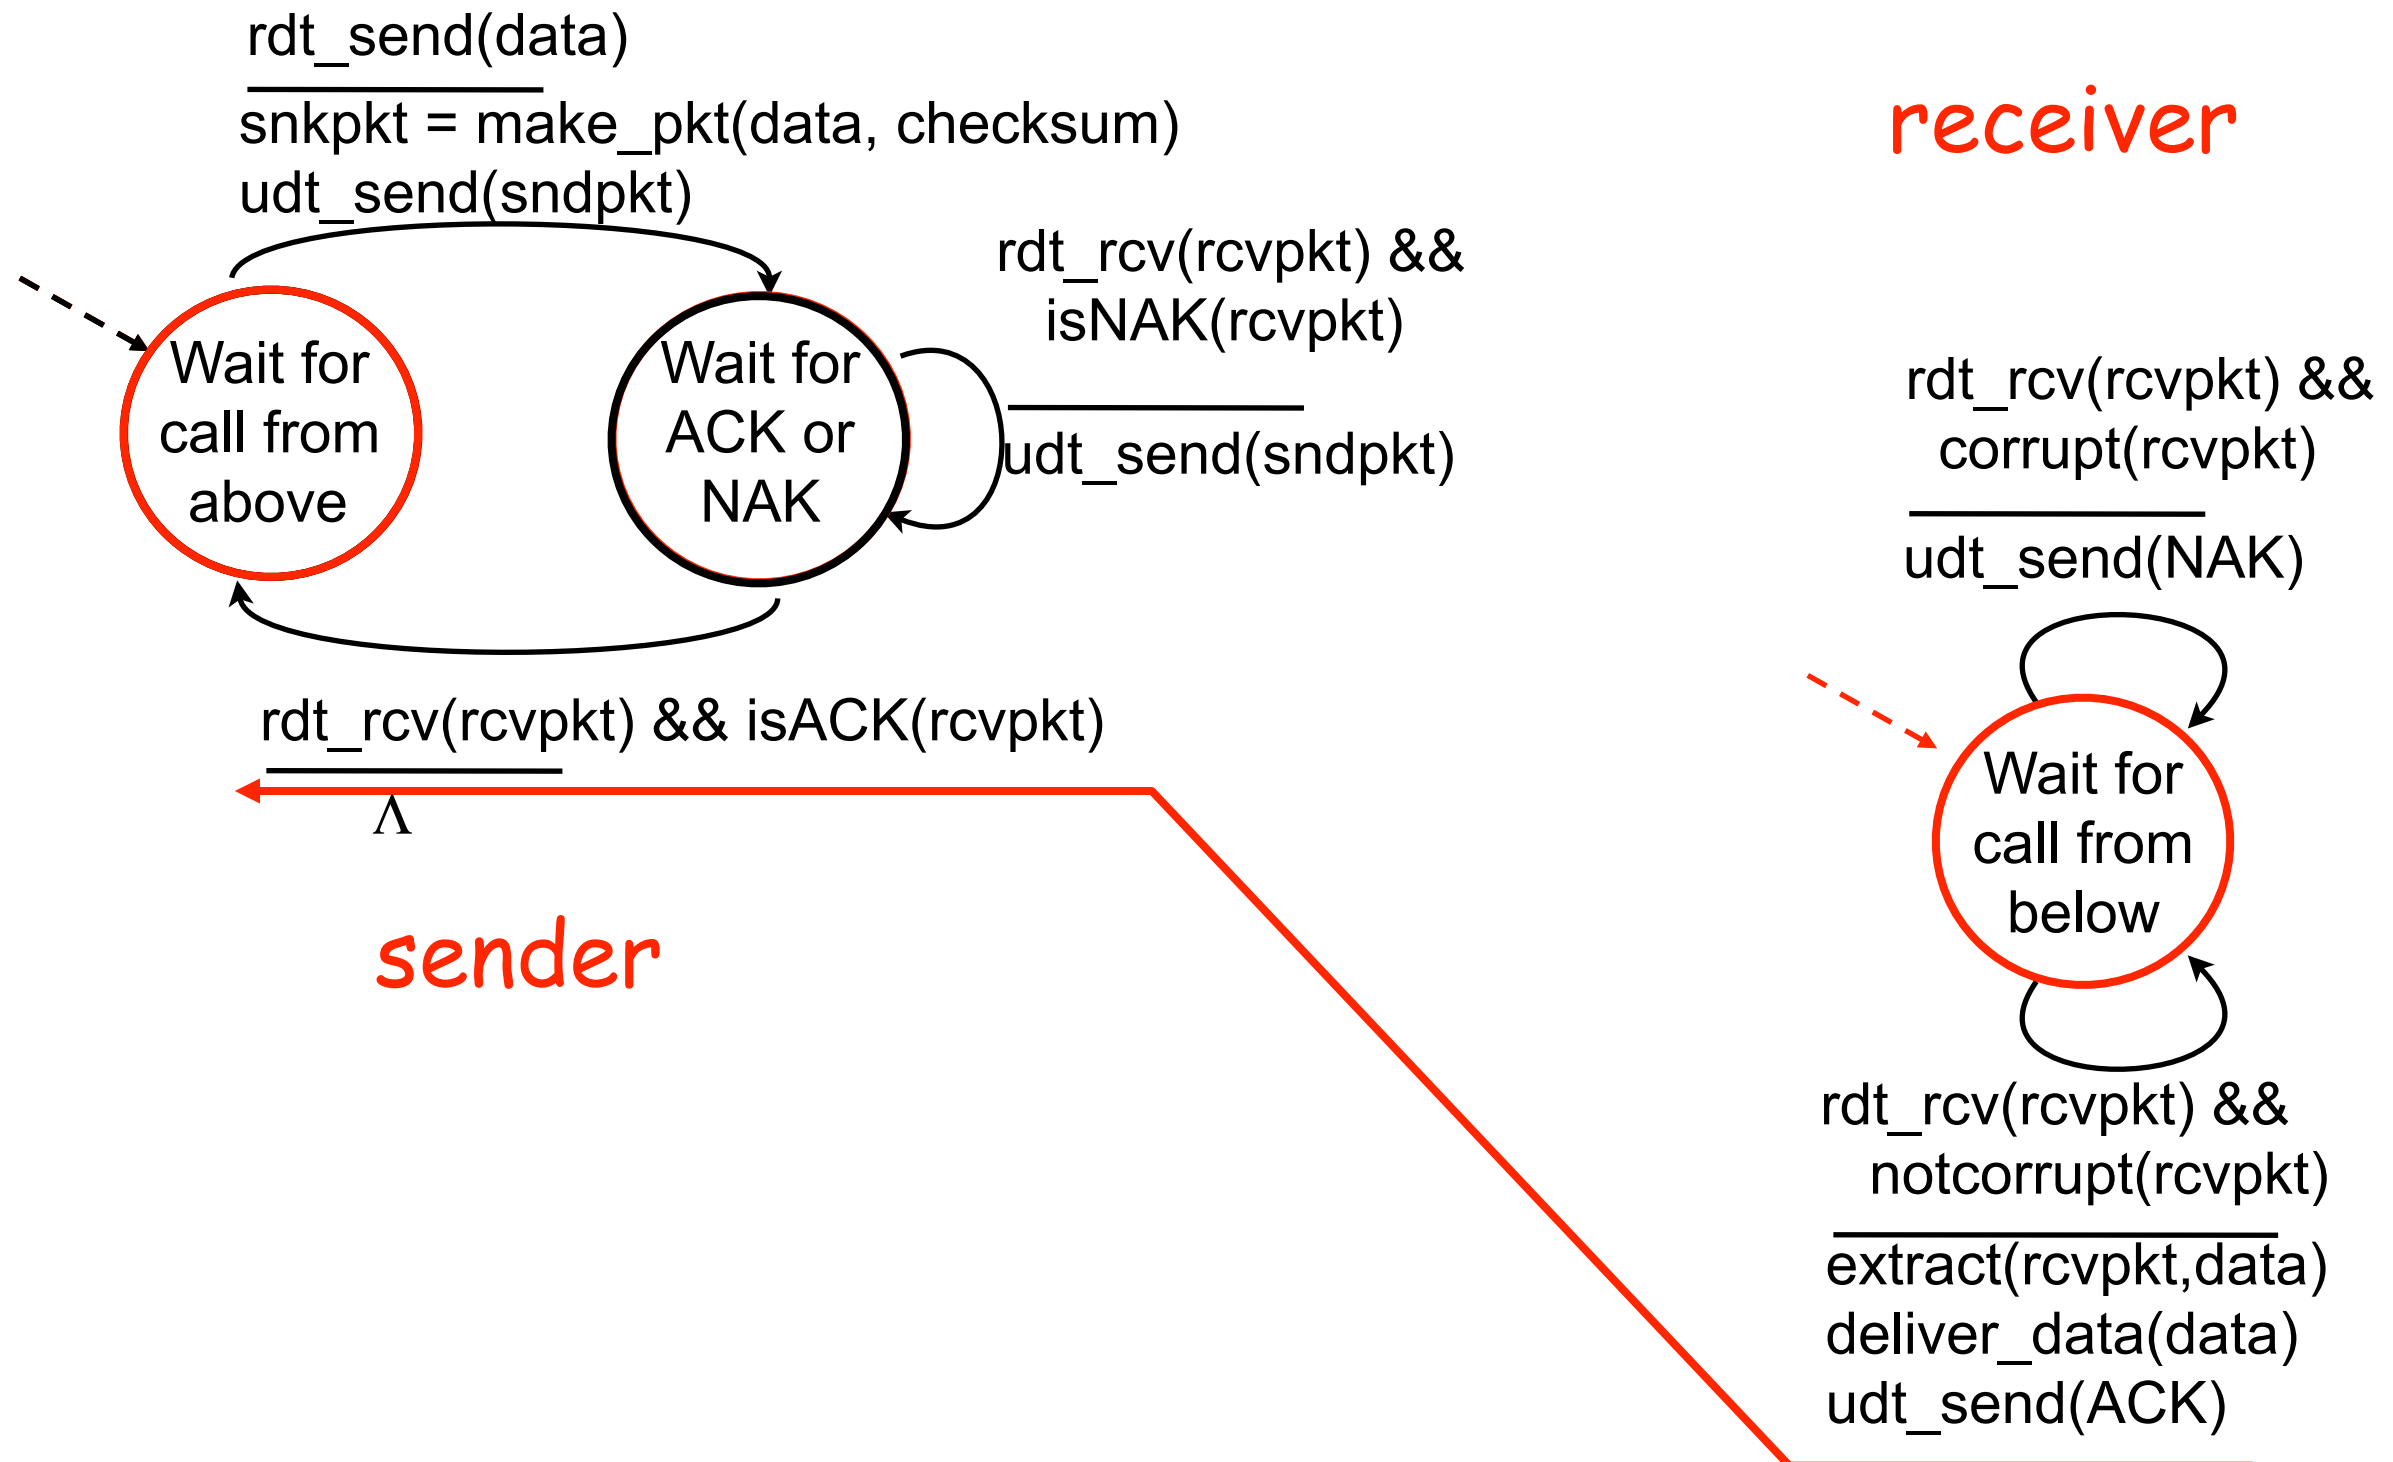
receiver



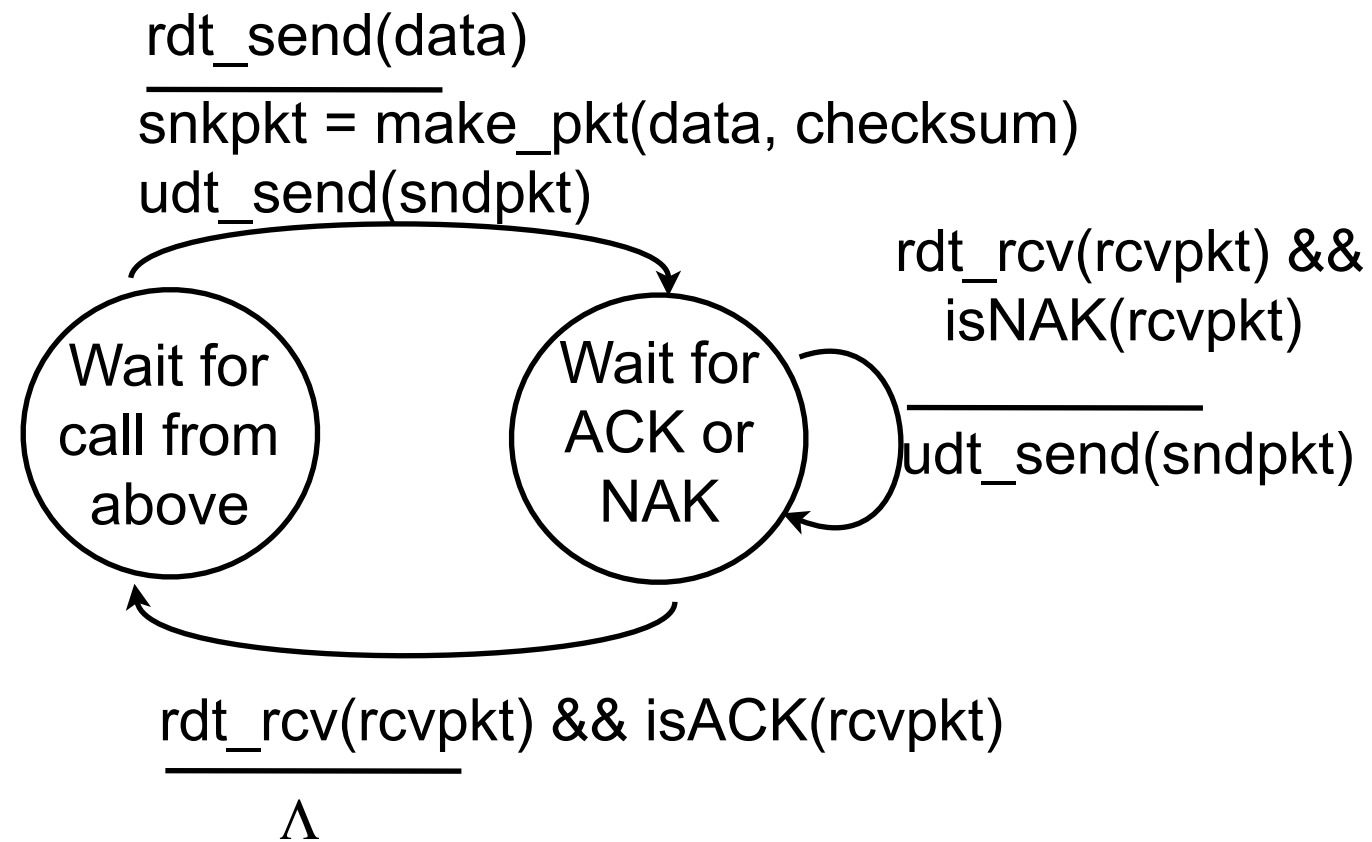
rdt2.0: operation with no errors



rdt2.0: operation with no errors

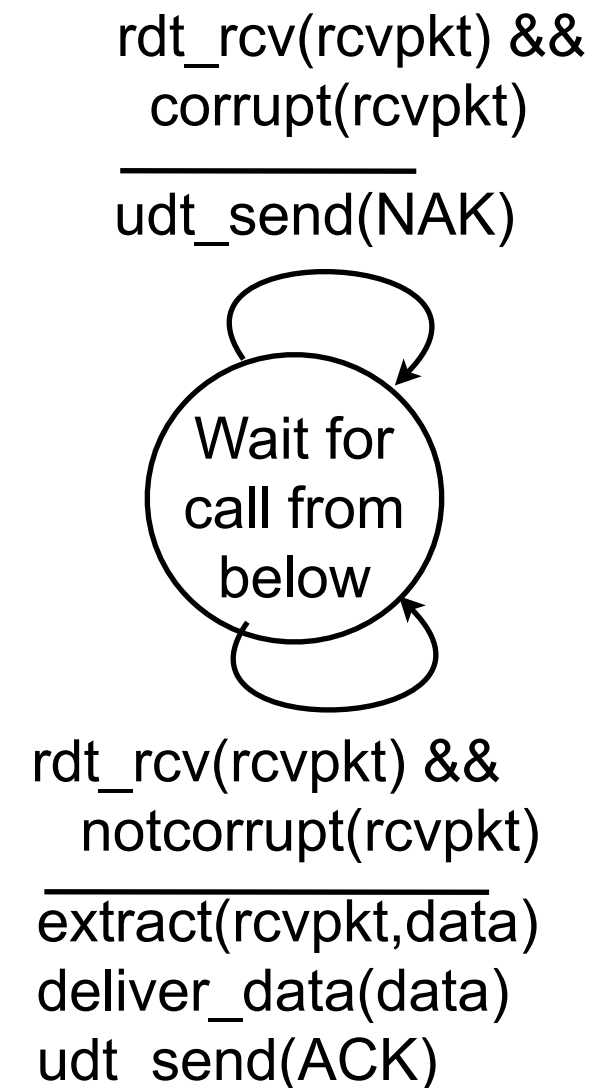


rdt2.0: error scenario

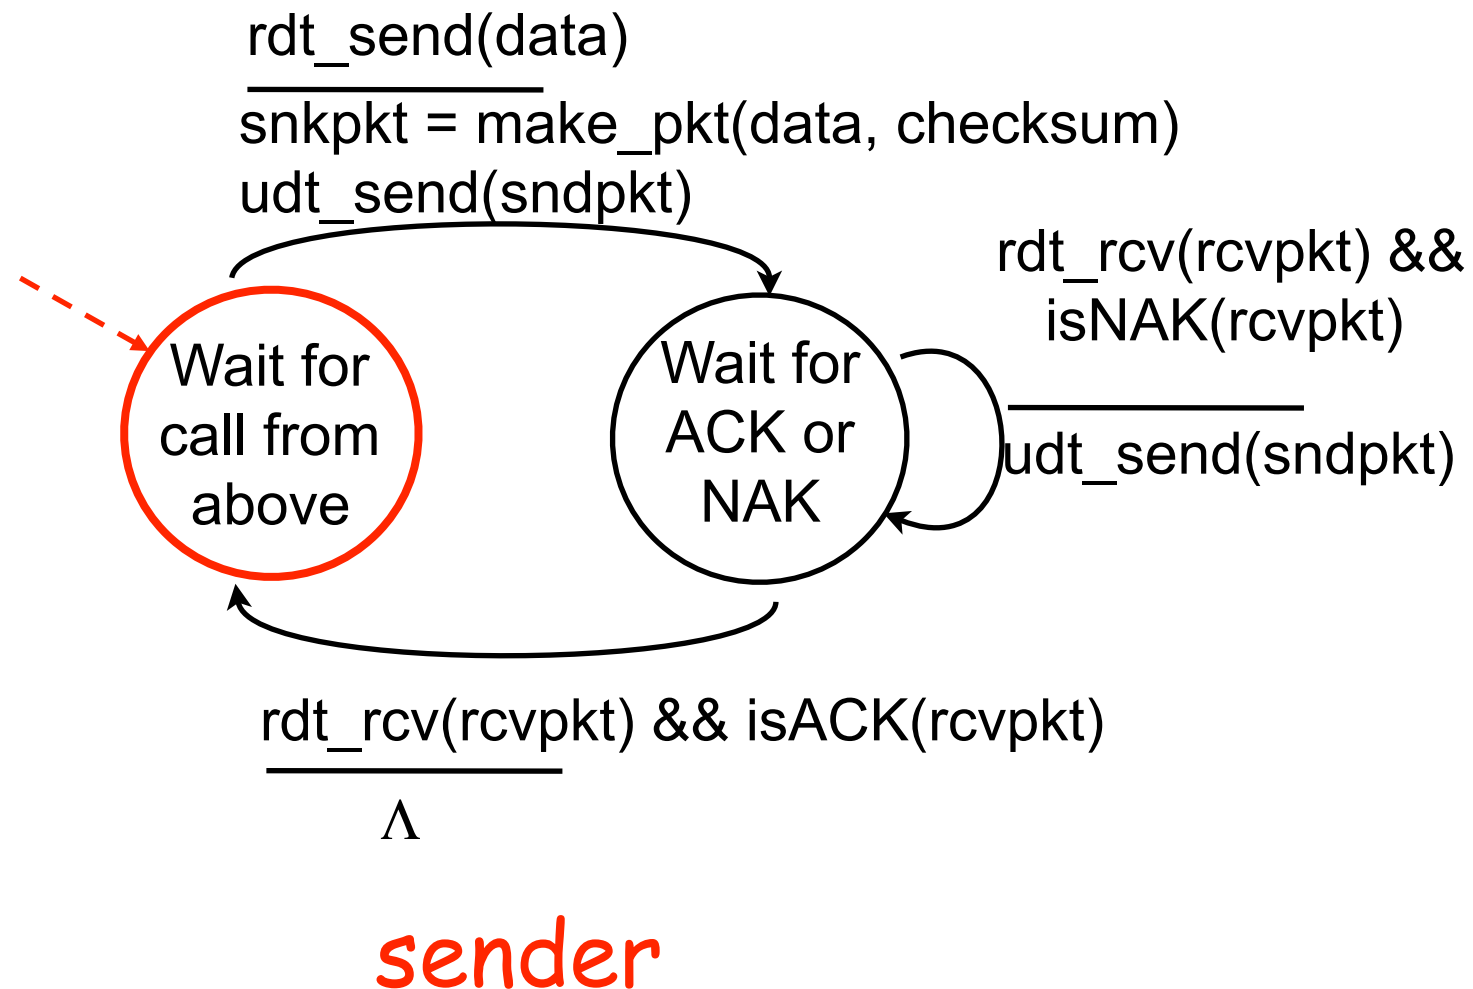


sender

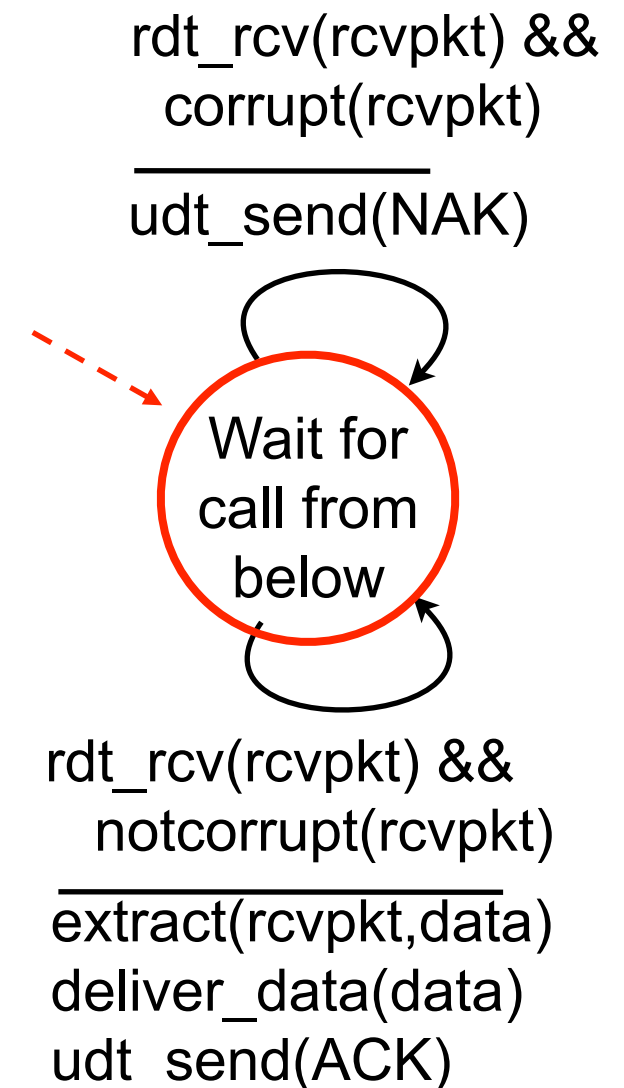
receiver



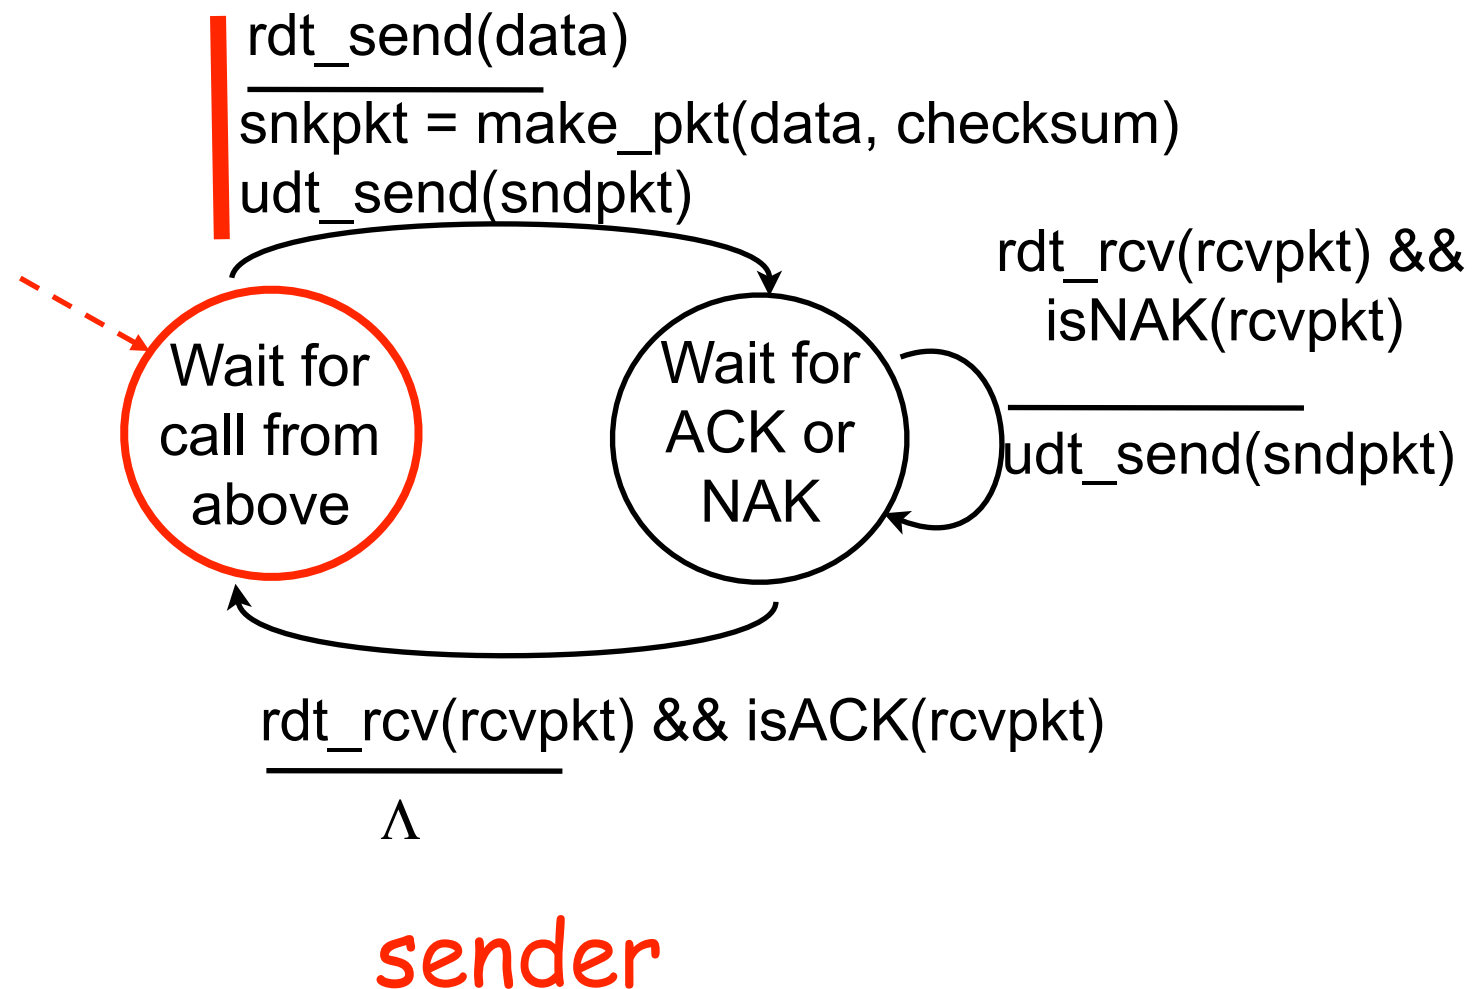
rdt2.0: error scenario



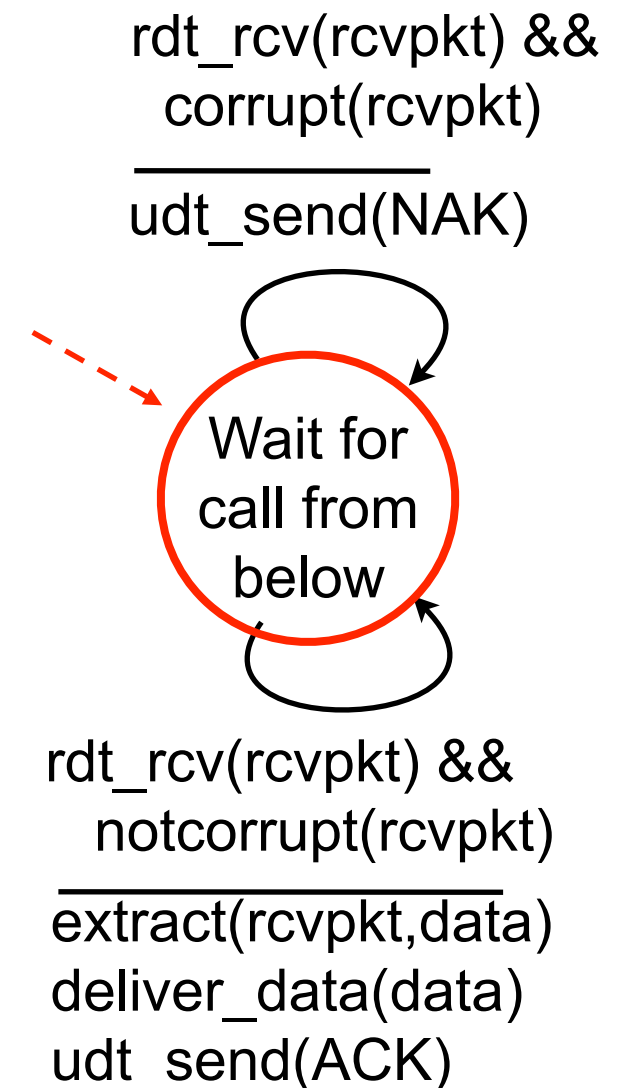
receiver



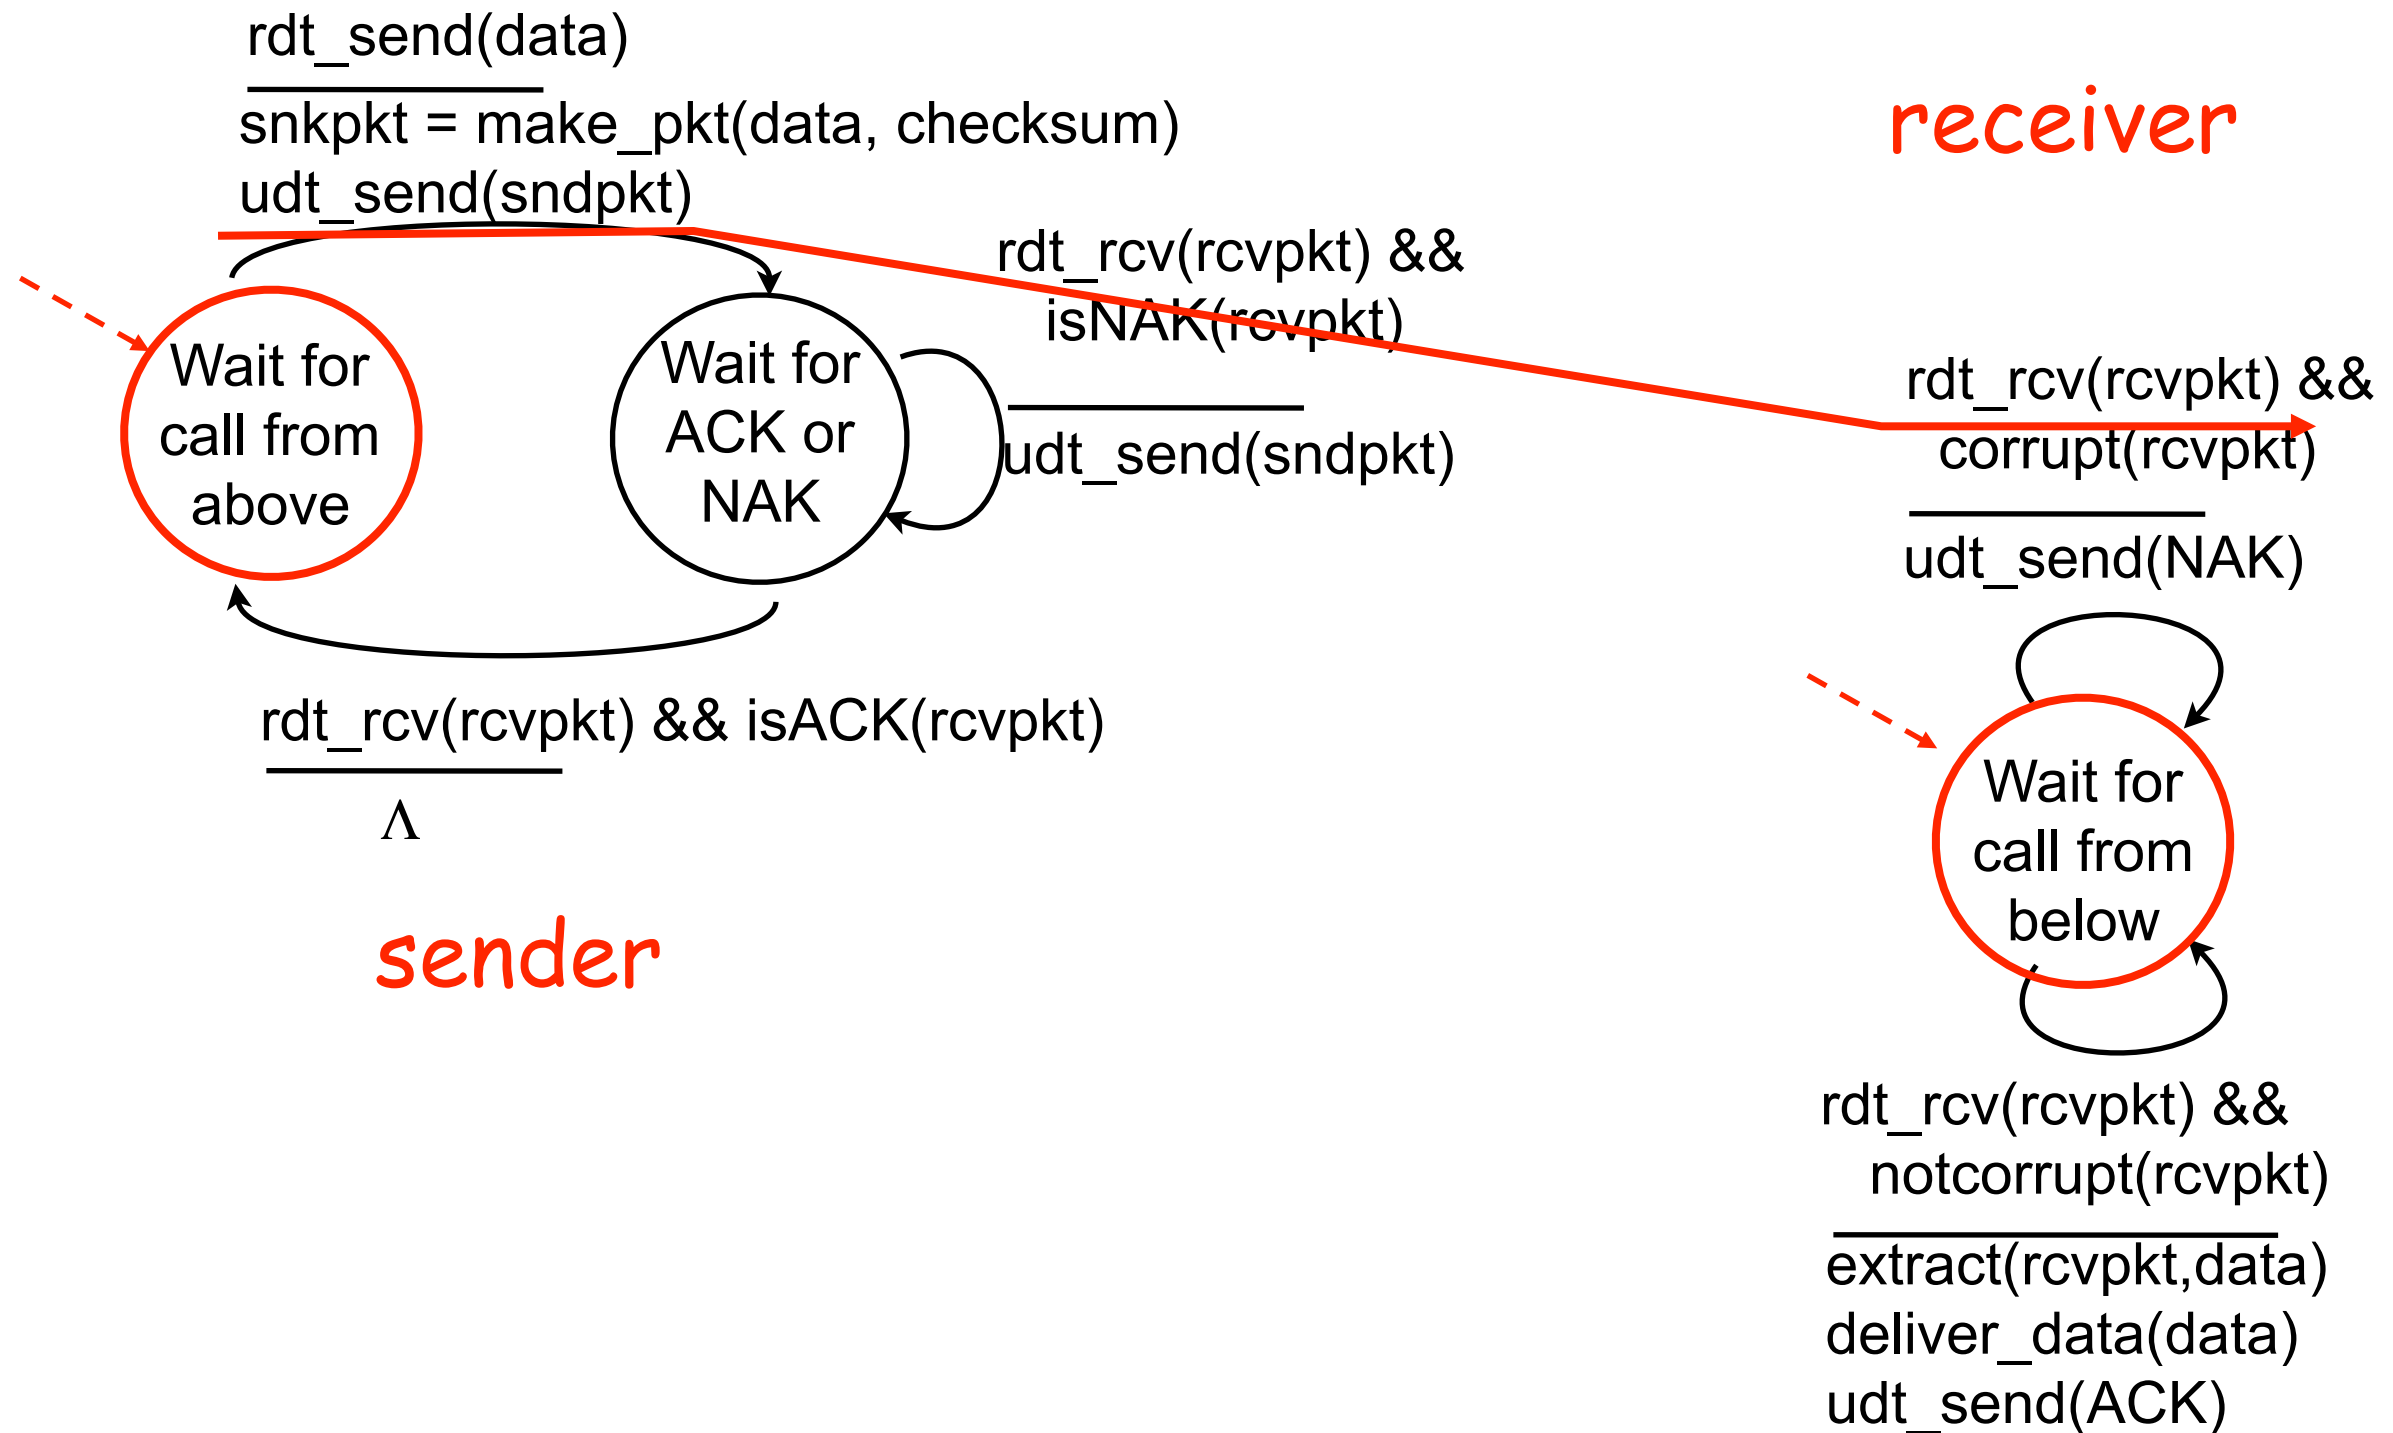
rdt2.0: error scenario



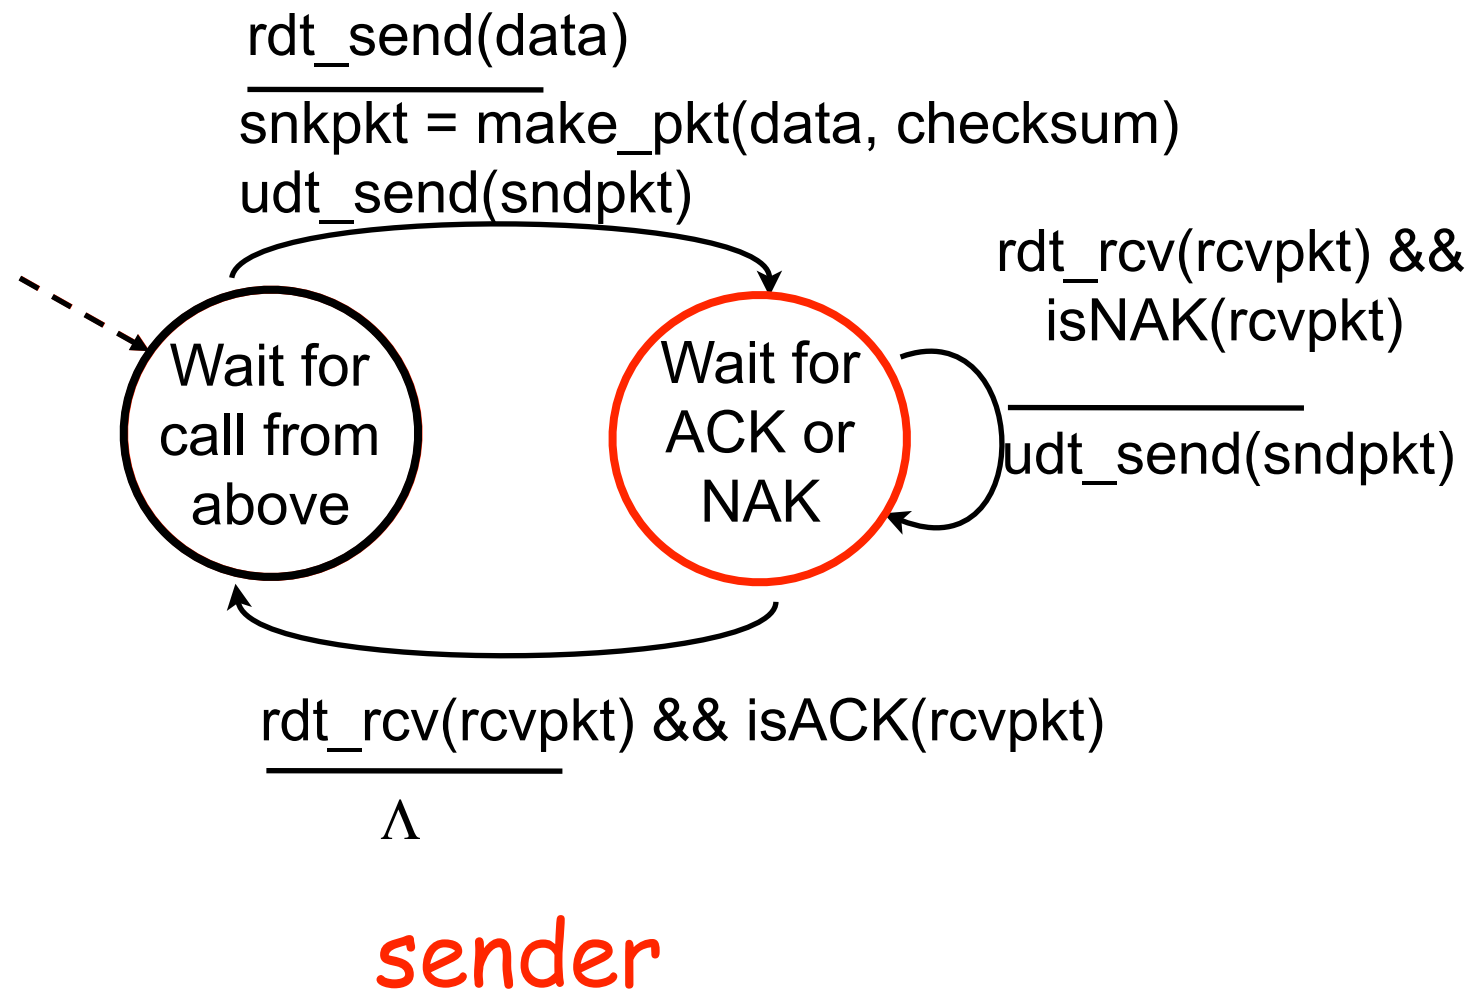
receiver



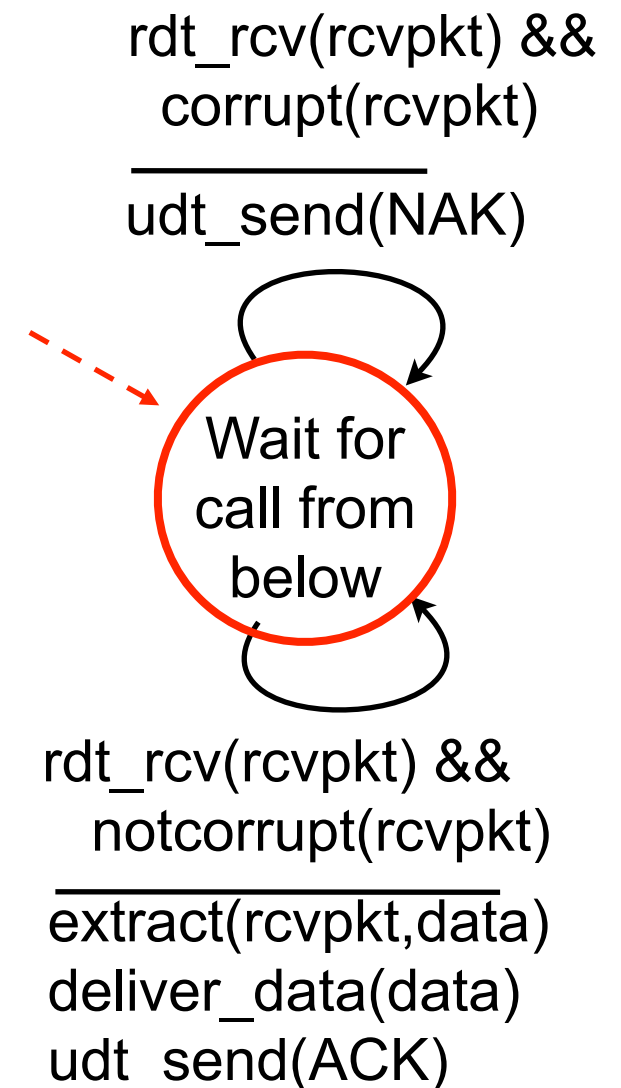
rdt2.0: error scenario



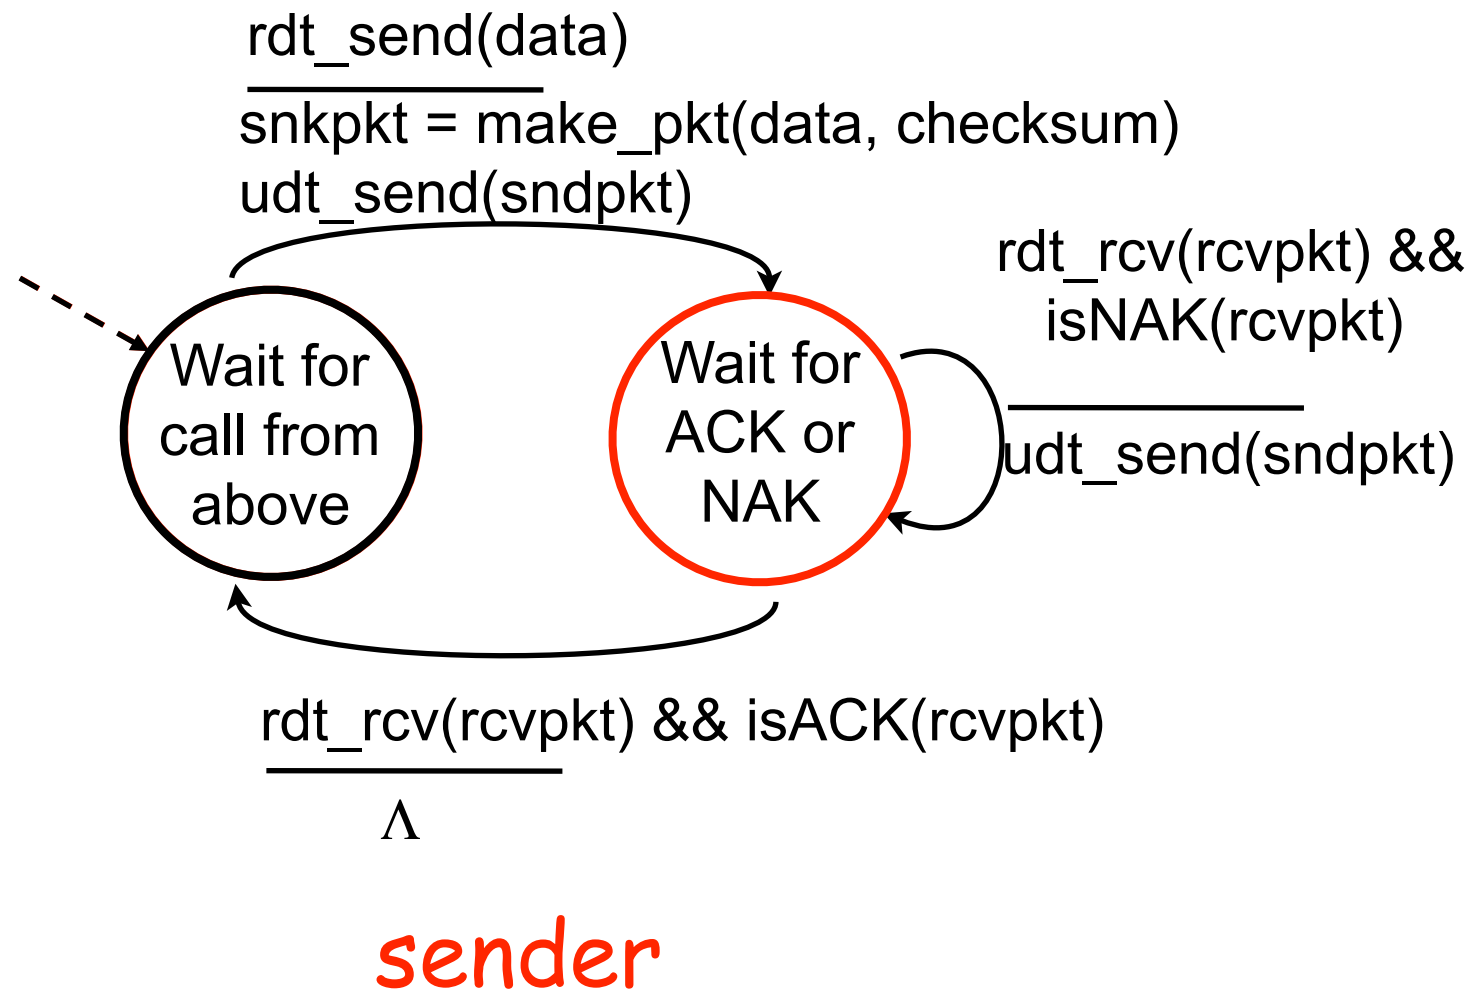
rdt2.0: error scenario



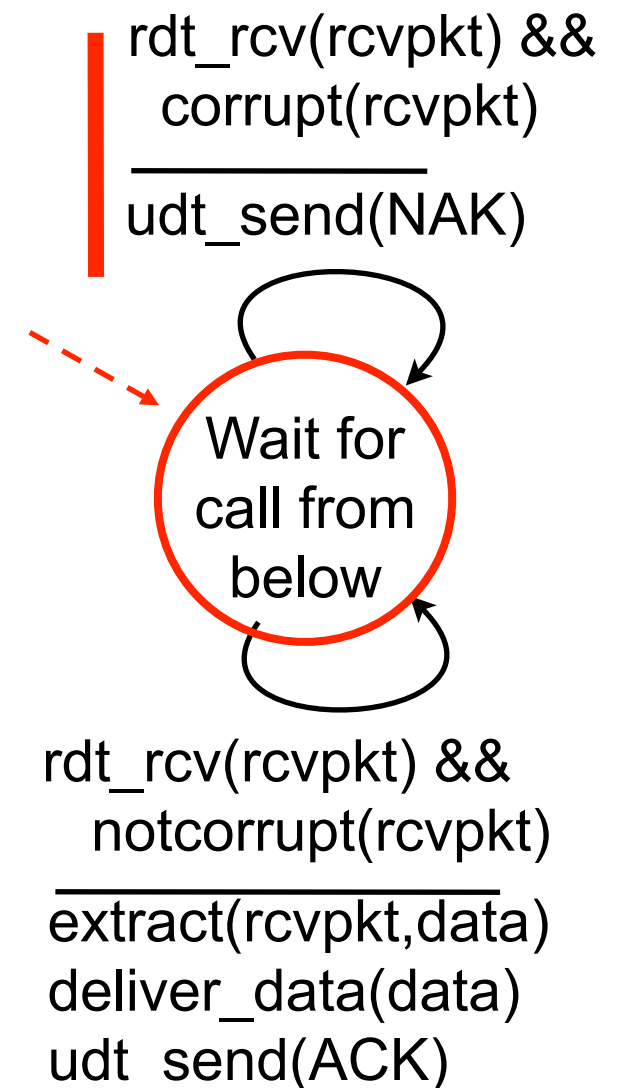
receiver



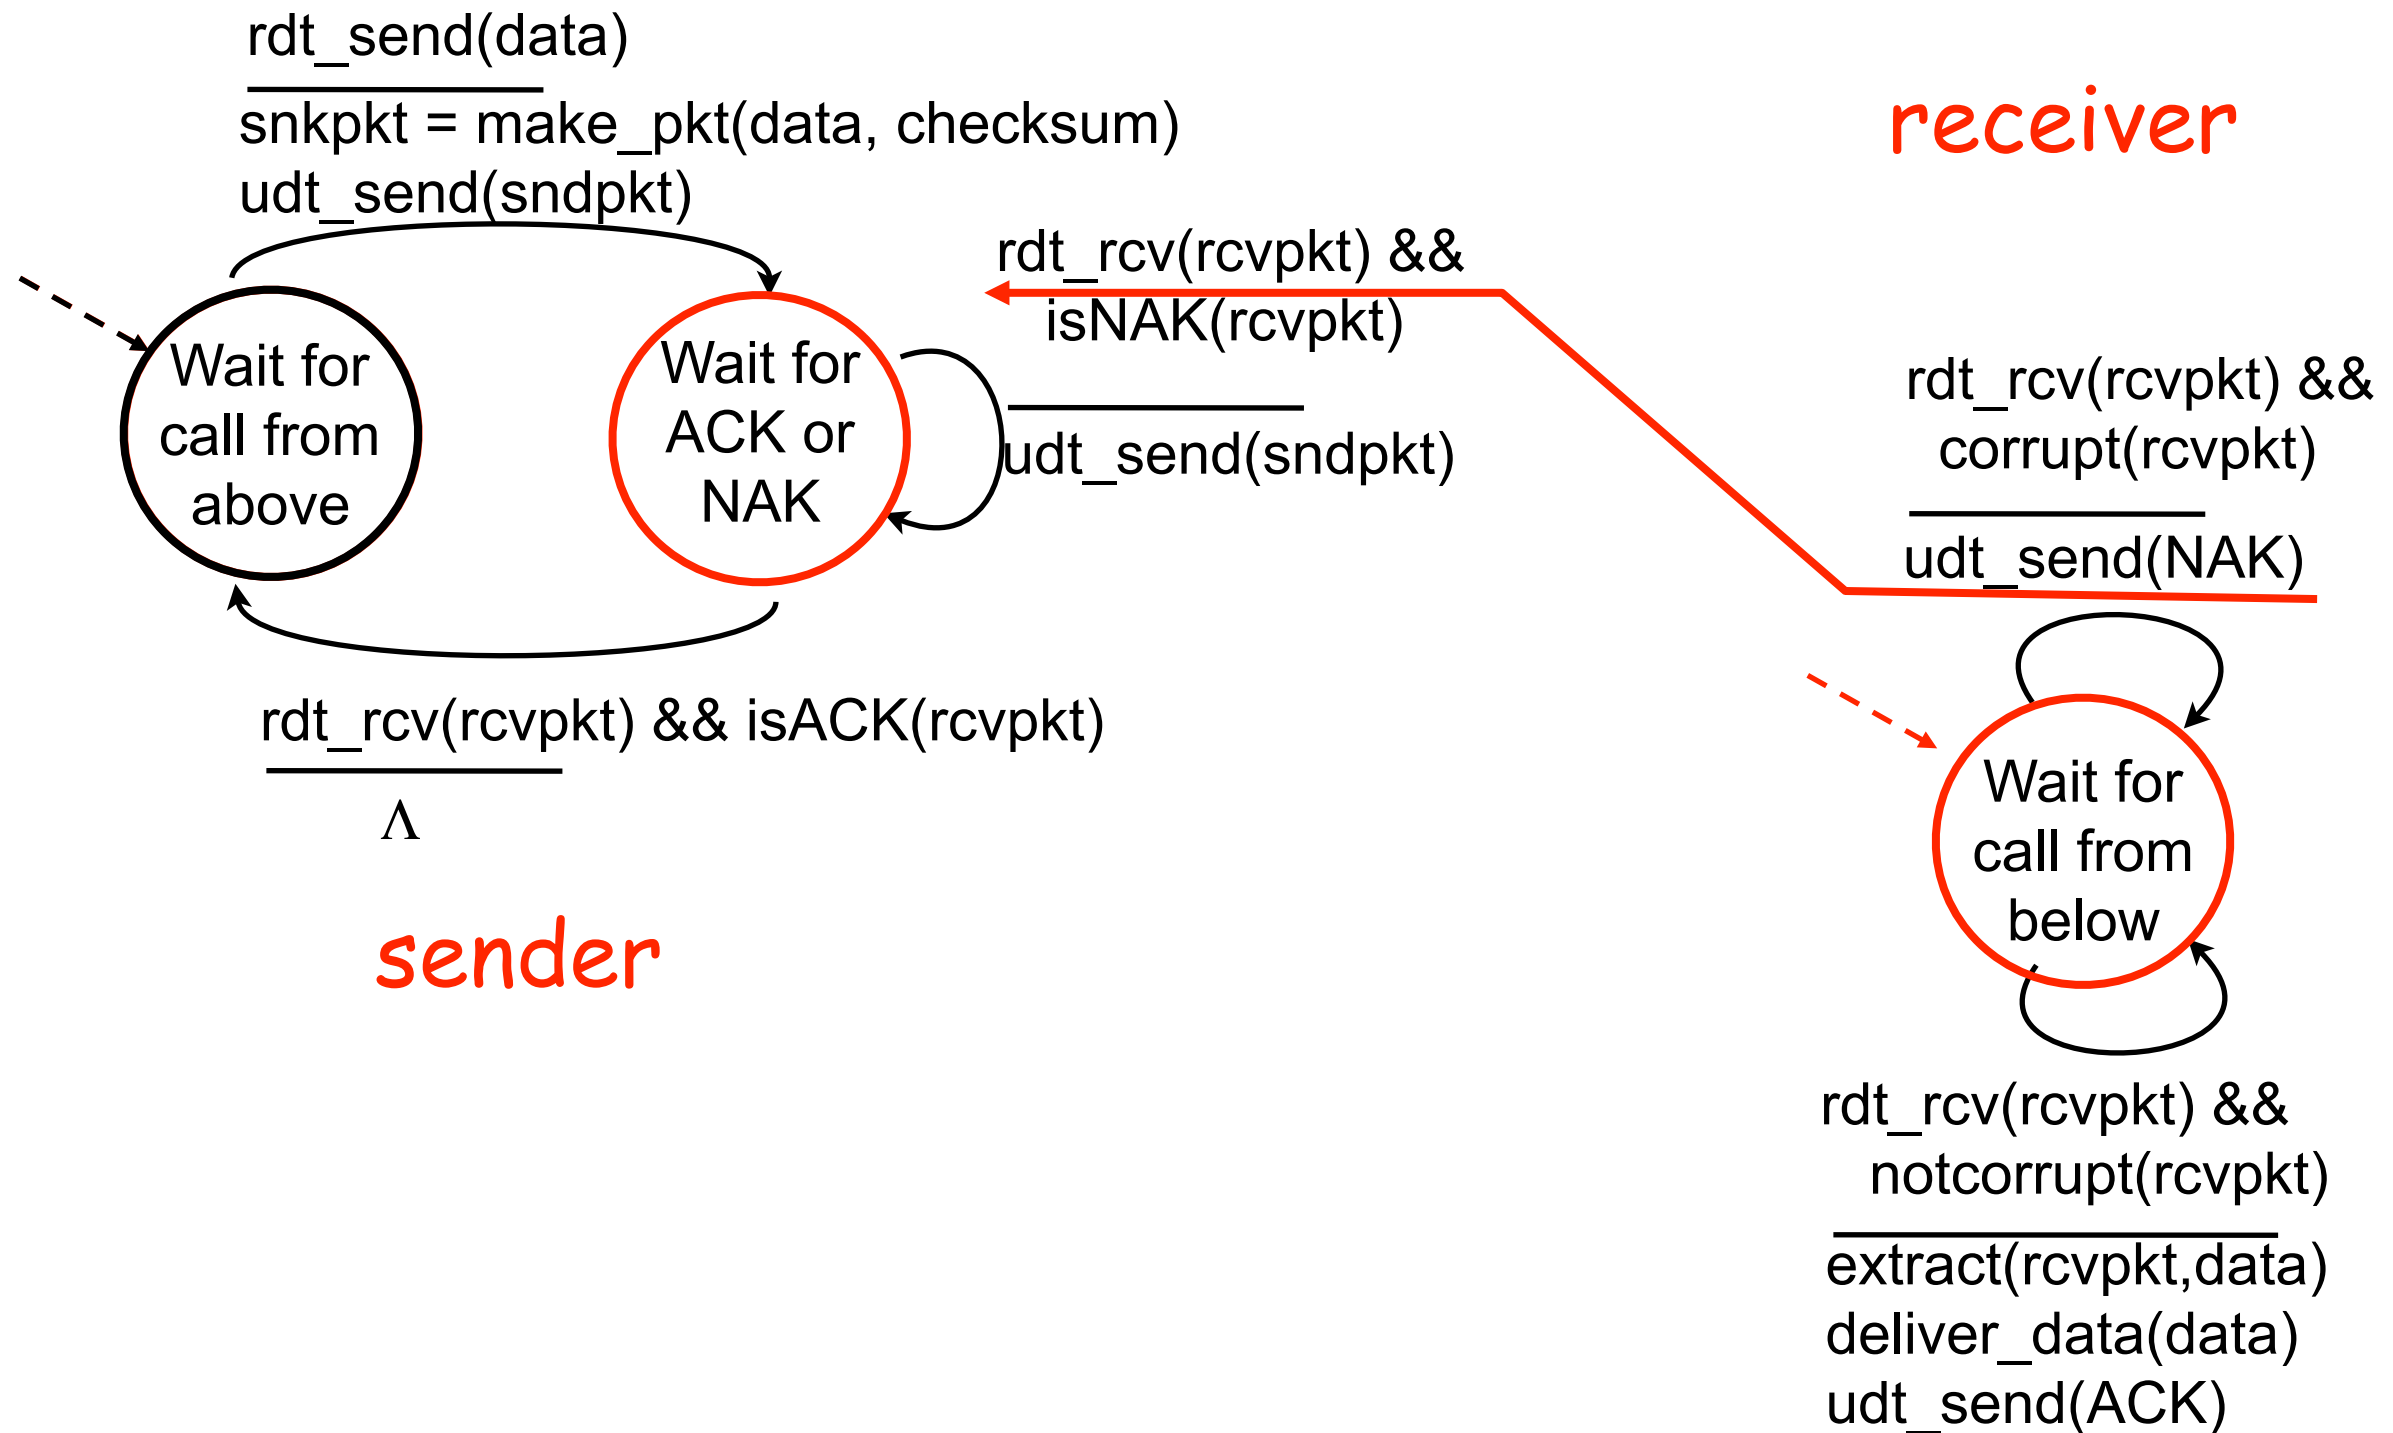
rdt2.0: error scenario



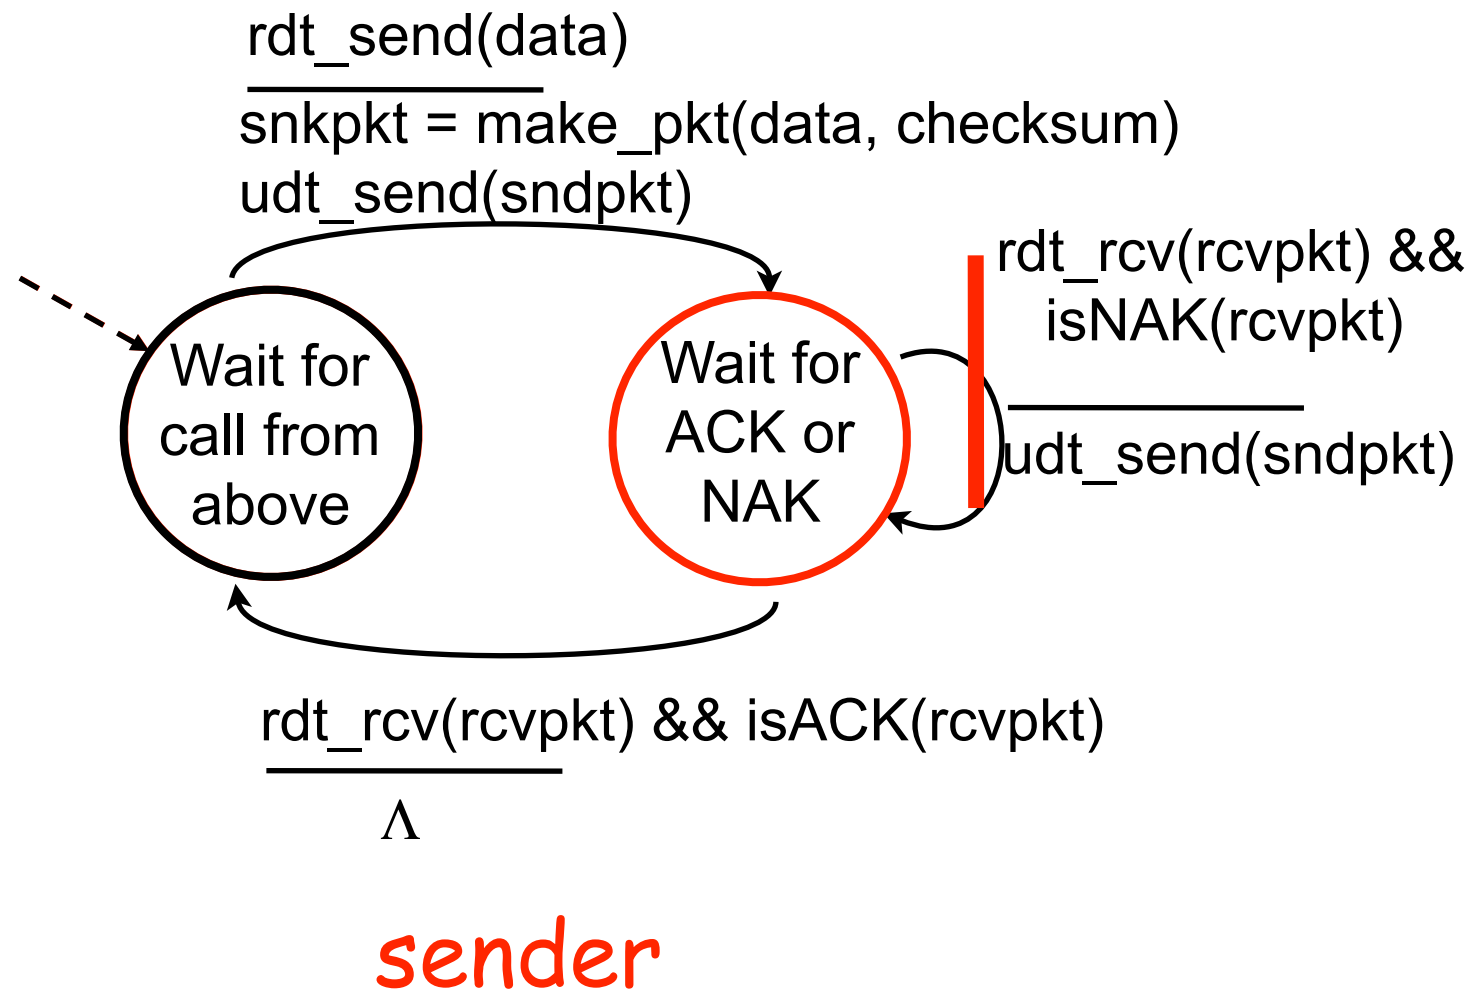
receiver



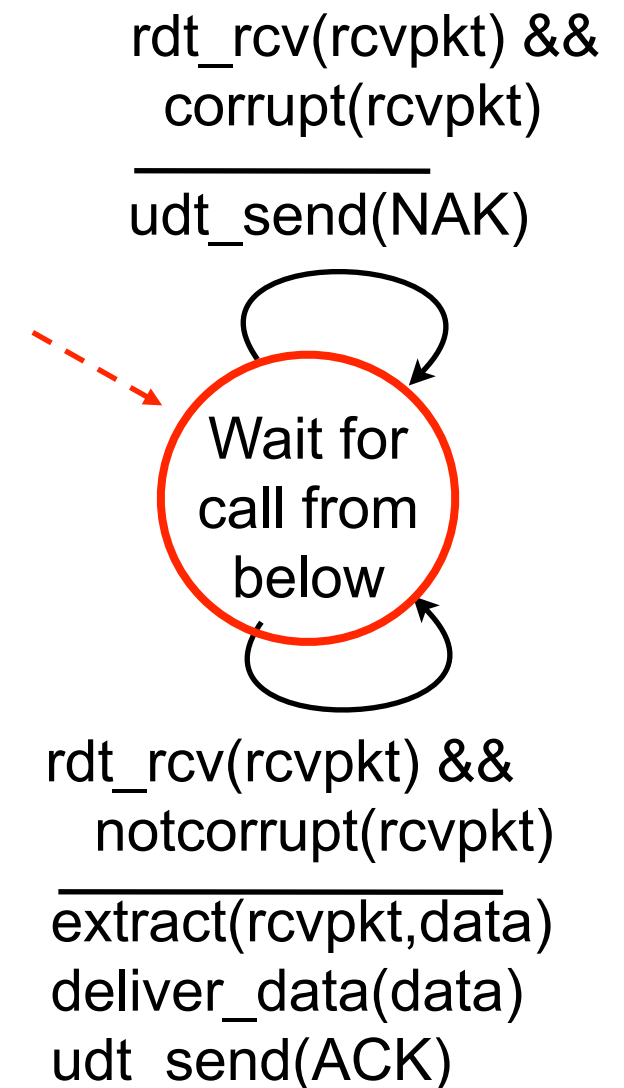
rdt2.0: error scenario



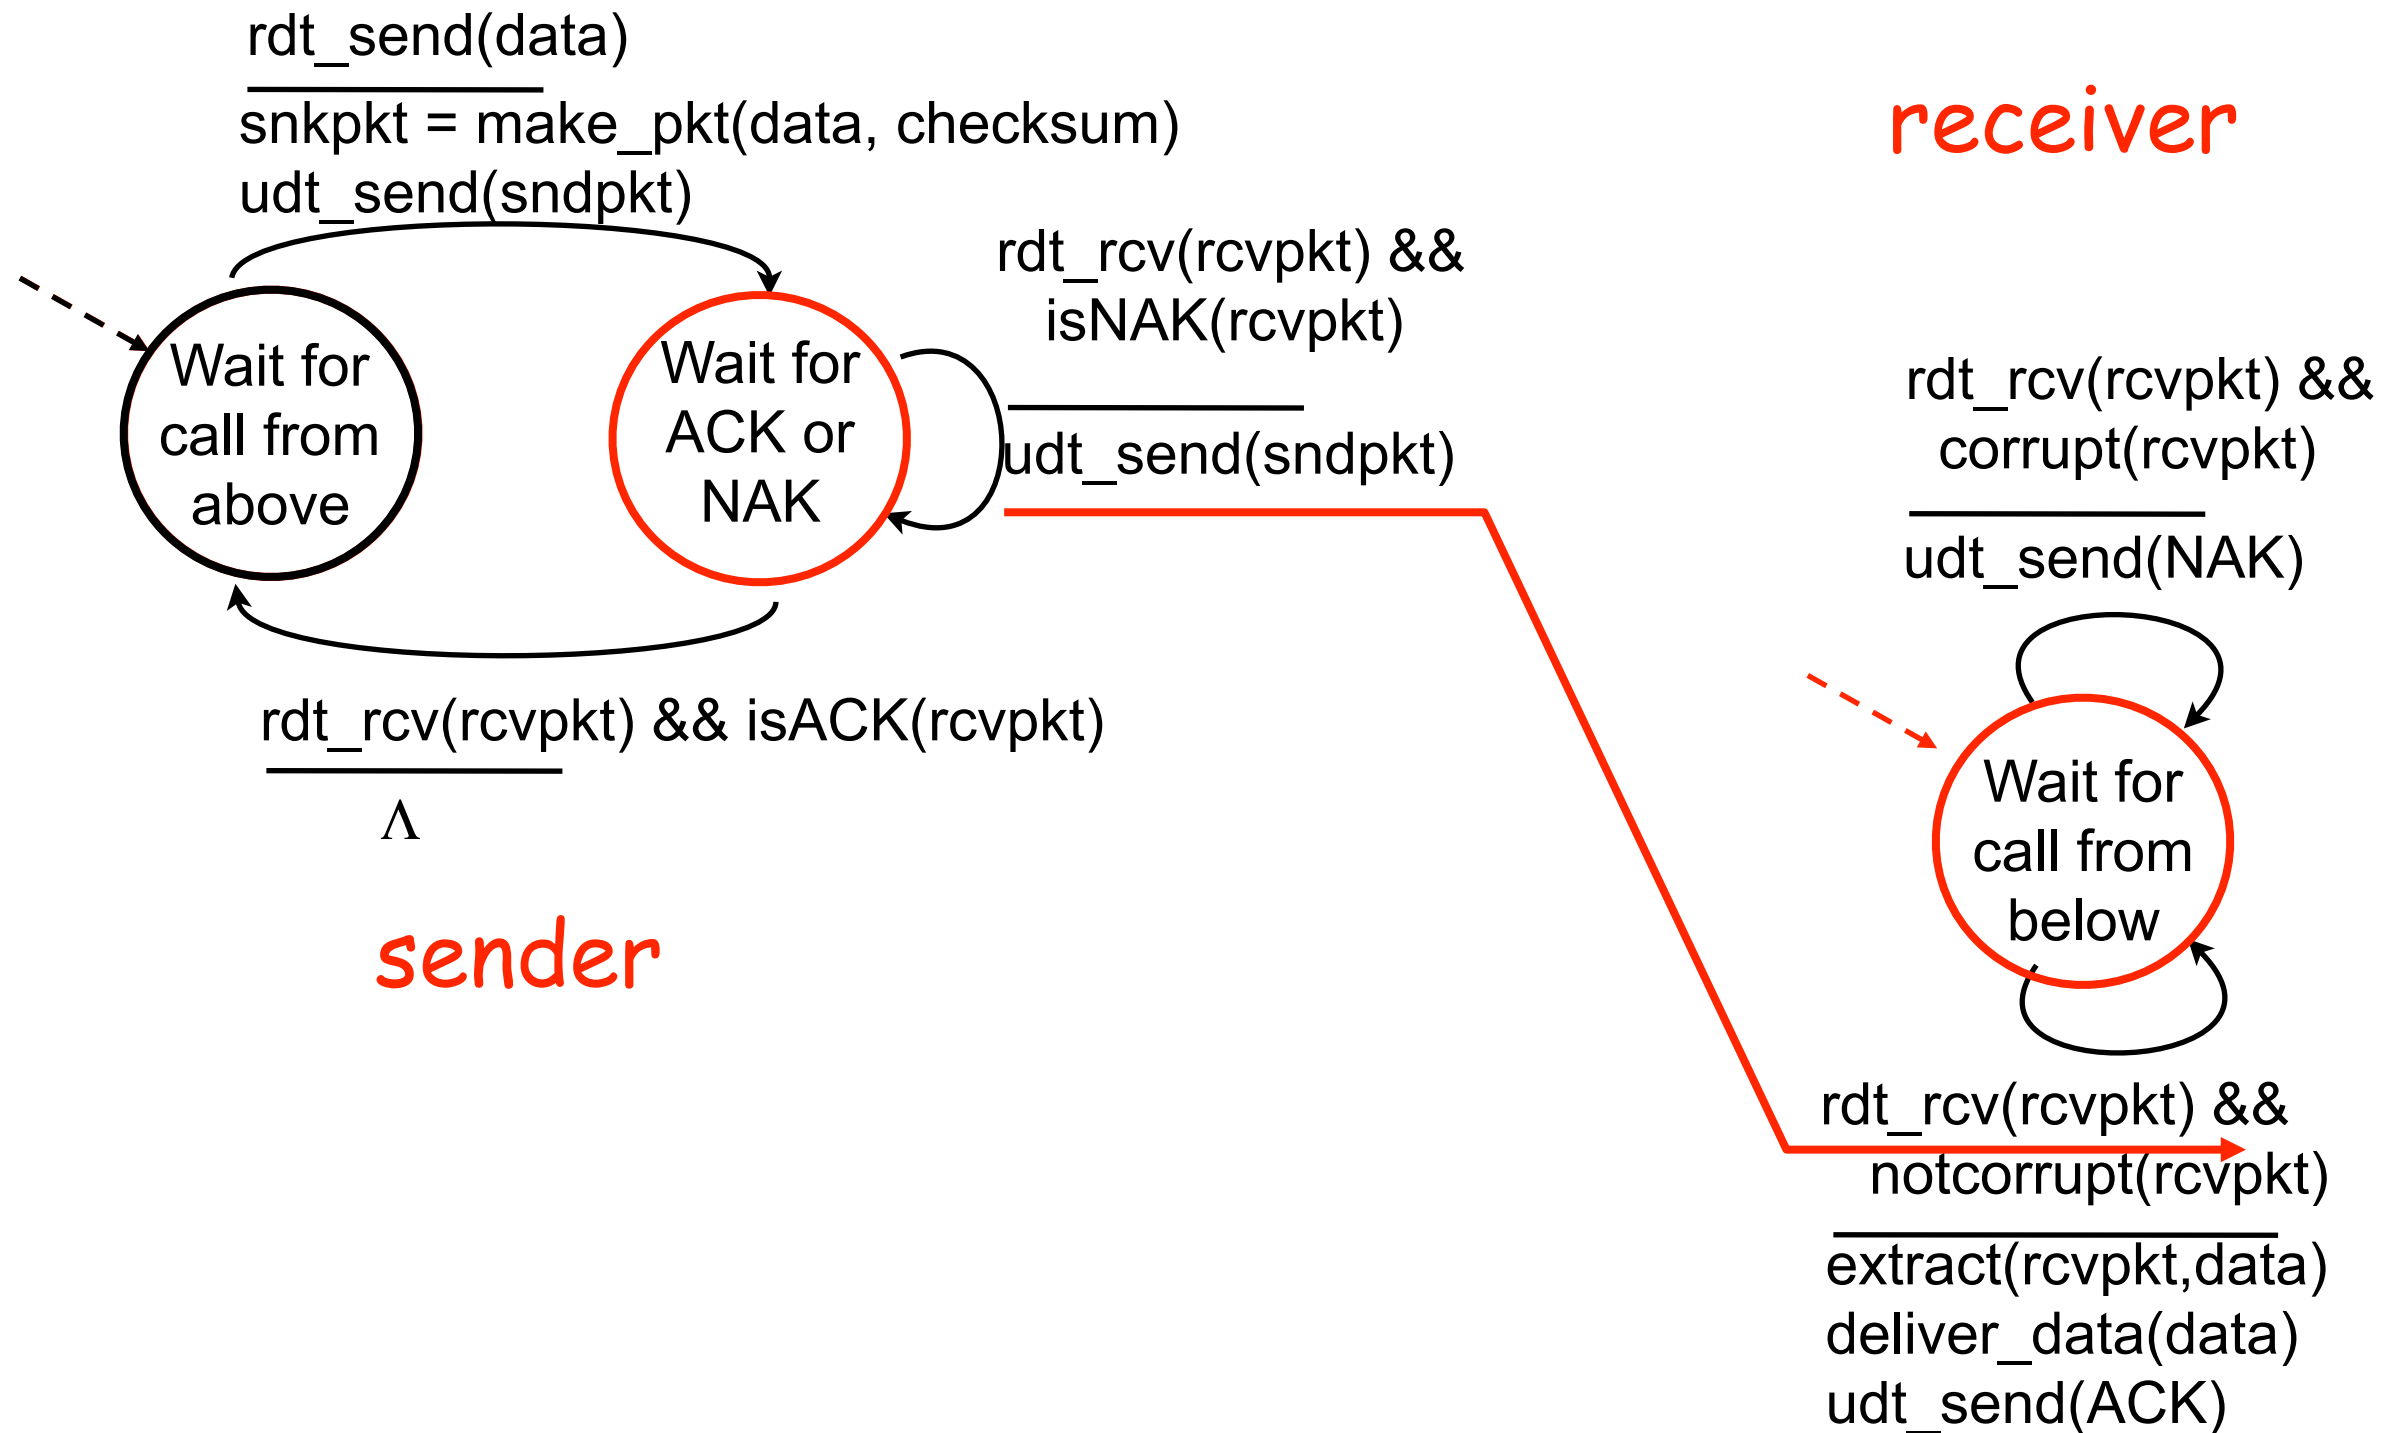
rdt2.0: error scenario



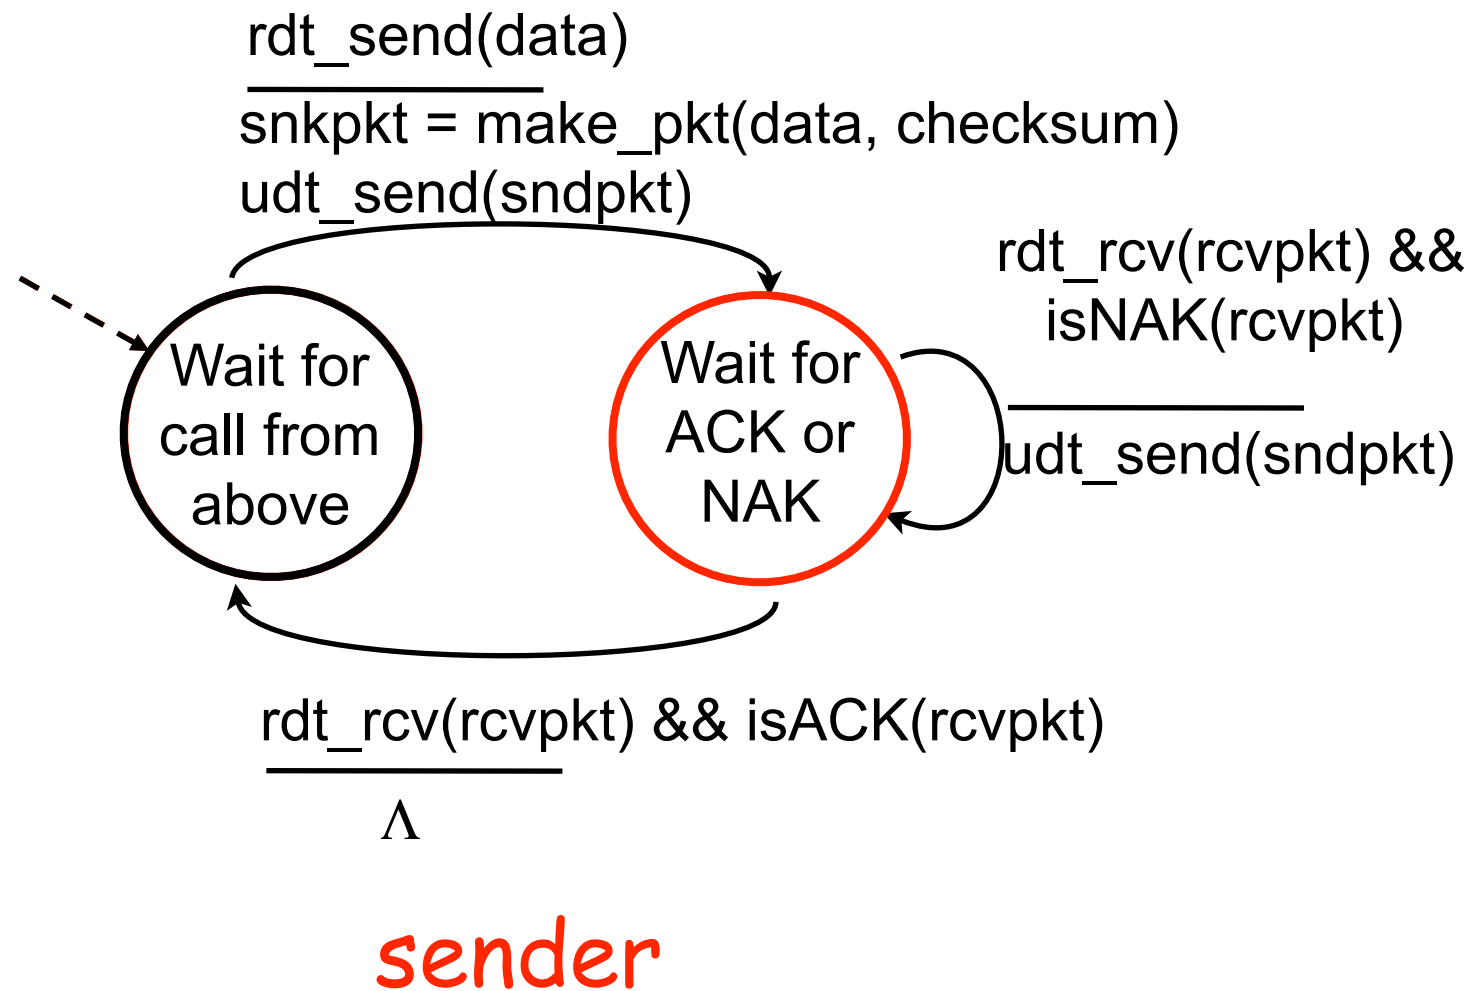
receiver



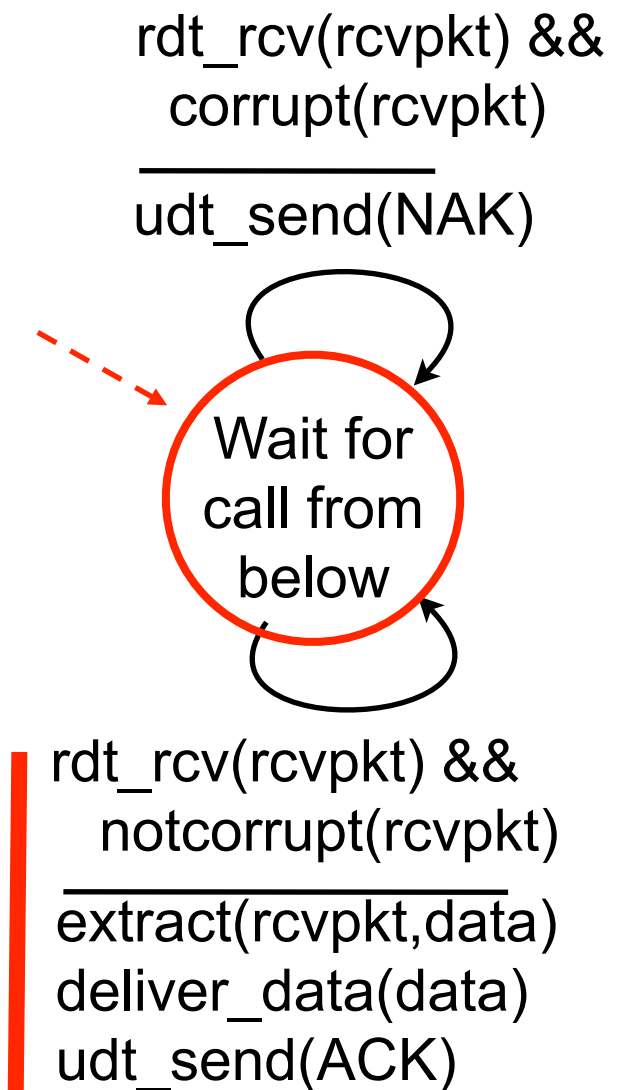
rdt2.0: error scenario



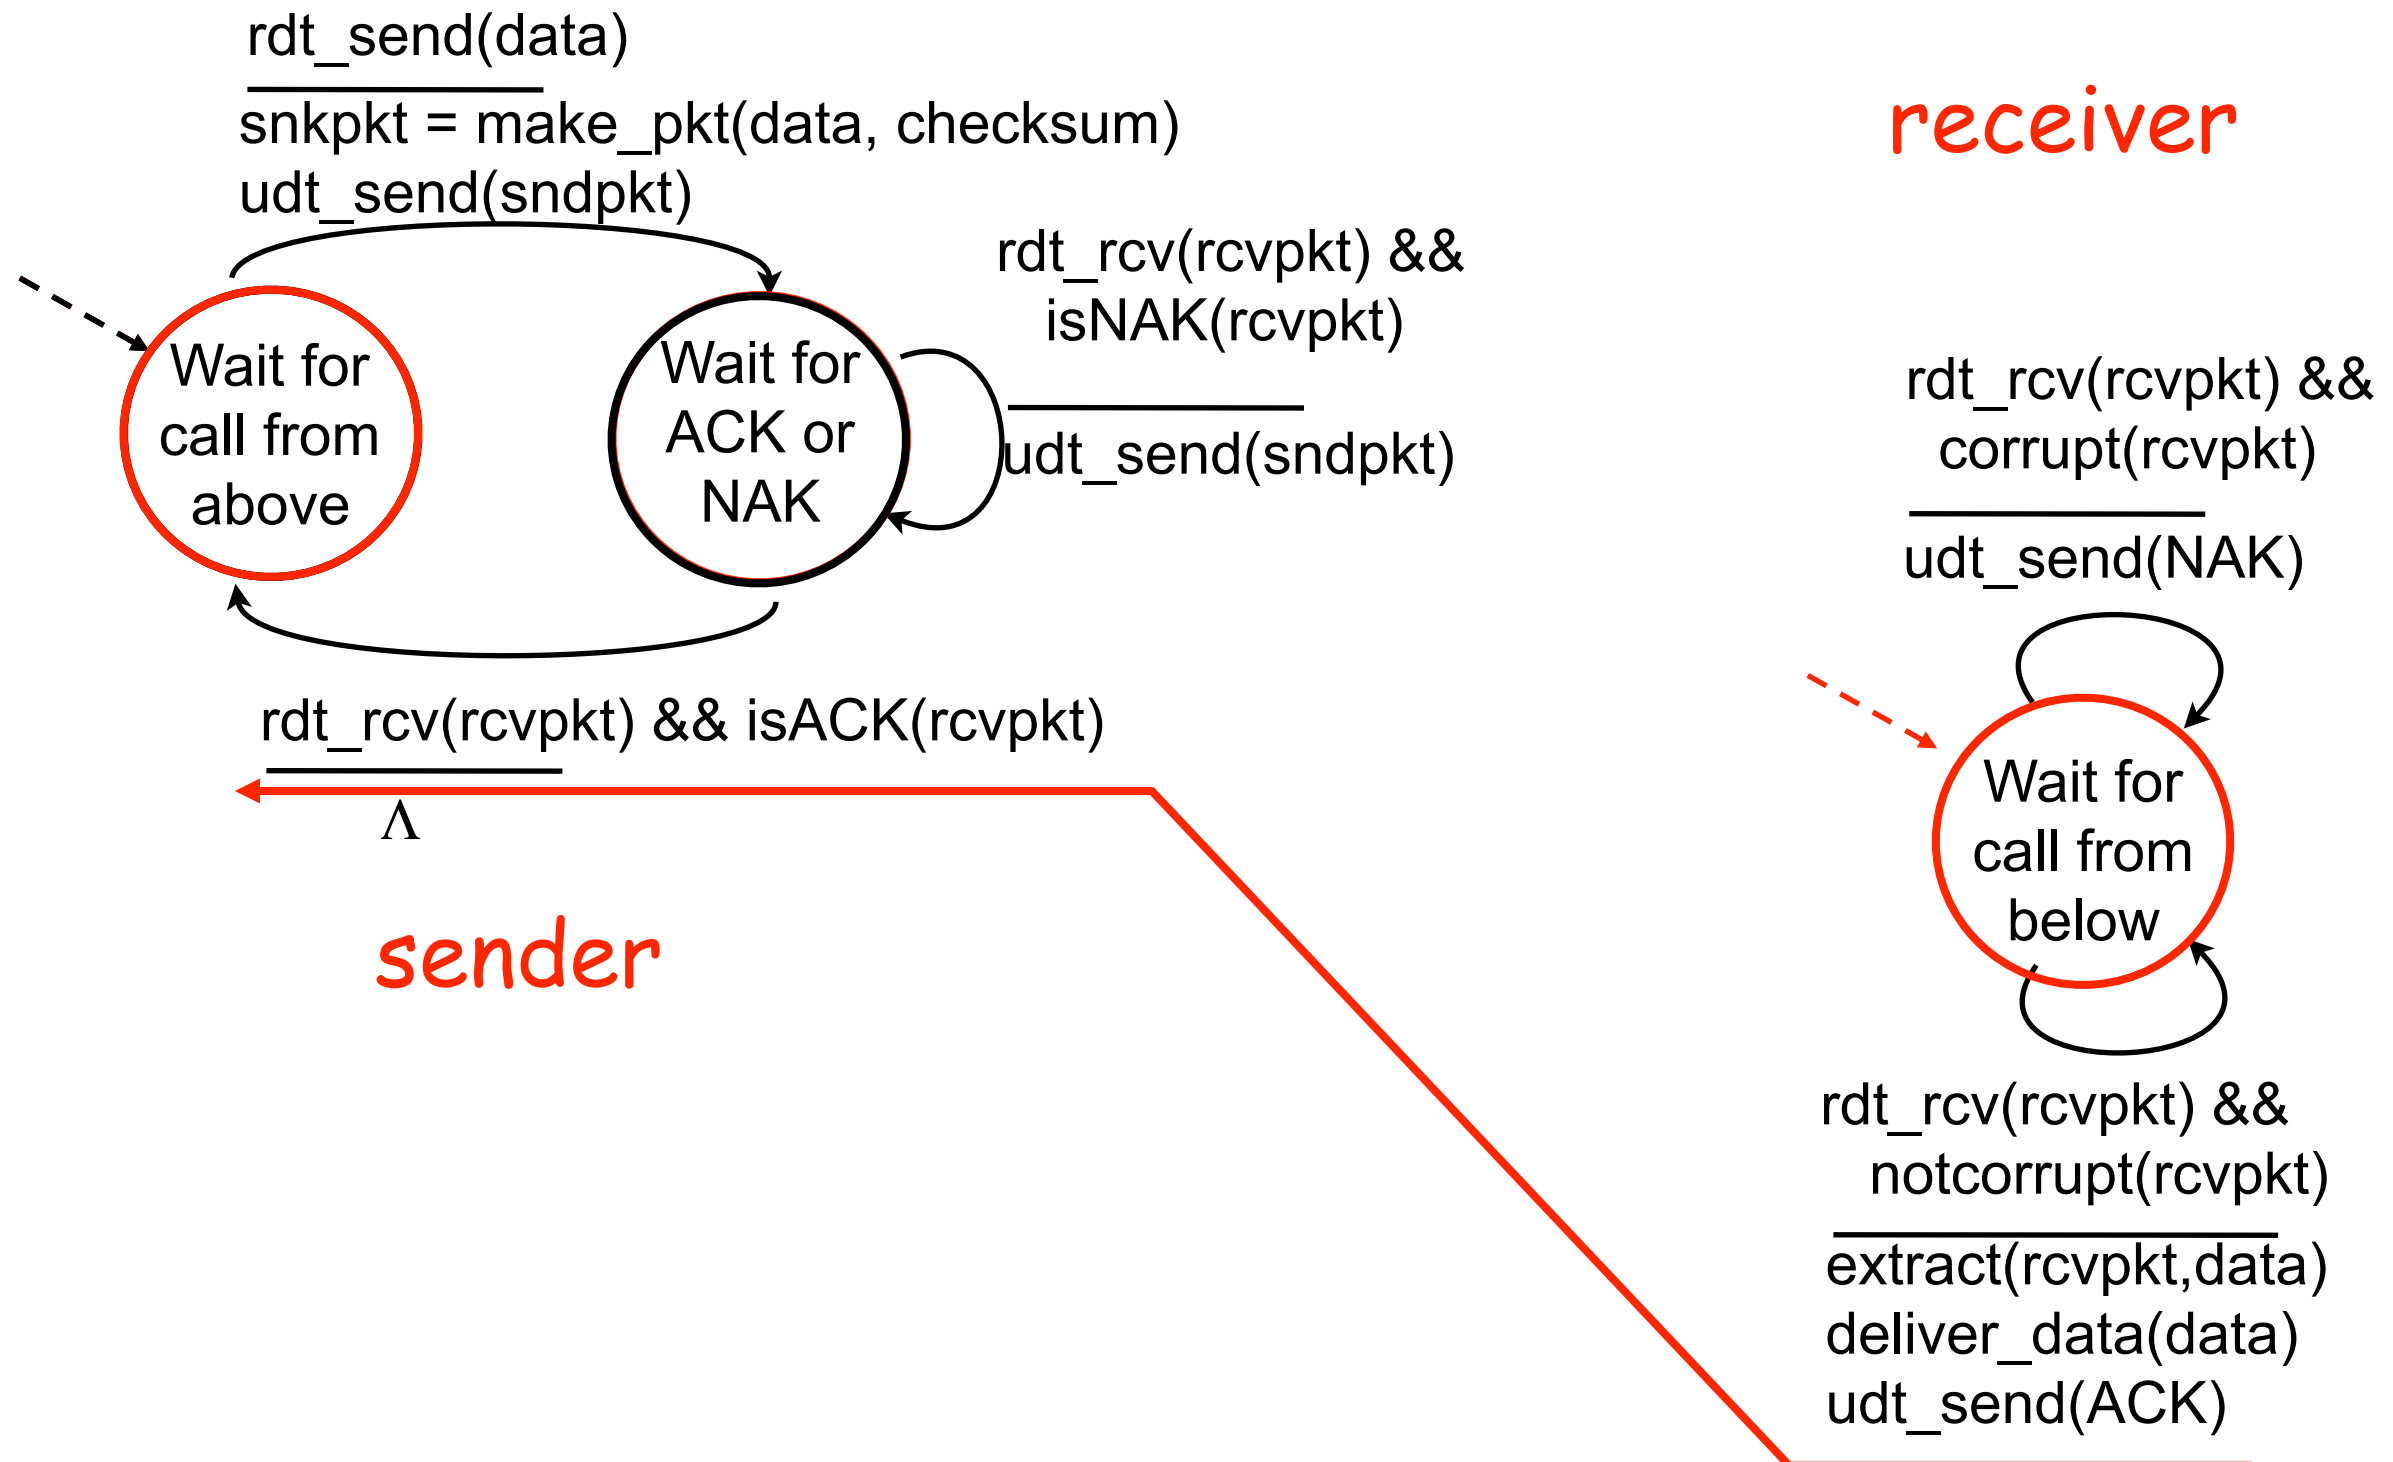
rdt2.0: error scenario



receiver



rdt2.0: error scenario



rdt2.0 has a fatal flaw!

rdt2.0 has a fatal flaw!

What happens if ACK/
NAK corrupted?

- ❖ sender doesn't know what happened at receiver!

rdt2.0 has a fatal flaw!

What happens if ACK/
NAK corrupted?

- ❖ sender doesn't know what happened at receiver!
- ❖ can't just retransmit:
possible duplicate