



Network Layer

Part 7

Mark Allman
mallman@case.edu

EECS 325/425
Fall 2018

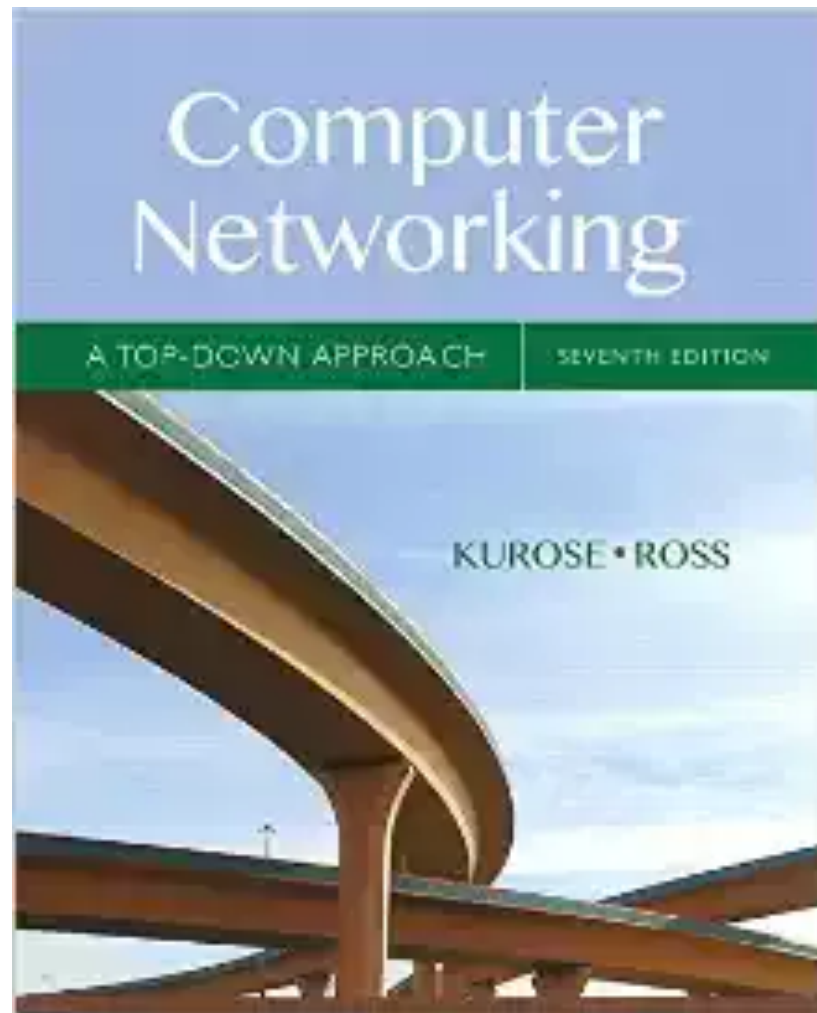
*“I was working in the lab late one night
When my eyes beheld an eerie sight”*

These slides are more-or-less directly from the slide set developed by Jim Kurose and Keith Ross for their book “Computer Networking: A Top Down Approach, 5th edition”.

The slides have been lightly adapted for Mark Allman’s EECS 325/425 Computer Networks class at Case Western Reserve University.

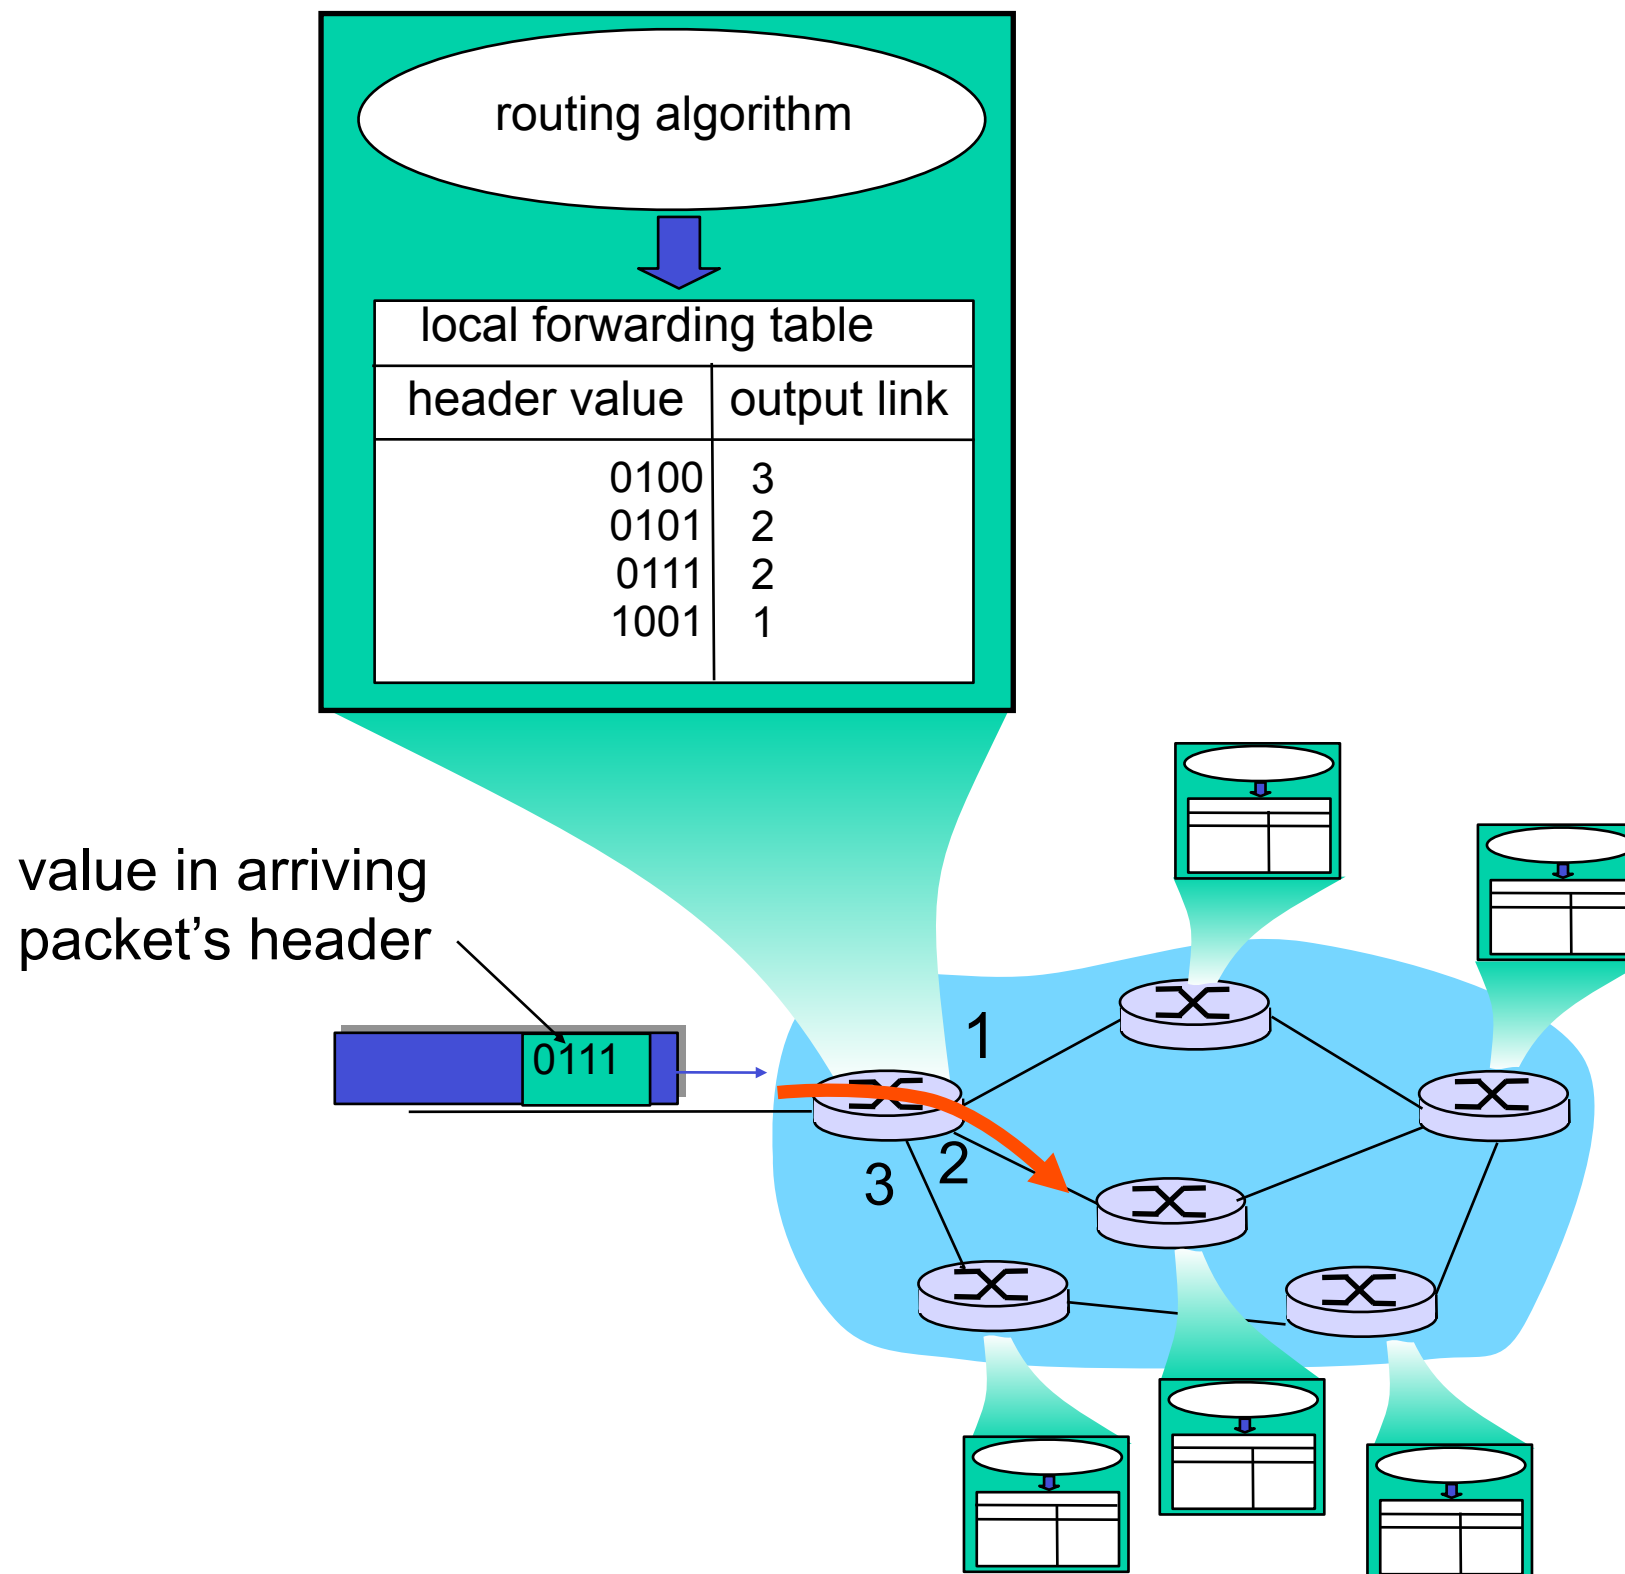
All material copyright 1996-2010
J.F Kurose and K.W. Ross, All Rights Reserved

Reading Along ...




- Network layer is chapters 4 & 5
- Routing Algorithms

Interplay between routing, forwarding



Graph abstraction




Graph: $G = (N, E)$

N = set of routers = $\{ u, v, w, x, y, z \}$


E = set of links = $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

Graph abstraction

Graph abstraction




Graph abstraction




- $c(x,x') = \text{cost of link } (x,x')$
 - e.g., $c(w,z) = 5$

Graph abstraction



- $c(x,x') = \text{cost of link } (x,x')$
 - e.g., $c(w,z) = 5$
- cost could always be 1, or inversely related to bandwidth, or inversely related to congestion

Graph abstraction



- $c(x, x') = \text{cost of link } (x, x')$


- e.g., $c(w, z) = 5$

- cost could always be 1, or inversely related to bandwidth, or inversely related to congestion

Cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

E.g., cost of $(u, w, z) = c(u, w) + c(w, z) = 5 + 5 = 10$

Graph abstraction



- $c(x, x') = \text{cost of link } (x, x')$

- e.g., $c(w, z) = 5$


- cost could always be 1, or inversely related to bandwidth, or inversely related to congestion

Cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

E.g., cost of $(u, w, z) = c(u, w) + c(w, z) = 5 + 5 = 10$

Question: What path to use between u and z ?

Graph abstraction



- $c(x, x') = \text{cost of link } (x, x')$

- e.g., $c(w, z) = 5$

- cost could always be 1, or inversely related to bandwidth, or inversely related to congestion


Cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

E.g., cost of $(u, w, z) = c(u, w) + c(w, z) = 5 + 5 = 10$

Question: What path to use between u and z ?

Question: What's the least-cost path between u and z ?

Graph abstraction



- $c(x, x') = \text{cost of link } (x, x')$

- e.g., $c(w, z) = 5$

- cost could always be 1, or inversely related to bandwidth, or inversely related to congestion

Cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

E.g., cost of $(u, w, z) = c(u, w) + c(w, z) = 5 + 5 = 10$

Question: What path to use between u and z ?

Question: What's the least-cost path between u and z ?

Routing algorithm: algorithm that finds least-cost path

Routing Algorithm classification

Routing Algorithm classification


Global or decentralized
information?

Routing Algorithm classification

Global or decentralized
information?

Global:

- ❖ all routers have complete topology, link cost info
- ❖ "link state" algorithms




Routing Algorithm classification

Global or decentralized information?


Global:

- ❖ all routers have complete topology, link cost info
- ❖ "link state" algorithms



Decentralized:

- ❖ router knows physically-connected neighbors, link costs to neighbors
- ❖ iterative process of computation, exchange of info with neighbors
- ❖ "distance vector" algorithms



Routing Algorithm classification

Static or dynamic?

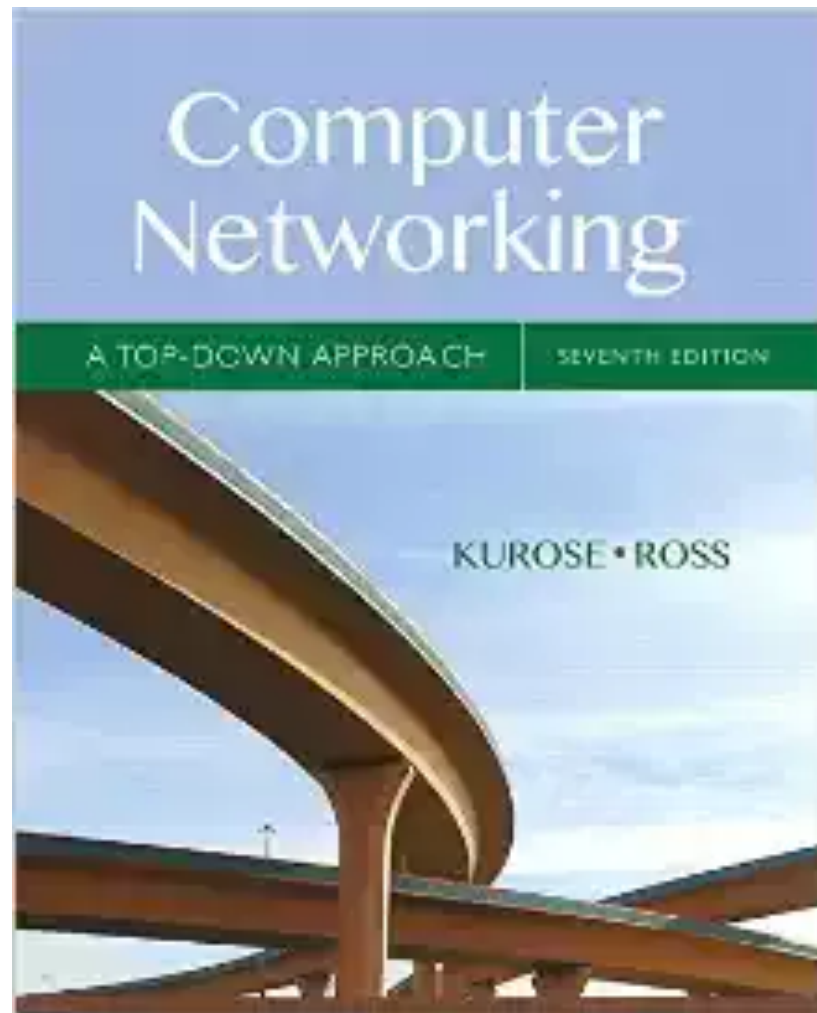
Static:

- ❖ routes change slowly over time

Dynamic:

- ❖ routes change more quickly
 - periodic update
 - in response to link cost changes

Reading Along ...



- Network layer is chapters 4 & 5
- Routing Algorithms
Link state routing

A Link-State Routing Algorithm

Dijkstra's algorithm

- ❖ net topology, link costs known to all nodes
 - accomplished via "link state broadcast"
 - all nodes have same info
- ❖ computes least cost paths from one node ('source') to all other nodes
 - gives forwarding table for that node
- ❖ iterative: after k iterations, know least cost path to k dest.'s

Dijkstra's Algorithm

1 **Initialization:**

2 $N' = \{a\}$

3 for all nodes b

4 if b adjacent to a

5 then $D(b) = c(a,b)$

6 else $D(b) = \infty$

7

8 **Loop**

9 find e not in N' such that $D(e)$ is a minimum

10 add e to N'

11 update $D(f)$ for all f adjacent to e and not in N' :

12 $D(f) = \min(D(f), D(e) + c(e,b))$

13 /* new cost to f is either old cost to f or known

14 shortest path cost to e plus cost from e to f */

15 **until all nodes in N'**

Dijkstra's algorithm, discussion

Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic


Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic



initially


Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic



initially


Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic




Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic





Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic




Dijkstra's algorithm, discussion

Algorithm complexity: n nodes


- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic



initially



... recompute
routing



Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic





Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic






Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic






Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic







Dijkstra's algorithm, discussion

Algorithm complexity: n nodes

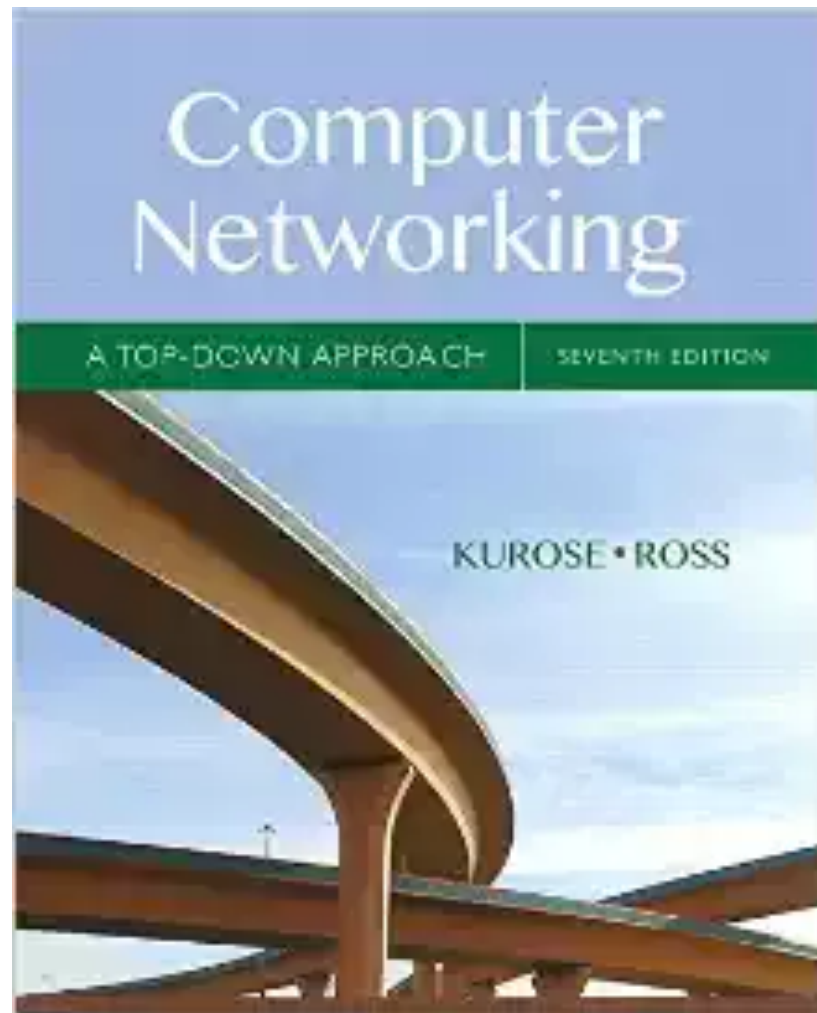
- ❖ each iteration: need to check all nodes, w , not in N
- ❖ $n(n+1)/2$ comparisons: $O(n^2)$
- ❖ more efficient implementations possible: $O(n \log n)$

Oscillations possible:

- ❖ e.g., link cost = amount of carried traffic



Reading Along ...



- Network layer is chapters 4 & 5
- Routing Algorithms
Distance vector routing

Distance Vector Algorithm

Distance Vector Algorithm

Bellman-Ford Equation (dynamic programming)

Define


$d_x(y) :=$ cost of least-cost path from x to y

Then

$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$


where min is taken over all neighbors v of x

Bellman-Ford example



Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$


Bellman-Ford example



Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

Bellman-Ford example




Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$d_u(z) = \min \{ c(u,v) + d_v(z),$$

Bellman-Ford example




Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$d_u(z) = \min \{ c(u,v) + d_v(z), \\ c(u,x) + d_x(z),$$

Bellman-Ford example




Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$d_u(z) = \min \{ c(u,v) + d_v(z), \\ c(u,x) + d_x(z), \\ c(u,w) + d_w(z) \}$$

Bellman-Ford example




Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \end{aligned}$$

Bellman-Ford example




Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \end{aligned}$$

Bellman-Ford example




Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$d_u(z) = \min \{ c(u,v) + d_v(z), \\ c(u,x) + d_x(z), \\ c(u,w) + d_w(z) \}$$

$$= \min \{ 2 + 5, \\ 1 + 3, \\ 5 + 3 \} = 4$$

Bellman-Ford example



Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

Node that achieves minimum is next
hop in shortest path → forwarding table

Distance Vector Algorithm

- ❖ $D_x(y)$ = estimate of least cost from x to y
 - x maintains distance vector $D_x = [D_x(y): y \in N]$

❖ node x :

- knows cost to each neighbor v : $c(x,v)$
- maintains its neighbors' distance vectors.

For each neighbor v , x maintains

$$D_v = [D_v(y): y \in N]$$

Distance vector algorithm

Basic idea:

- ❖ from time-to-time, each node sends its own distance vector estimate to neighbors
- ❖ when x receives new DV estimate from neighbor, it updates its own DV using B-F equation:
$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \quad \text{for each node } y \in N$$
- ❖ under minor, natural conditions, the estimate $D_x(y)$ converge to the actual least cost $d_x(y)$

Distance Vector Algorithm

Iterative, asynchronous:

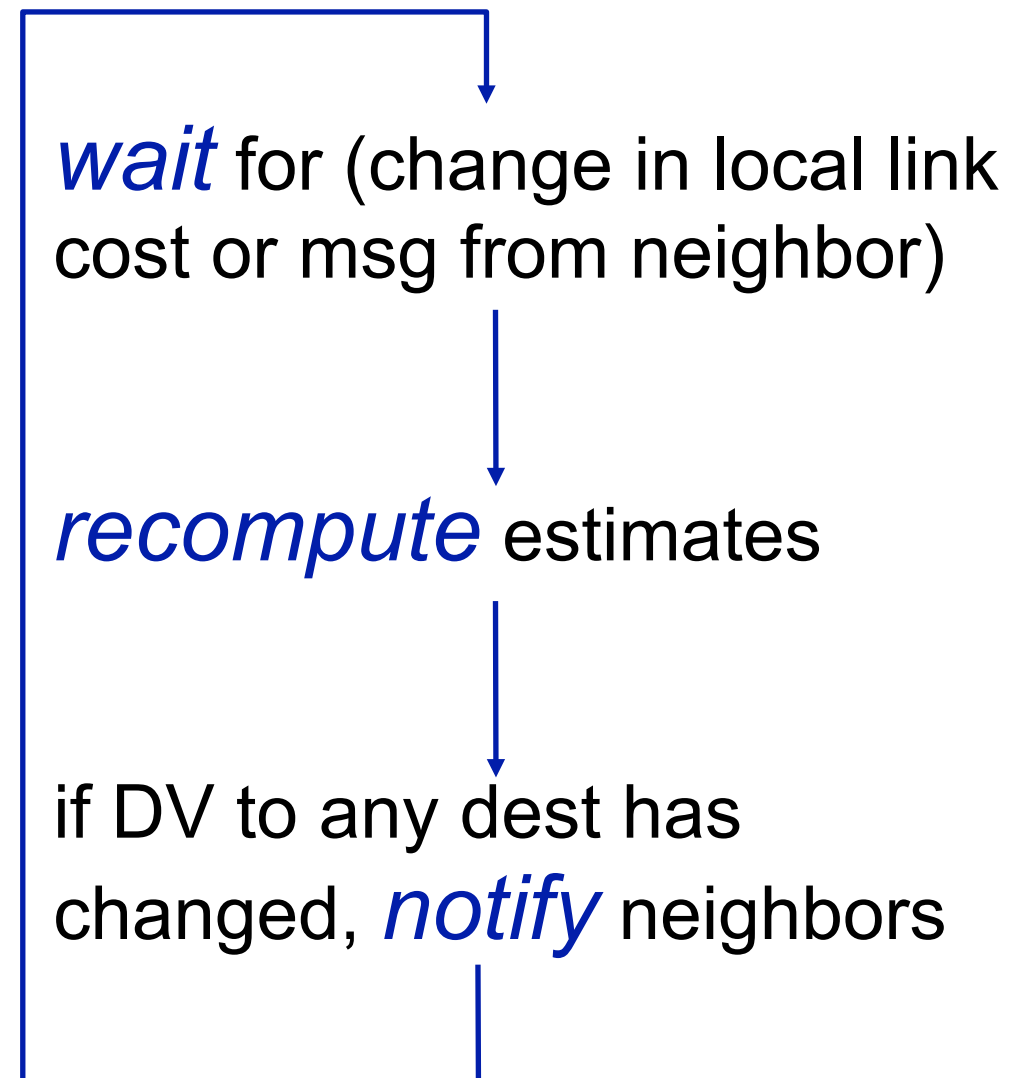
each local iteration caused by:

- ❖ local link cost change
- ❖ DV update message from neighbor

Distributed:

- ❖ each node notifies neighbors only when its DV changes
 - neighbors then notify their neighbors if necessary


Each node:



Distance Vector: link cost changes

Link cost changes:


- ❖ node detects local link cost change
- ❖ updates routing info, recalculates distance vector
- ❖ if DV changes, notify neighbors
- ❖ good news travels fast



Distance Vector: link cost changes

Link cost changes:

- ❖ good news travels fast
- ❖ bad news travels slow - "count to infinity" problem!
- ❖ 44 iterations before algorithm stabilizes: see text



Comparison of LS and DV algorithms

Comparison of LS and DV algorithms

Message complexity

- ❖ LS: with n nodes, E links,
 $O(nE)$ msgs sent
- ❖ DV: exchange between
neighbors only
 - convergence time varies

Comparison of LS and DV algorithms

Message complexity

- ❖ LS: with n nodes, E links,
 $O(nE)$ msgs sent
- ❖ DV: exchange between
neighbors only
 - convergence time varies

Speed of Convergence

- ❖ LS: $O(n^2)$ algorithm requires
 $O(nE)$ msgs
 - may have oscillations
- ❖ DV: convergence time varies
 - may be routing loops
 - count-to-infinity problem

Comparison of LS and DV algorithms

Message complexity

- ❖ LS: with n nodes, E links, $O(nE)$ msgs sent
- ❖ DV: exchange between neighbors only
 - convergence time varies

Speed of Convergence

- ❖ LS: $O(n^2)$ algorithm requires $O(nE)$ msgs
 - may have oscillations
- ❖ DV: convergence time varies
 - may be routing loops
 - count-to-infinity problem

Robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect **link** cost
- each node computes only its own table

DV:


- DV node can advertise incorrect **path** cost
- each node's table used by others
 - error propagate thru network

Routing Algorithms


Link State

Distance Vector

Routing Algorithms




Link State




Distance Vector

Routing Algorithms



Link State

❖ Both can work in small networks




Distance Vector

❖ Neither great for dealing with big networks

- millions of destinations


Routing Algorithms

- ❖ We must apply divide and conquer
- ❖ Aggregate routers within organizations together and deal with them as a single unit
 - An "autonomous system"




Routing Algorithms


- ❖ Must have two routing algorithms:
 - Intra-AS or Interior routing algorithms to manage routing inside an AS
 - Inter-AS or Exterior routing algorithms to manage routing between ASes



Interconnected ASes



Interconnected ASes



❖ forwarding table configured by both intra- and inter-AS routing algorithm

- intra-AS sets entries for internal dests
- inter-AS & intra-As sets entries for external dests