



# Congestion Control

Mark Allman  
*mallman@case.edu*

EECS 325/425  
Fall 2018

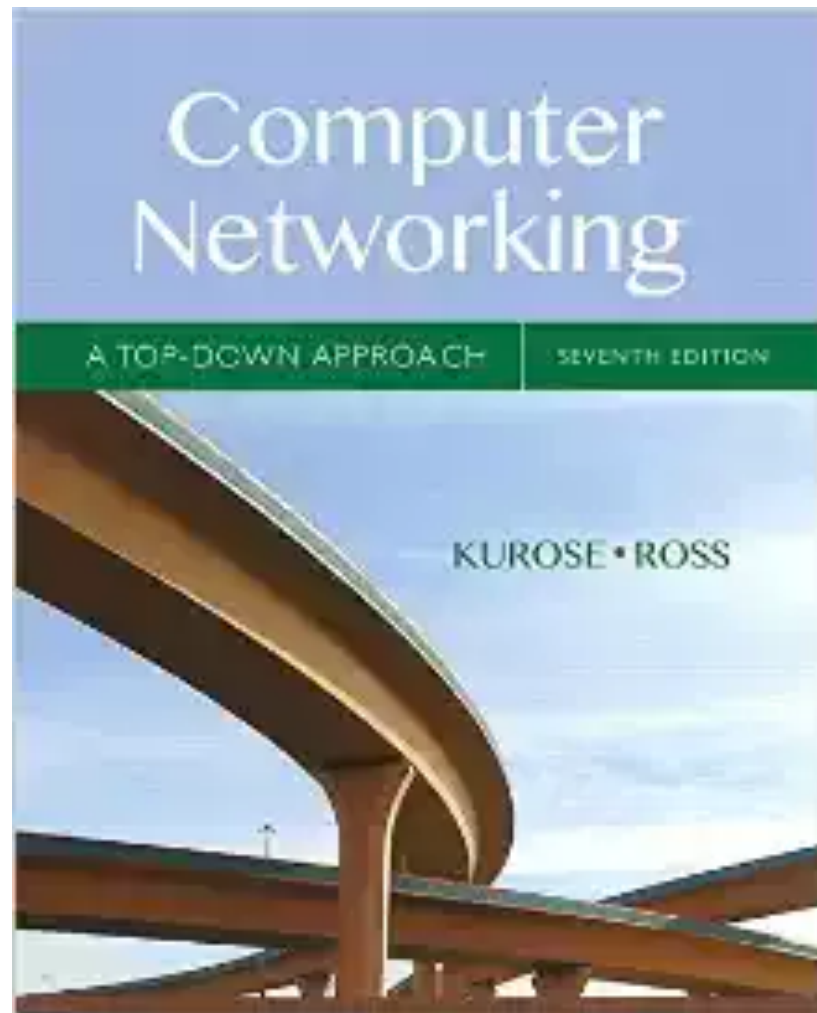
*“Well, we busted outa class, had to get away from those fools,  
We learned more from a three minute record than we ever learned in school.”*

These slides are more-or-less directly from the slide set developed by Jim Kurose and Keith Ross for their book “Computer Networking: A Top Down Approach, 5th edition”.

The slides have been lightly adapted for Mark Allman’s EECS 325/425 Computer Networks class at Case Western Reserve University.

All material copyright 1996-2010  
J.F Kurose and K.W. Ross, All Rights Reserved

# Reading Along ...



- 3.6: Principles of Congestion Control

# Congestion Control


# Congestion Control

❖ What is "congestion"?

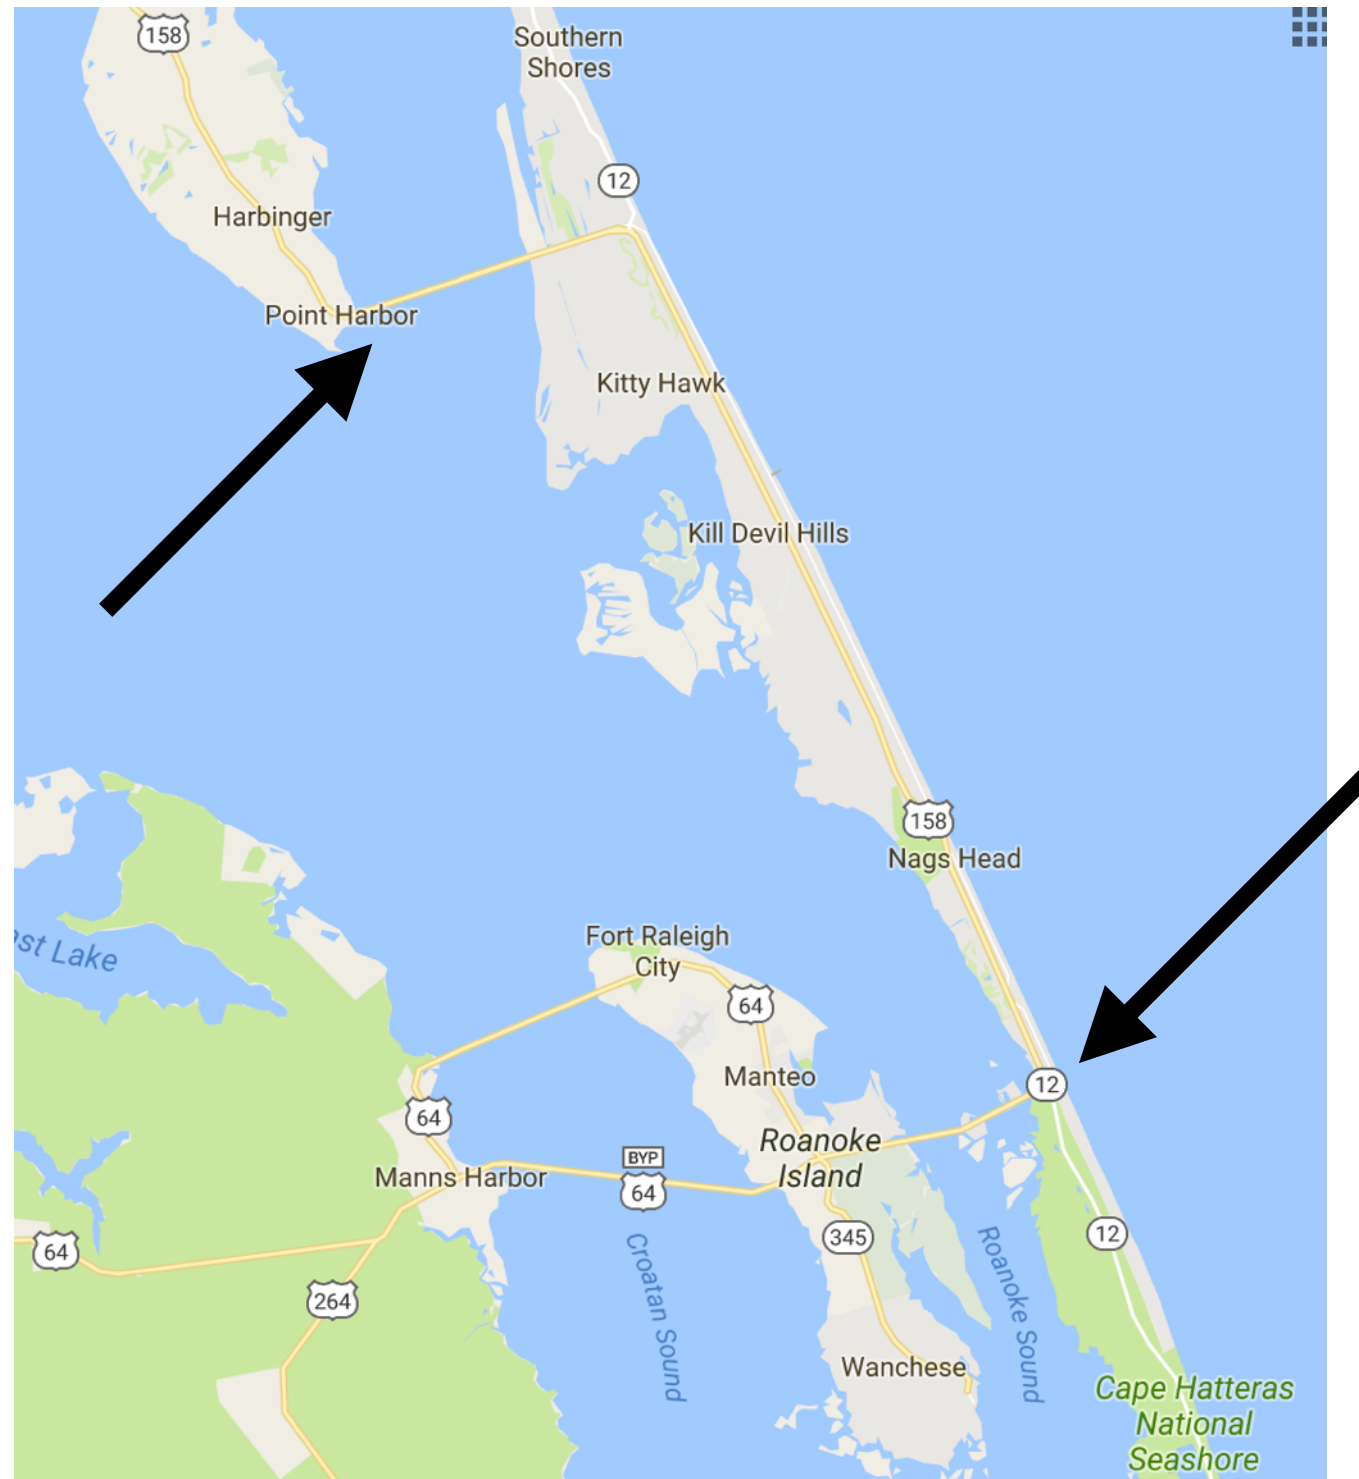
# Congestion Control

- ❖ What is "congestion"?
- ❖ Why does it need "controlled"?

# Congestion




# Congestion






# Congestion



A map of the Albemarle-Pamlico region in North Carolina, showing the Albemarle Sound and Pamlico River. The map includes major cities and towns such as Jamestown, Newport News, Hampton, Norfolk, Virginia Beach, Chesapeake, Suffolk, Windsor, Courtland, Franklin, Boykins, Murfreesboro, Winton, Gatesville, South Mills, Moyock, Knotts Island, Currituck, Corolla, Elizabeth City, Winfall, Edenton, Duck, Kitty Hawk, and Kill Devil Hills. Major highways are marked with route numbers (e.g., 95, 460, 58, 13, 17, 158, 11, 42, 32, 343, 12). Seven arrows indicate the locations of the study sites: three black arrows point to sites in the northern part of the region (near Jamestown, Windsor, and Norfolk), and four red arrows point to sites in the southern part of the region (near Winton, South Mills, Currituck, and Duck/Kitty Hawk/Kill Devil Hills).

# Congestion












# Worth the Effort!



# A Networking Example




Receiver

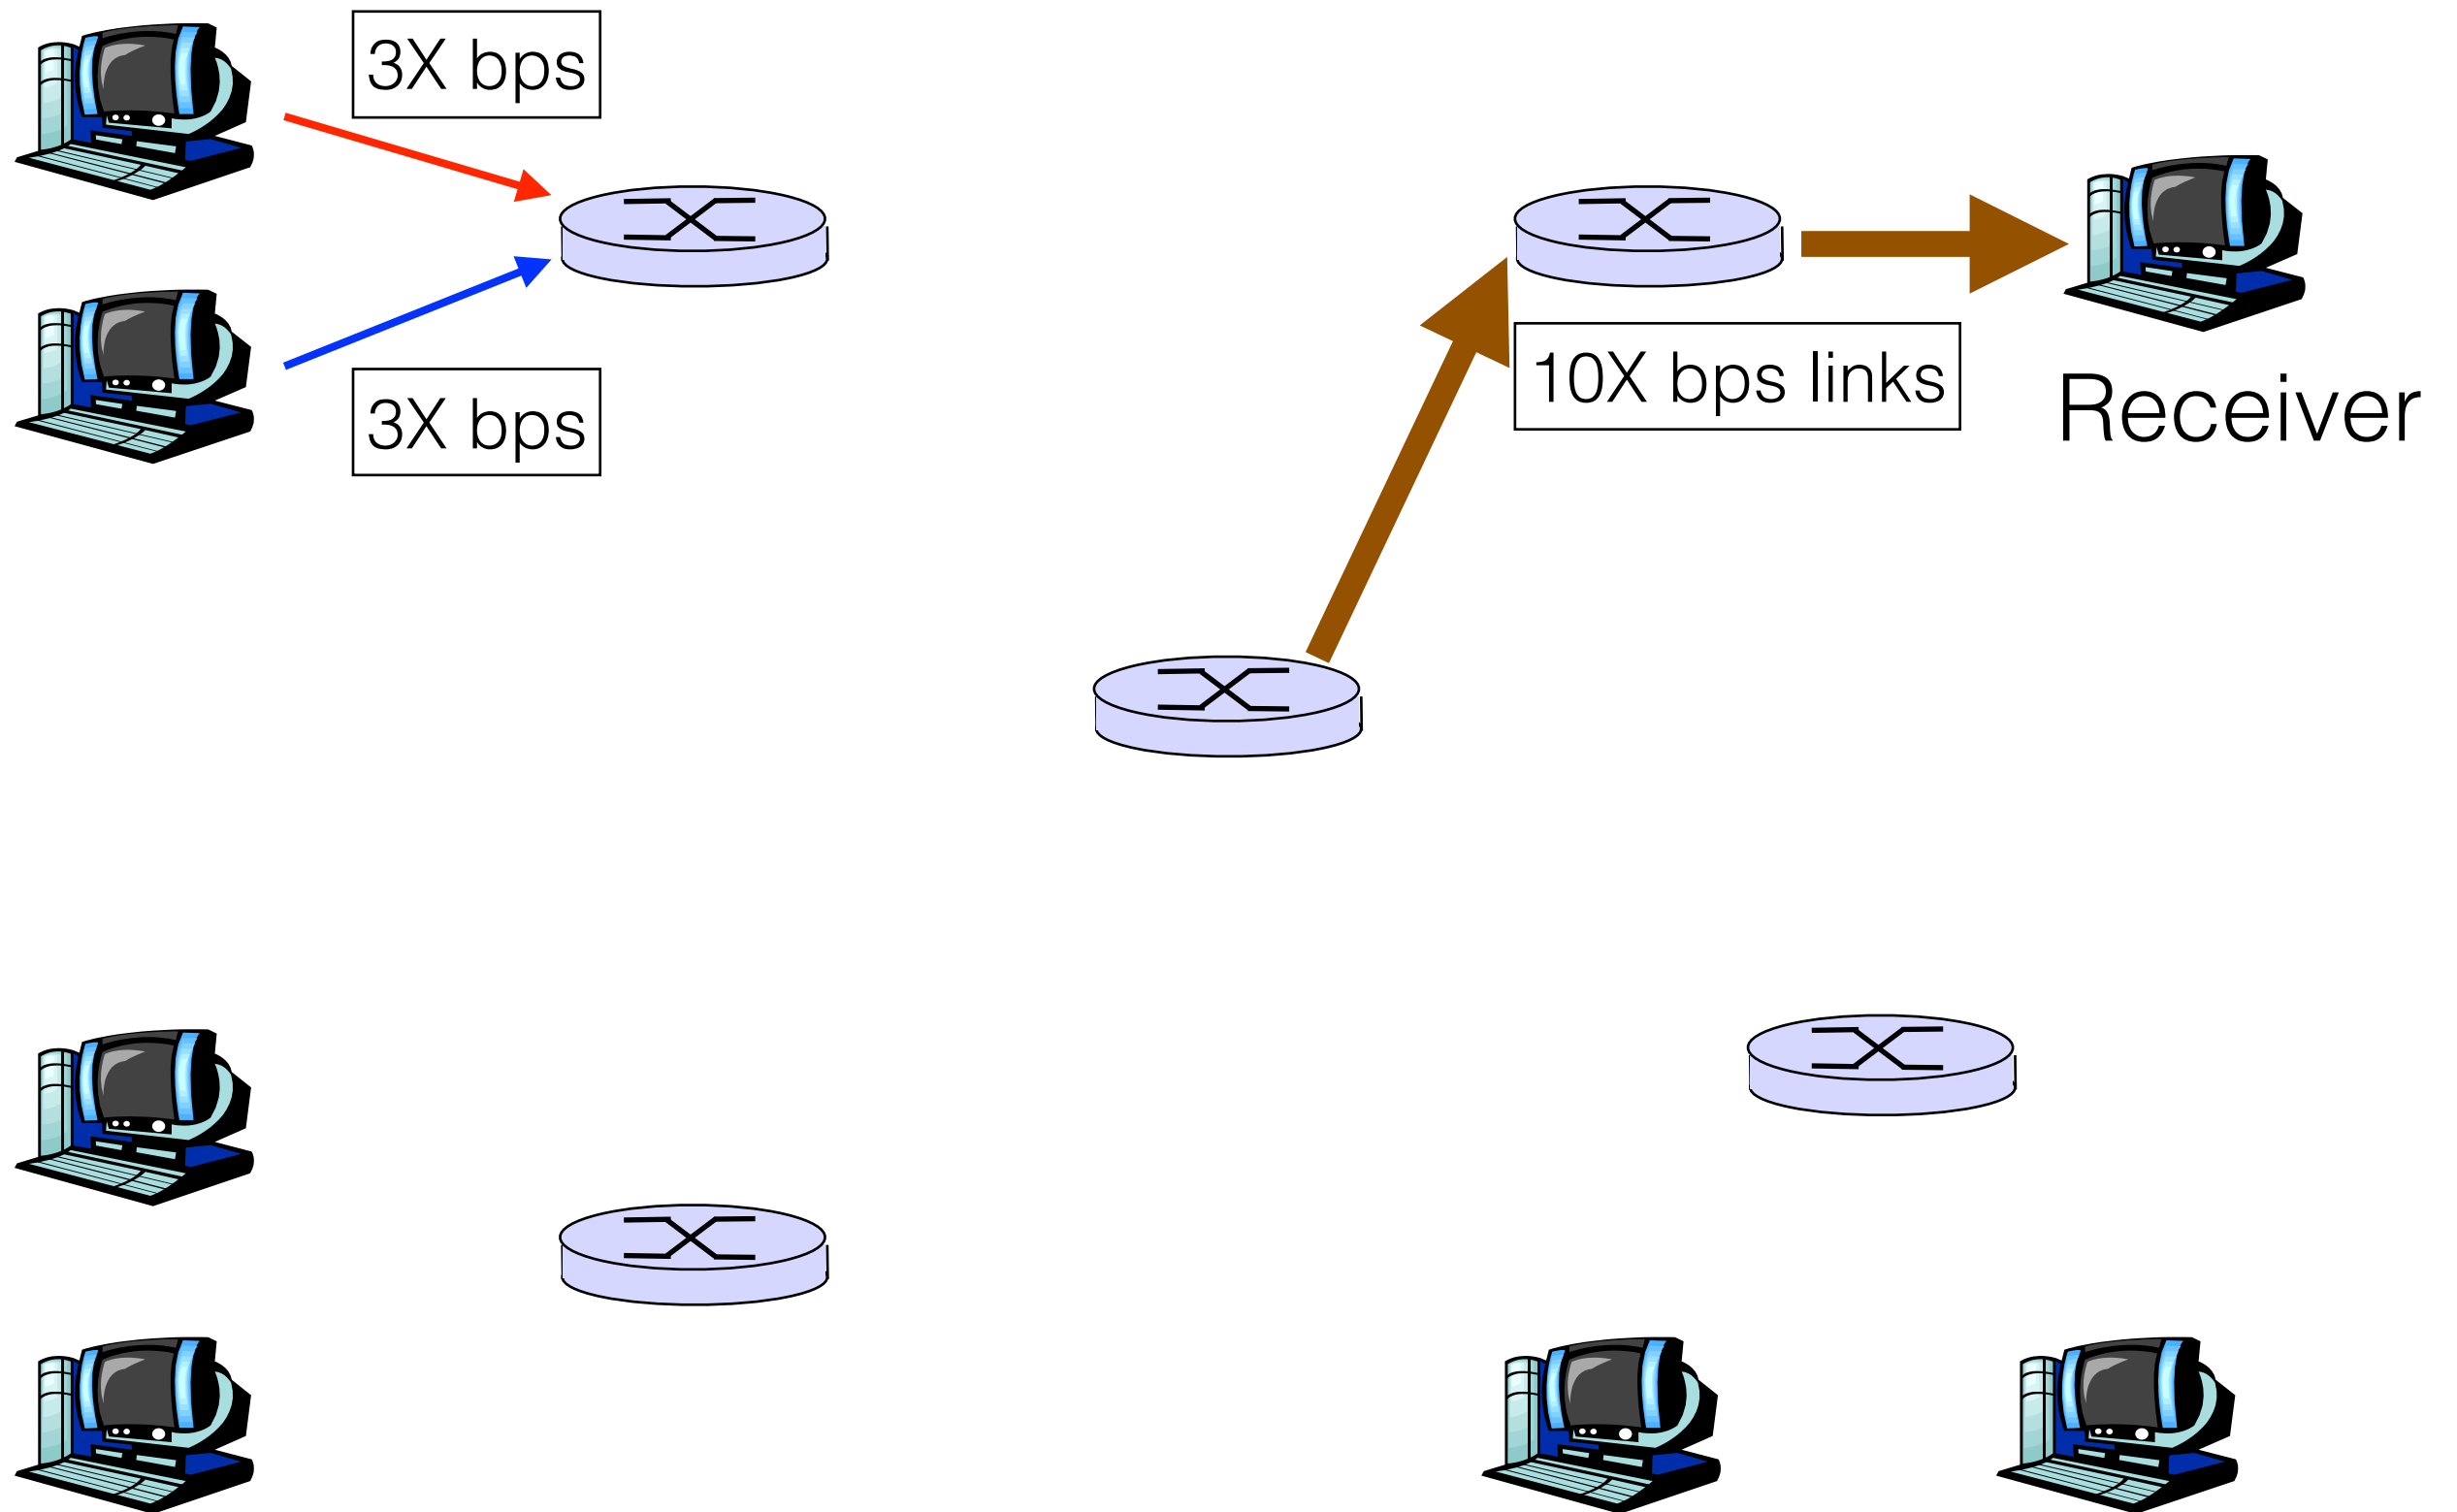




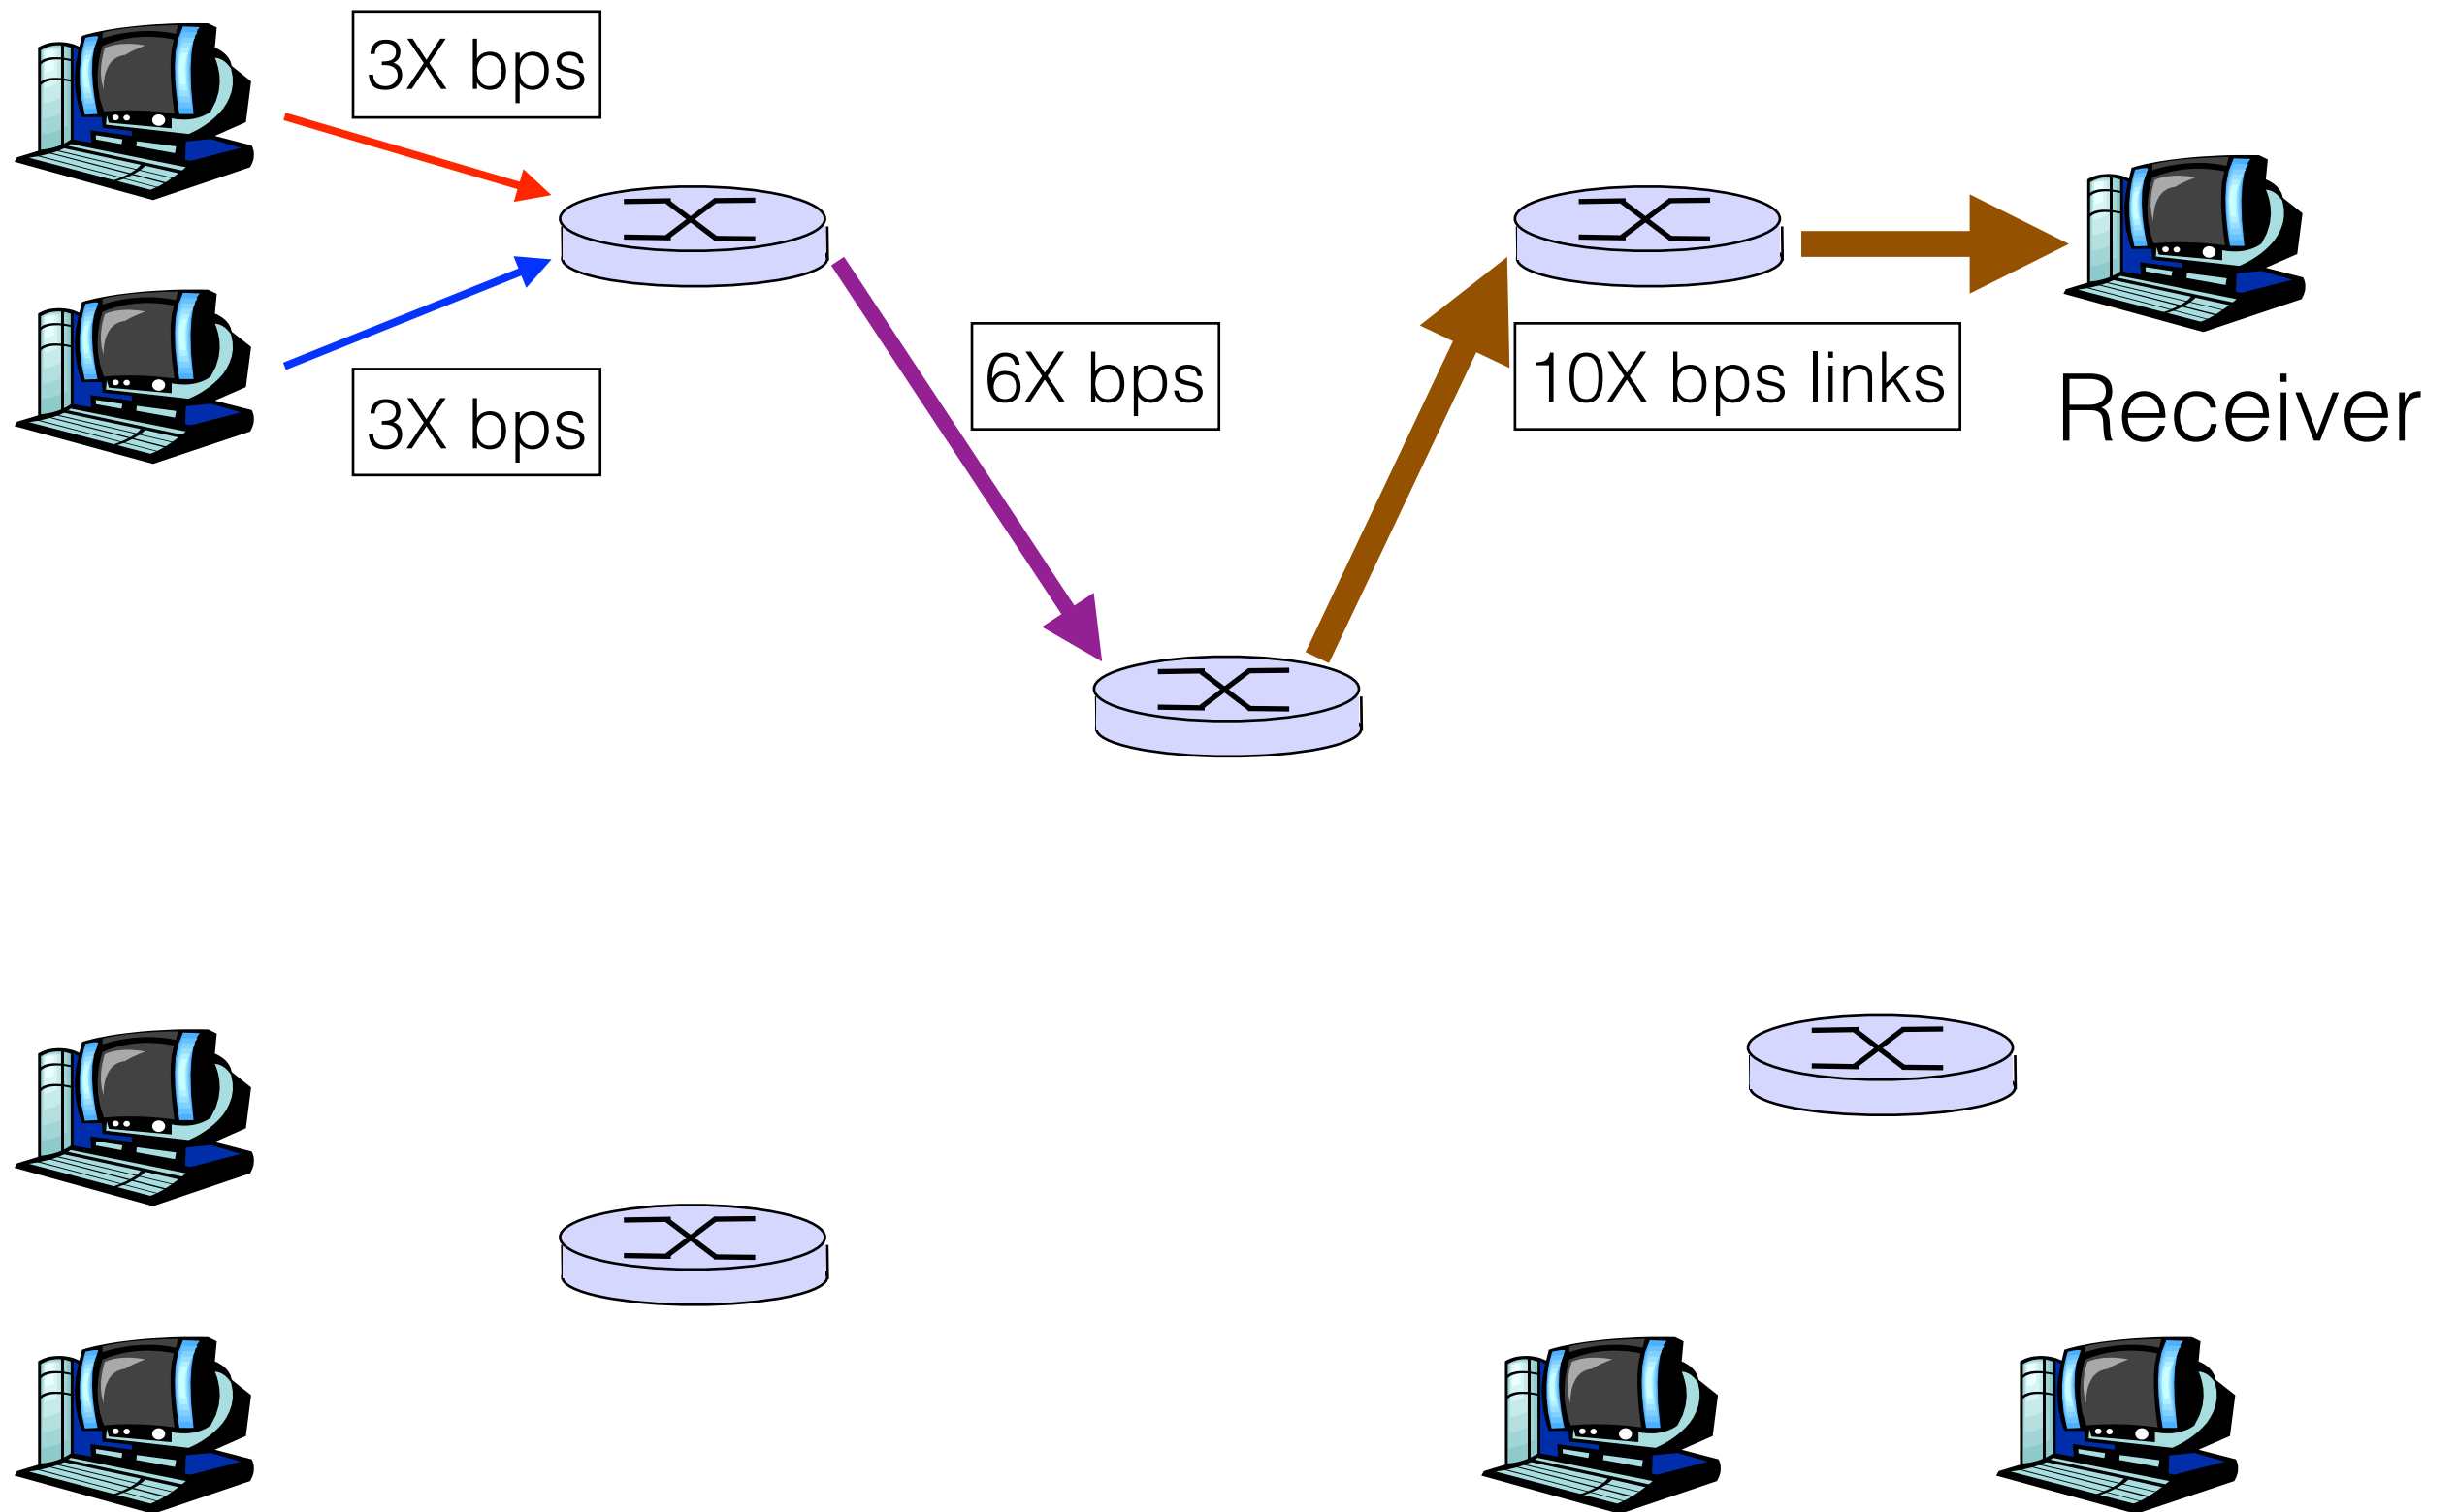
# A Networking Example



# A Networking Example

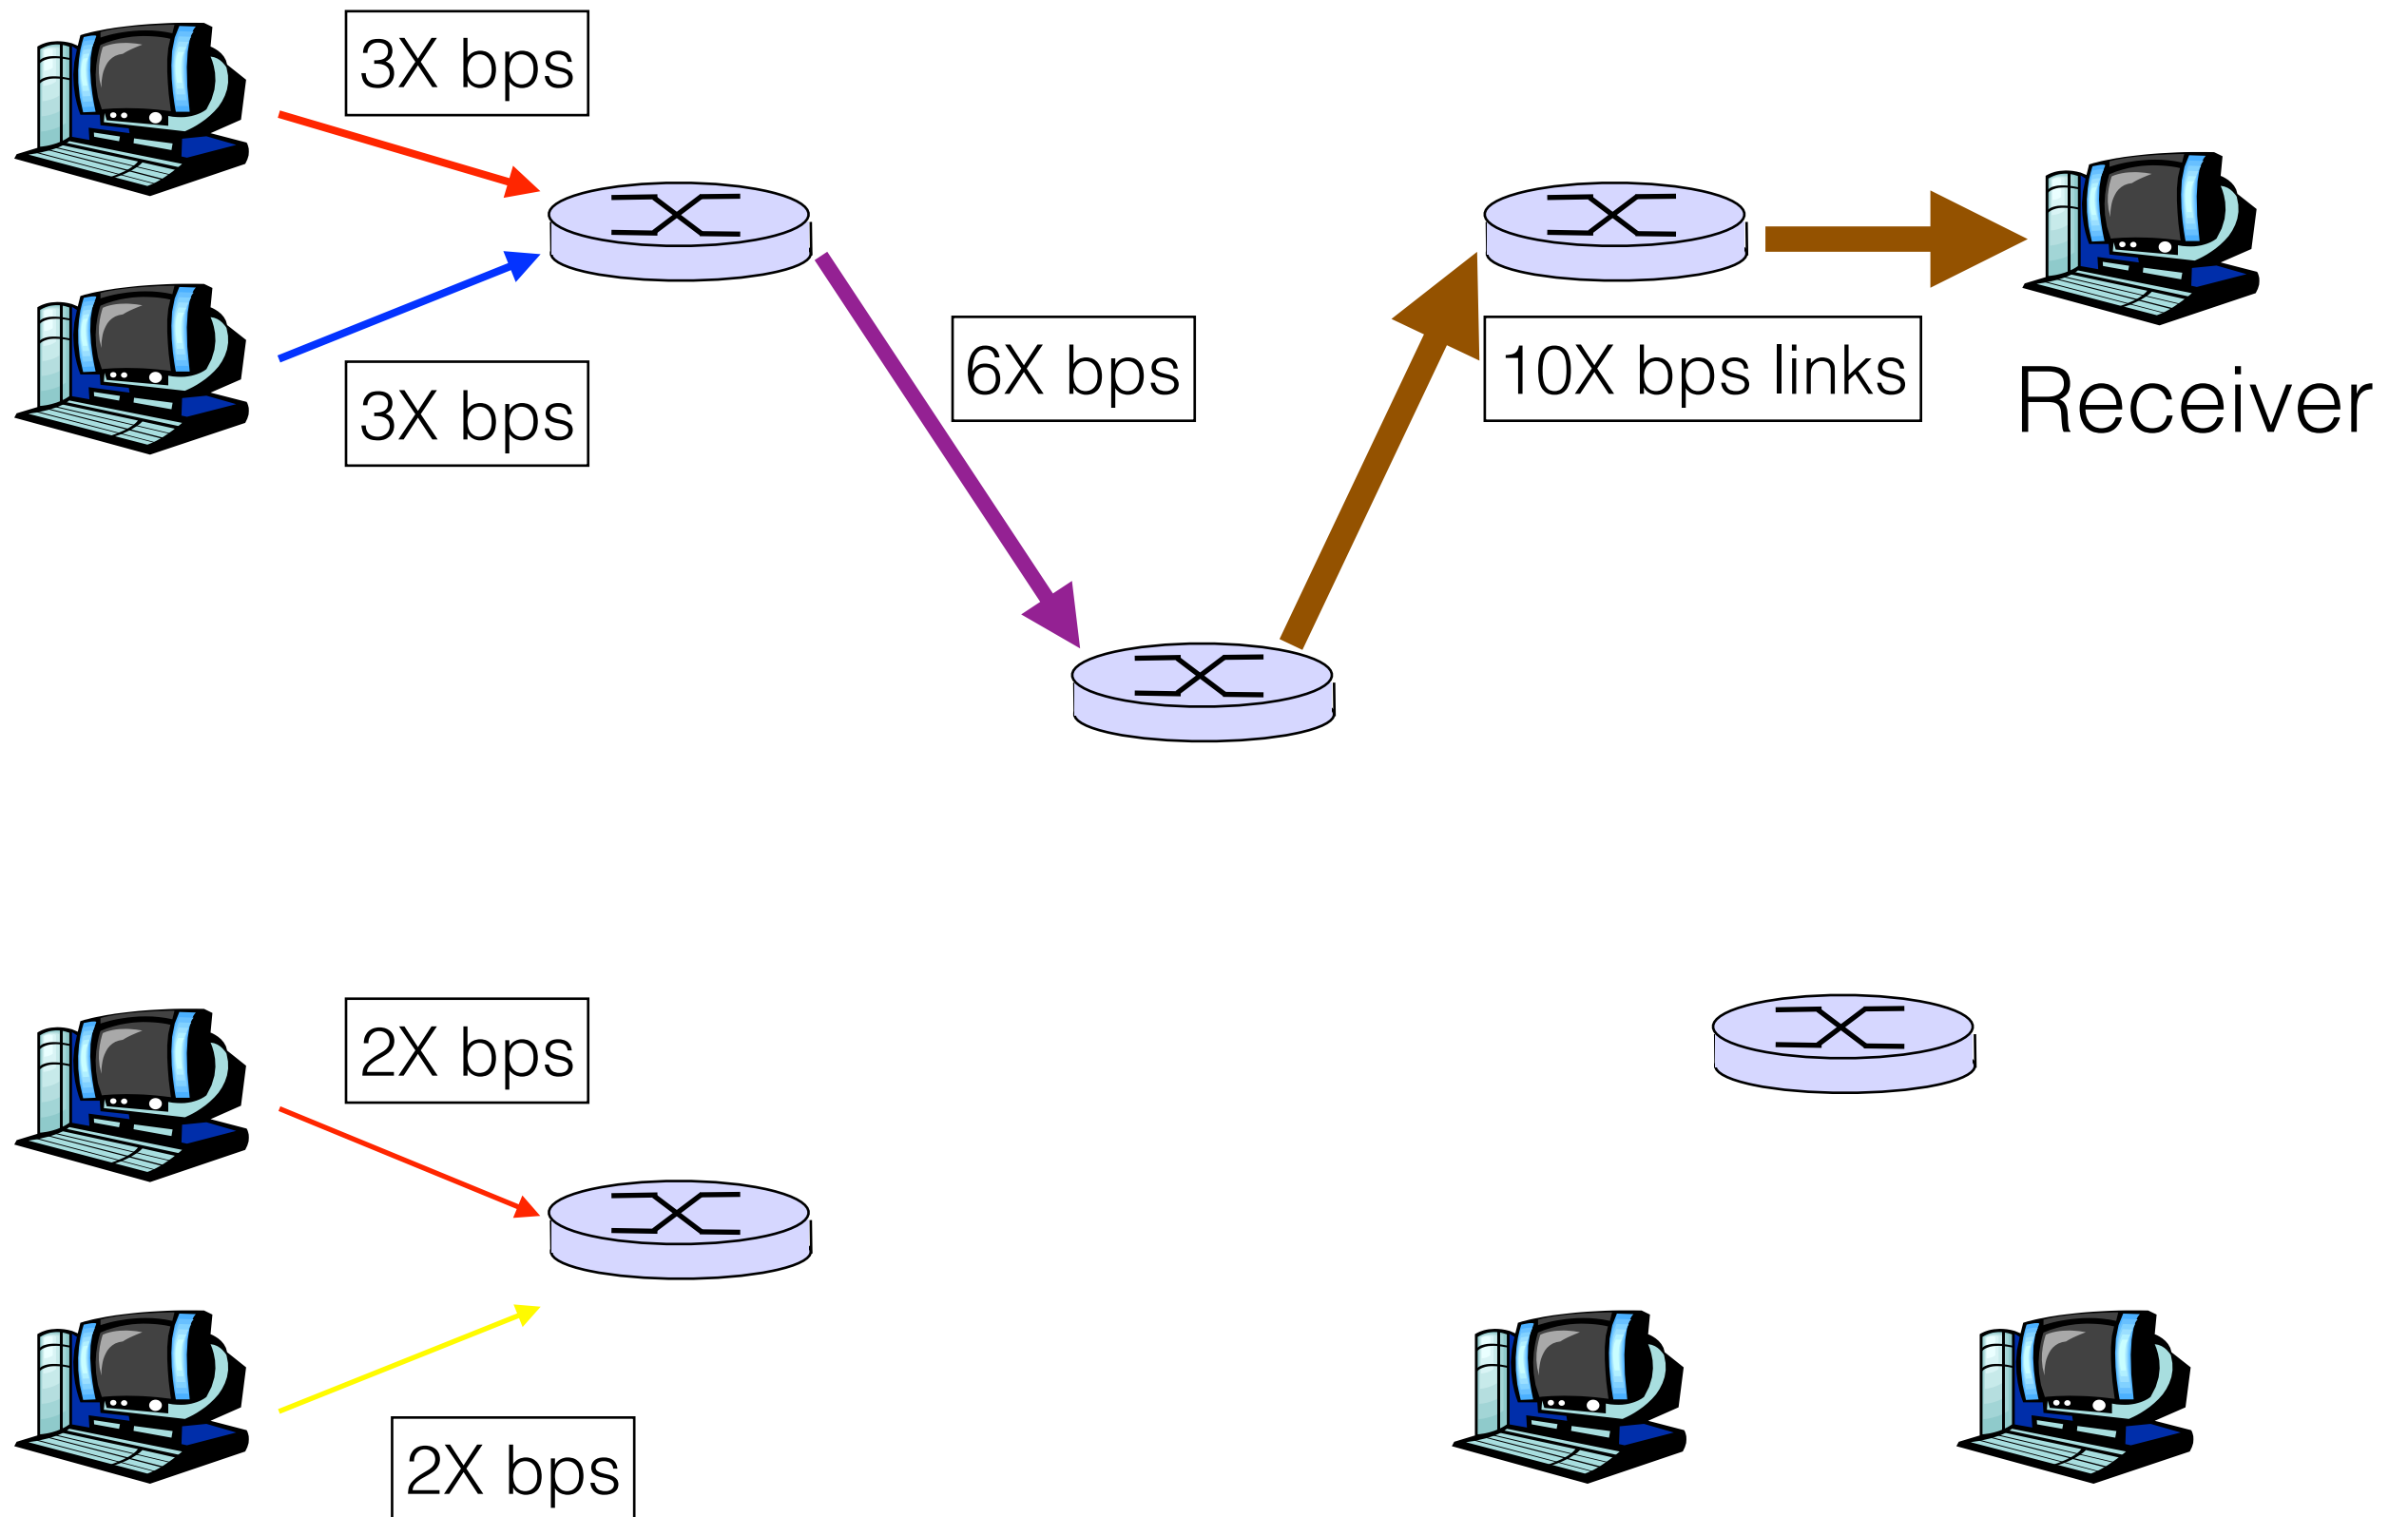


# A Networking Example

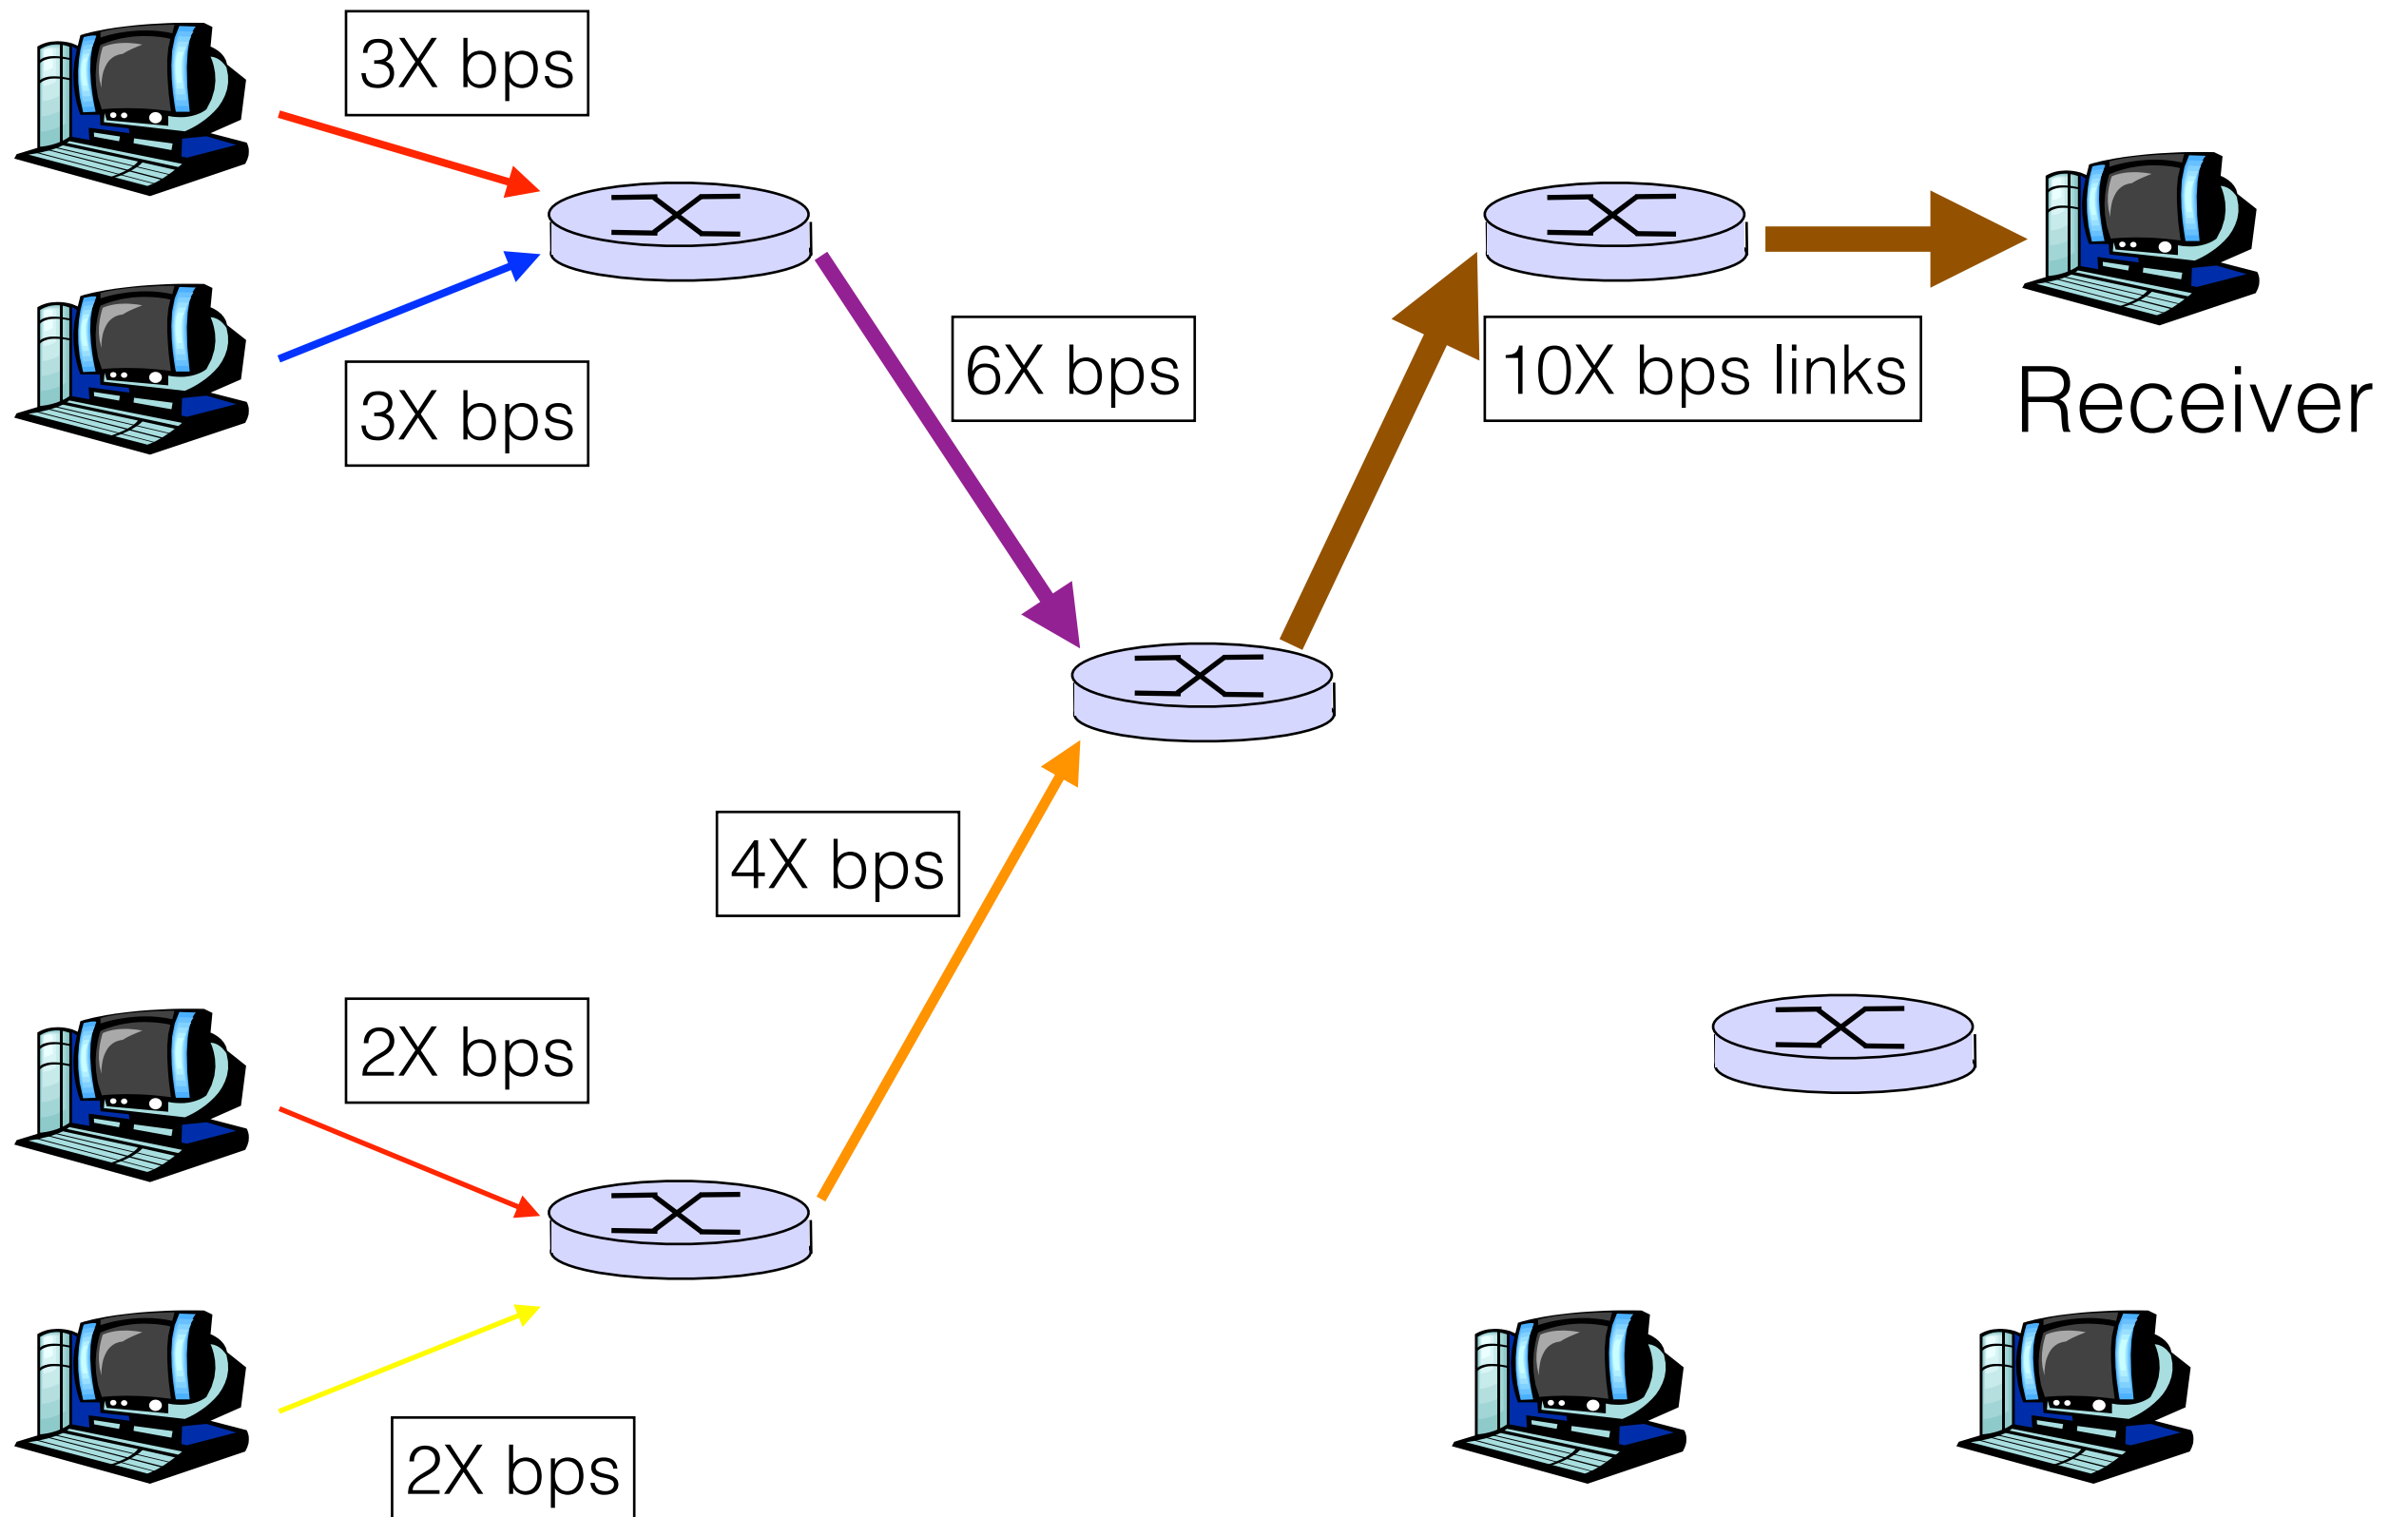




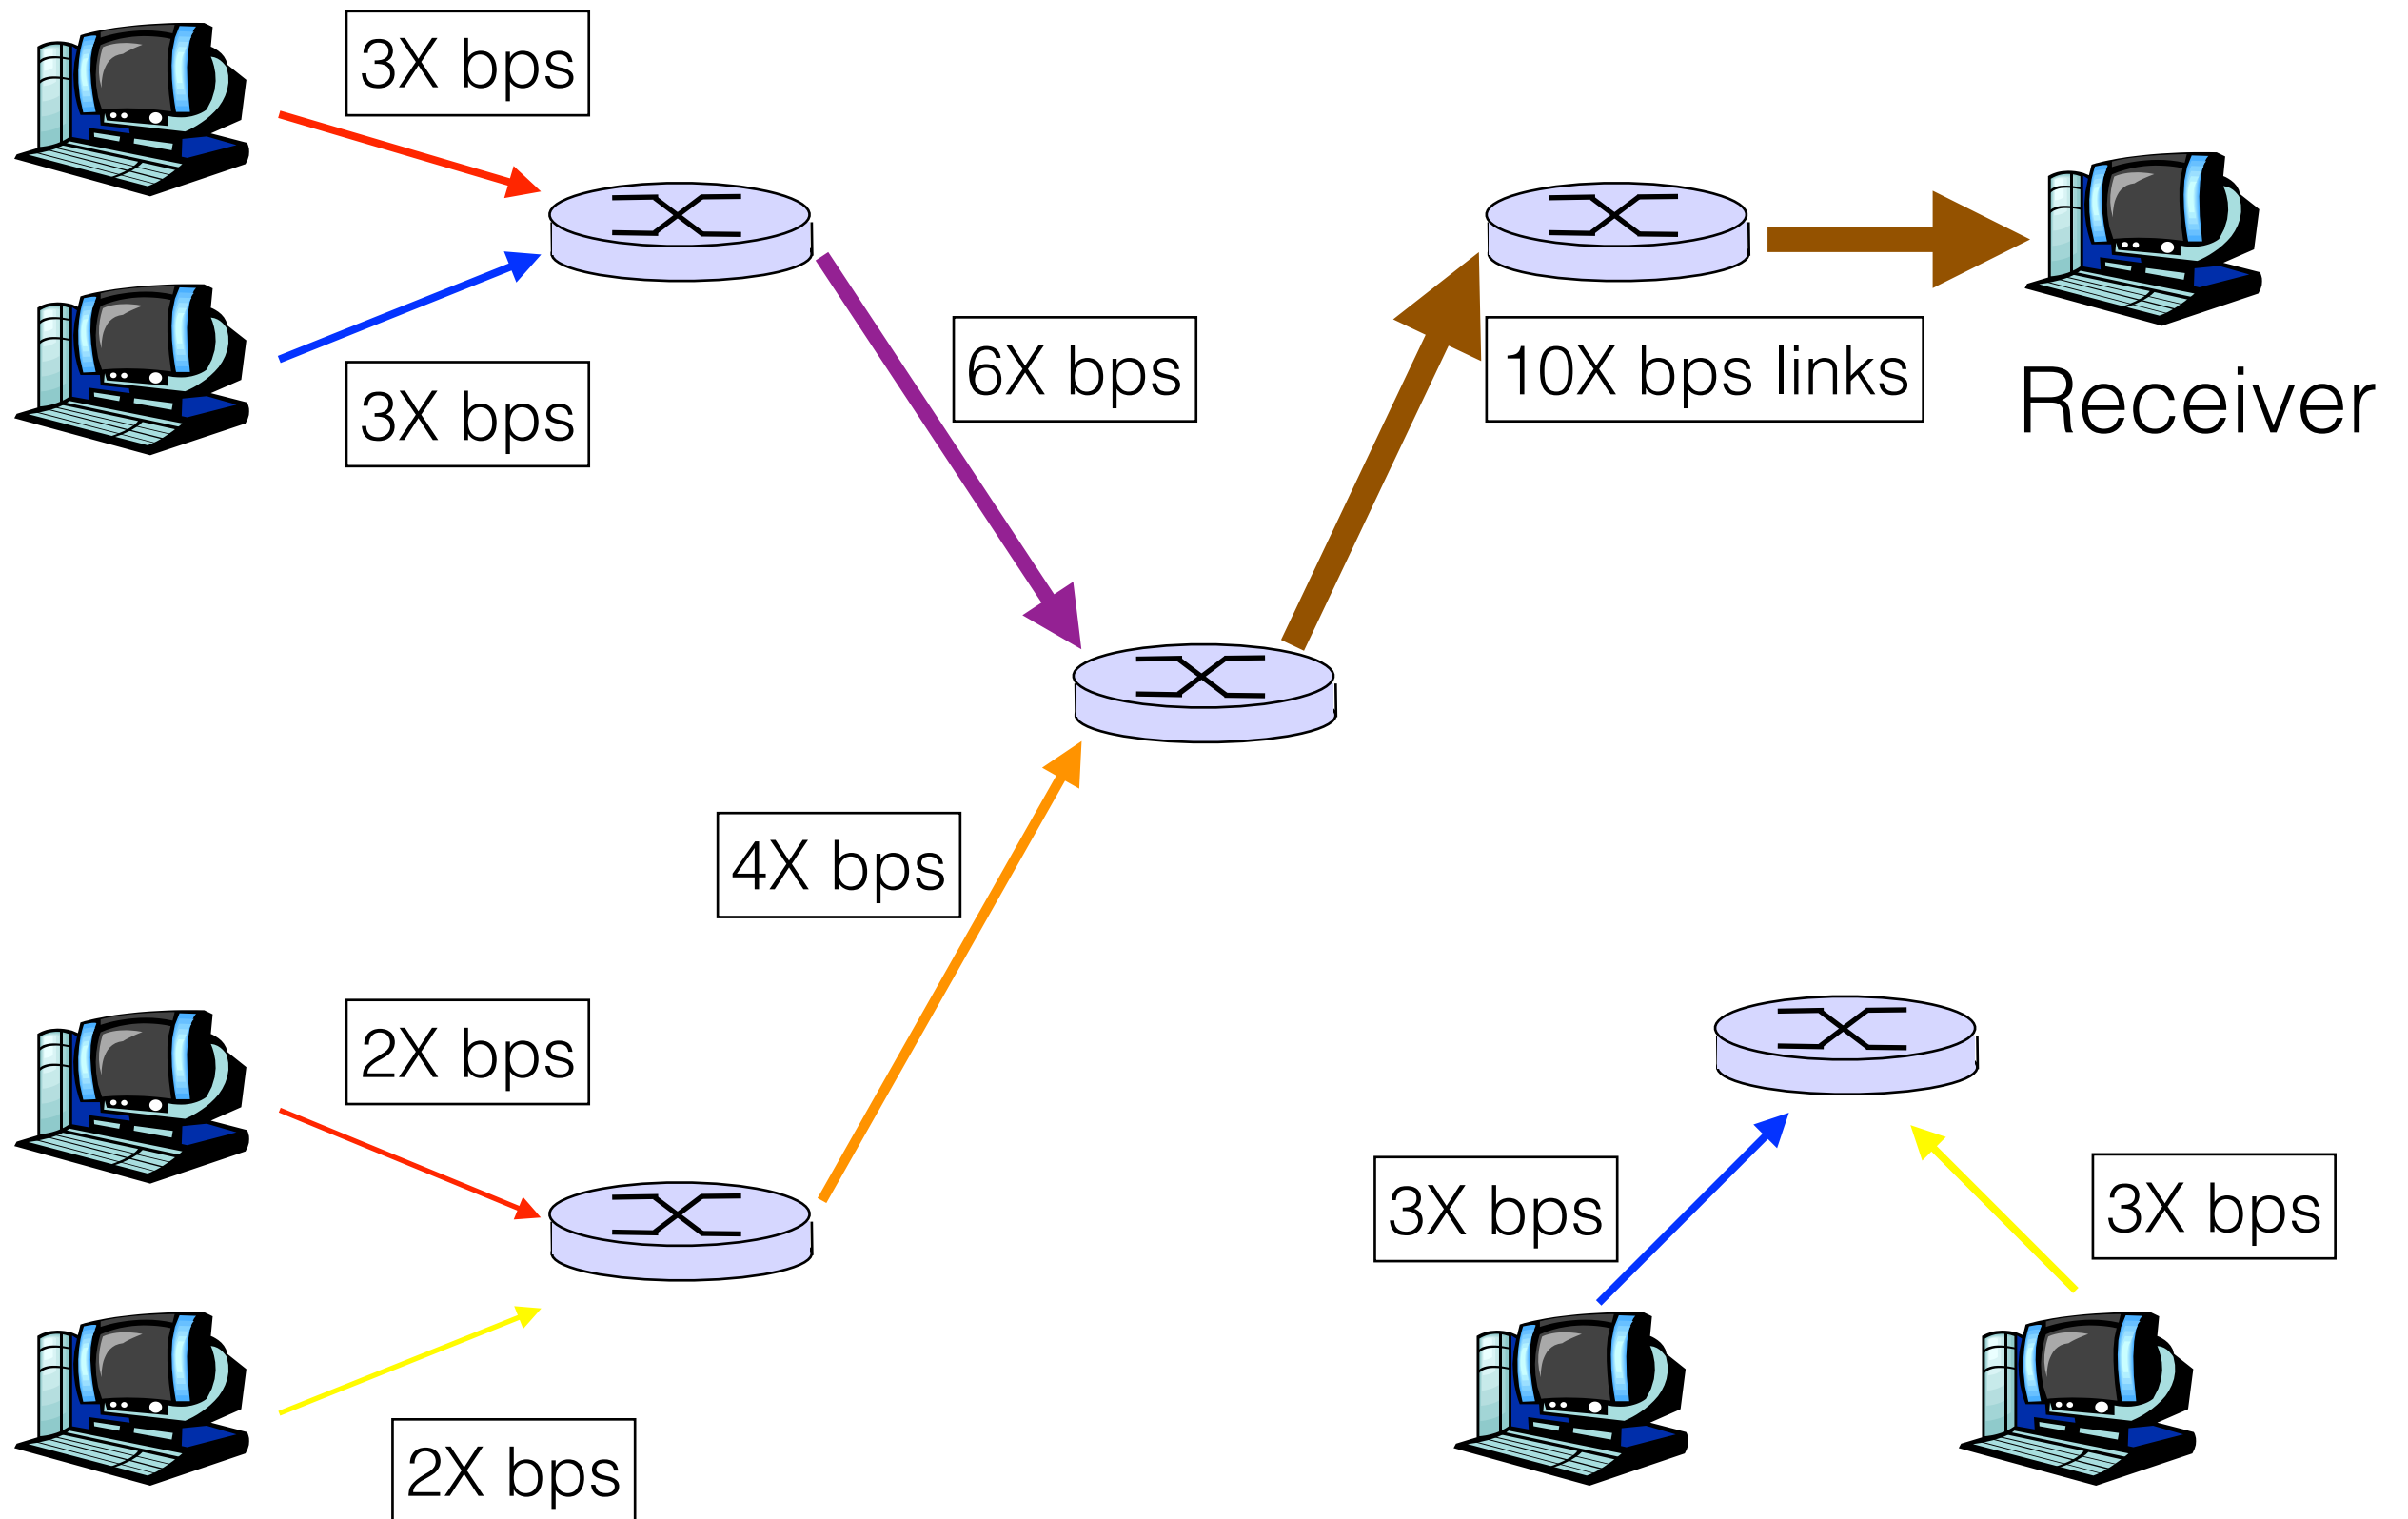
# A Networking Example



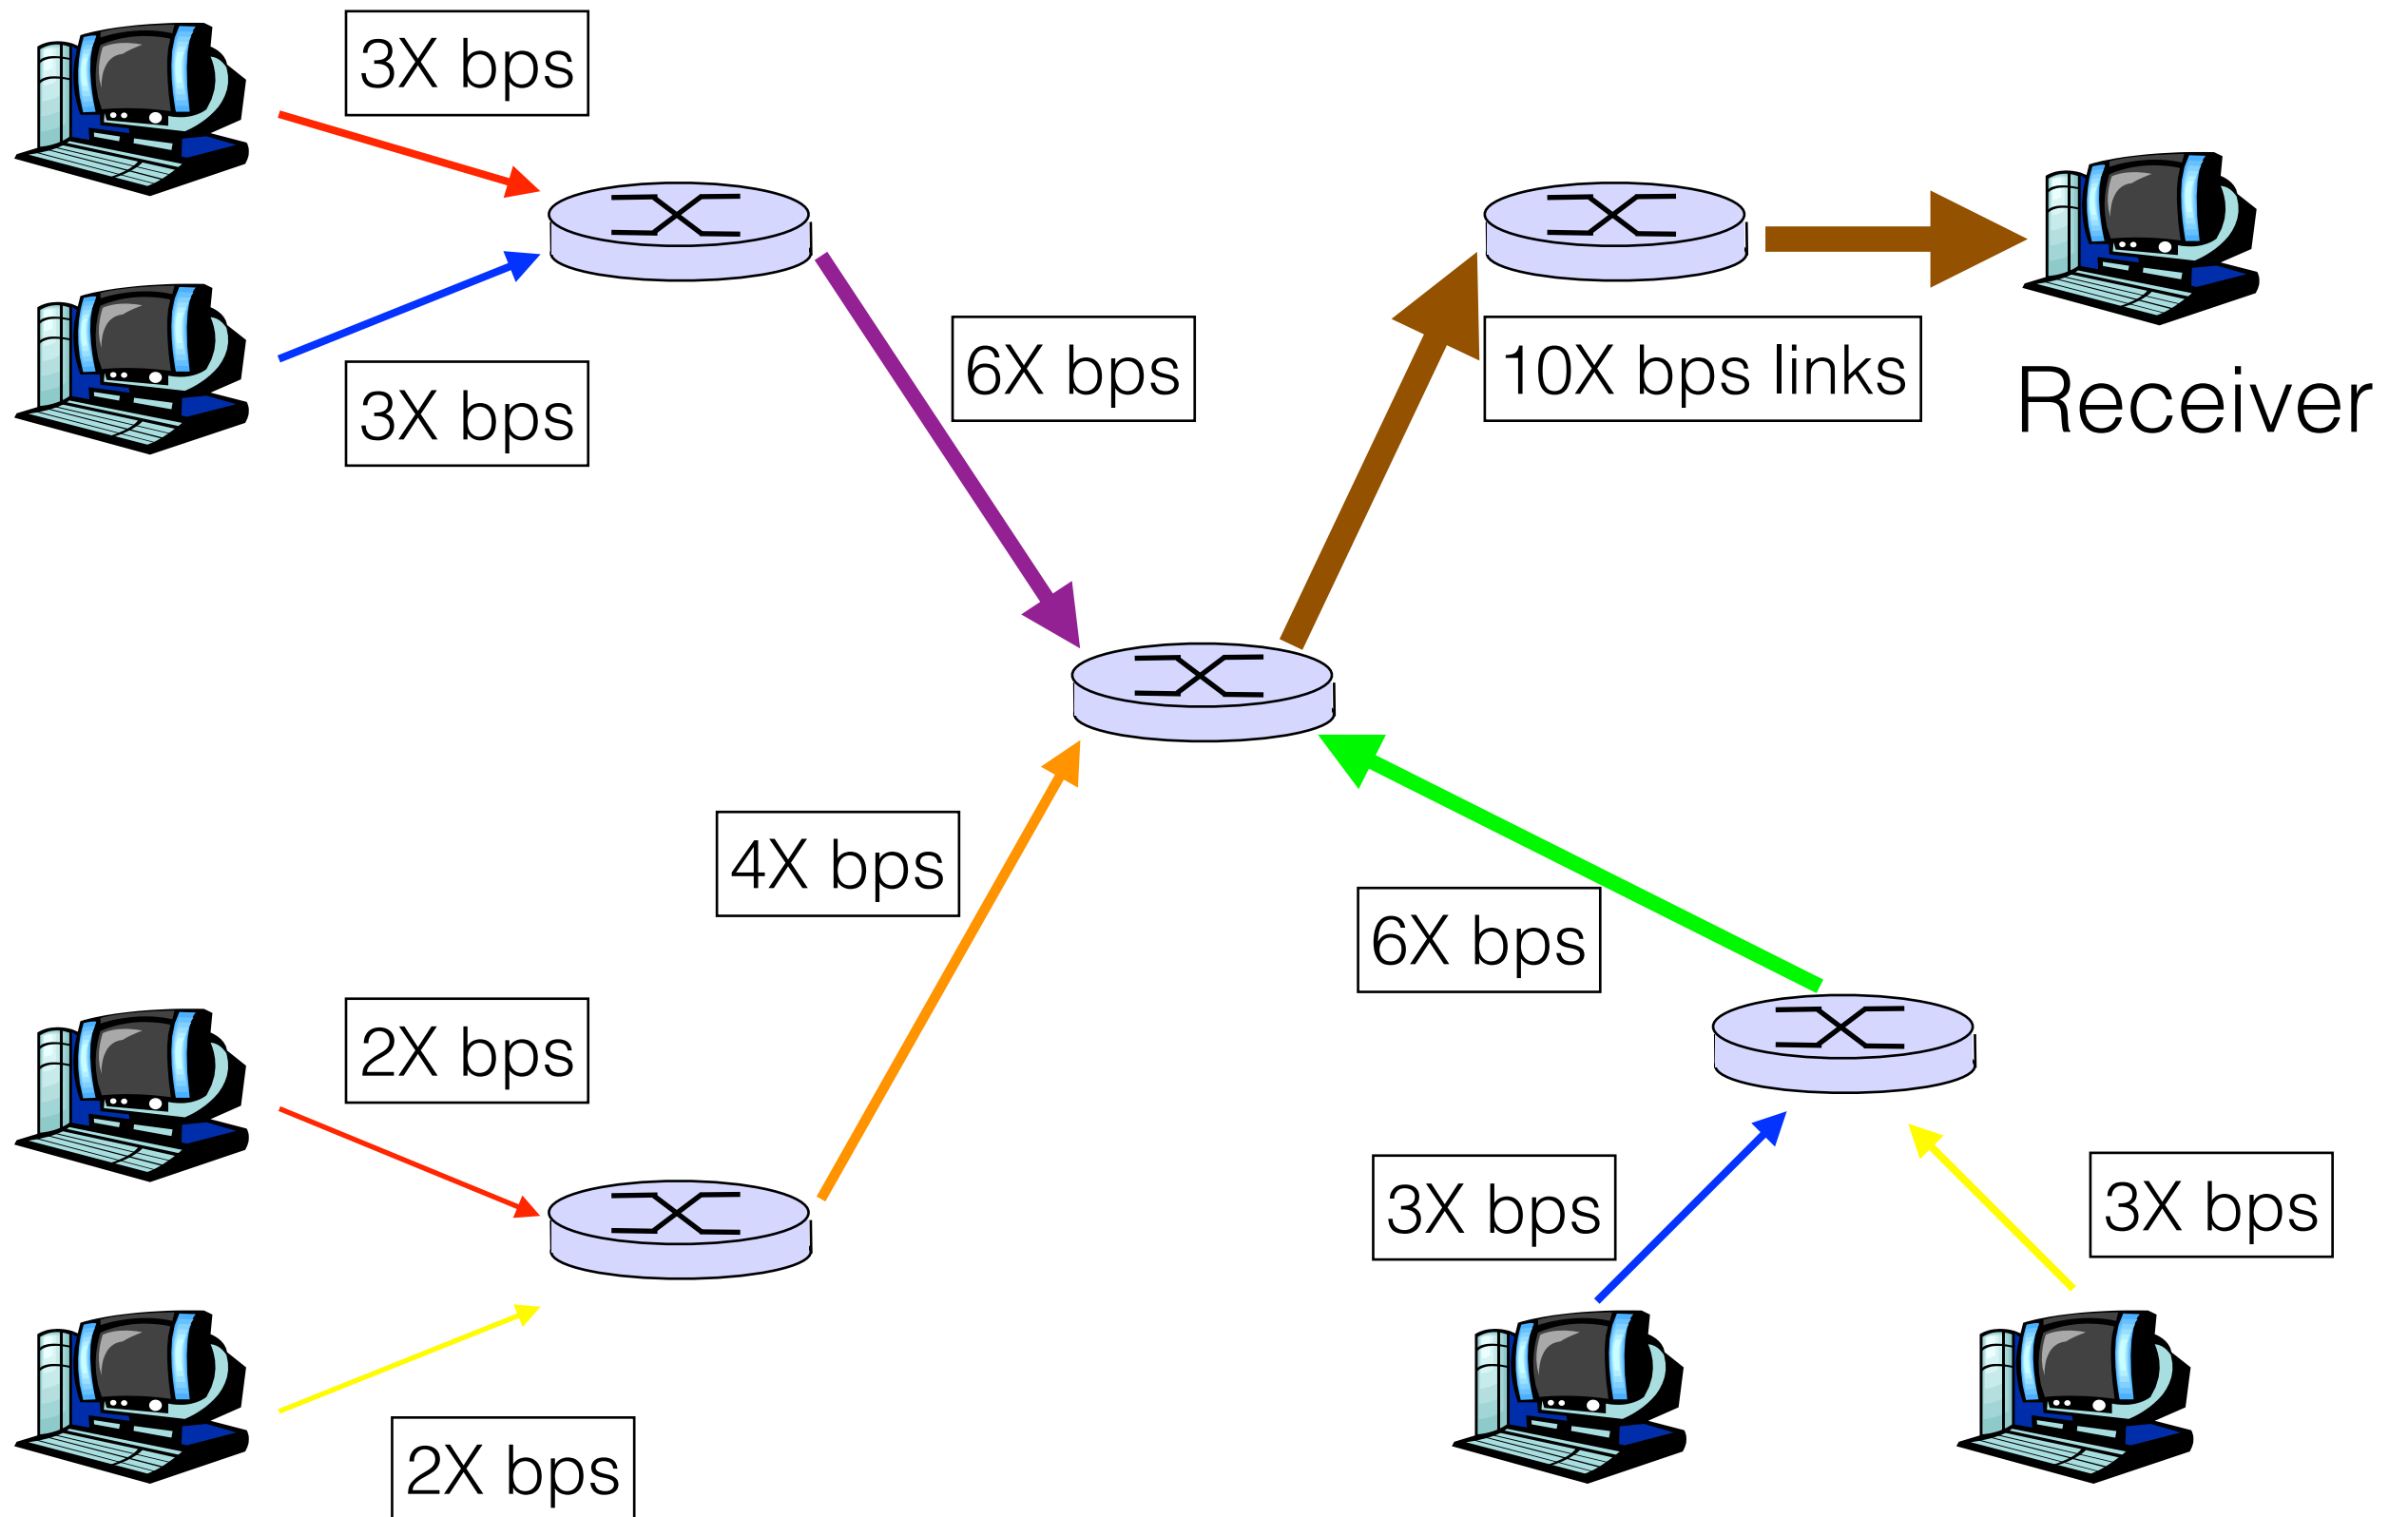
# A Networking Example



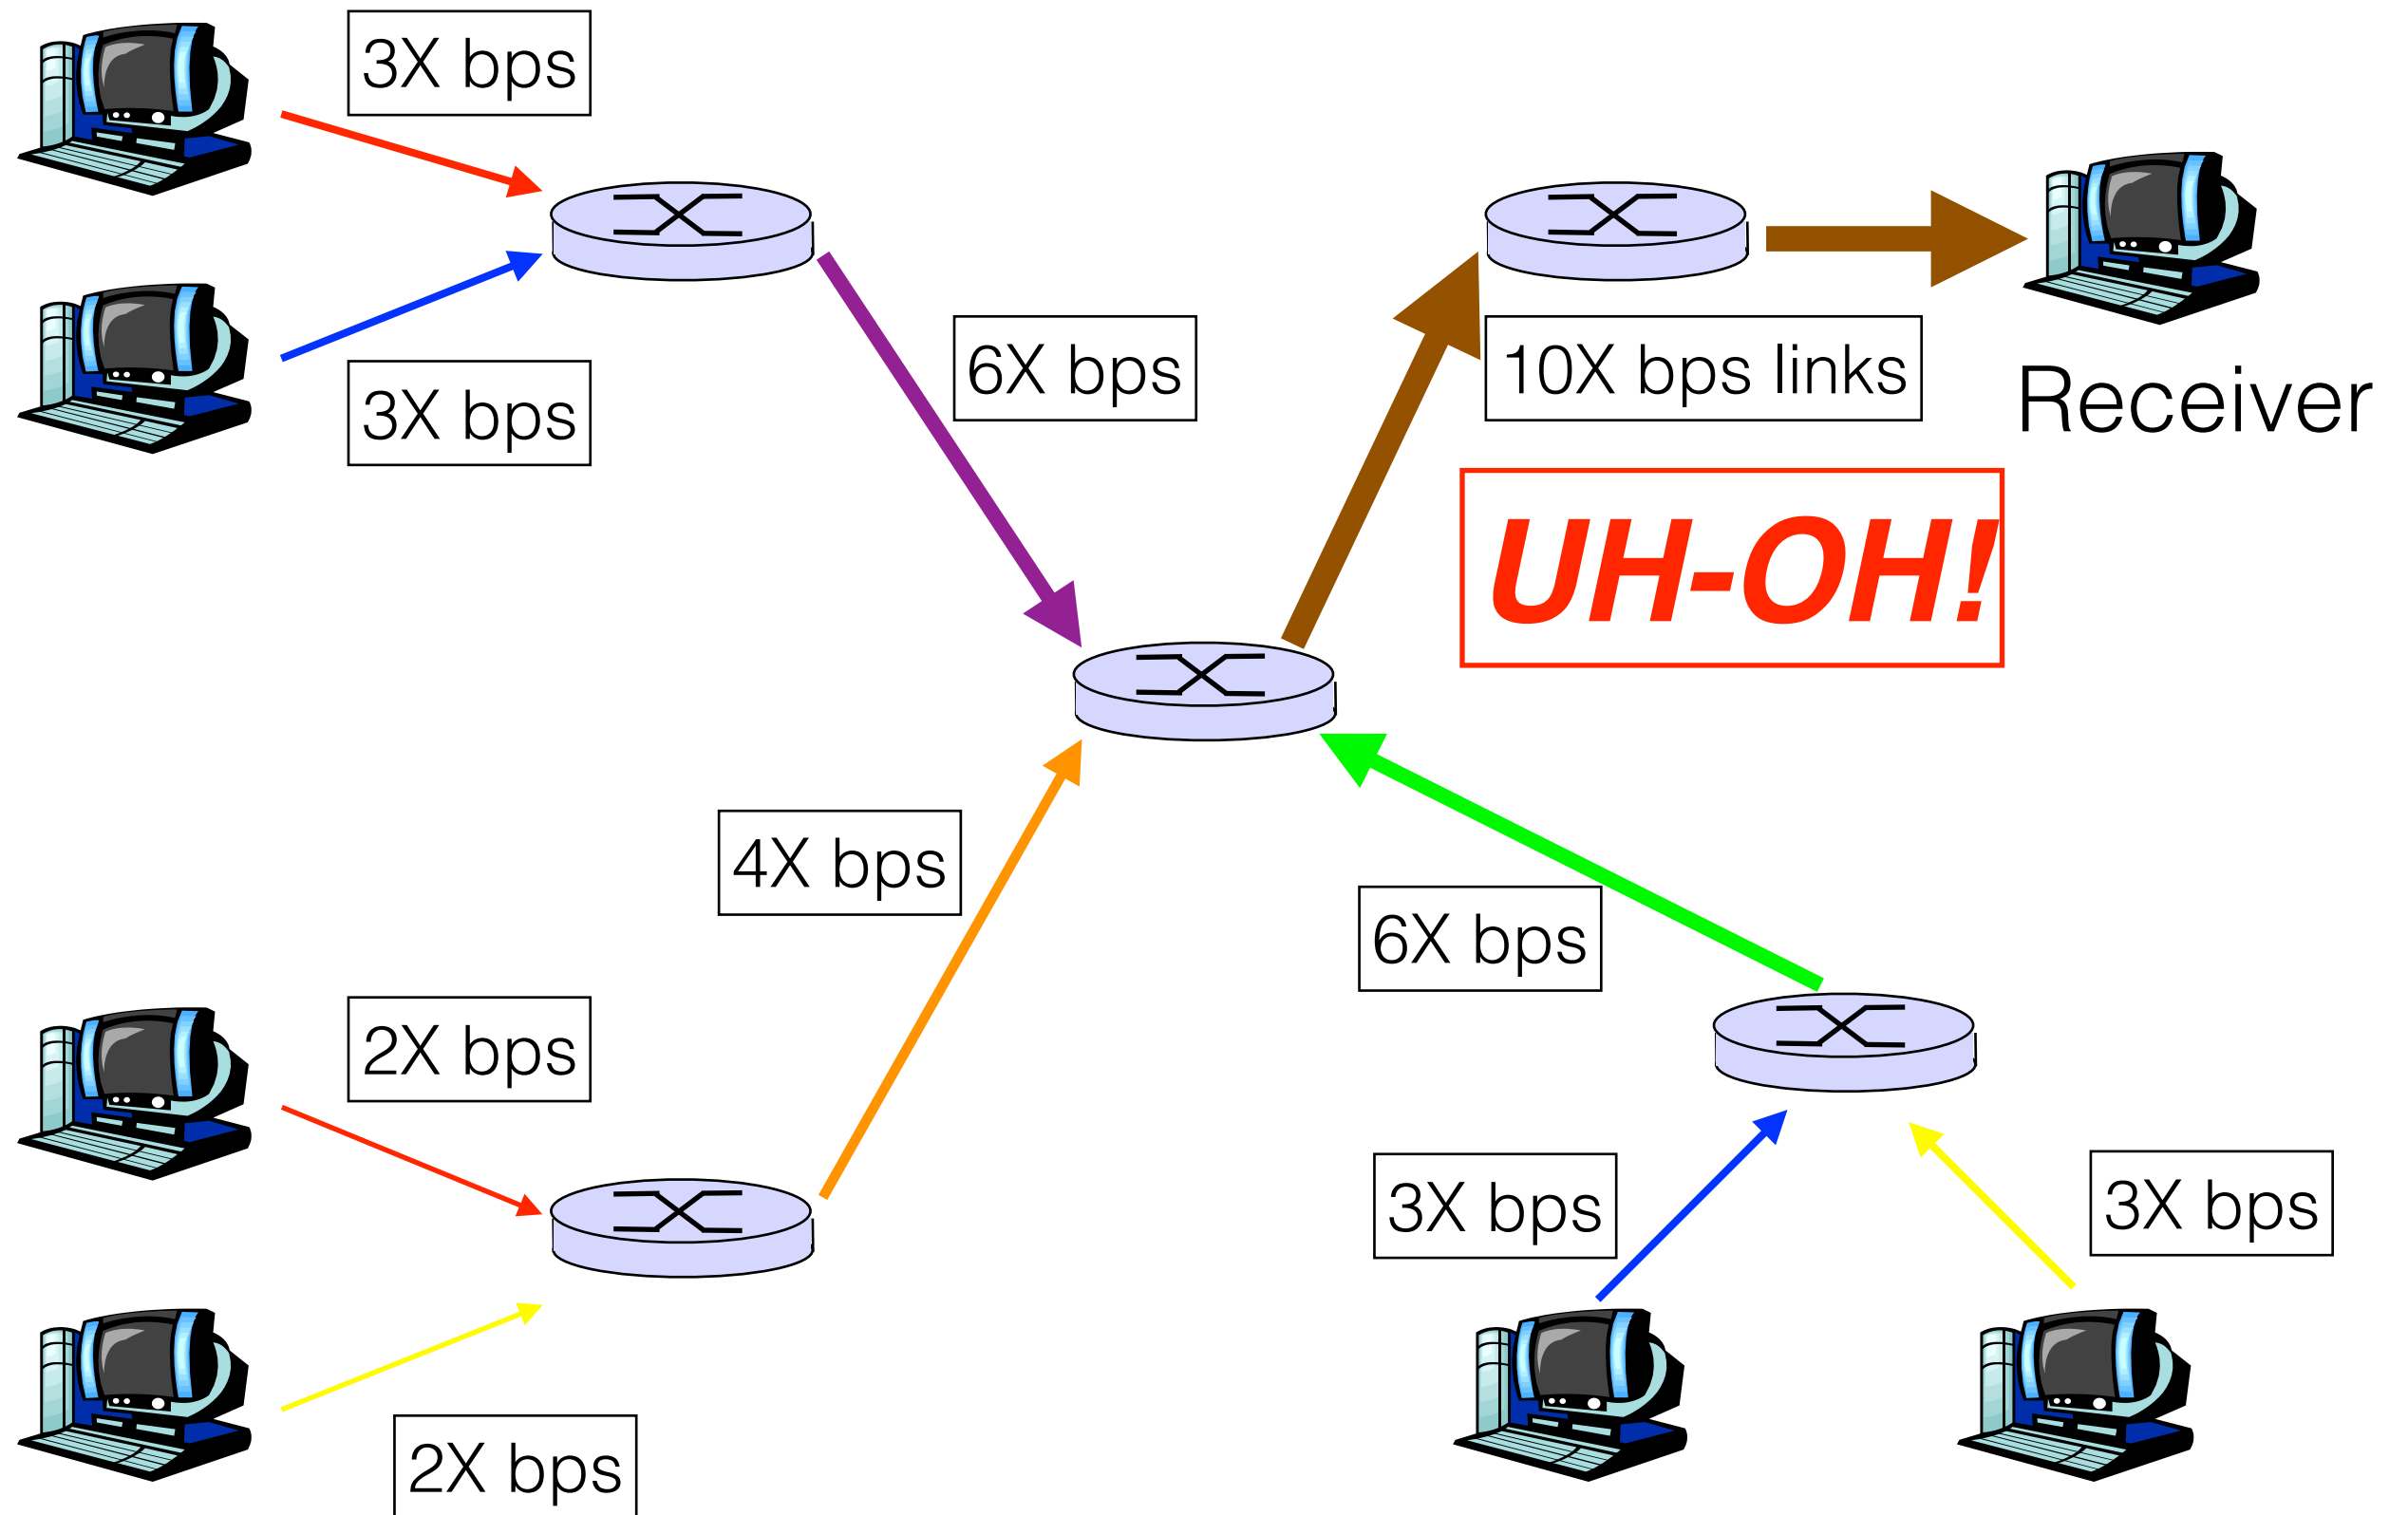
# A Networking Example



# A Networking Example



# A Networking Example



# Principles of Congestion Control

# Principles of Congestion Control

## Congestion:

- ❖ informally: "too many sources sending too much data too fast for **network** to handle"
- ❖ different from flow control!

# Principles of Congestion Control

## Congestion:

- ❖ informally: "too many sources sending too much data too fast for **network** to handle"
- ❖ different from flow control!
- ❖ manifestations:



# Principles of Congestion Control

## Congestion:

- ❖ informally: "too many sources sending too much data too fast for **network** to handle"
- ❖ different from flow control!
- ❖ manifestations:
  - lost packets (buffer overflow at routers)
  - long delays (queueing in router buffers)

# Principles of Congestion Control

## Congestion:

- ❖ informally: "too many sources sending too much data too fast for **network** to handle"
- ❖ different from flow control!
- ❖ manifestations:
  - lost packets (buffer overflow at routers)
  - long delays (queueing in router buffers)
- ❖ a top-10 problem!

# Costs of Congestion

# Costs of Congestion

- ❖ more work for a given sending rate
  - i.e., more retransmissions

# Costs of Congestion

- ❖ more work for a given sending rate
  - i.e., more retransmissions
- ❖ retransmissions take network resources

# Costs of Congestion

- ❖ more work for a given sending rate
  - i.e., more retransmissions
- ❖ retransmissions take network resources
- ❖ less timely data delivery

# Costs of Congestion


- ❖ more work for a given sending rate
  - i.e., more retransmissions
- ❖ retransmissions take network resources
- ❖ less timely data delivery
- ❖ but, are these annoyance or fundamental issues?!

# Costs of Congestion

- ❖ more work for a given sending rate
  - i.e., more retransmissions
- ❖ retransmissions take network resources
- ❖ less timely data delivery
- ❖ but, are these annoyance or fundamental issues?
  - left unchecked congestion can lead to “collapse” whereby a network expends scarce resources on packets that will ultimately be dropped
    - i.e., accomplish little—if any—meaningful work




# Congestion: Single Bottleneck




All xmit  
@ 10pps

# Congestion: Single Bottleneck




All xmit  
@ 10pps

# Congestion: Multiple Bottlenecks




All xmit  
@ 10pps

# Congestion: Multiple Bottlenecks




All xmit  
@ 10pps

# Congestion: Multiple Bottlenecks



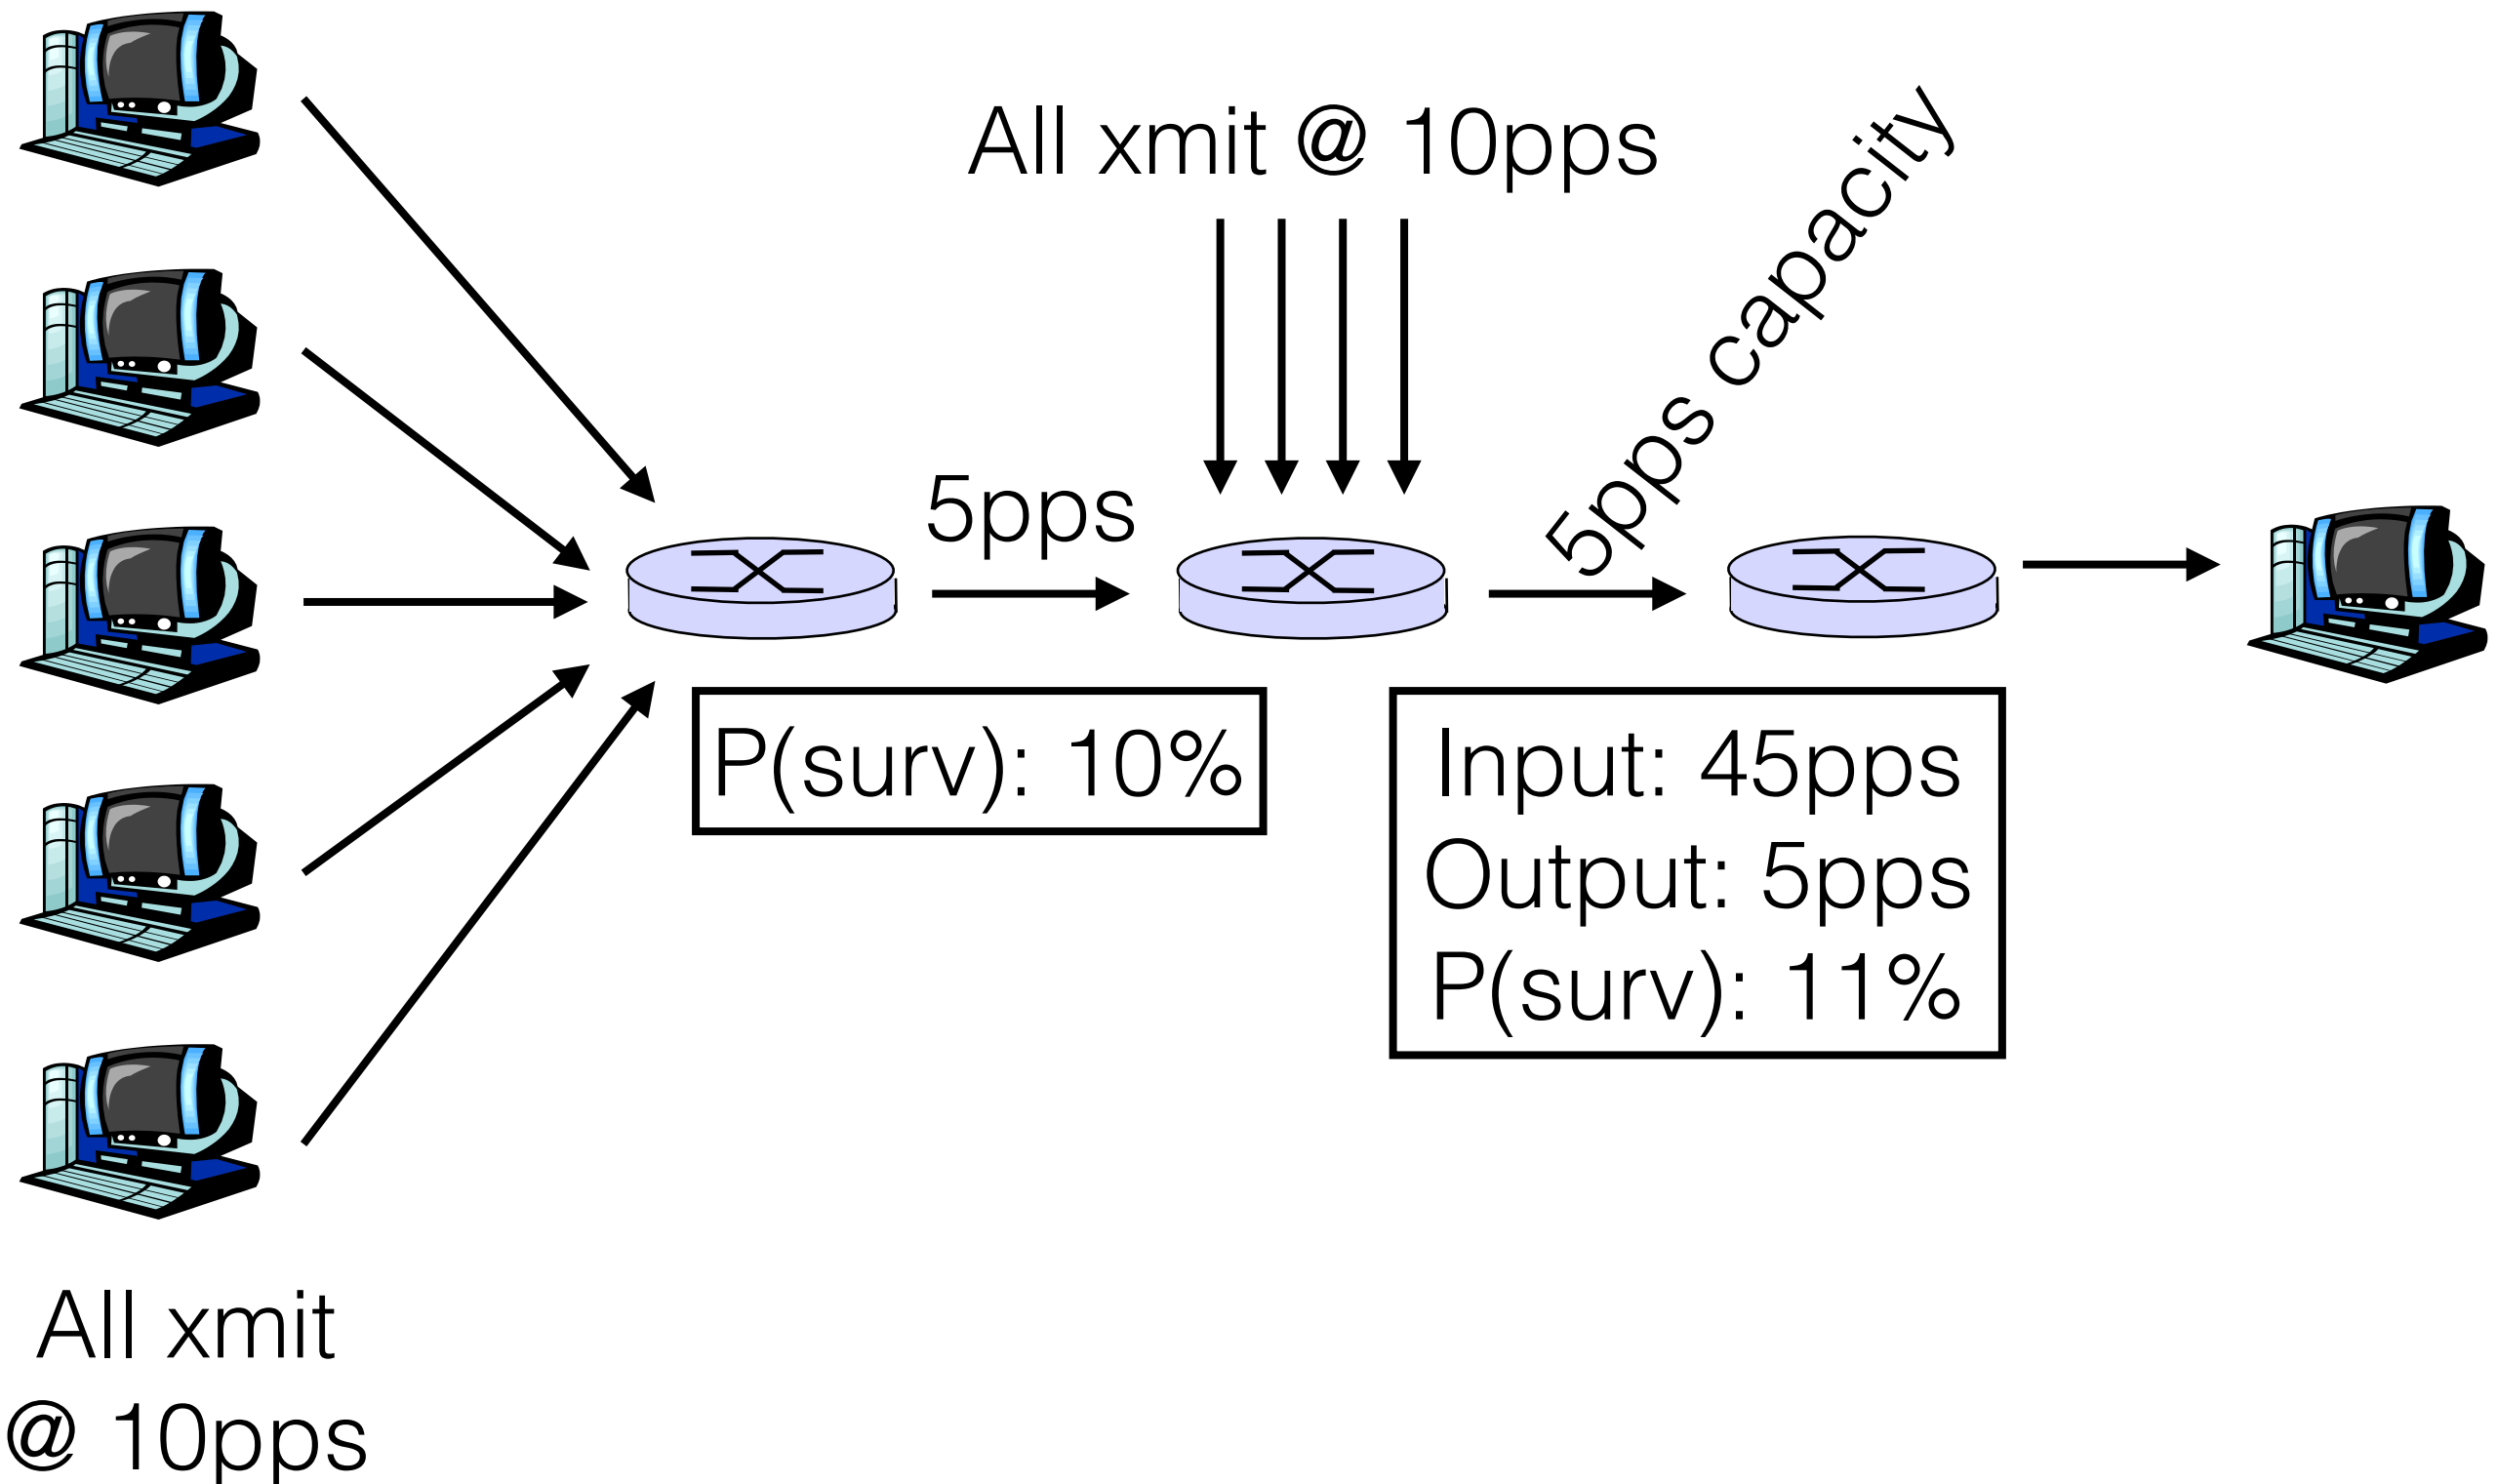
All xmit  
@ 10pps

# Congestion: Multiple Bottlenecks




All xmit  
@ 10pps

# Congestion: Multiple Bottlenecks




# Congestion: Multiple Bottlenecks




All xmit  
@ 10pps




# Congestion: Multiple Bottlenecks



# Congestion: Multiple Bottlenecks

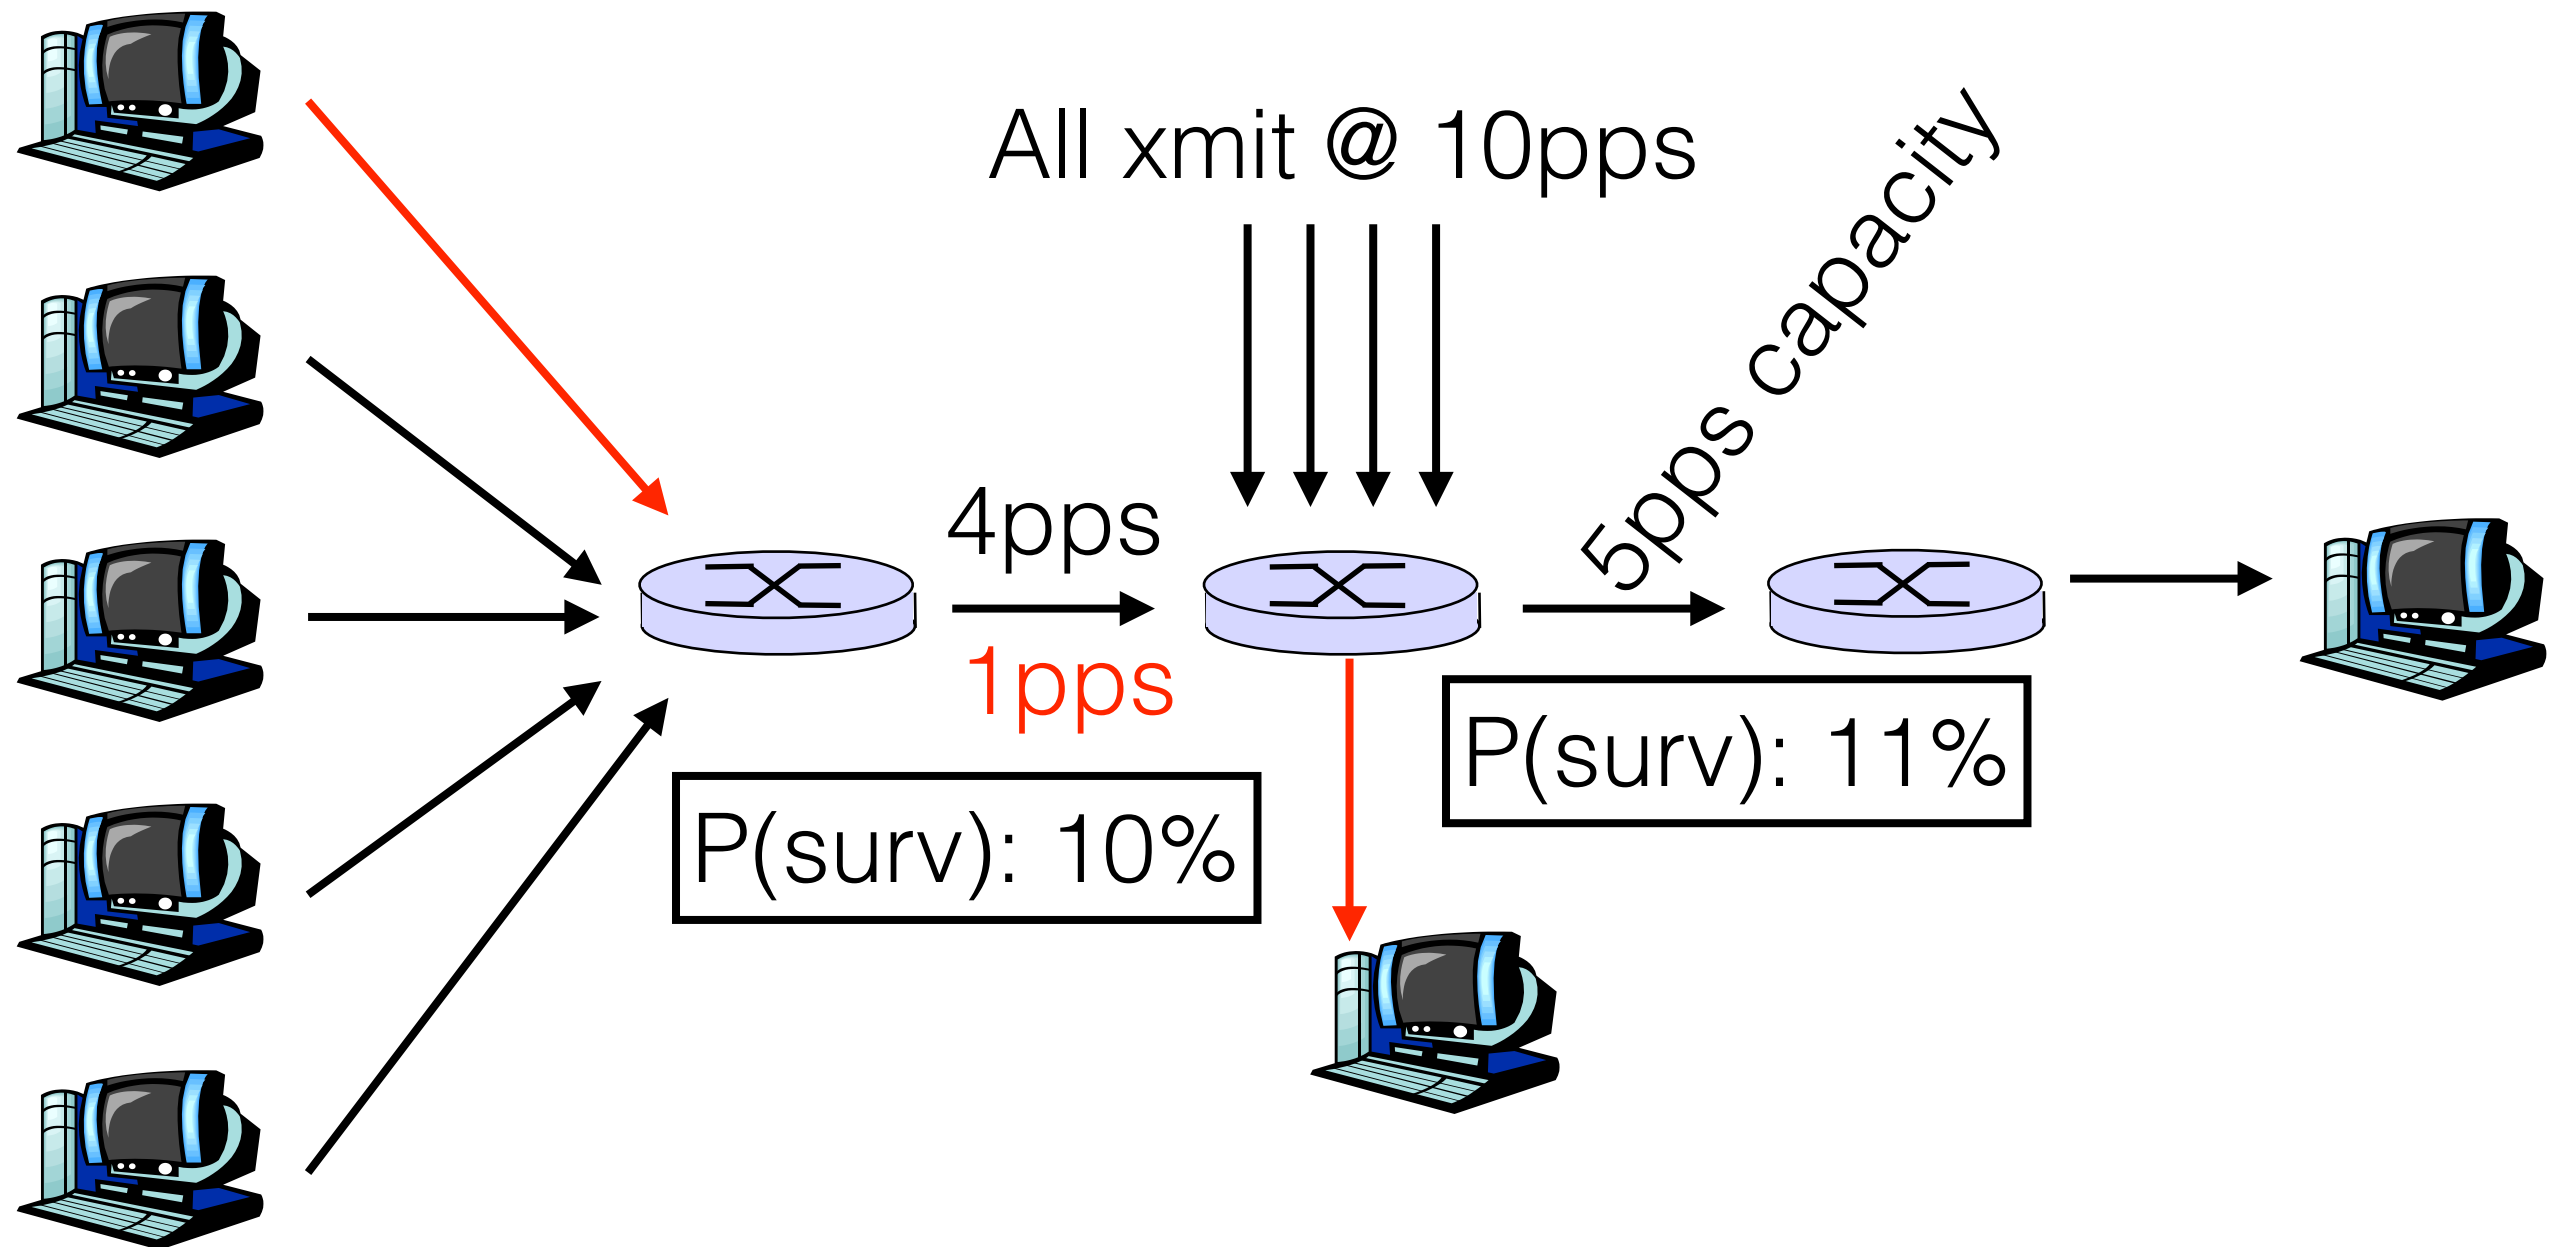


# Congestion: Multiple Bottlenecks



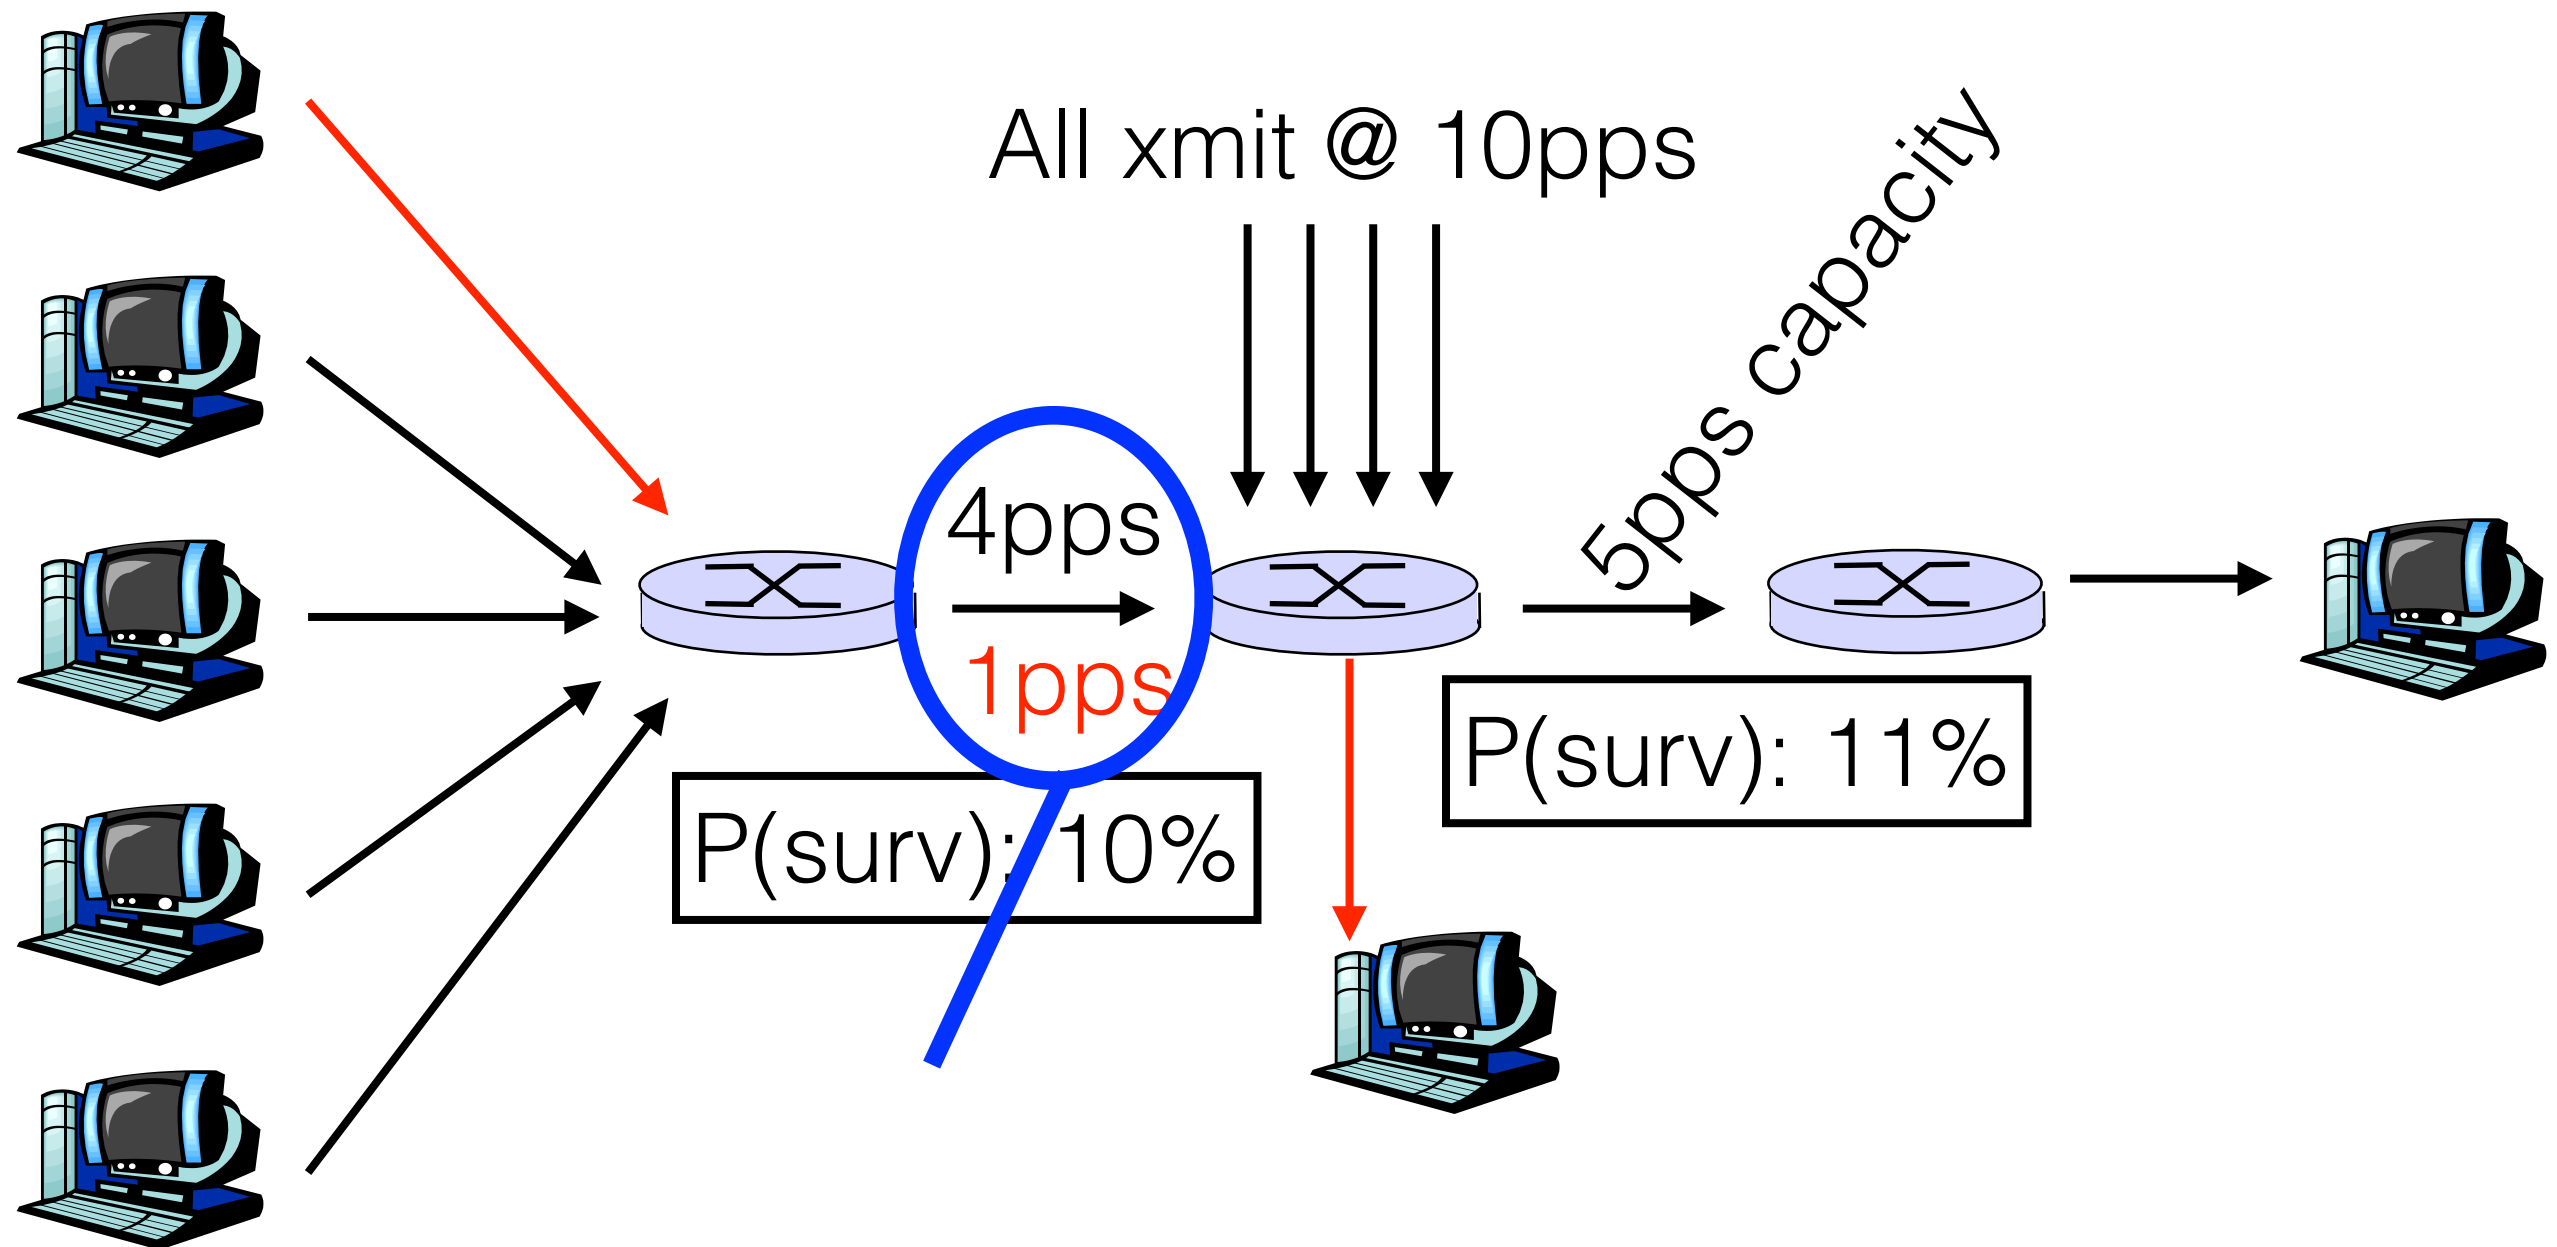
All xmit  
@ 10pps

# Congestion: Multiple Bottlenecks



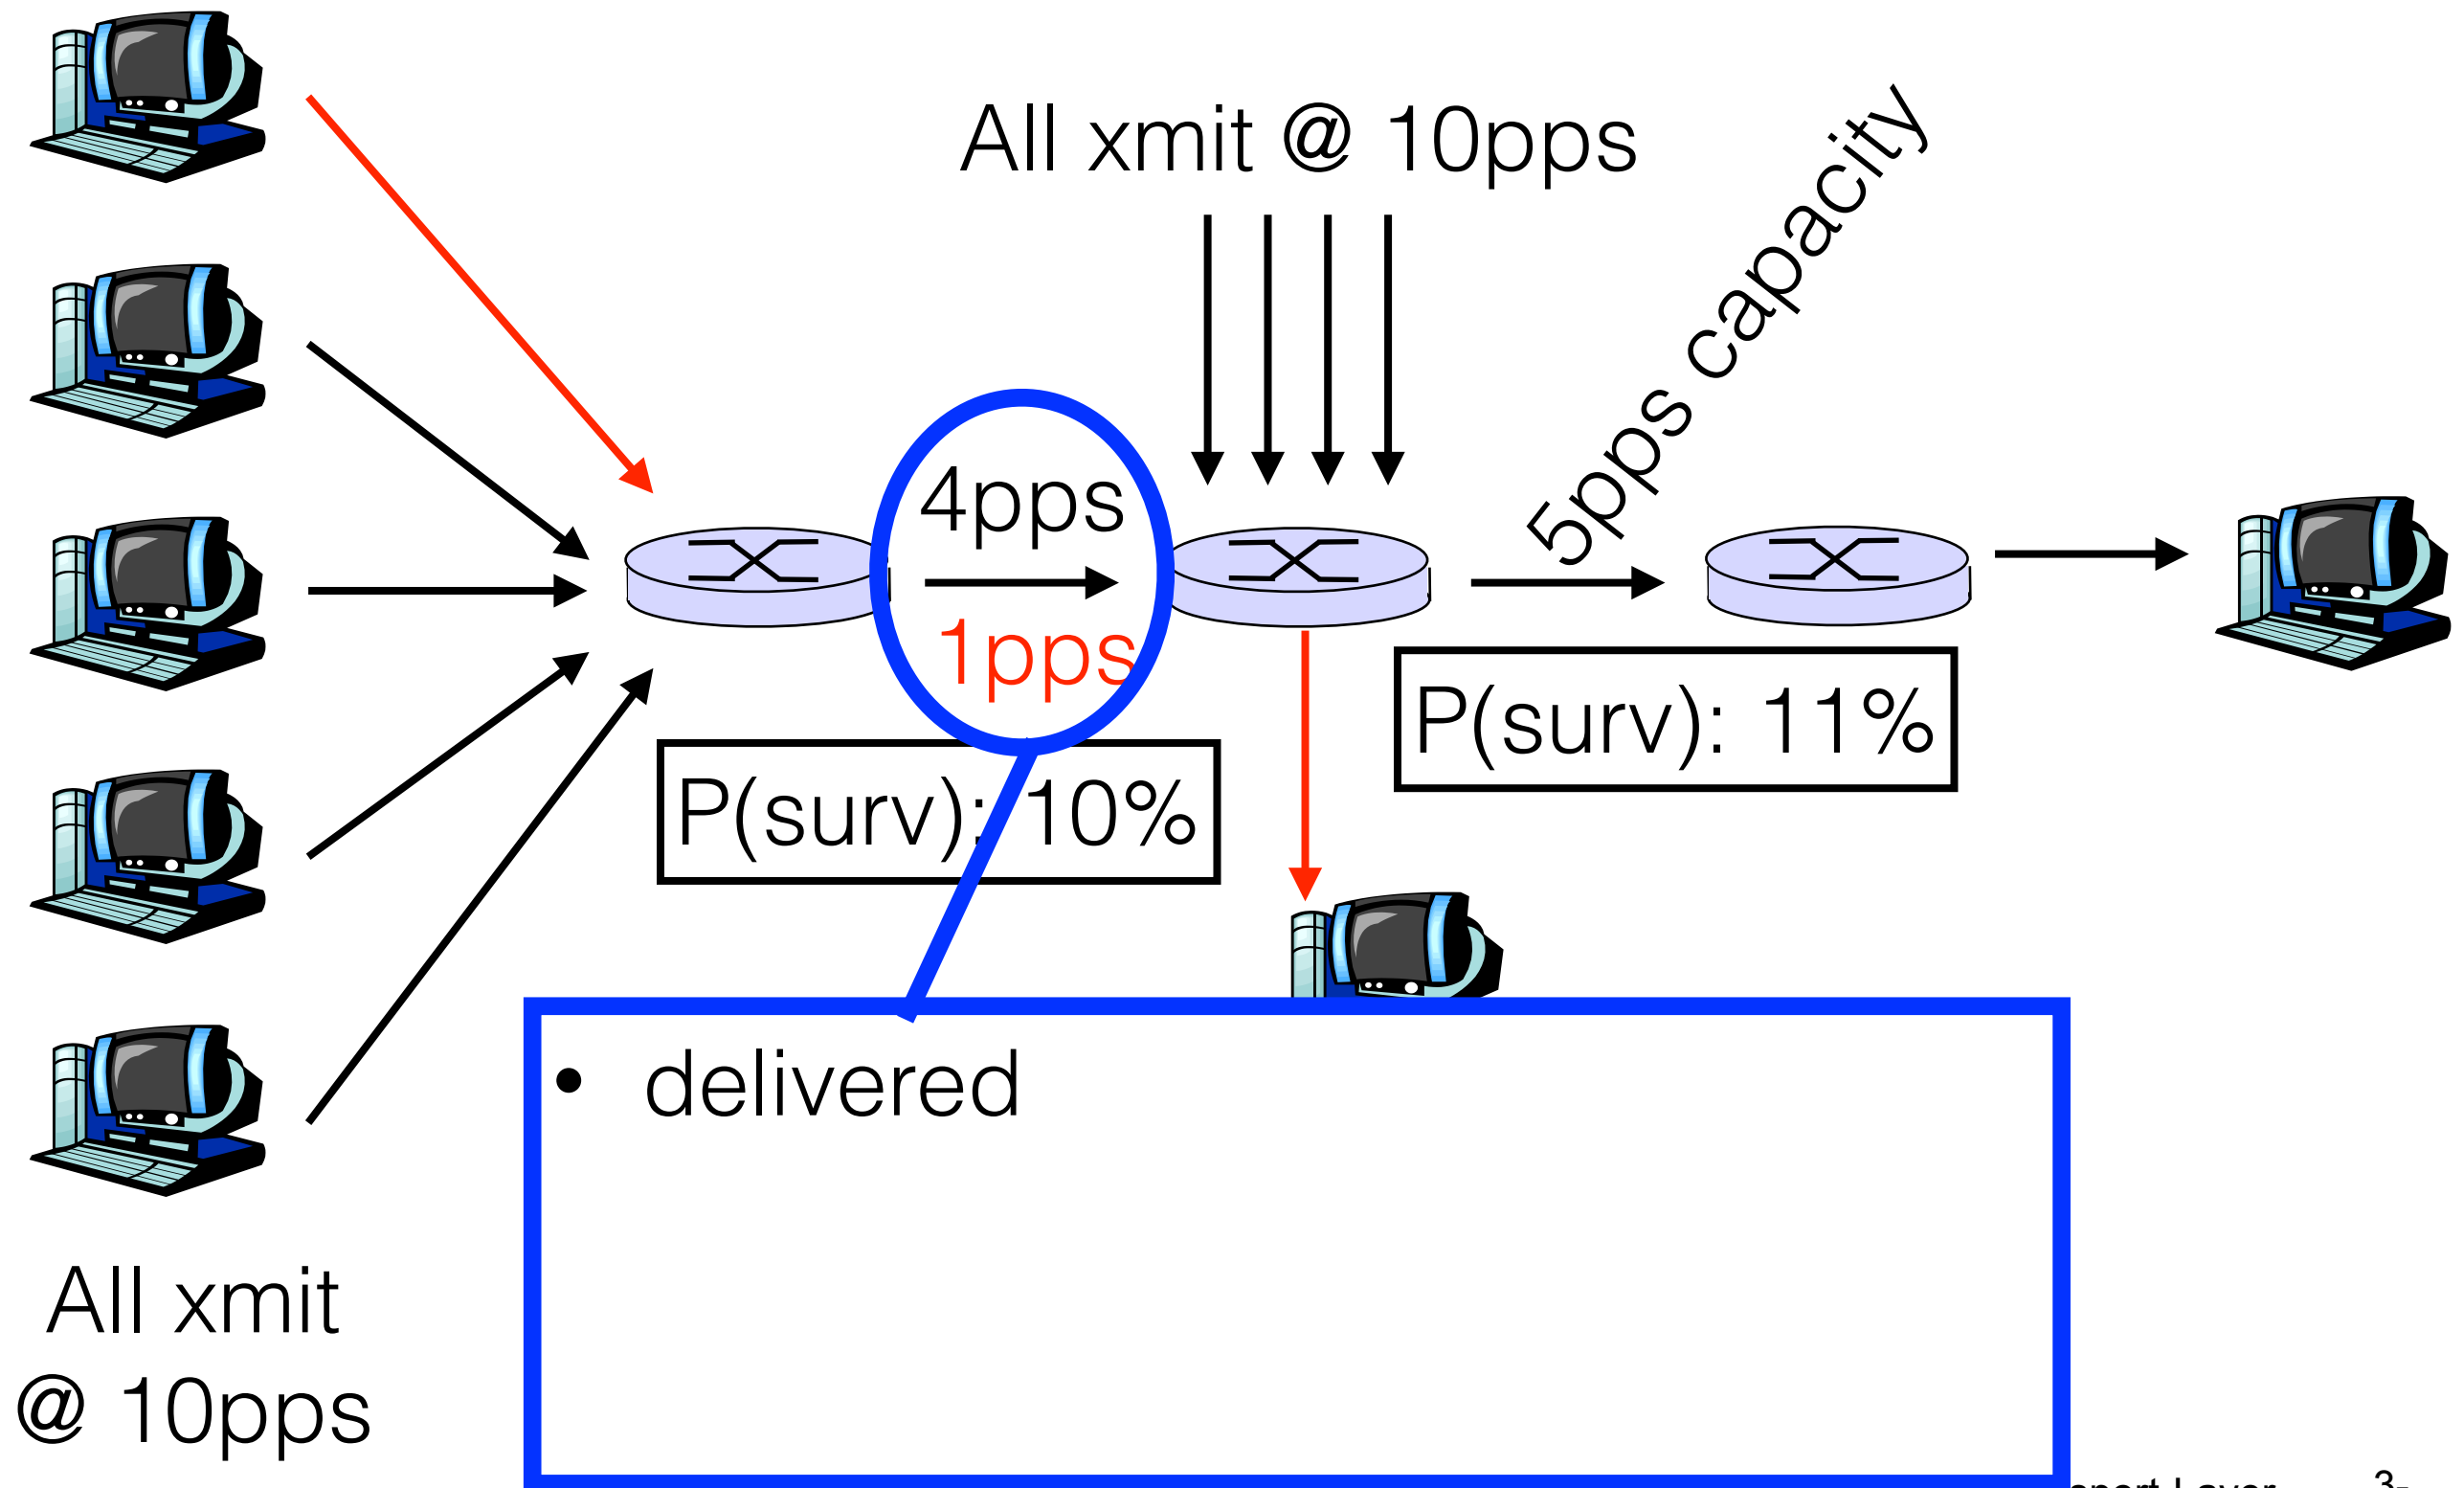
All xmit  
@ 10pps

# Congestion: Multiple Bottlenecks




All xmit  
@ 10pps


# Congestion: Multiple Bottlenecks



# Congestion: Multiple Bottlenecks




# Congestion: Multiple Bottlenecks





# Congestion: Multiple Bottlenecks



# Approaches towards congestion control

# Approaches towards congestion control

Application
Transport
Network
Data Link
Physical

# Approaches towards congestion control

# Approaches towards congestion control

❖ Fundamental problem ...

# Approaches towards congestion control


- ❖ Fundamental problem ...
  - congestion happens at a location that has no control over the imposed load

# Approaches towards congestion control

## ❖ Fundamental problem ...

- congestion happens at a location that has no control over the imposed load
- the devices that impose load have no explicit understanding of congestion

# Congestion: Multiple Bottlenecks



All xmit  
@ 10pps



# Approaches towards congestion control

# Approaches towards congestion control

## ❖ Fundamental problem ...

- congestion happens at a location that has no control over the imposed load
- the devices that impose load have no explicit understanding of congestion

# Approaches towards congestion control

## ❖ Fundamental problem ...

- congestion happens at a location that has no control over the imposed load
- the devices that impose load have no explicit understanding of congestion
- rather, congestion is generally caused by aggregate behavior from many hosts
  - so, it's hard to blame one entity
  - and, it's hard for one entity to mitigate

# Approaches towards congestion control

## ❖ Fundamental problem ...

- congestion happens at a location that has no control over the imposed load
- the devices that impose load have no explicit understanding of congestion
- rather, congestion is generally caused by aggregate behavior from many hosts
  - so, it's hard to blame one entity
  - and, it's hard for one entity to mitigate

## ❖ "tragedy of the commons" problem

# Approaches towards congestion control

# Approaches towards congestion control

Two broad approaches towards congestion control:

## end-end congestion control:

- ❖ no explicit feedback from network
- ❖ congestion inferred from end-system observed loss, delay
- ❖ approach taken by TCP

# Approaches towards congestion control

Two broad approaches towards congestion control:


## end-end congestion control:

- ❖ no explicit feedback from network
- ❖ congestion inferred from end-system observed loss, delay
- ❖ approach taken by TCP

## network-assisted congestion control:

- ❖ routers provide feedback to end systems
  - single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
  - explicit rate sender should send at

# Reading Along ...



- 3.7: TCP Congestion Control



# TCP Congestion Control

# TCP Congestion Control

❖ RFCs 3390, 3517, 3782, 5681, etc.

# TCP Congestion Control

- ❖ RFCs 3390, 3517, 3782, 5681, etc.
- ❖ Sender defines another window:
  - the "congestion window" or cwnd
  - limits amount of unACKed data

# TCP Congestion Control

- ❖ RFCs 3390, 3517, 3782, 5681, etc.
- ❖ Sender defines another window:
  - the "congestion window" or cwnd
  - limits amount of unACKed data
  - strictly limited by the flow control window

# TCP Congestion Control

- ❖ RFCs 3390, 3517, 3782, 5681, etc.
- ❖ Sender defines another window:
  - the "congestion window" or cwnd
  - limits amount of unACKed data
  - strictly limited by the flow control window
  - dynamic
  - adjusts based on the estimated state of the network

# TCP Congestion Control

- ❖ RFCs 3390, 3517, 3782, 5681, etc.
- ❖ Sender defines another window:
  - the "congestion window" or cwnd
  - limits amount of unACKed data
  - strictly limited by the flow control window
  - dynamic
  - adjusts based on the estimated state of the network
  - cwnd is sender state and not transmitted to the receiver

# TCP Congestion Control

# TCP Congestion Control

- ❖ Assume loss means the network is congested
  - not always a good assumption



# TCP Congestion Control

- ❖ Additive increase, multiplicative decrease (AIMD)
  - when no congestion detected, increase the rate (cwnd) linearly
  - when congestion detected, decrease the rate (cwnd) by half

# TCP Congestion Control: details

❖ sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

❖ roughly,

$$\text{rate} = \frac{\text{cwnd}}{\text{RTT}} \text{ Bytes/sec}$$

How does sender  
perceive congestion?

❖ loss event = timeout or  
3 duplicate acks

❖ cwnd is dynamic, function of  
perceived network congestion

# Getting Started

- ❖ OK, so, send at most cwnd bytes
- ❖ But, how do we get started?


# TCP Slow Start

# TCP Slow Start

- ❖ when connection begins, increase rate exponentially until first loss event:
  - initially  $cwnd = 1 \text{ MSS}$
  - double  $cwnd$  every RTT
  - done by incrementing  $cwnd$  for every ACK received
- ❖ summary: initial rate is slow but ramps up exponentially fast

# TCP Slow Start

- ❖ when connection begins, increase rate exponentially until first loss event:
  - initially  $cwnd = 1 \text{ MSS}$
  - double  $cwnd$  every RTT
  - done by incrementing  $cwnd$  for every ACK received
- ❖ summary: initial rate is slow but ramps up exponentially fast



# Evolving the Initial cwnd

# Evolving the Initial cwnd

❖ Late 80s: 1 segment (RFC 2001)



# Evolving the Initial cwnd

- ❖ Late 80s: 1 segment (RFC 2001)
  - why?

# Evolving the Initial cwnd

- ❖ Late 80s: 1 segment (RFC 2001)
  - why?
- ❖ Late 90s: 2-4 segments (RFCs 2414 & 3390)
  - depends on segment size

# Evolving the Initial cwnd

- ❖ Late 80s: 1 segment (RFC 2001)
  - why?
- ❖ Late 90s: 2-4 segments (RFCs 2414 & 3390)
  - depends on segment size
  - why?

# Evolving the Initial cwnd

- ❖ Late 80s: 1 segment (RFC 2001)
  - why?
- ❖ Late 90s: 2-4 segments (RFCs 2414 & 3390)
  - depends on segment size
  - why?
- ❖ 2013: 10 segments (RFC 6928)

# Evolving the Initial cwnd

- ❖ Late 80s: 1 segment (RFC 2001)
  - why?
- ❖ Late 90s: 2-4 segments (RFCs 2414 & 3390)
  - depends on segment size
  - why?
- ❖ 2013: 10 segments (RFC 6928)
  - why?

# Inferring congestion

# Inferring congestion

- ❖ after 3 dup ACKs:
  - cwnd is cut in half
  - window then grows linearly

# Inferring congestion

❖ after 3 dup ACKs:

- cwnd is cut in half
- window then grows linearly

❖ but after timeout event:

- cwnd instead set to 1 MSS;
- window then grows exponentially
- to a threshold, then grows linearly



# Inferring congestion


- ❖ after 3 dup ACKs:
  - cwnd is cut in half
  - window then grows linearly
- ❖ but after timeout event:
  - cwnd instead set to 1 MSS;
  - window then grows exponentially
  - to a threshold, then grows linearly

## Philosophy:


- ❖ 3 dup ACKs indicates network capable of delivering some segments
- ❖ timeout indicates a "more alarming" congestion scenario




# Congestion Window Evolution




# Congestion Window Evolution




# Congestion Window Evolution




# Congestion Window Evolution



# Congestion Window Evolution




# Congestion Window Evolution







# Congestion Window Evolution




# Congestion Window Evolution




# Congestion Window Evolution




# Congestion Window Evolution



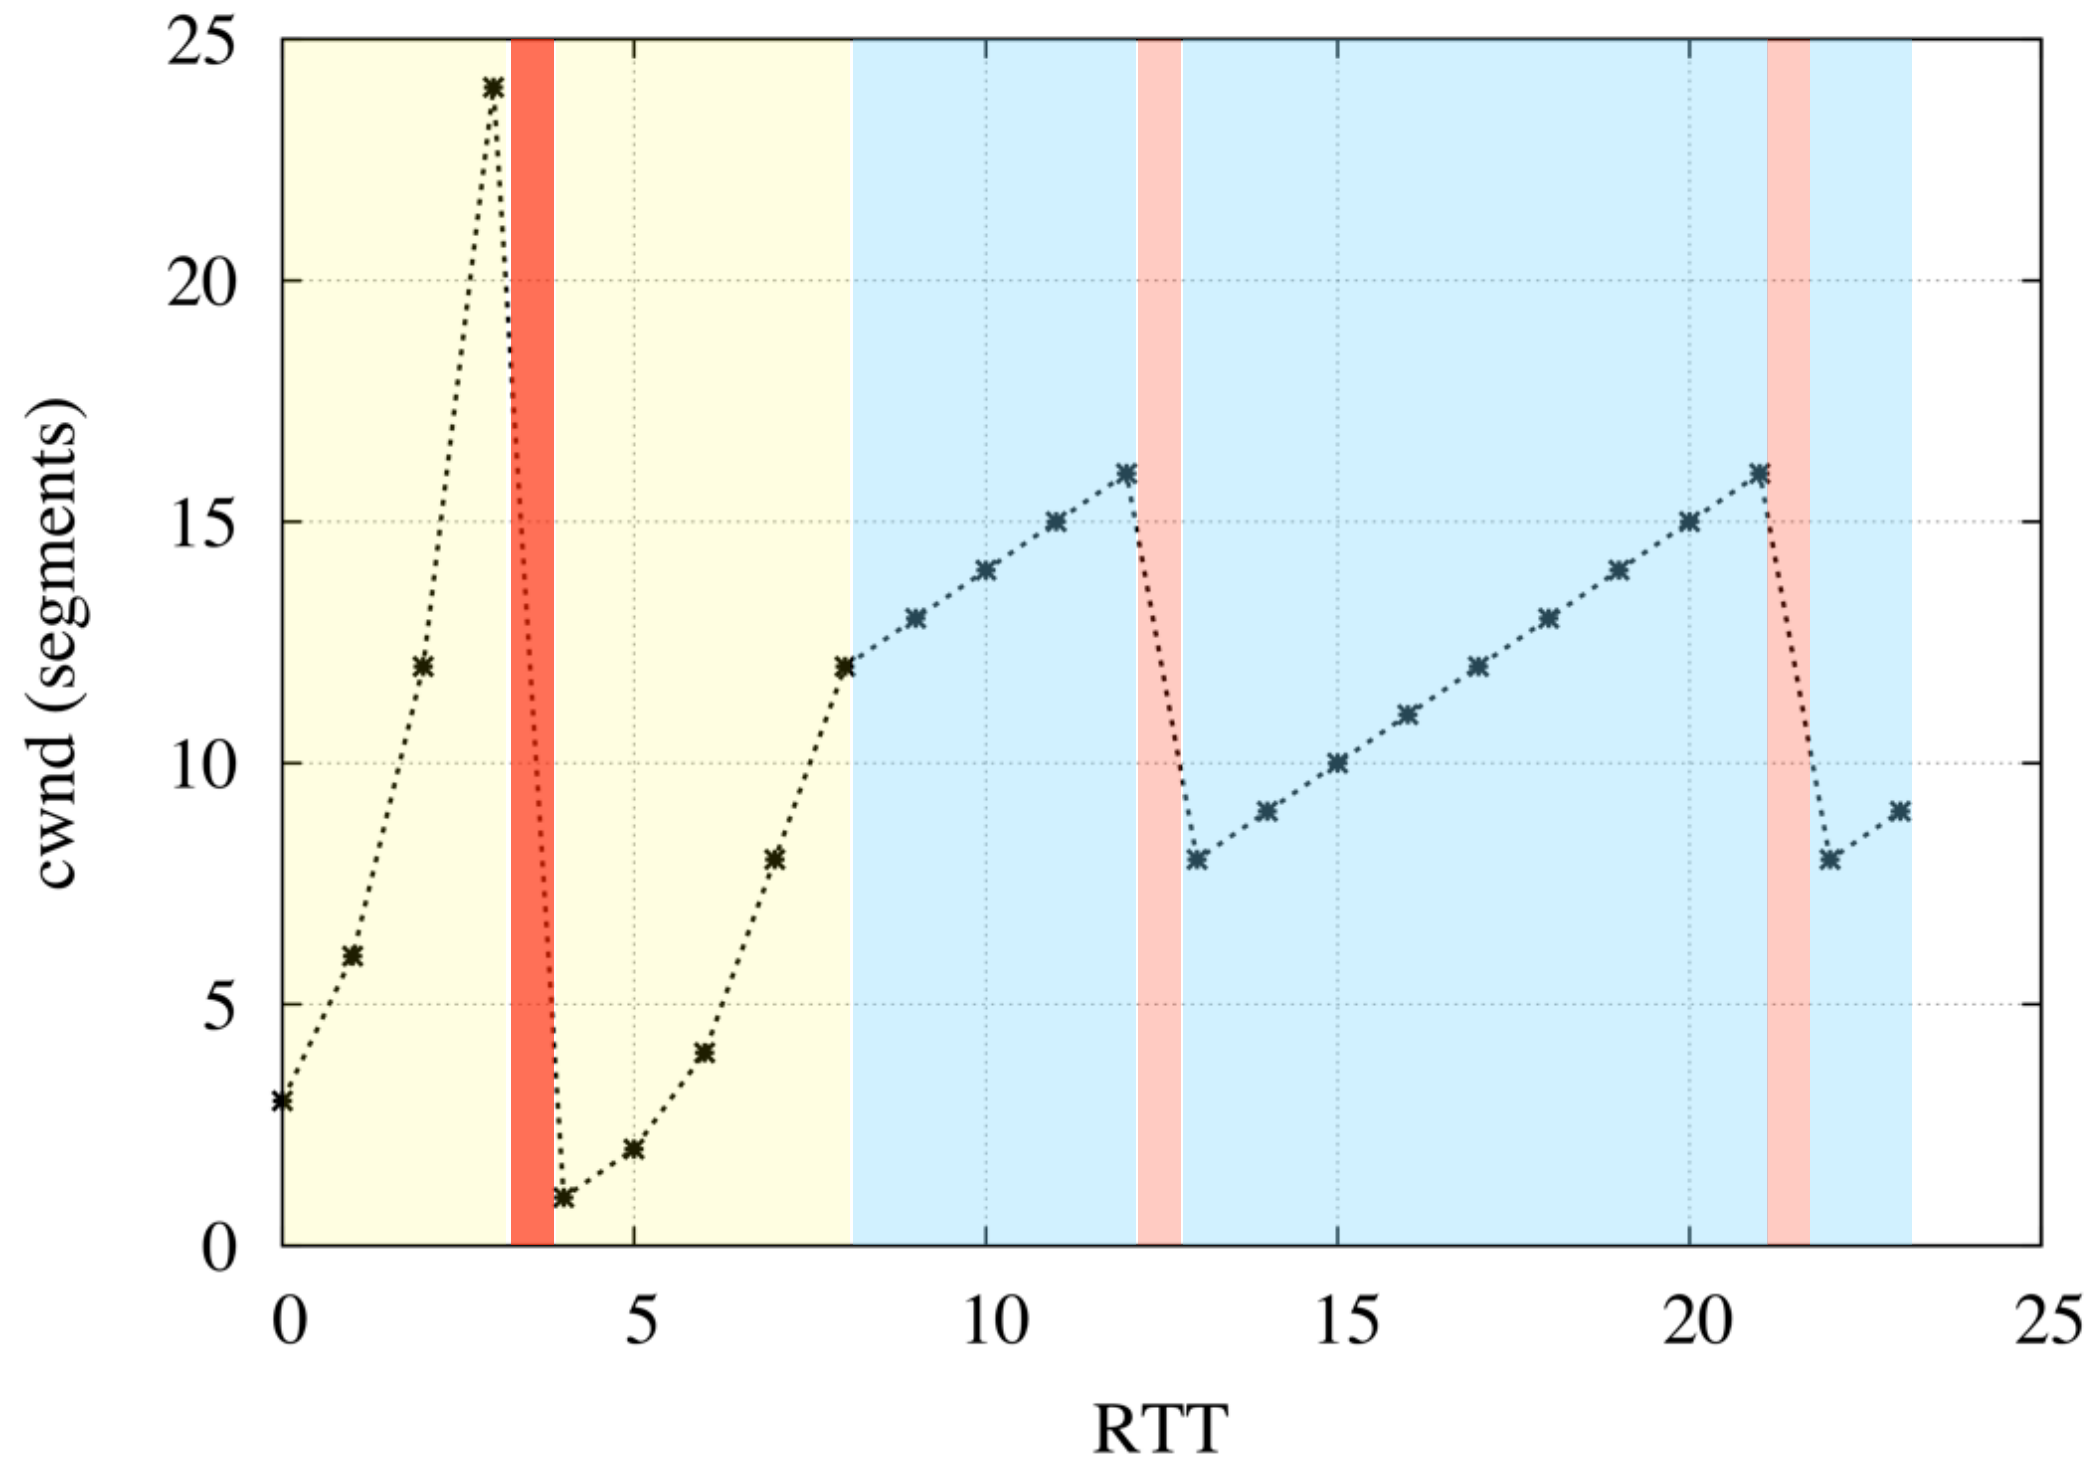
# Congestion Window Evolution




# Congestion Window Evolution



# cwnd Evolution



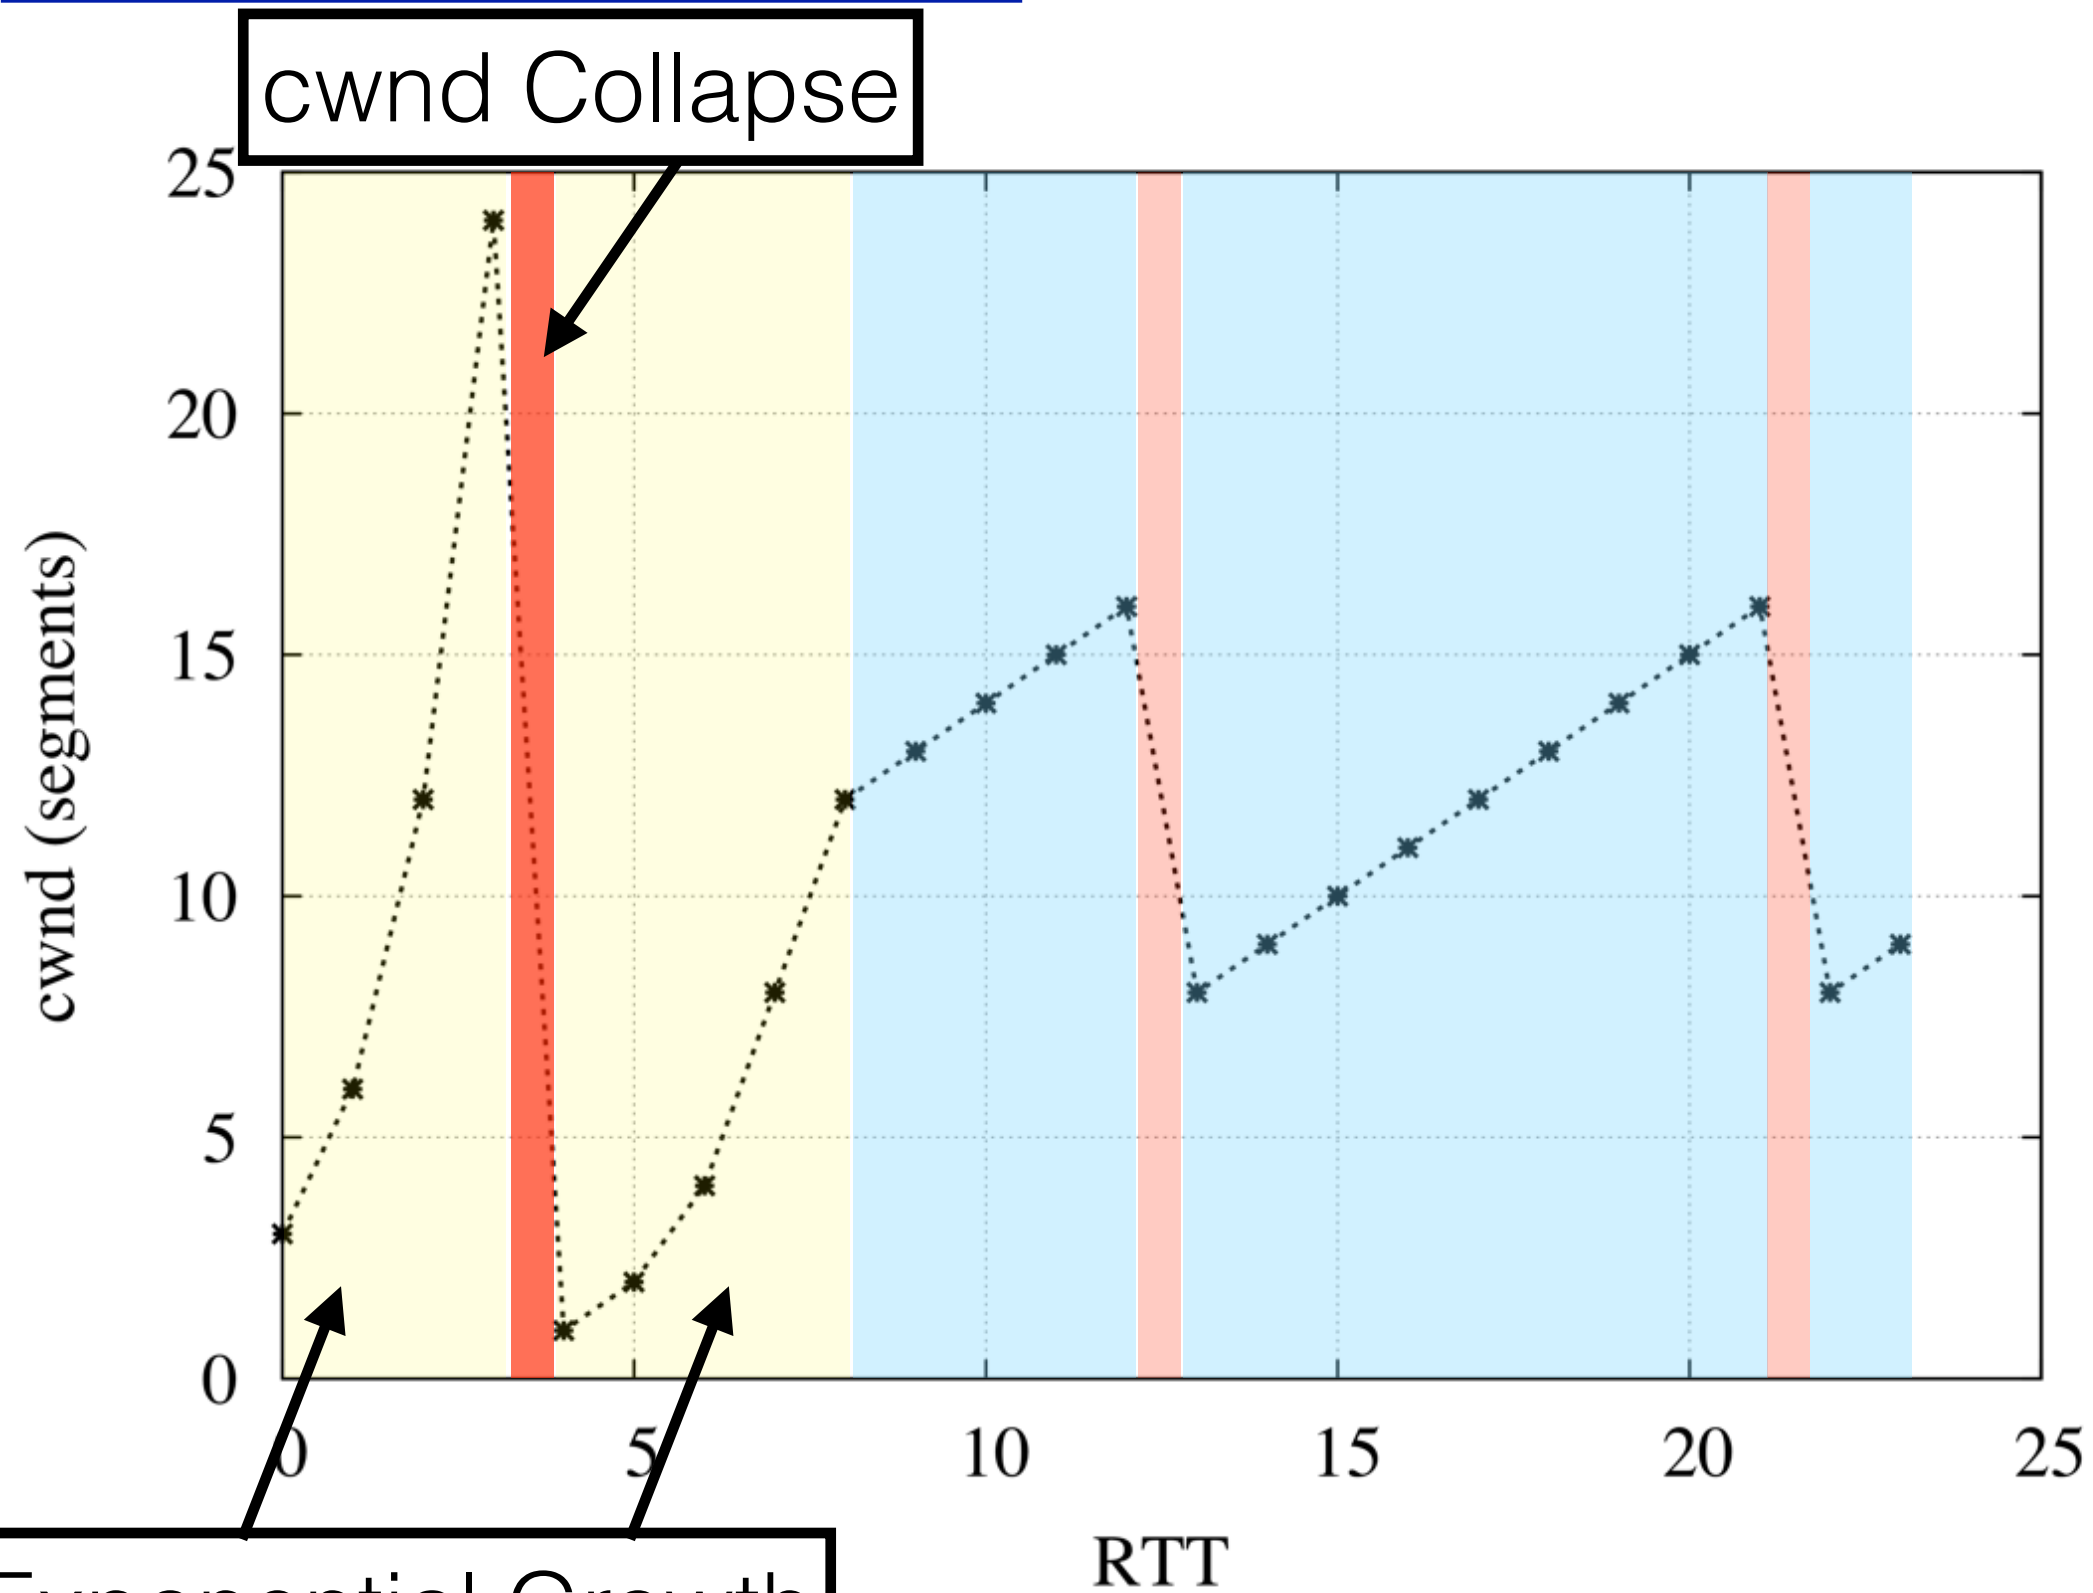
# cwnd Evolution




Exponential Growth



# cwnd Evolution



# cwnd Evolution



# cwnd Evolution

