



# Congestion Control

## Part 2

Mark Allman  
*[mallman@case.edu](mailto:mallman@case.edu)*


EECS 325/425  
Fall 2018

These slides are more-or-less directly from the slide set developed by Jim Kurose and Keith Ross for their book “Computer Networking: A Top Down Approach, 5th edition”.

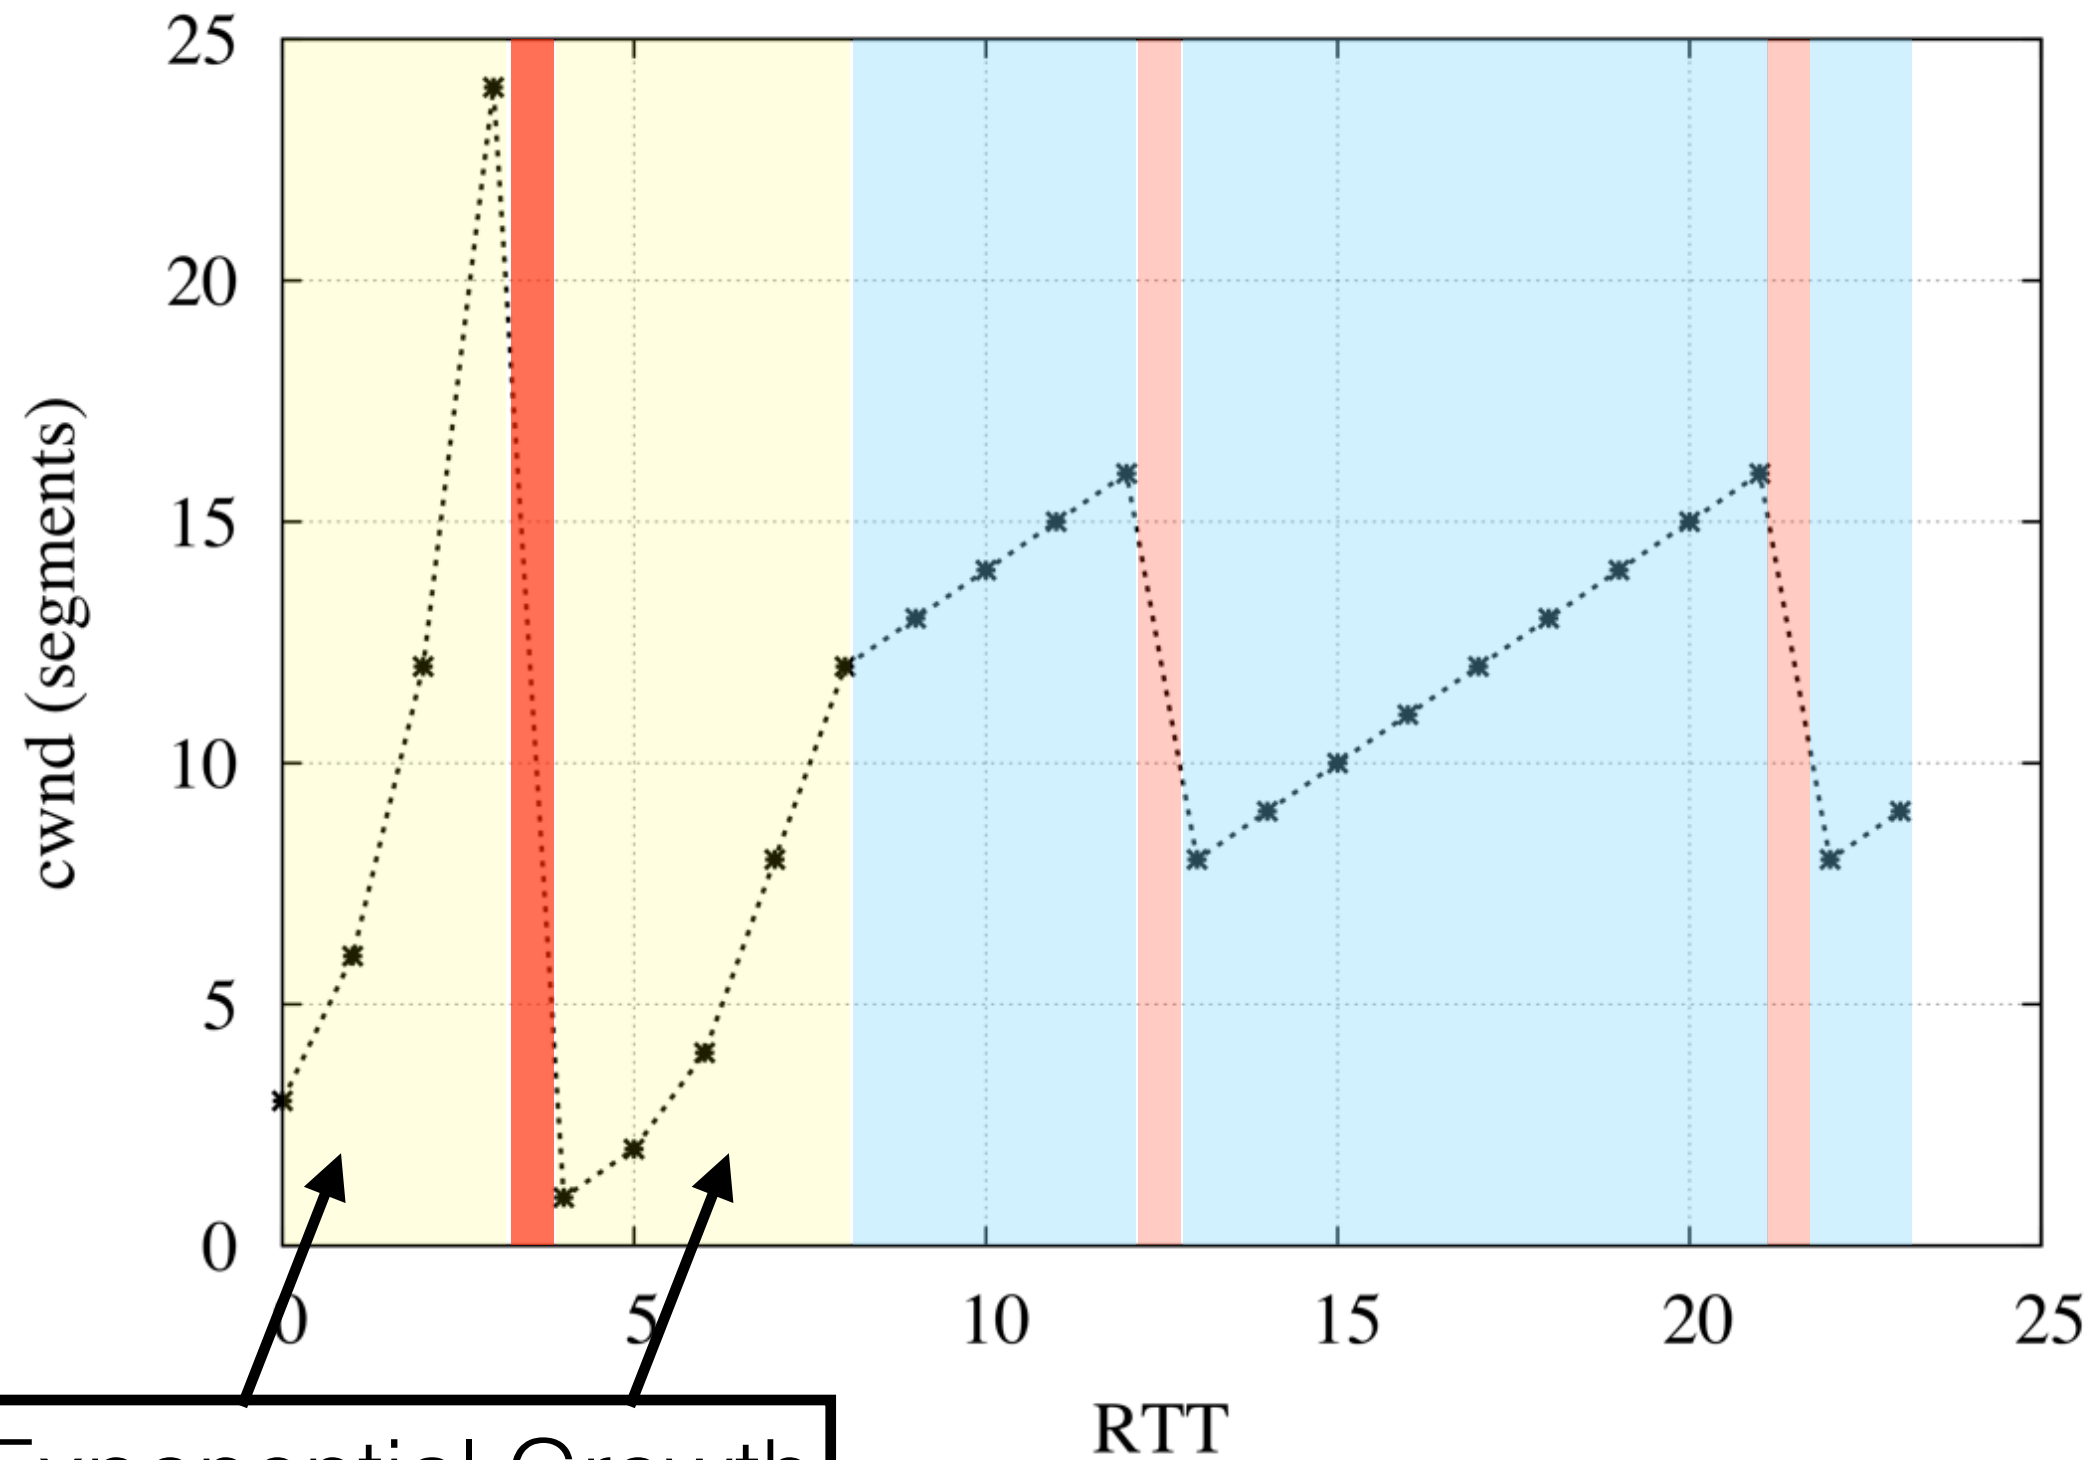
The slides have been lightly adapted for Mark Allman’s EECS 325/425 Computer Networks class at Case Western Reserve University.

All material copyright 1996-2010  
J.F Kurose and K.W. Ross, All Rights Reserved

# cwnd Evolution




# cwnd Evolution

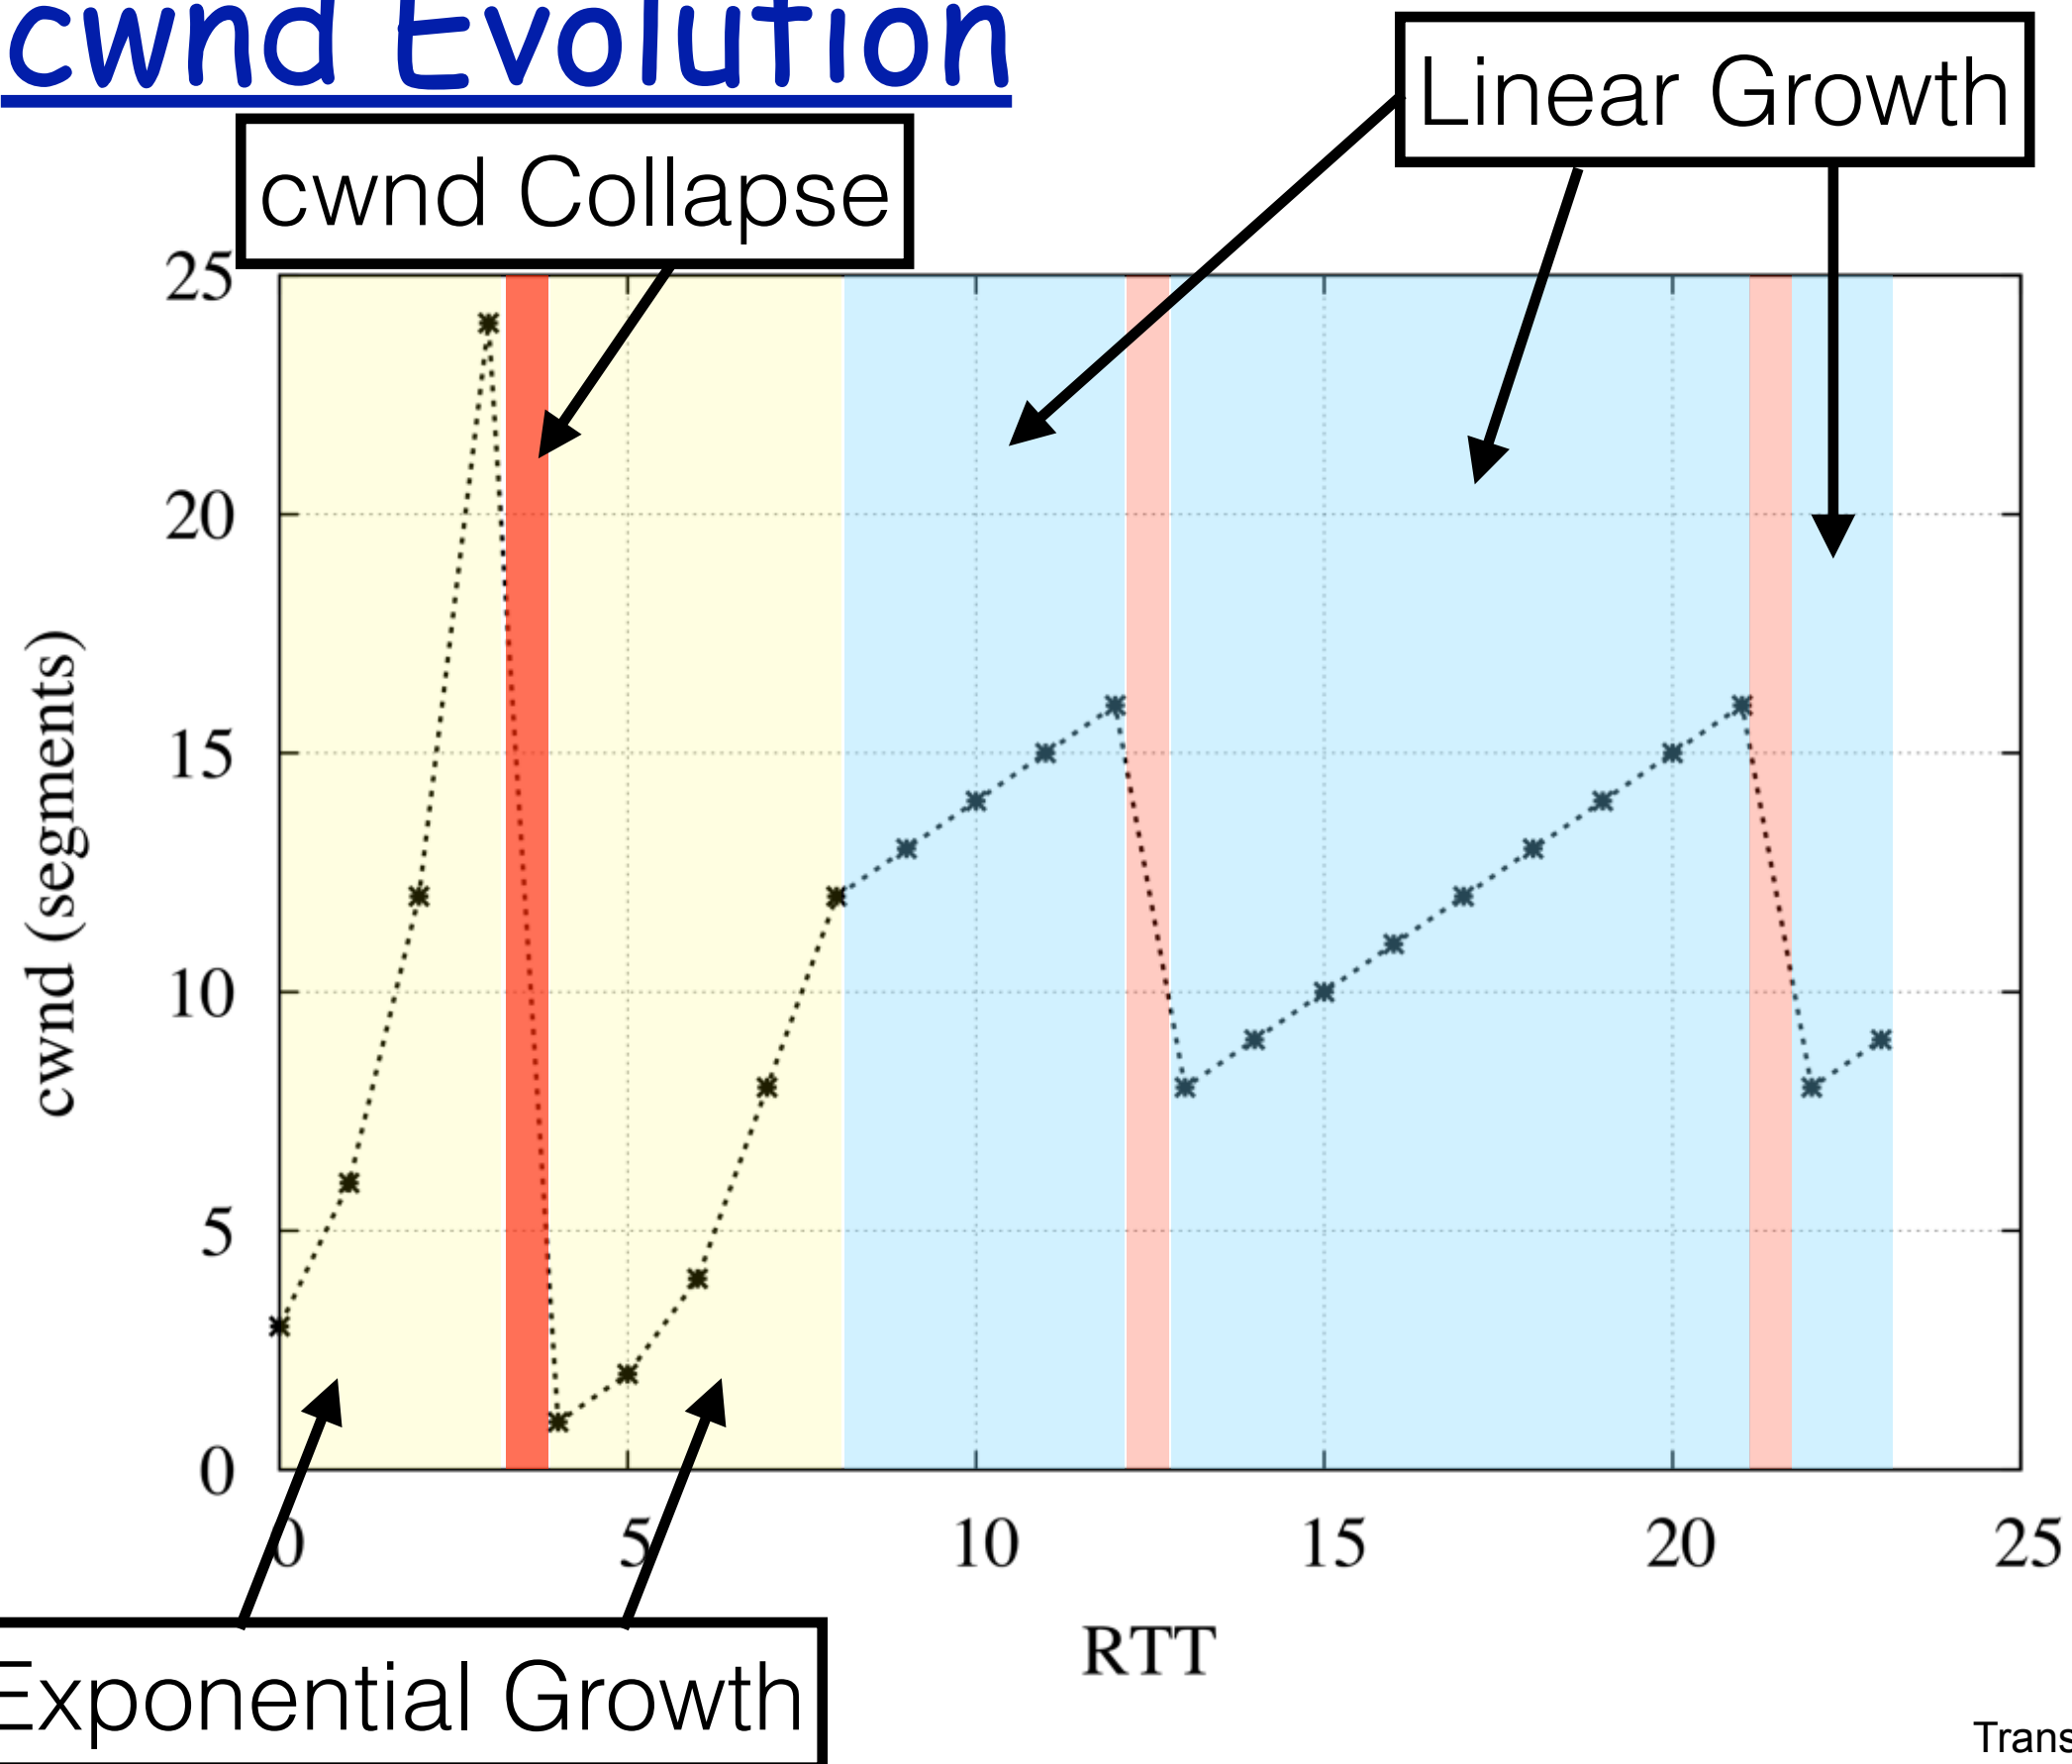


Exponential Growth


# cwnd Evolution




# cwnd Evolution



# cwnd Evolution




# Asymmetry & CC






# Asymmetry & CC




# Asymmetry & CC




$$\text{Net Asym} = 25 / 5 = 5$$

# Asymmetry & CC




$$\text{Net Asym} = 25 / 5 = 5$$

# Asymmetry & CC



$$\text{Net Asym} = 25 / 5 = 5$$


# Asymmetry & CC



$$\text{Net Asym} = 25 / 5 = 5$$

$$\text{TCP Asym} = 1518 / 58 \approx 26$$

# Asymmetry & CC




$$\text{Net Asym} = 25 / 5 = 5$$


$$\text{TCP Asym} = 1518 / 58 \approx 26$$

OK as long as  $\text{TCP Asym} > \text{Net Asym}$

# Asymmetry & CC



# Asymmetry & CC




$$\text{Net Asym} = 25 / 0.5 = 50$$

$$\text{TCP Asym} = 1518 / 58 \approx 26$$

OK as long as TCP Asym > Net Asym



# Asymmetry & CC




$$\text{Net Asym} = 25 / 0.5 = 50$$

$$\text{TCP Asym} = 1518 / 58 \approx 26$$

OK as long as  $\text{TCP Asym} > \text{Net Asym}$

# Asymmetry & CC



$$\text{Net Asym} = 25 / 0.5 = 50$$


$$\text{TCP Asym} = 1518 / 58 \approx 26$$

OK as long as  $\text{TCP Asym} > \text{Net Asym}$

# Asymmetry & CC

- What to do in highly asymmetric networks?
  - dynamically set the rate of ACKs

# cwnd Evolution



# Performance

- We know that network performance is *bounded* by a simple equation:

$$\text{window} = \text{bw} \times \text{delay}$$

$$\text{bw} = \text{window} \div \text{delay}$$

- But, the congestion window is dynamic ....

# TCP Performance

- Refine:

$$\text{rate} \approx (1.22 \times \text{MSS}) \div (\text{RTT} \times \text{sqrt}(\text{L}))$$

- example: 1500 byte segments, 100ms RTT, want 10 Gbps throughput
- $L = 2 \cdot 10^{-10}$     **Wow – a very small loss rate!**

# TCP Performance

$$\text{rate} \approx \frac{(1.22 \times \text{MSS})}{(\text{RTT} \times \text{sqrt}(\text{L}))}$$

# TCP Performance

$$\text{rate} \approx (1.22 \times \text{MSS}) \div (\text{RTT} \times \text{sqrt}(\text{L}))$$

- Shorter RTTs get preference



# TCP Performance

$$\text{rate} \approx (1.22 \times \text{MSS}) \div (\text{RTT} \times \text{sqrt}(\text{L}))$$

- Shorter RTTs get preference
  - consider 2 connections: with RTTs of 10msec and 50msec, but all else the same

# TCP Performance

$$\text{rate} \approx (1.22 \times \text{MSS}) \div (\text{RTT} \times \text{sqrt}(\text{L}))$$

- Shorter RTTs get preference
  - consider 2 connections: with RTTs of 10msec and 50msec, but all else the same
  - the connection with a 10msec RTT will get 5x the performance of the other connection

# TCP Performance

$$\text{rate} \approx (1.22 \times \text{MSS}) \div (\text{RTT} \times \text{sqrt}(\text{L}))$$

- Shorter RTTs get preference
  - consider 2 connections: with RTTs of 10msec and 50msec, but all else the same
    - the connection with a 10msec RTT will get 5x the performance of the other connection
- TCP is fair as long as the MSS, RTT and L are the same

# Endpoint CC

- Why do we use endpoint CC?
  - What are the pros? Cons?
- Should we involve the network?
  - Why or why not?

# Approaches towards congestion control

Two broad approaches towards congestion control:

## end-end congestion control:

- ❖ no explicit feedback from network
- ❖ congestion inferred from end-system observed loss, delay
- ❖ approach taken by TCP

## network-assisted congestion control:

- ❖ routers provide feedback to end systems
  - single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
  - explicit rate sender should send at

# The Network Side of CC

- Active queuing
  - e.g., RED
- Fair queuing
- Traffic shaping

# Active Queuing

# Active Queuing

- Drop segments *before* the queue fills to invoke endpoint CC



# Active Queuing

- Drop segments *before* the queue fills to invoke endpoint CC
- E.g., Random Early Detection (RED)
  - (lots of others)

# Active Queuing

- Drop segments *before* the queue fills to invoke endpoint CC
- E.g., Random Early Detection (RED)
  - (lots of others)
- Goal: the network works with endpoints, but *implicitly*
  - (a little explicitness has been added ...)

# RED

# RED

- Track *average* queue occupancy,  $q$

# RED

- Track *average* queue occupancy,  $q$
- Set minimum and maximum thresholds:  $min^{th}$  &  $max^{th}$

# RED

- Track *average* queue occupancy,  $q$
- Set minimum and maximum thresholds:  $min^{th}$  &  $max^{th}$
- When  $q < min^{th}$ , forward all packets

# RED

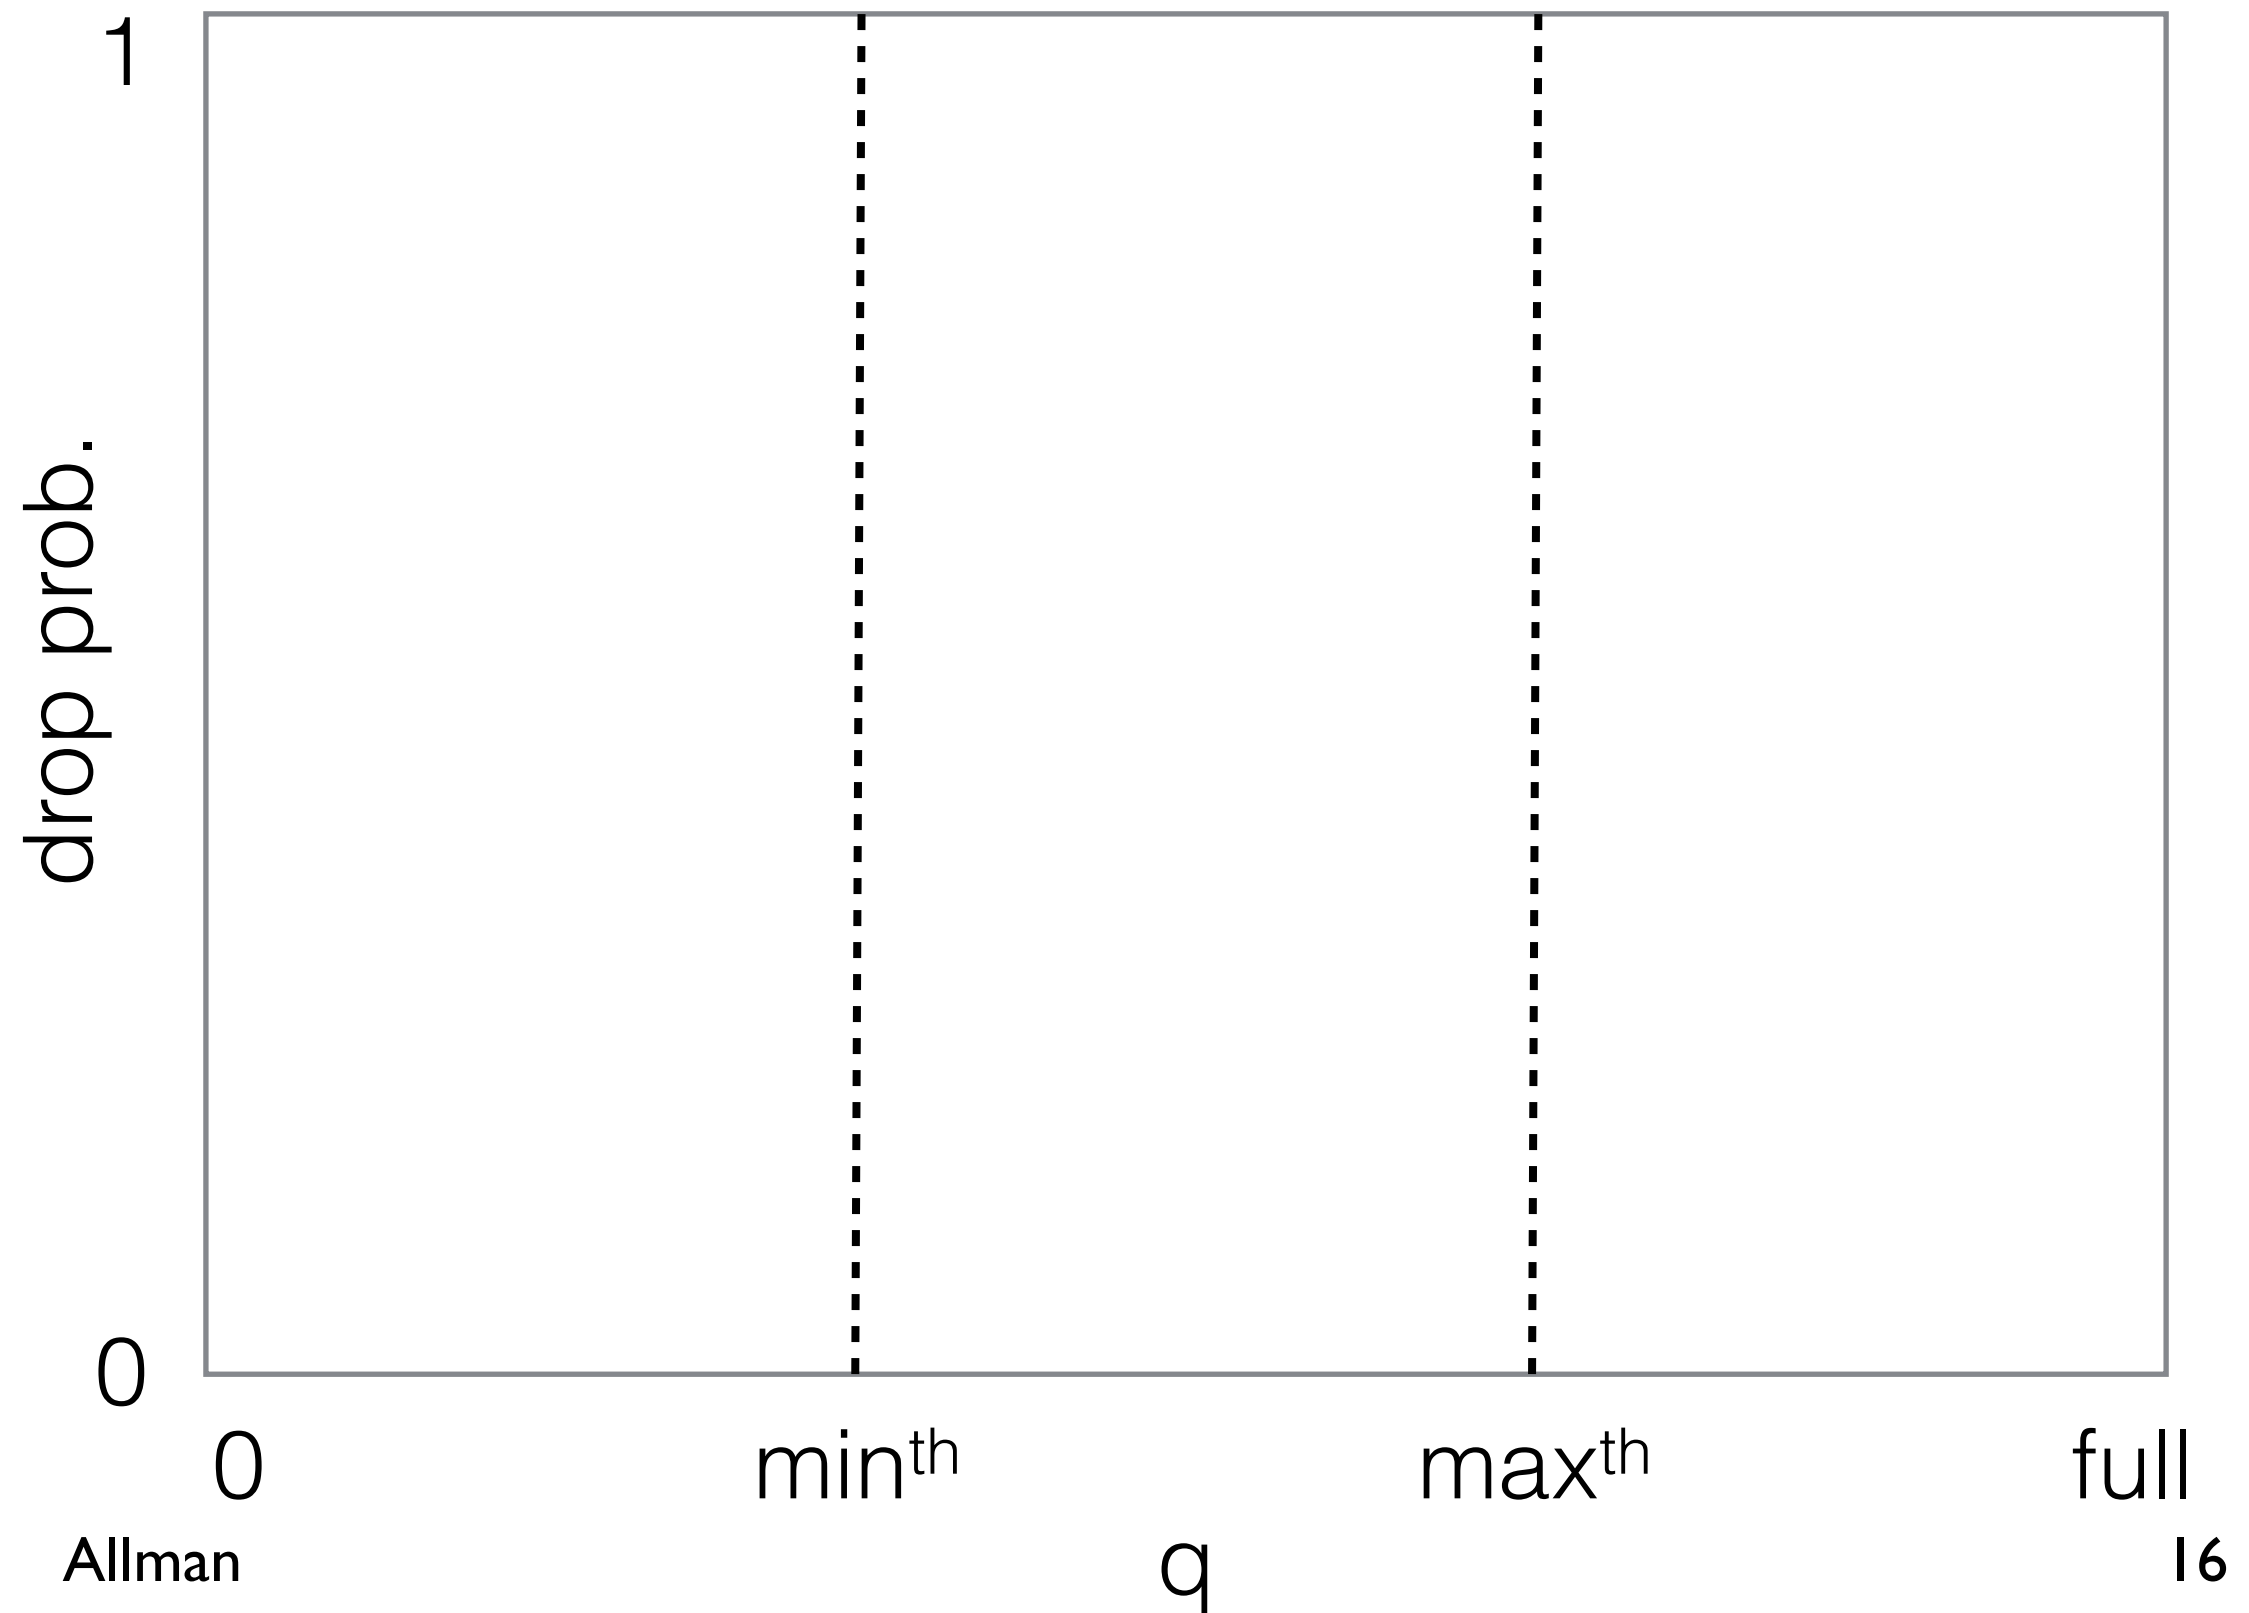
- Track *average* queue occupancy,  $q$
- Set minimum and maximum thresholds:  $min^{th}$  &  $max^{th}$
- When  $q < min^{th}$ , forward all packets
- When  $q > max^{th}$ , drop all packets

# RED


- Track *average* queue occupancy,  $q$
- Set minimum and maximum thresholds:  $min^{th}$  &  $max^{th}$
- When  $q < min^{th}$ , forward all packets
- When  $q > max^{th}$ , drop all packets
- When  $min^{th} \leq q \leq max^{th}$ , drop probabilistically




# RED



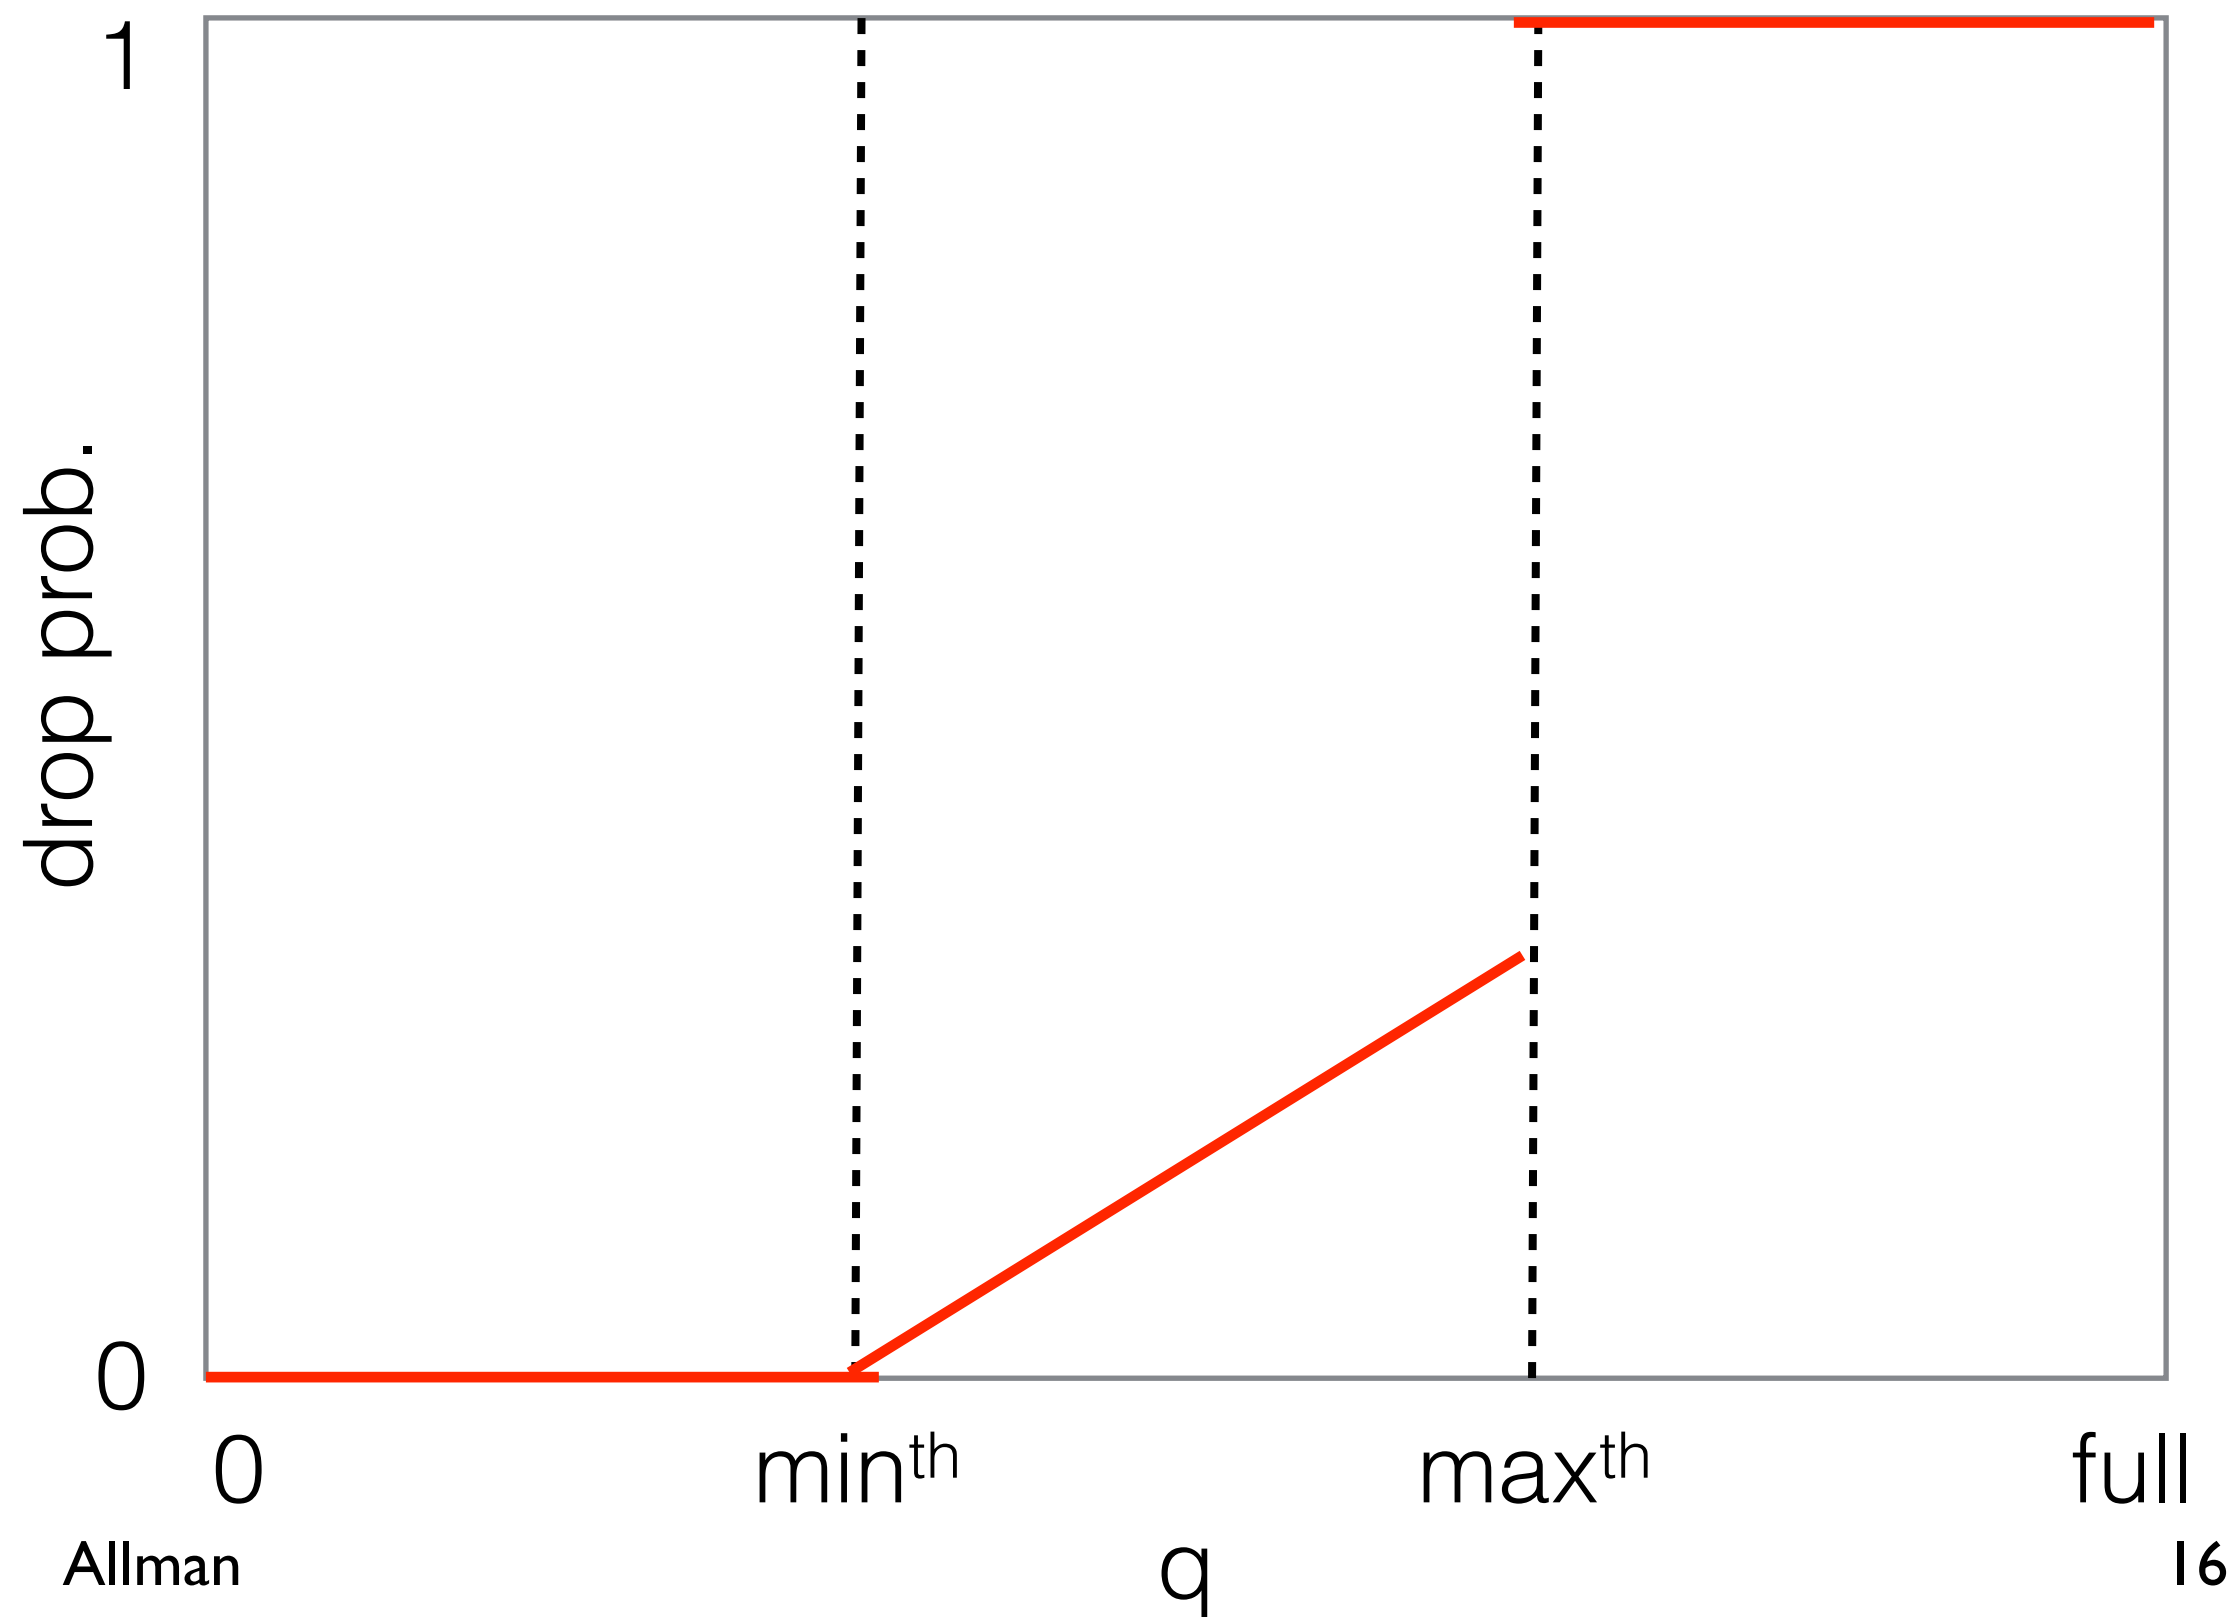
# RED




# RED



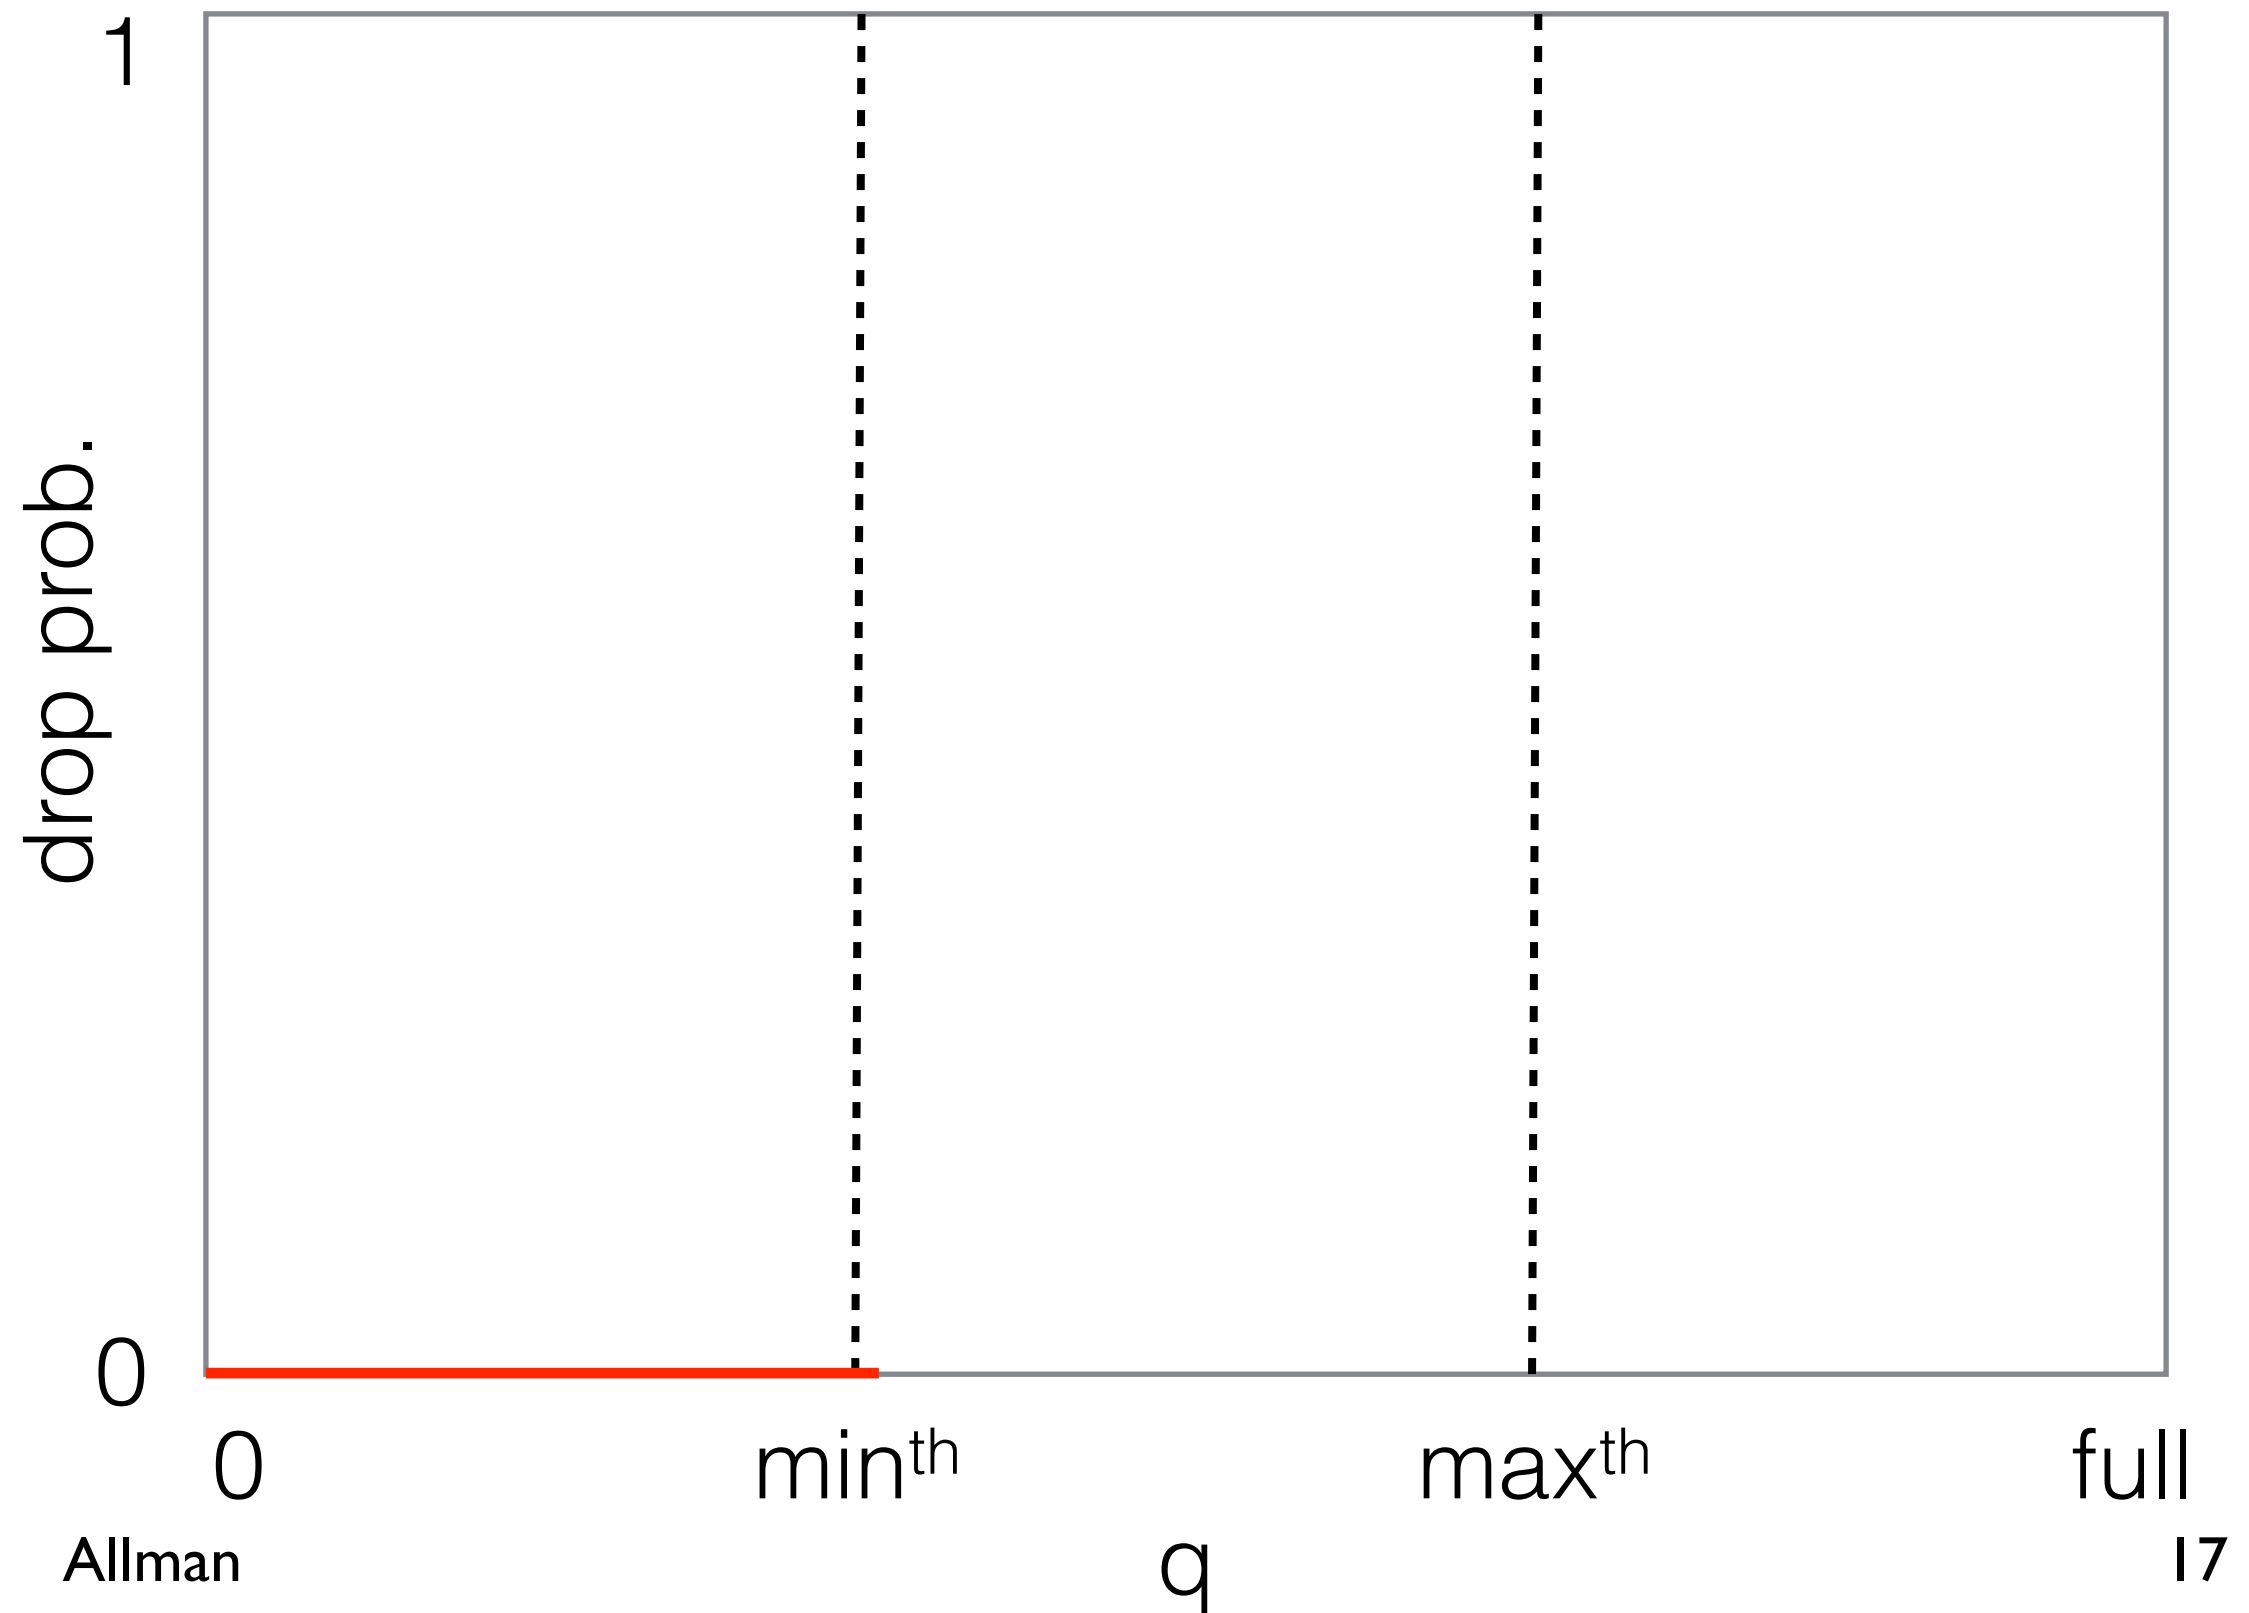
# RED




# Gentle RED




# Gentle RED



# Gentle RED



# Gentle RED

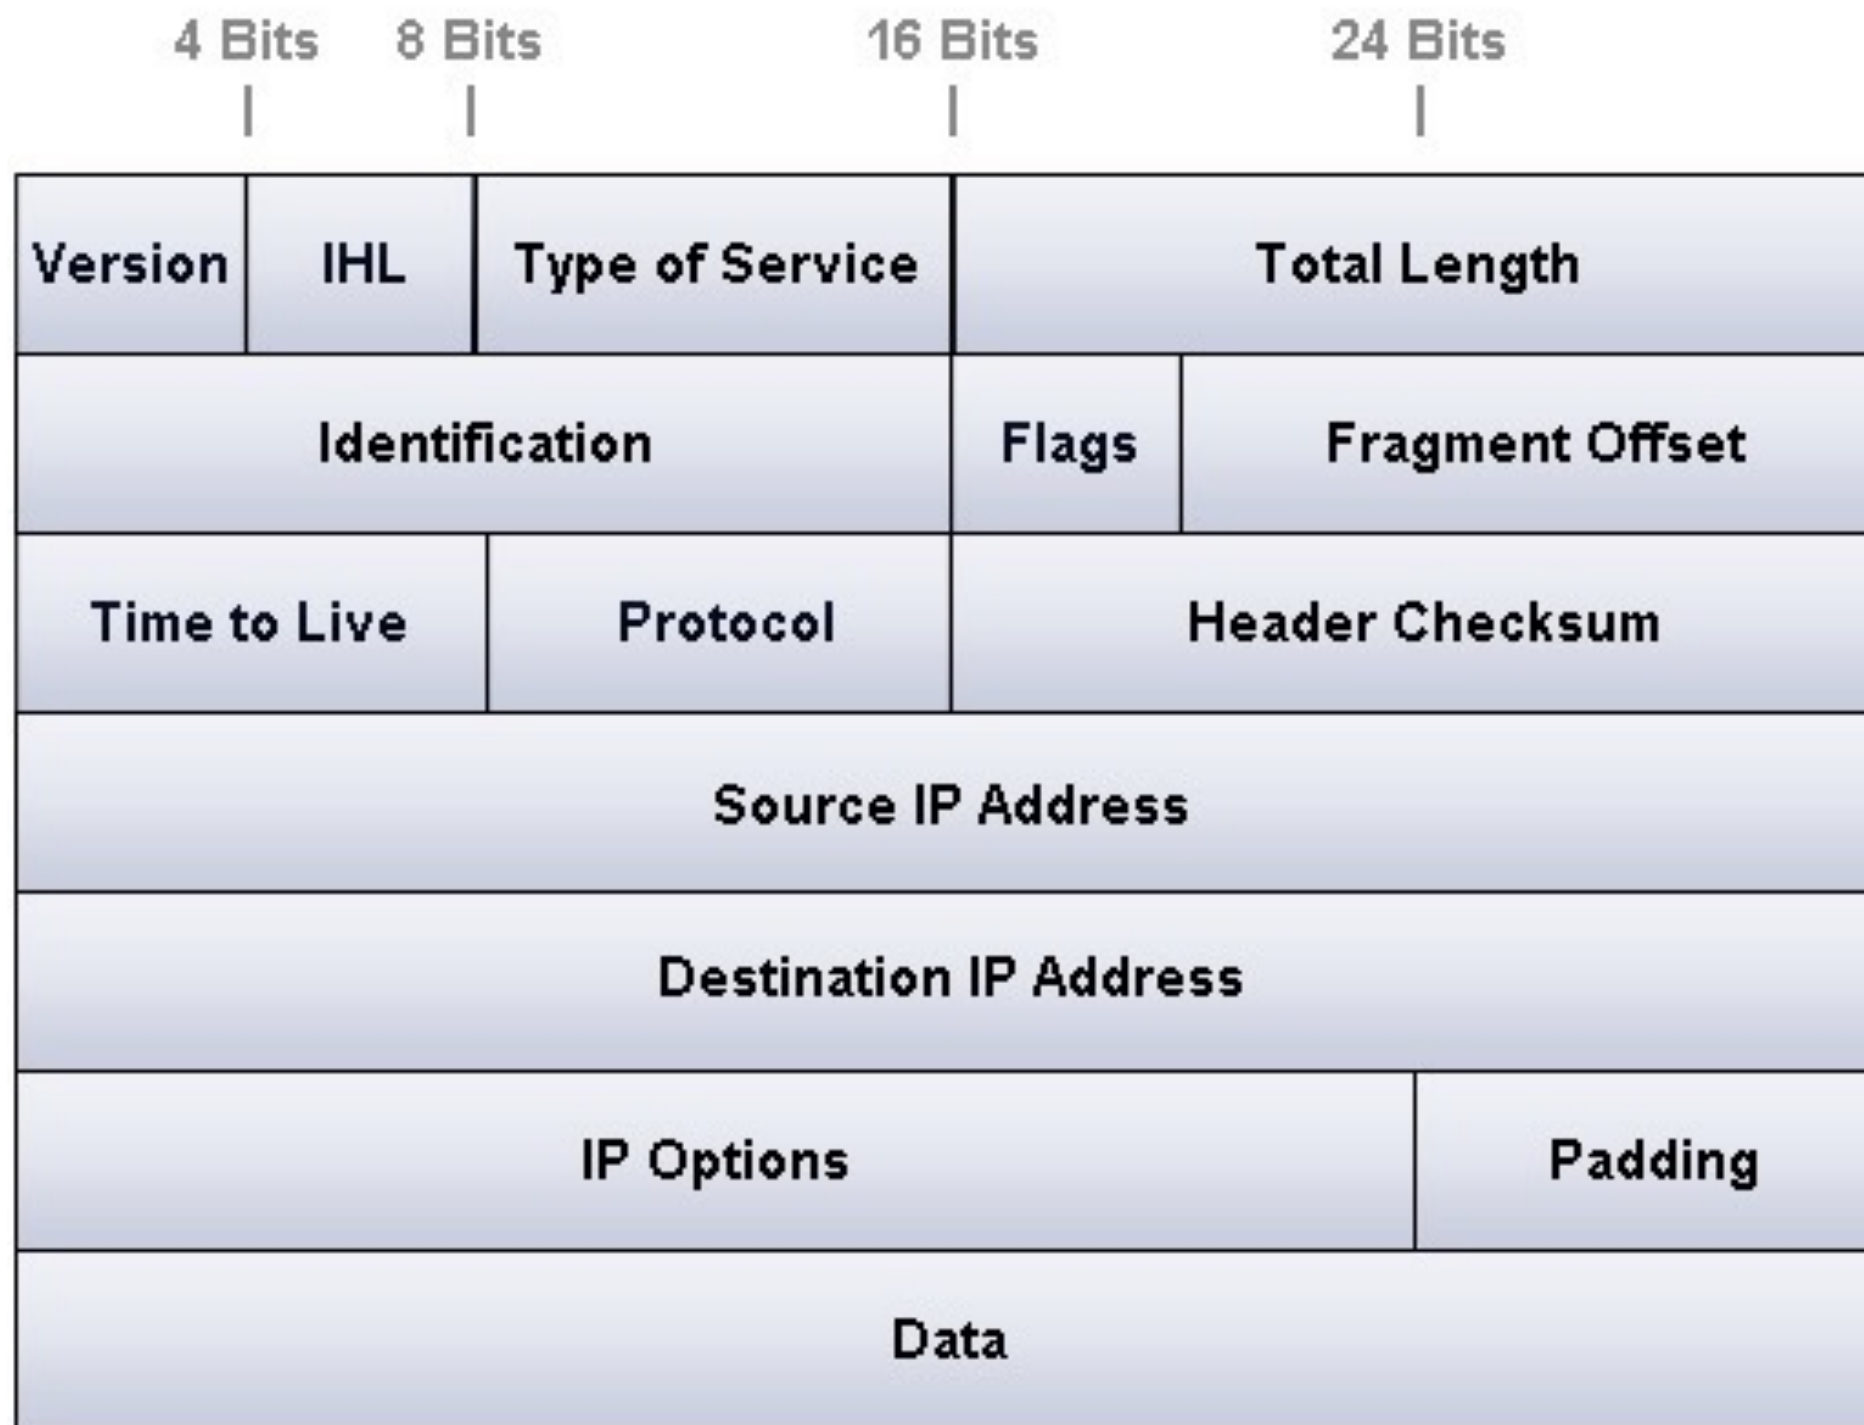




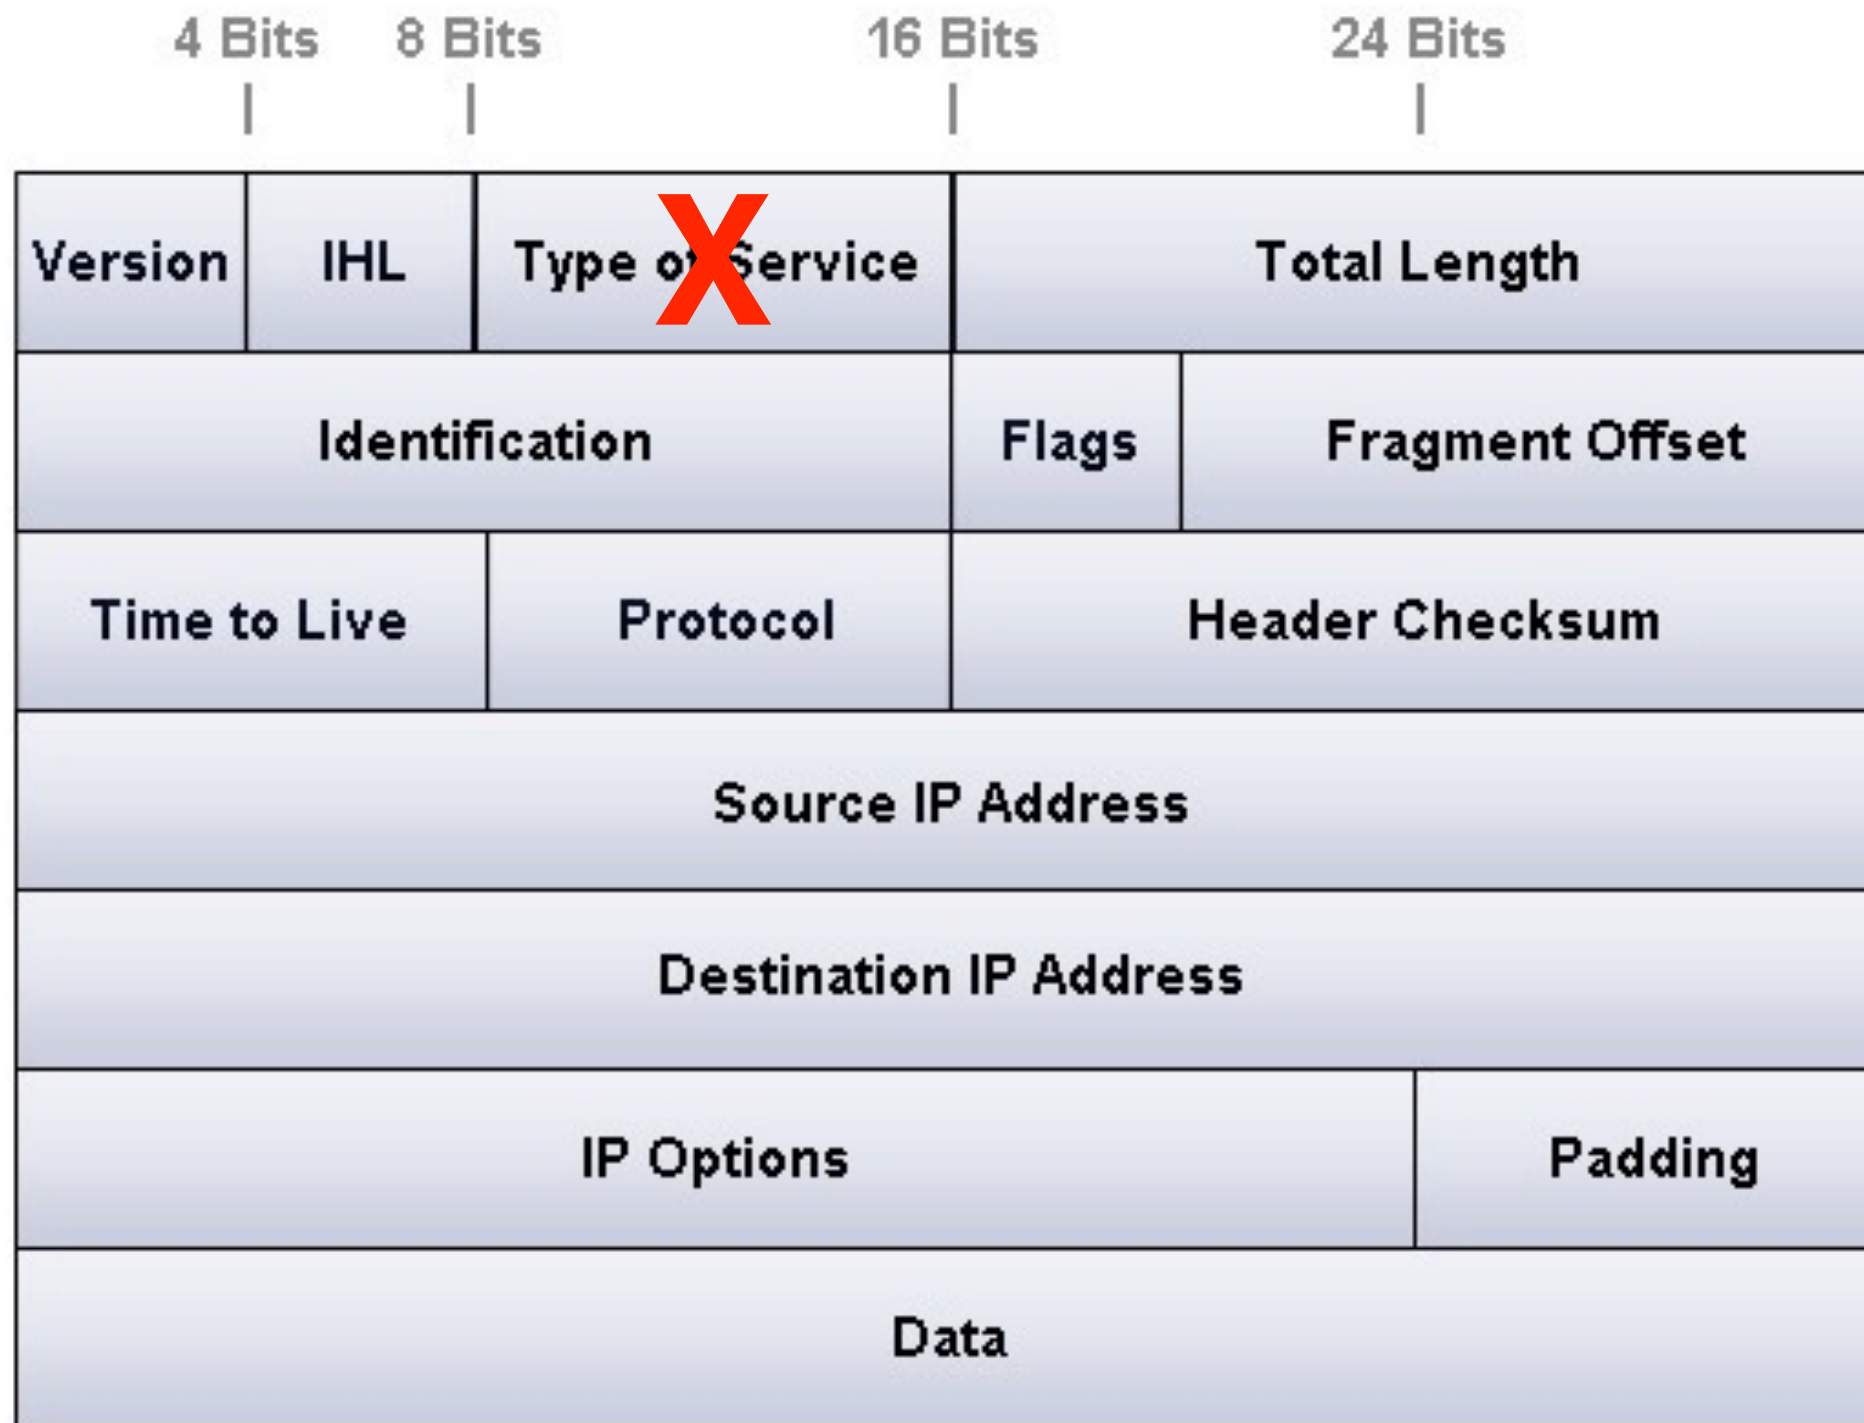
# Active Queuing

- Operators are somewhat hesitant to use these schemes .... why?!

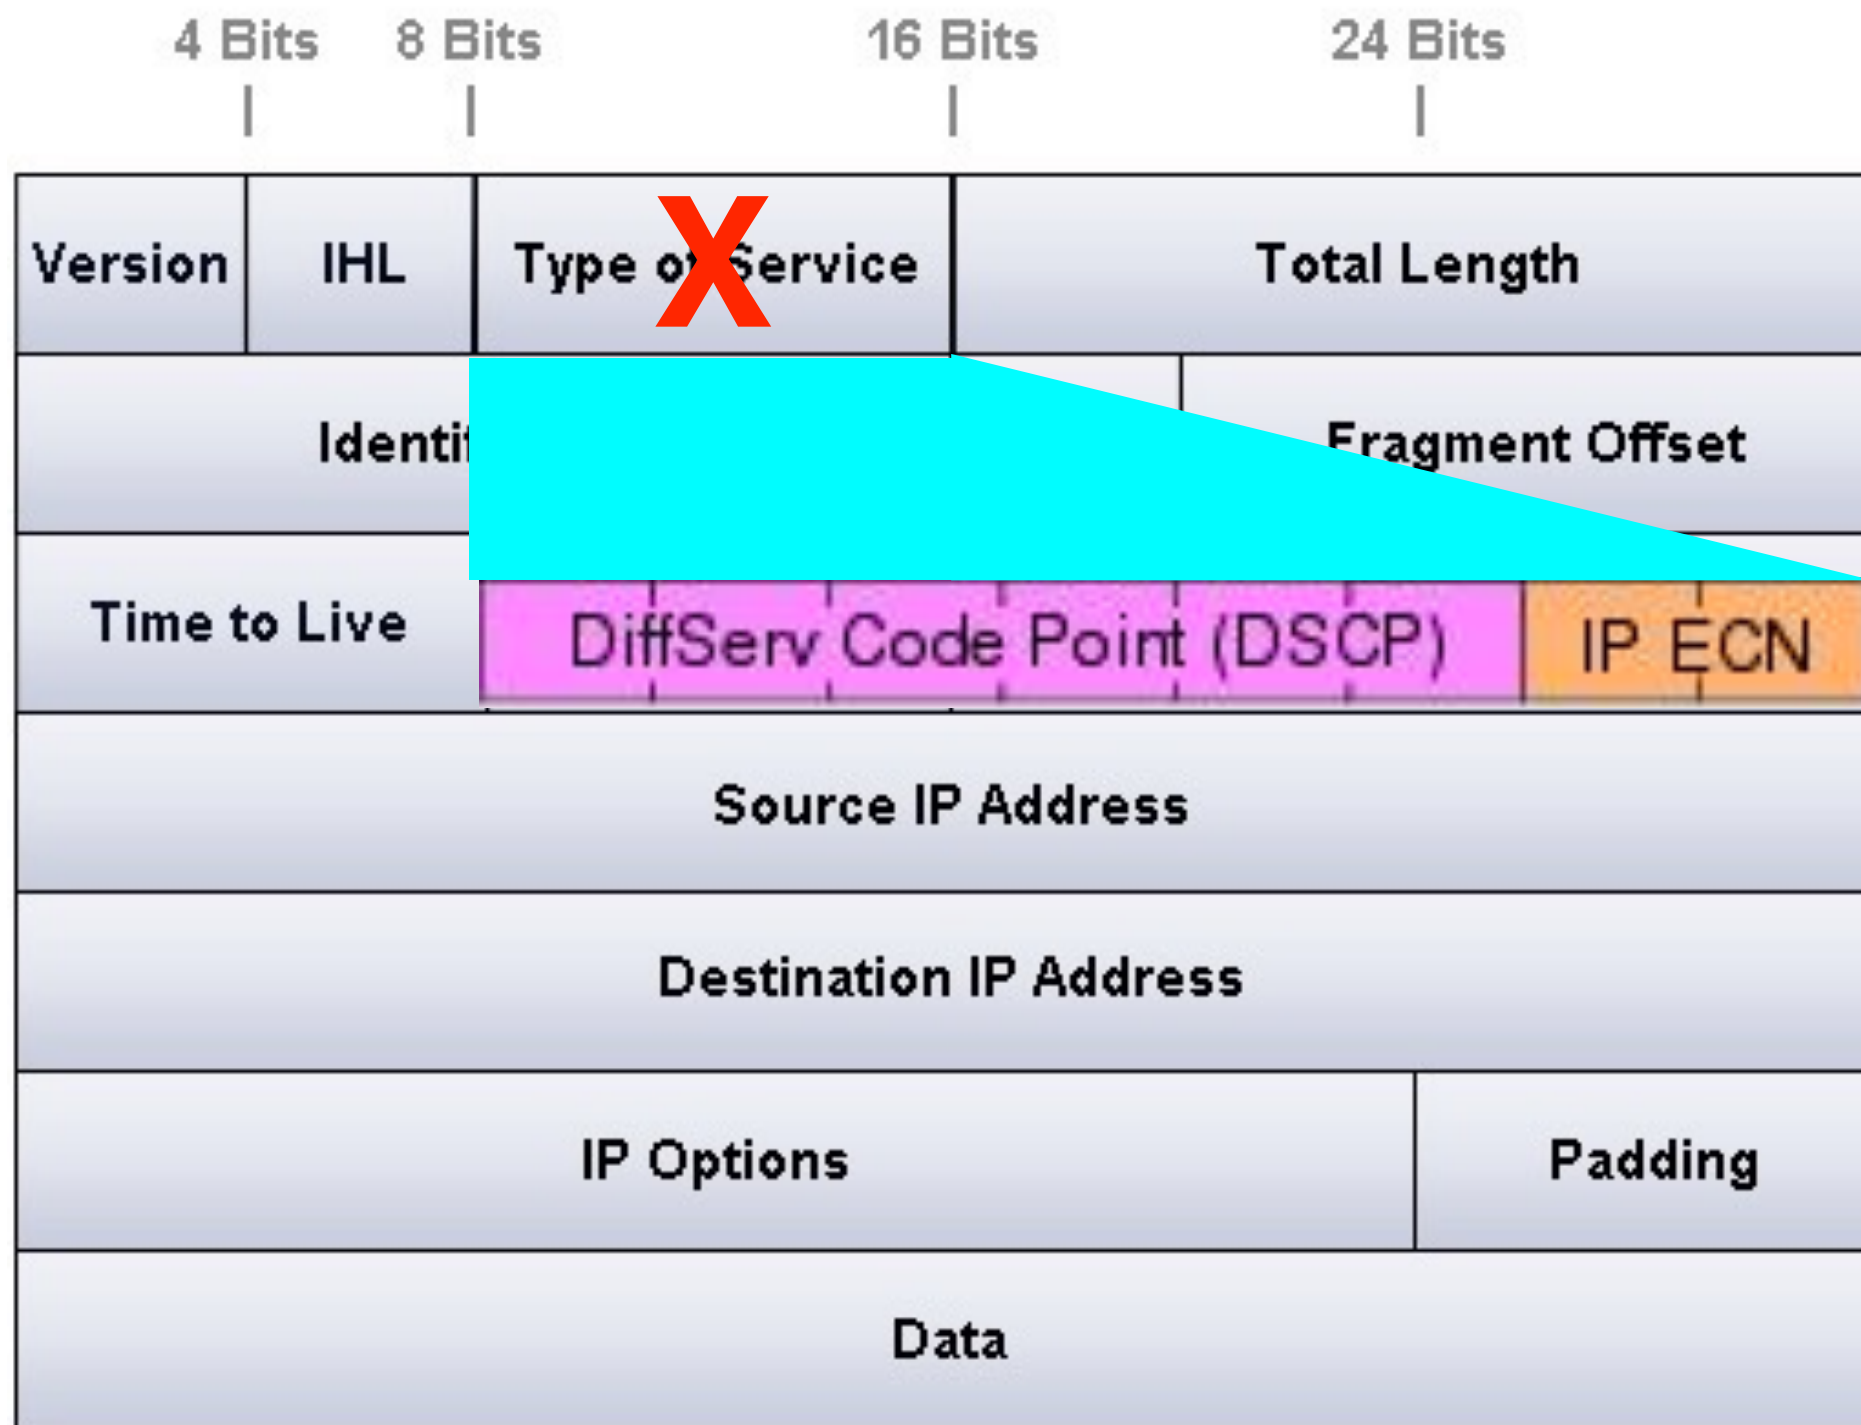
# Explicit Congestion Notification



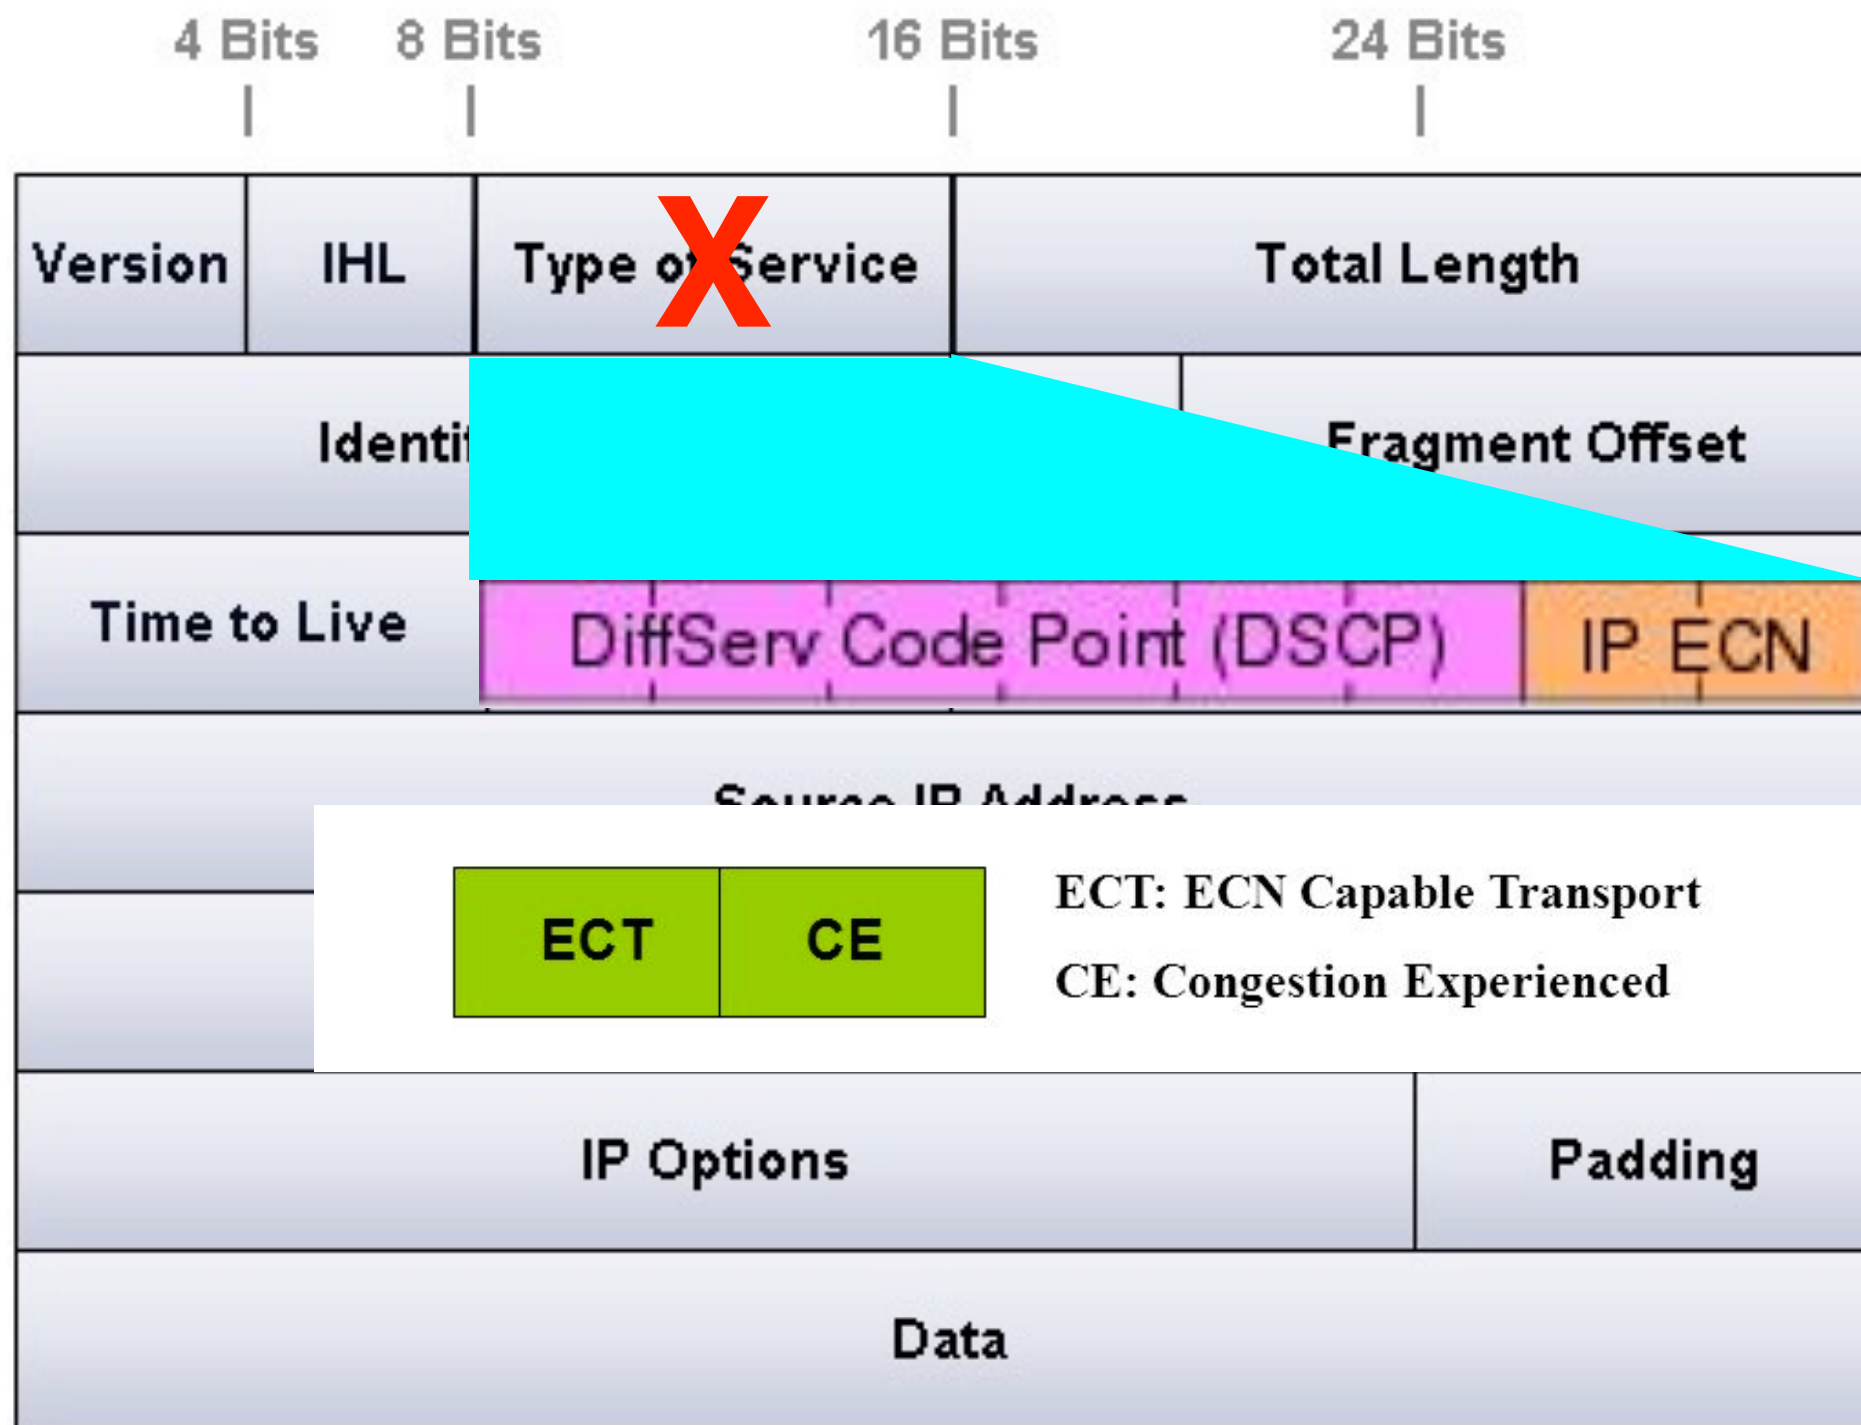
# Explicit Congestion Notification




# Explicit Congestion Notification



# Explicit Congestion Notification



# Gentle RED + ECN



# Fair Queuing

# Fair Queuing

- Independent logical queues for different traffic
  - e.g., for different source IPs
  - e.g., for different TCP connections
  - e.g., for different transport protocols



# Fair Queuing

- Independent logical queues for different traffic
  - e.g., for different source IPs
  - e.g., for different TCP connections
  - e.g., for different transport protocols
- I.e., protect traffic from competition

# Fair Queuing

- Independent logical queues for different traffic
  - e.g., for different source IPs
  - e.g., for different TCP connections
  - e.g., for different transport protocols
- I.e., protect traffic from competition
- We have figured out the hardware to do this at line rate

# Traffic Shaping

- Lots of techniques to manage traffic at various points in the network according to policy
  - expensive!
- Does this obviate the need for endpoint CC?